

VINCENT WANG-MASCIANICA

STRING DIAGRAMS FOR TEXT

Contents

0 Synopsis	7
0.1 What this thesis is about	8
0.2 Question: What is the practical value of studying language when Large Language Models exist?	10
0.3 First Reply: I don't know. Maybe explainability, maybe something else.	11
0.3.1 Objection: You're forgetting the bitter lesson.	13
0.3.2 Objection: GOFAI? GO-F-yourself.	14
0.3.3 Point of Information: What is computational irreducibility?	14
0.3.4 Objection: How does any of this improve capabilities?	15
0.3.5 Objection: Aren't string diagrams just graphs?	15
0.3.6 Objection: Pictures are heuristics, not math	16
0.4 Second Reply: LLMs don't help us understand language; how might string diagrams help?	17
0.4.1 Objection: Isn't the better theory the one with better predictions?	17
0.4.2 Objection: What's wrong with λ -calculus and sequent calculi and graphs?	19
0.5 Synopsis of the thesis	20
1 Bibliography	23
2 Background	27
2.1 Process Theories	28
2.1.1 What does it mean to copy and delete?	31
2.1.2 What is an update?	33
2.1.3 Pregroup diagrams and correlations	35
2.1.4 Equational Constraints and Frobenius Algebras	35
2.1.5 Processes, Sets, and Computers	35
2.2 An introduction to weak n-categories for formal linguists	38
2.2.1 String-rewrite systems as 1-object-2-categories	39

2.2.2	Tree Adjoining Grammars	42
2.2.3	Tree adjoining grammars with local constraints	47
2.2.4	Braiding, symmetries, and suspension	49
2.2.5	TAGs with links	54
3	String Diagrams for Text	61
3.1	Previously, on DisCoCat	62
3.1.1	Lambek’s Linguistics	62
3.1.2	Coecke’s Composition	66
3.1.3	Categorical quantum mechanics	68
3.1.4	Enter computational linguistics	71
3.1.5	I killed DisCoCat, and I would do it again.	75
3.2	Internal wirings redux	85
3.3	How do we communicate using language?	85
3.3.1	Grammars of speakers and listeners	86
3.3.2	Discrete Monoidal Fibrations	92
3.3.3	Strictified diagrams for monoidal categories	98
3.3.4	Monoidal Cofunctors	100
3.4	Communicative constraints as a cofunctor from productive to parsing grammar	101
3.5	Relating circuits and CFGs fibrationally	108
3.6	Where internal wirings come from	109
3.6.1	Discrete monoidal fibrations for grammatical functions	109
3.7	Discussion and Limitations	110
3.8	A generative grammar for text circuits	115
3.8.1	Text Circuits	115
3.8.2	A circuit-growing grammar	118
3.8.3	Sentences	119
3.8.4	Simple sentences	120
3.8.5	Modifiers	121
3.8.6	Rewriting to circuit-form	122
3.8.7	Extensions I: relative and reflexive pronouns	122
3.8.8	Extensions II: grammar equations	123
3.8.9	Extensions III: higher-order modifiers	123
3.8.10	Equivalence to internal wirings	123
3.8.11	Text circuit theorem	123
3.8.12	Related work	123

4 Continuous relations: a palette for toy models	125
4.1 Continuous Relations: A concept-compliant setting for text circuits	126
4.1.1 Why not use something already out there?	127
4.2 Continuous Relations	130
4.3 ContRel diagrammatically	131
4.3.1 Relations that are always continuous	131
4.4 Continuous Relations by examples	134
4.5 Populating space with shapes using sticky spiders	140
4.5.1 When does an object have a spider (or something close to one)?	140
4.6 Mathematician's endnotes	164
4.6.1 The category ContRel	164
4.6.2 Symmetric Monoidal structure	164
4.6.3 Rig category structure	165
4.6.4 ContRel and Rel are related by a free-forgetful adjunction	166
4.6.5 Why not Span(Top) ?	168
4.6.6 Why not a Kleisli construction on Top ?	168
4.6.7 Where is the topology coming from?	169
4.6.8 Why are continuous relations worth the trouble?	169
5 Sketches of the shape of language	171
5.1 Topological concepts in flatland via ContRel	172
5.1.1 Shapes and places	172
5.1.2 The unit interval	175
5.1.3 Displacing shapes	178
5.1.4 Moving shapes	181
5.1.5 Rigid motion	185
5.1.6 Modelling linguistic topological concepts	188
5.1.7 States, actions, manner	193
5.2 On entification, general anaphora, computers, and lassos.	200
5.3 Neural nets, and (im)possibility results for learning text circuits from data	212
5.3.1 Deterministic Neural Nets and Closed Monoidal Categories	212
5.3.2 Approximating Text Circuits with deterministic neural nets	215
5.3.3 Text circuits with unbounded depth and width	222
5.3.4 Text circuits of unbounded width in noncartesian settings	224
5.3.5 A value proposition for quantum machine learning	225
5.4 Towards learning gates that satisfy First-Order Logic specifications over boundedly finite models	227

5.4.1	Discussion	231
5.5	Modelling metaphor	233
5.5.1	Orders, Temperature, Colour, Mood	233
5.5.2	Complex conceptual structure	233
5.5.3	234

(Acknowledgements will go in a margin note here.)

0

Synopsis

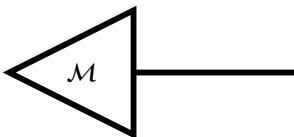


Figure 1: Let's say that the meaning of text is how it updates a model. So we start with some model of the way things are, modelled as data on a wire.

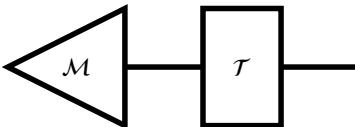


Figure 2: Text updates that model; like a gate updates the data on a wire.

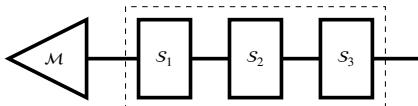


Figure 3: Text is made of sentences; like a circuit is made of gates and wires.

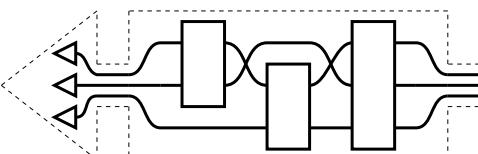


Figure 4: Let's say that *The meaning of a sentence is how it updates the meanings of its parts*. As a first approximation, let's say that the parts of a sentence are the nouns it contains or refers to. Noun data is carried by wires. Collections of nouns are related by gates, which play the roles of verbs and adjectives.

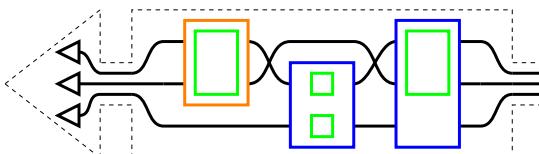


Figure 5: Gates can be related by higher order gates, which play the roles of adverbs, adpositions, and conjunctions; anything that modifies the data of first order gates like verbs.

0.1 What this thesis is about

THIS THESIS IS ABOUT STUDYING LANGUAGE USING STRING DIAGRAMS.

I am interested in using contemporary mathematical tools as a fresh approach to modelling some features of natural language considered as a formal object. Specifically, I am concerned with the compositional aspect of language, which I seek to model with the compositionality of string diagrams. Insofar as compositionality is the centrepiece of 'knowledge of language', I share a common interest with linguists, but I will not hold myself hostage to their methods, literature, nor their concern with empirical capture. I will make all the usual simplifying assumptions that are available to theoreticians, such that an oracular machine will decide on lexical disambiguation and the appropriate parse using whatever resources it wants, so that I am left to work with lexically disambiguated words decorating some formal grammatical structure. It is with this remaining disambiguated mathematical structure that I seek to state a general framework for *meaningful compositional representations of text*, in the same way we humans construct rich and interactable representations of things-going-on in our minds when we read a storybook. So if you are interested in understanding language, this thesis is an invitation to a conception of formal linguistics that's maybe worth a damn in a world where large language models exist.

OBJECTION: ISN'T THAT REINVENTING THE WHEEL?

Yes, to an extent. I am not interested in the human language faculty *per se*, so my aims differ. There are several potential practical and theoretical benefits that a fresh mathematical perspective on language enables. First, the mathematics of applied category theory allows us to unify different views of syntax, and conservatively generalise formal semantics to aspects of language that may have seemed beyond the reach of rigour, such as metaphor. Practically, the same mathematics allows us to construct interfaces between syntax/structure and semantics/implementation in such a way that we can control the former and delegate the latter by providing specifications without explicit implementation, which (for historical reasons I will explain shortly) is possibly the least-bad idea for getting at natural language understanding in computers from the bottom-up. Second, there are probably benefits to expressing linguistics in the same mathematical lingua franca that can be used to reason about linear and affine algebra CITE, first order logic CITE, causal networks CITE, signal flow CITE, electrical circuits CITE, database operations CITE, spatial relations CITE, game theory CITE, petri nets CITE, hypergraphs CITE, probability theory CITE, machine learning CITE, and quantum theory CITE, to name a few applications. At the moment, the practical achievements of language algorithms de-emphasise the structure of language, and there is no chance of reintroducing the study of structure with dated mathematics.

POINT OF INFORMATION: WHAT DO YOU MEAN BY NATURAL LANGUAGE?

Natural language is a human superpower, and the foundation of our collective achievements and mistakes as a species. By *natural language* I mean a non-artificial human language that some child has grown up speaking. English is a natural language, while Esperanto and Python are constructed languages. If you are still reading then you probably know a thing or two already about natural language. Insofar as there are rules for natural languages, it is probable that like most natural language users, you obey the rules of language intuitively without knowing what they are formally; for example while you may not know what adpositions are, you know where to place words like *to*, *for*, *of* in a sentence and how to understand those sentences. At a more complex level, you understand idioms, how to read between the lines, how to flatter, insult, teach, promise, wager, and so on. There is a dismissive half-joke that "engineering is just applied physics", which we might analyse to absurdity as "law is just applied linguistics"; in its broadest possible conception, linguistics is the foundational study of everything that can possibly be expressed.

POINT OF INFORMATION: WHAT ARE STRING DIAGRAMS?

String diagrams are a pictorial syntax for interacting processes. They are compositional blueprints that we can give semantics to – i.e. instantiate – in just about any system with a notion of sequential and parallel composition of processes. In particular, this means string diagrams may be interpreted as program specifications on classical or quantum computers, or as neural net architectures. Moreover, we can devise equations between string diagrams to govern the behaviour of each process without having to spell out a bottom-up implementation, by asking processes to interact with other processes in certain ways. The mathematical foundation of string diagrams – applied category theory – is in short, the mathematics of compositionality.

Many fields of study have developed string diagrams as informal calculational aids, unaware of their common usage across disciplines and the rather new mathematics that justifies their use; everybody knows, but it isn't common knowledge. Why is that so? Because just as crustaceans independently converge to crab-like shapes within their own ecological niches by what is called *carcinisation*, formal notation for formal theories of "real world" problem domains undergo "string-diagrammatisation" in similar isolation. Why is that so? Because our best formal theories of the real world treat complexity as the outcome of composing simple interacting parts; perhaps nature really works that way, or we cannot help but conceptualise in compositional terms. When one has many different processes sending information to each other via channels, it becomes tricky to keep track of all the connections using one-dimensional syntax; if there are N processes, there may be on the order of $\mathcal{O}(N^2)$ connections, which quickly becomes unmanageable to write down in a line, prompting the development of indices in notation to match inputs and outputs. In time, probably by doodling a helpful line during calculation to match indices, link-ed indices become link-ing wires, and string-diagrammatisation is complete.

I will demonstrate how they are used in Section 2.1.5 and define them formally in Section ???. String diagrams are a heuristically natural yet mathematically formal syntax for representing complex, composite systems. I say *mathematically formal* to emphasise that string diagrams are not merely heuristic tools backed by

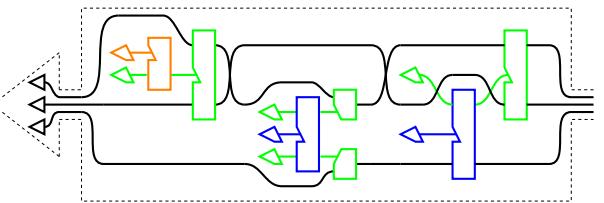


Figure 6: In practice, higher order gates may be implemented as gates that modify parameters of other gates. Grammar, and *function words* – words that operate on meanings – are in principle absorbed by the geometry of the diagram. These diagrams are natural vehicles for *dynamic semantics*, broadly construed, where states are prior contexts and sentences-as-processes update prior contexts.

a handbook of standards decided by committee: they are unambiguous mathematical objects that you can bet your life on [9? , 12, 11, 18].

0.2 *Question: What is the practical value of studying language when Large Language Models exist?*

Large language models raise questions of existential concern for the field of linguistics. More narrowly, they demand justification as to why I am writing a thesis about theoretical approaches to basic linguistics as a computer scientist in current year. Although this thesis is pure theory, I wish to address the question of practical value early because I imagine practical people are impatient.

LET ME OUTLINE THE TERMS AND STAKES. Because the field is developing so quickly, assume that everything about LLMs is prefaced with "at the time of writing". Large Language Models are programs trained – using a lot of data and a lot of compute time – to predict the next word in text, computational techniques for which have evolved from Markov n-grams to transformers [23]. This sounds unimpressive, but – in tandem with fine-tuning from human feedback in the case of chatGPT [16] – it is enough to tell and explain jokes [1], pass the SAT [20] and score within human ranges on IQ tests [21]. There is an aspect of genuine scientific and historical surprise that text-prediction can do this kind of magic. On the account of [13], computational linguistics began in a time when compute was too scarce to properly attempt rationalist, knowledge-based and theoretically-principled approaches to modelling language. Text-prediction as a task arose from a deliberate pursuit of "low-hanging fruit" as a productive and knowledge-lean alternative to doing nothing in an increasingly data-rich environment. Some observers [5] expressed concern that all this fruit would be picked bare in a generation to force a return to knowledge-based methods, but those concerns appear now to be unfounded.

I'm sure there will be many further notable developments, and to be safe I won't make any claims about what machines can't do if we keep making them bigger and feed them more data or have them interact with one another in clever ways. Nonetheless there remain limitations that seem persistent for the foreseeable future, not in terms of *capabilities*, but in terms of *explainability and safety*. These models have a tendency to hallucinate facts and (ironically, for a computer) bad arithmetic [7], and I imagine that the cycle of discovering limitations and overcoming them will continue. Despite whatever limitations exist in the state-of-the-art, it is evident to all sane observers that this is an important technology, for several reasons.

1. LLMs are a civilisational milestone technology. A force-multiplication tool for natural language – the universal interface CITE – built from abundant data and compute in the silicon age may have comparably broad, deep, and lasting impact to the conversion of abundant chemical fuel to physical energy by steam engines in the industrial revolution CITE .
2. LLMs represent a paradigm shift for humanity because they threaten our collective self-esteem, in a more pointed manner than losing at chess or Go to a computer; modifying a line of thinking from [6], LLMs

demonstrate that language and (the appearance of) complex thought that language facilitates is not a species-property for humans, and this stings on par with Darwin telling us we are ordinary animals like the rest, or Galileo telling us our place in the universe is unremarkable.

3. LLMs embody the latest and greatest case study of the bitter lesson [19]. The tragedy goes like this: there's a group of people who investigate language – from syntax and semantics to pragmatics and analogies and storytelling and slang – who treat their subject with formal rigour and have been at it for many centuries. Their role in the story of LLMs is remarkable because it doesn't exist. They were the only qualified contestants in a "let's build a general-purpose language machine" competition, and they were a no-show. Now the farce: despite the fact that all of their accumulated understanding and theories of language were left out of the process, the machine is not only built but also far exceeds anything we know how to build in a principled way out of all their hard-earned insight. That is the bitter lesson: dumb methods that use a lot of data and compute outperform clever design and principled understanding.

I will note in passing that I have an ugly duckling problem, in that I am not strictly aligned with machine learning, nor linguists broadly construed, nor mathematical linguists. I am unfortunately placed in that I feel enough affinity to have defensive instincts for each camp, but I am distanced enough from each that I am sure to suffer attacks from all sides. Perhaps a more constructive metaphor than war is that I am writing in a cooperative spirit between domains, or that I am an arbitrageur of ideas between them. With that in mind, I am for the moment advocating on behalf of pen-and-paper-and-principles linguists in formulating a two-part reply to the devastating question, and I will switch sides later for balance. First a response that answers with practical values in mind, and then a response that asserts and rests upon the distinct values of linguists.

0.3 *First Reply: I don't know. Maybe explainability, maybe something else.*

EXPRESSING GRAMMAR AS COMPOSITION OF PROCESSES MIGHT YIELD PRACTICAL BENEFITS. MOREOVER, WE WANT ECONOMY, GENERALITY, AND SAFETY FOR LANGUAGE MODELS, AND WE CAN POTENTIALLY DO THAT WITH FEW TRADEOFFS IF WE USE THE RIGHT FRAMEWORK.

Simplified, half of the problem of learning language is learning the meaning of words. The meanings change over time and are context-dependent, and the words are always increasing in number. Encoding these meanings by hand is a sisyphean task. Data-driven learning methods are a good fit: the patterns to be learnt are complex and nebulous, and there is a lot of data. However, data-driven methods may be weaker at the second half of the problem: learning and executing the composition of meanings according to syntax. We can see just how much weaker when we consider the figures involved in 'the poverty of the stimulus'.

In short, this famous problem is the observation that humans learn language despite having very little training data, in comparison to the complexity of the learned structure. It is on the basis of this observation

– alongside many others surrounding language acquisition and use – that Chomsky posits [3] that language is an innate human faculty, the development of which is less like effortfully going to the gym and more like effortlessly growing arms you were meant to have. The explanation goes like this: we can explain how a complex structure like grammar gets learnt from a small amount of data if everyone shares an innate Universal Grammar with a small number of free parameters to be learned. Whether or not the intermediate mechanism is a species-property of humans, the point is that we humans get a very small amount of input data, that data interacts with the mechanism in some way, and then we know a language. So, now that there are language-entities that are human-comparable in competence, we can make a back-of-the-envelope approximation of how much work the intermediate mechanism is doing or saving by comparing the difference in how much data and compute is required for both the human and for the machine to achieve language-competence. Humans get about 1.5 megabytes of data [14], 90 billion neurons [8], and an adult human consumes around 500 calories per day for thinking, for let's say 20 years of language learning. Rounding all values *up* to the closest order of magnitude, this comes to a cost metric of 10^{29} bits \times joules \times neurons. PaLM – an old model which is by its creators' account the first language model to be able to reason and joke purely on the basis of linguistic ability and without special training [4, 15] – required 780 billion training tokens of natural language (let's discount the 198 gigabytes of source code training data), which we generously evaluate at a rate of 4 characters per token [10] and 5 bits per character. The architecture has 540 billion neurons, and required 3.2 million kilowatt hours of energy for training [22]. Rounding values for the three units down *down* to the nearest order of magnitude comes to a cost metric of 10^{41} bit-joule-neurons. Whatever the human mechanism is, it is responsible for an order of magnitude in efficiency *give or take an order of magnitude of orders of magnitude*. It's possible that over time we can explain this difference away by various factors such as the efficiency of meat over minerals, separating knowledge of the world from knowledge of language, more efficient model architectures, or the development of efficient techniques to train new language models using old ones CITE . One thing is clear: if it is worth hunting a fraction of a percent of improvement on a benchmark, forget your hares, a 10{10 factor is a stag that's worth cooperating to feast on.

THE LINGUISTIC STRATEGY FOR HUNTING THE STAG STARTS WITH WHAT WE KNOW ABOUT HOW THE MECHANISM BETWEEN OUR EARS WORKS WITH LANGUAGE. The good news is that the chief methodology of armchair introspection is egalitarian and democratic. The bad news is that it is also anarchistic and hard-by-proximity; we are like fish in water, and it is hard for fish to characterise the nature of water. So the happy observations are difficult to produce and easily verified, and that means there are just a few that we know of that are unobjectionably worth taking into account. One, or *the* such observation is *systematicity*. Systematicity CITE refers to when a system can (generate/process) infinitely many (inputs/outputs/expressions) using finite (means/rules/pieces) in a "consistent" (or "systematic") manner – some tautologies require geniuses like Fodor to put into words. Like pornography, examples are easier than definitions. We know finitely many words but we can produce and understand infinitely many texts; we can make infinitely many lego sculp-

tures out of finitely many types of pieces; we can describe infinite groups and other mathematical structures using finitely many generators and relations; in the practical domain of computers, systematicity is synonymous with programmability and expressibility.

The concepts of systematicity and compositionality are deeply linked, because the only way we know how to achieve systematicity in practice is composition. Frege's initial conception of compositionality CITE was borne of meditations on language, and states that a whole is the sum of its parts – some tautologies require geniuses like Frege to put into words. Later conceptions of compositionality CITE , the most notable deviation arising from meditations on quantum theory, are the same as the original, modulo variations on the formal definitions of parts and the method of summation. So there is our starting point: language is systematic and systematicity is the empirical surface of compositionality as far as we know, so compositionality is probably part of the solution to the problem of the stimulus, if not most of it.

The reasoning I report above should clarify why some folks don't think large language models have anything to do with language. The issue with purely data-driven architectures is either that we know immediately that they cannot be operating upon their inputs in a compositional way, or even if they do appear to be doing compositional things, their innards are too large and their workings too opaque to tell with confidence. Insofar as the task of learning language splits between learning meanings and learning the compositional rules of syntax that give rise to systematicity, I hope the framework I present in this thesis can be a proposal to split the cake sensibly between the two halves of the problem: meanings for the machines, compositionality for the commons. Syntax is still difficult and vast, but the rules are finite and relatively static. We can break the black-box by reexpressing syntax as the composition of smaller black-boxes. We all stand to benefit: we may give machines an easier time – now they only have to learn the meanings of words well – and we can have confidence that the internal representations of the machine – their "mind's eye" – contains something we can probe and understand.

0.3.1 *Objection: You're forgetting the bitter lesson.*

The bitter lesson is so harsh and often-enough repeated that this viewpoint is worth addressing proactively. The caveat that saves us is that the curse of expertise applies only to the object-language of the problem to be solved, not model architectures. We agree that qualitative improvements in problem-solving ability rarely if ever arise from encoding expert knowledge of the problem domain. Instead – and we see this historically CITE [deep, decision, qlearn, GAN, transformers] – these improvements come from architectural innovations, which means altering the parts and internal interactions of a model: changing *how* it thinks rather than *what* it thinks, to paraphrase Sutton's original prescription. These structural changes are motivated by understandings (at varying degrees of formality) of the "geometry of the problem" CITE . The value proposition here is that with an appropriate mathematical lingua franca for structure, composition, and interaction, we can mindfully design rather than stumble upon the "meta-methods" Sutton calls for, allowing experts to en-

code *how* machines think and discover rather than *what*. I hope to demonstrate in Section ?? how importing compositional and structural understanding from linguistics to machine learning via string diagrams might allow us to cheat the bitter lesson in spirit while adhering to the letter.

0.3.2 *Objection: GOFAI? GO-F-yourself.*

Hostility (or at least indifference) to symbolic approaches is a stance espoused by virtually all of modern machine learning, and for good reasons. This stance is worth elaborating and steelmanning for pen-and-paper-people in the context of language. First, many linguistic phenomena are nebulous CITE : the boundary of a simile is like that of a cloud, not sharp like the boundary of a billiard ball. Second, linguistic phenomena are complex, dynamic, and multifactorial: there are so many interacting mechanisms and forces in the production and comprehension of language that even if we do have crisp mathematical models for all the constituent processes we are still left with something computationally irreducible. These two points together weakly characterise the kinds of problem domains where machine learning shines. It is just fact that LLM outputs today conform to any reasonable understanding of syntax, semantics, pragmatics, conversational implicature, and whatever else we have theorised. It is just a fact that they produce better poetry and humor than anything we could explicitly program according to our current understanding. It is almost a natural law that they will only get better from here. So what good are pen-and-paper theories as far as practical applications are concerned? To borrow terms from concurrency, there is already plenty of liveness, what is needed is more safety; liveness is when the program does something good, and safety is a guarantee it won't do something bad. For example, there is ongoing work in integrating LLMs with structured databases for uses where facts and figures matter CITE ; there is still a need for safeguards to prevent harmful outputs CITE and adversarial attacks like prompt injection CITE ; while LLMs give a very convincing impression of reasoned thought, we would like to be sure if ever we decide to use such a machine for anything more than entertainment, like deciding on a medical course of treatment CITE or making decisions with financial consequences CITE . The good news is that symbolic-compositional theories are the right shape for safety concerns, because they can be picked apart and reasoned about CITE . It is clear however that symbolic-compositional approaches by themselves are nowhere near achieving the kind of liveness LLMs have. Therefore, the direction of progress is synthesis.

0.3.3 *Point of Information: What is computational irreducibility?*

Computational irreducibility CITE , elsewhere called "type 2" problems CITE refers to a special kind of computational difficulty where the explanation of a system amounts to a total computational simulation of it. This is best understood by example. We have a closed-form mathematical expression to shortcut the computation of the evolution of a system of two point masses under gravity, but we have no such shortcut for the three-body problem; the best we can do is simulate the system's evolution, and the lesson is that even for very simple

systems, it is possible that no amount of causal-mechanistic understanding will simplify computational simulation.

0.3.4 *Objection: How does any of this improve capabilities?*

It's not meant to. The core value proposition for synthesis is explainable AI, which operates in a manner we can analyse, and if appropriate, constrain. For this purpose, merely knowing *what* a deep-learning model is thinking is not enough¹: solving symbol-grounding alone is a necessary but insufficient component. For instance, merely knowing what the weights of subnetworks of an image classification model represent CITE does not meet our requirement of an understanding of the computations that manipulate those representations. Moreover, purely data-driven methods to control the computation may incur ethical costs CITE , to say nothing of the potential harm that may result from a poorly safeguarded model CITE . Add to this the ever-growing dirty laundry lists of AI models failing CITE in inhuman ways, and the task of incorporating compositionality – a formal understanding of *how* models learn and reason – gains urgency.

The work in this thesis builds upon a research programme – DisCoCat, elaborated in Section ?? – which lies somewhere in the middle of a duality of approaches to merging connectionism and symbolic-composition². It is, to the best of my knowledge, the only approach that explicitly incorporates mathematically rigorous compositional structures from the top-down alongside data-driven learning methods from the bottom-up. Fortifying this bridge across the aisle requires a little give from both sides; I ask only that reader entertain some pretty string diagrams.

0.3.5 *Objection: Aren't string diagrams just graphs?*

Yes and no! This point is best communicated by a mathematical koan. Consider the following game between two players, you and me. There are 9 cards labelled 1 through 9 face up on the table. We take turns taking one of the cards. The winner is whoever first has three cards in hand that sum to 15, and the game is a draw if we have taken all the cards on the table and neither of us have three cards in hand that sum to 15. I will let you go first. Can you guarantee that you won't lose? Can you spell out a winning strategy? If you have never heard this story, give it an honest minute's thought before reading on.

The usual response is that you don't know a winning strategy. I claim that you probably do. I claim that even a child knows how to play adeptly. I'll even wager a drink that you have played this game before. The game is Tic-Tac-Toe, also known as Naughts-and-Crosses: it is possible to arrange the numbers 1 to 9 in a 3-by-3 magic square, such that every column, row, and diagonal sums to 15.

The lesson here is that choice of representations matter. In the mathematical context, representations matter because they generalise differently. On the surface, here is an example of two representations of the same platonic mathematical object. However, Tic-Tac-Toe is in the same family as Connect-4 or 5-in-a-row on an unbounded grid, while the game with numbered cards generalises to different variants of Nim. That they

¹ I recount the following from CITE , which argues that symbol-grounding is solvable from data alone, and in the process surveys the front of the symbol-grounding problem in AI: the issue of whether LLMs encode what words refer to and mean. On the account of [2], the performance of current LLMs is a form of Chinese Room [17] phenomenon, so no amount of linguistic competence can be evidence that LLMs solve the symbol-grounding problem. However, the available evidence appears to suggest otherwise – large models converge on word embeddings for geographical place names that are isomorphic to their physical locations CITE . Since we know that brain activity patterns encode words in a manner that facilitates analogical reasoning in an abstract conceptual space CITE , extrapolating the ability of LLMs to encode analogical representations would in the limit suggest that LLMs encode meanings in a way isomorphic to how we do – at least for individual tokens.

² The investigation of the common ground between symbolic-composition and connectionism takes on, I suggest, essentially two, dual forms. The first kind uses connectionist methods to simulate symbolic-composition, which we can see the beginnings of in LLMs by examples such as chain-of-thought reasoning . The second kind is the inverse, where connectionist architectures are organised and reasoned with by symbolic-compositional means. Some examples of the first include implementing data structures as operations on high-dimensional vectors, taking advantage of the idiosyncrasies of linear algebra in very high dimension CITE , or work that explores how the structure of word-embeddings in latent space encode semantic relationships between tokens CITE . Some examples of the second include reasoning about the capability of graph neural networks by identifying their underlying compositional structure CITE , or architectures explicitly designed to instantiate symbolic-compositional structures using neural nets as constituent parts, such as GANs CITE and gradient boosted decision trees CITE .

coincide in one instance is a fork in the path. In the same way, viewing string diagrams as "just graphs" is taking the wrong path, just as it would be a mistake to dismiss graphs as "just sets of vertices and edges". String diagrams are indeed "just" a special family of graphs, just as much as prime numbers are special integers and analytic functions are special functions. Throughout this thesis I will be re-presenting familiar and unfamiliar things in string diagrams, so I request the reader to remember the koan and keep an open mind.

In a broader context, representations matter for the sake of improved human-machine relations. These two representations are the same as far as a computer or a formal symbol-pusher is concerned, but they make world of difference to a human native of meatspace. We ought to swing the pendulum towards incorporating human-friendly representations in language models, so that we may audit those representations for explainability concerns. As it stands, there is something fundamentally inhuman and behavioural about treating the production of language as a string of words drawn from a probability distribution. I don't know about you, but I tend to use language to express pre-existing thoughts in my head that aren't by nature linguistic. Even if we grant that the latent space of a data-driven architecture is an analog for the space of internal mental states of a human user of language, how can we know whether the spaces are structurally analogous to the extent that human-machine partnership through the interface of natural language is safe? So here again is the possible solution: by composing architectures in the shape of language from the start, we can attempt guarantees that the latent-space representation of the machine is built up in the same way we build up a mental representation when we read a book or watch a film. I'll sketch how to approach this in Section ??.

0.3.6 *Objection: Pictures are heuristics, not math*

A deeper objection here is that diagrams do not look like serious mathematics. Later I will give ample space to show how they are serious, but the reasons behind this rather common prejudice are worth elaborating. This is the wound Bourbaki has inflicted. Nicolas Bourbaki is a pseudonym for a group of French mathematicians, who wrote a highly influential series of textbooks. It is difficult to overstate their influence. The group was founded in the aftermath of the First World War, around the task of writing a comprehensive and rigorous foundations of mathematics from the ground up. The immediate *raison-d'être* for this project was that extant texts at the time were outdated, and the oral tradition and living history of mathematics in institutions of learning were decimated by the deaths of mathematicians at war. In a broader historical context CITE, Bourbaki was a reactionary response to the crisis in the foundations of mathematics at the beginning of the century, elicited by Russell's paradox. Accordingly, their aims were rationalist, totalitarian, and high-modernist CITE, favouring abstraction and disdaining visualisation, in line with their contemporary artistic and musical fashions. Consequently, Bourbaki's Definition-Proposition-Theorem style of mathematical exposition is an evolved form of Euclid's that eschews intuition and example, a format pretending at timelessness that requires years of initiation to effectively read and write, and remains *de rigueur* for rigour today in dry mathematics textbooks. The deeper objection arises from the supposition that serious mathematics ought

to look arcane and difficult, as most mathematics exposition after Bourbaki is. The reply is that it need not be so, and that it was not always so! The Bourbaki format places emphasis and prestige upon the deductive activity that goes into proving a theorem, displacing other aspects of mathematical activity such as constructions, algorithms, and taxonomisation [CITE](#). These latter aspects are better suited for the nebulous subject matter of natural language, which not lend itself well to theorems, but is a happy muse for mathematical play.

0.4 Second Reply: LLMs don't help us understand language; how might string diagrams help?

Another way to deal with the devastating question of LLMs is to reject it, on the basis that using or understanding LLMs is completely different from understanding language, and language is worth understanding in its own right. To illustrate this point by a thought experiment, what would linguistics look like if it began today? LLMs would appear to us as oracles; wise, superhumanly capable at language, but inscrutable. Similarly, most people effortlessly use language without a formal understanding which they can express. So the fundamental mystery would remain unchanged. Understanding how an LLM works cannot help: to borrow a thought from [CITE](#), suppose you knew the insides of a mechanical calculator by heart. Does that mean you *understand* arithmetic? At best, obliquely: implementing a computer for ideal arithmetic means compromises; the calculator is full of inessentialities and tricks against the constraints of physics. You would not know where the tricks begin and the essence ends. Similarly, suppose you knew every line of code and every piece of data used to train an LLM; does that mean you understand how language works? How does one delineate what is essential to language, and what is accidental? So the value proposition to establish is how string diagrams come into the picture for the linguist who is (definitionally) concerned with understanding how language works. Let's entertain one more objection from the practical reader and one objection from the theoretical reader before formulating a reply.

0.4.1 Objection: Isn't the better theory the one with better predictions?

Whether LLMs are even a theory of language is a best debatable. There are various criteria – not all independent – that are arguably necessary for something to qualify as an explanatory theory, and while LLMs suffice (or even excel) at some, they fail at others. Empirical adequacy – the ability of theory to account for the available empirical data and make good predictions about future observations – is one such criterion [CITE](#), and here LLMs excel. In contrast to the idealised and partial nature of formal theories, the nature of LLMs is that they are trained on empirical data about language that captures the friction of the real world. So, in terms of raw predictive power, we should naturally expect the LLMs to have an advantage over principled theories. They are so good at empirical capture that to some degree they automatically satisfy the related criteria of coherence – consistency with other established linguistic theories – and scope – the ability

To illustrate the insufficiency of empirical capture to make a theory, consider the historical case study of models of what we now call the solar system. The Ptolemaic geocentric model of the solar system was more empirically precise than the heliocentric Copernican, even though the latter was "more correct" CITE . This should not be surprising, because Ptolemaic epicycles can overfit to approximate any observed trajectory of planets. It took until Einstein's relativity to explain the precession of perihelion of mercury, which at last aligned theoretical understanding with empirical observation. But Newton's theory of gravity was undeniably worthwhile science, even if it was empirically outperformed by its contemporaries. Consider just how divorced from reality Newton was: Aristotelian physics is actually correct on earth, where objects don't continue moving unless force is continually supplied, because friction exists. It took a radical departure from empirical concerns to the frictionless environment of space in order obtain the simplified and idealised model of gravity that is the foundation of our understanding of the solar system and beyond. The lessons as I see them are as follows. First, aimed towards some advocates of theory-free approaches, we should bely the order to evacuate linguistics departments because performance is to some degree orthogonal to understanding. In fact, the scientific route of understanding involves simplified and idealised models that ignore friction, and will necessarily suffer in performance while maturing, so one must be patient. Second, aimed towards some theoreticians, haphazard gluing together of different theories and decorating them with bells-and-whistles for the sake of fitting empirical observation is no different than adding epicycles; one must either declare a foundational or philosophical justification apart from empirical capture (which machines are better at anyway), or state outright that it's just a fun hobby. Third, while explainability studies in terms of analysing the behaviour of subnetworks is practically important, the whole approach seems fundamentally misguided; imagine an epicyclist explaining the precession of mercury's perihelion by pointing at a collection of epicycles and calling it a "distributed representation".

to capture a wide range of phenomena. But while empirical capture is necessary for explanatory theories, it is insufficient.

There are several criteria where the adequacy of LLMs is unclear or debatable. Fruitfulness is a sociological criterion for goodness of explanatory theories, in that they should generate new predictions and lead to further discoveries and research CITE . While they are certainly a potent catalyst for research in many fields even beyond machine learning, it is unclear for now how they relate to the subject matter of linguistics. Whether they satisfy Popper's criterion of falsifiability is as of yet not determined, because it is not settled how to go about falsifying the linguistic predictions of LLMs, or even express what the content of a theory embodied by an LLM is. The closest examples to falsifiability that come to mind are tests of LLM fallibility for reasoning and compositional phenomena CITE , or their weakness to adversarial prompt-injections CITE , but these weaknesses do not shed light on their linguistic competence and "understanding" directly.

Now the disappointments. LLMs are far from simple, and simplicity (Occam's Razor) is an ancient criterion for the goodness of explanation. Moreover, they fail at providing explanatory mechanisms CITE , and they do not unify or subsume our prior understandings CITE . The first two points are unobjectionable, so I will briefly elaborate on the criterion of unification and subsumption of prior understandings, borrowing a framework from cognitive neuroscience. A common methodology for investigating cognitive systems is Marr's 3 levels CITE (poorly named, since they are not hierarchical, but more like interacting domains.) Level 1 is the computational theory, an extensional perspective that concerns tasks and functions: at this level on asks what the contents and aims of a system are, to evaluate what the system is computing and why, respectively. Level 2 is representation and algorithm, an intensional perspective that concerns the representational format of the contents within the system, and the procedures by which they are manipulated to arrive at outcomes and outputs. Level 3 is hardware, which concerns the mechanical execution of the system, as gears in a mechanical calculator or as values, reads, and writes in computer memory. To be fair, in the case of LLMs, we understand well the nature of computational theory level, at least in their current incarnation as next-token-predictors, which is a narrow and clear task. Furthermore, we understand the hardware level well, from the silicon going up through the ladder of abstraction to software libraries and the componentwise activity of neural nets. There is something deeply wrong about our understanding – if we can call it that – of level 2. We have working understandings of several aspects about this level. We know something about the nature of internal representations in neural nets both in terms of semantic encoding within weight distributions CITE and of token-embeddings in latent space CITE . We can explain how transformer models work in terms of attention mechanisms and lookback CITE , which serves as working understanding of the procedural aspect of LLMs. We also understand mathematically how it is that these models are trained using data to produce the outputs they do. The deep problem is that in spite of these understandings which should jointly cover all of level 2, we only obtain explanations at the wrong level of abstraction for the purposes we care about CITE . Level 2 is in a sense the important level to get right for the purposes of explainability, auditability, and extension, since it is at the level of representation and procedure that we can investigate internal

structure and match levels of abstraction across domains. Since the three levels interact, the challenge is to slot in a story about level 2 that coheres with what we already know about levels 1 and 3. I claim that we can go about this challenge using string diagrams as a lingua franca for mathematical linguists and machines.

0.4.2 *Objection: What's wrong with λ -calculus and sequent calculi and graphs?*

There is nothing wrong with punchcard machines either, as far as computability is concerned. String diagrams, and applied category theory more broadly, are a good metalanguage for formal linguistics. The usual choice of set theory is not well-suited for complex and interacting moving parts. The chief drawback is that set theory requires bottom-up specifications, so for instance if you want to specify a function, you have to spell out how it behaves on the domain and codomain, which means spelling out what the innards of the domain and codomain are; to specify a set theoretic model necessitates providing complete detail of how every part looks on the inside³. As you may already know, this representation-dependency is a bureaucratic nightmare when dealing with a complex system. This leads to at least three problems.

1. The sociological problem is that this makes things difficult to understand unless you have invested a lot of time into mathematics in general.
2. Interoperability is tricky. When a programmer wants to use a data structure or algorithm, they do not always write it from scratch or copy code from stackoverflow or get their LLM to do it for them; they may use a library that provides them structures and methods they can call without worrying about how those structures and methods are implemented all the way down. However, if you building a complex theory by spelling out implementations set-theoretically from the start, incorporating a new module from elsewhere becomes difficult if that module has encoded things in sets differently. A lot of busywork goes into translating foundations of formalisms at an analogous level to machine code, which is time better spent building upwards and outwards. A computer scientist might say that some abstraction is needed, and being one, I say so.
3. Third, and related to the second, is that set-theory is not the native language for the vast majority of practical computation. Often in the design of complex theories, we do not care about how precisely representations are implemented, instead we only care about placing constraints or guarantees on the behaviour of interacting interactions – that is, we care about operational semantics.

The problems I have mentioned above are obstacles, and I hope to show that using applied category theory as a metalanguage may be a solution. A broad theme of this thesis is to illustrate the economy and reach of applied category theory for dealing with compositional phenomena. Our capacity for language is one of the oldest and sophisticated pieces of compositional technology, maybe even the foundation of compositional thought. So, linguists are veteran students of compositionality and modularity. How does syntax

³ This is an innate feature of set theory. Consider the case of the cartesian product of sets, one of the basic constructions. $A \times B$ is the "set of ordered pairs" (a, b) of elements from the respective sets, but there are many ways of encoding that are equivalent in spirit but not in syntax; a sign that the syntax is a hindrance. What we really want of the product is the property that $(a, b) = (c, d)$ just when $a = c$ and $b = d$. Now here is a small sampling of different ways to encode an ordered pair. Kuratowski's definition is

$$A \times B := \left\{ \{\{a\}, \{a, b\}\} \mid a \in A, b \in B \right\}$$

Which could have just as easily been:

$$A \times B := \left\{ \{a, \{a, b\}\} \mid a \in A, b \in B \right\}$$

And here is Wiener's definition:

$$A \times B := \left\{ \{\{a, \emptyset\}, b\} \mid a \in A, b \in B \right\}$$

Novel contributions:

- Section ?? demonstrates how string-diagrammatic reasoning allows for graphical proofs of strong equivalences between typological, string-production, and further strong equivalence to a fragment of tree-adjoining grammars.
- Text diagrams and text circuits lie at the heart of the above correspondences and of this thesis, which we introduce and investigate in Section ?? in an abridged re-presentation of CITE , culminating in a proof relating the expressive capacity of text circuits to a controlled fragment of English that serves as evidence that text circuits are a natural metalanguage for grammatical relationships that make no extraneous distinctions.
- In Section ??, moving towards applications, I introduce the category of continuous relations, to set a mathematical stage upon which we can build toy models, expanding upon my previous work on linguistically compositional spatial relations CITE towards modelling mechanical systems and containers.
- I mathematically investigate the possibilities and limitations of textual modelling with text circuits on classical and quantum computers in Section ?? by examining the limitations of cartesian monoidal categories for modelling text circuits, taking the universal approximation theorem into account.
- In Section ??, I extend the string-diagrammatic techniques used to prove correspondences between different syntactic theories to text circuits provides a framework for the formal, conceptually-compliant modelling of textual metaphor.
- I demonstrate a formal connection between tame topologies and tensed language in Section ??, which extends to a formal framework to model narratives as database rewrites in Section ??.

compose meaning? How do the constraints and affordances of language interact? The discipline embodies a encyclopaedic record of how compositionality works "in the field", just as botanists record flowers, early astronomers the planetary motions, or stamp collectors stamps. But a disparate collection of observations encoded in different formats does not a theory make; we will inevitably wish to bring it all together. I show that we can progress towards this aim using formal diagrams that our visual cortices have built-in rules to manipulate, and that allow us to work at the level of abstraction we choose, so that we may easily incorporate other modules and find implementations in a variety of settings.

To summarise the first value proposition, string diagrams are an aesthetic, intuitive, flexible, and rigourous metalanguage syntax that gives agency to the modeller by operating at a level of abstraction of their choice. In this vein, theories of syntax expressed in terms of string diagrams makes it easier to reason about expressive equivalence between theories at a compositional level. More precisely, a theory of syntax is expressed as a finitely presented symmetric monoidal category, and relationships between theories are expressed as symmetric monoidal functors, which are generalised structure-preserving maps. The upshot of reasoning in this way is that equivalences are established at a structural level between the atomic components of corresponding theories, which lets us do some crazy things.

TL;DR OF INTRODUCTION: LLMs do not explain language, and formal linguists "explain" language using the mathematical equivalent of punchcard machines. Both sides stand to gain from synthesis, but synthesis requires a shared mathematical metalanguage. I propose string diagrams and applied category theory as a candidate, and the rest of the thesis is about justifying the proposal.

0.5 *Synopsis of the thesis*

Chapter 2 provides the relevant background and foundations for category theory, machine learning, and formal syntax for this thesis, which lives at the intersection. The ideas required from the parent fields will be basic, so the exposition is meant to get readers across disciplines on the same page, not impress experts. For string diagrams I will first provide a primer for how process-theoretic reasoning with string diagrams work, by example. On the category theoretic end, I will recount symmetric monoidal categories as the mathematical objects that string diagrams are syntax for, as well as provide a working understanding of PROPs CITE and n-categories as formalised by CITE , which provide a metalanguage for specifying families of string diagrams. Once the reader is happy with string diagrams, for machine learning I will just introduce how deep neural nets and backpropagation work in string-diagrammatic terms to provide a foundation of formal understanding, and I will explain the mathematical and real-world reasons why deep learning is so powerful. For formal linguistics, I will sketch out a partial history of categorial linguistics in general, along the way briefly recasting CITE in more modern mathematical terminology to justify string diagrams as a generalisation of Montague's original conception of syntax and semantics.

Chapter 3 is about string diagrams for formal syntax. Here I recount context-free, pregroup, and tree-adjoining grammars to the reader, recast them string-diagrammatically, and relate them by means of discrete monoidal fibrations, a piece of mathematics I will develop. Then we (re)introduce text circuits as the common structure between those different theories of grammar that abstracts away differences in linear syntactic presentation while conserving a core set of grammatical relations. During my DPhil, I wrote a paper [CITE](#) in collaboration with Jonathon Liu and Bob Coecke which introduced text circuits in a pedestrian way, and characterised their expressive capacity with respect to a controlled fragment of English. I will just recover the main beats of that paper, this time using the metalanguage of n-categories.

Chapter 4 sets a mathematical stage for us to model and calculate using text circuits, for which purpose I introduce the category of continuous relations **ContRel**, a naïve generalisation of the category of continuous maps between topological spaces. **ContRel** is new, in the sense that the category-theoretically obvious approaches to obtaining such a category either do not work or yield something different. This section culminates in formal semantics for topological concepts such as `inside` which underpin the kinds of schematic doodle cartoons we might draw on paper to illustrate events occurring in space.

Chapter 5 is a formal invitation to playtime for the reader who gets that far. I don't expect that I've explored any novel linguistic phenomena, and I don't think I've invented any substantially new mathematics. All I've done is a form of intellectual arbitrage, putting tools from one field to work in another. To properly give weight to my claim that string diagrams and category theory are a good metalanguage for linguistics as a whole, it is necessary to demonstrate breadth. So, I model linguistic topological concepts; I give a mathematical setting for the study of generalised anaphora that reference any meaningful part of text; I provide formal semantics for the container metaphor in particular and textual metaphors in general; I sketch a formal correspondence between tensed language and tame topologies that extends to formally reckoning with narrative structure. All of this is to show that the methods I use are flexible and not doctrinal. I am not interested in whether these topics have been mathematicised more thoroughly and deeply before; what I care to demonstrate is that a little category theory and some imagination can go a long way.

Finally, I close with a discussion and prospectus. For the convenience of the reader, bibliographies are placed at the end of each chapter. Corrections, comments, and suggestions are welcome at vincentwangsemailaddress@gmail.com. I hope you enjoy the read, or if nothing else, I hope you like my diagrams!

1

Bibliography

- [1] Matthias Bastian. Google PaLM: Giant language AI can explain jokes, April 2022.
- [2] Emily M. Bender and Alexander Koller. Climbing towards NLU: On Meaning, Form, and Understanding in the Age of Data. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics*, pages 5185–5198, Online, 2020. Association for Computational Linguistics.
- [3] Noam Chomsky. *New Horizons in the Study of Language and Mind*. Cambridge University Press, Cambridge, 2000.
- [4] Aakanksha Chowdhery, Sharan Narang, Jacob Devlin, Maarten Bosma, Gaurav Mishra, Adam Roberts, Paul Barham, Hyung Won Chung, Charles Sutton, Sebastian Gehrmann, Parker Schuh, Kensen Shi, Sasha Tsvyashchenko, Joshua Maynez, Abhishek Rao, Parker Barnes, Yi Tay, Noam Shazeer, Vinodkumar Prabhakaran, Emily Reif, Nan Du, Ben Hutchinson, Reiner Pope, James Bradbury, Jacob Austin, Michael Isard, Guy Gur-Ari, Pengcheng Yin, Toju Duke, Anselm Levskaya, Sanjay Ghemawat, Sunipa Dev, Henryk Michalewski, Xavier Garcia, Vedant Misra, Kevin Robinson, Liam Fedus, Denny Zhou, Daphne Ippolito, David Luan, Hyeontaek Lim, Barret Zoph, Alexander Spiridonov, Ryan Sepassi, David Dohan, Shivani Agrawal, Mark Omernick, Andrew M. Dai, Thanumalayan Sankaranarayana Pillai, Marie Pellat, Aitor Lewkowycz, Erica Moreira, Rewon Child, Oleksandr Polozov, Katherine Lee, Zongwei Zhou, Xuezhi Wang, Brennan Saeta, Mark Diaz, Orhan Firat, Michele Catasta, Jason Wei, Kathy Meier-Hellstern, Douglas Eck, Jeff Dean, Slav Petrov, and Noah Fiedel. PaLM: Scaling Language Modeling with Pathways, October 2022. arXiv:2204.02311 [cs].
- [5] Kenneth Church. A Pendulum Swung Too Far. *Linguistic Issues in Language Technology*, 6, October 2011.
- [6] Luciano Floridi. *The Fourth Revolution: How the Infosphere is Reshaping Human Reality*. OUP Oxford, New York ; Oxford, June 2014.

- [7] Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring Mathematical Problem Solving With the MATH Dataset, November 2021. arXiv:2103.03874 [cs].
- [8] Suzana Herculano-Houzel. The remarkable, yet not extraordinary, human brain as a scaled-up primate brain and its associated cost. *Proceedings of the National Academy of Sciences of the United States of America*, 109 Suppl 1(Suppl 1):10661–10668, June 2012.
- [9] André Joyal and Ross Street. The geometry of tensor calculus, I. *Advances in Mathematics*, 88(1):55–112, July 1991.
- [10] Tabarak Khan. What are tokens and how to count them?, 2023.
- [11] Saunders Mac Lane. *Categories for the Working Mathematician*: 5. Springer, New York, NY, 2nd ed. 1978. softcover reprint of the original 2nd ed. 1978 edition edition, November 2010.
- [12] Saunders MacLane. Natural Associativity and Commutativity. *Rice Institute Pamphlet - Rice University Studies*, 49(4), October 1963. Accepted: 2011-11-08T19:13:47Z Publisher: Rice University.
- [13] Marjorie McShane and Sergei Nirenburg. *Linguistics for the Age of AI*. March 2021.
- [14] Francis Mollica and Steven T. Piantadosi. Humans store about 1.5 megabytes of information during language acquisition. *Royal Society Open Science*, 6(3):181393, March 2019.
- [15] Sharan Narang and Aakanksha Chowdhery. Pathways Language Model (PaLM): Scaling to 540 Billion Parameters for Breakthrough Performance, 2022.
- [16] OpenAI. ChatGPT: Optimizing Language Models for Dialogue, November 2022.
- [17] John R. Searle. Minds, brains, and programs. *Behavioral and Brain Sciences*, 3(3):417–424, September 1980. Publisher: Cambridge University Press.
- [18] Peter Selinger. A survey of graphical languages for monoidal categories. volume 813, pages 289–355. 2010. arXiv:0908.3347 [math].
- [19] Richard Sutton. The Bitter Lesson, 2019.
- [20] teddy [@teddynpc]. I made ChatGPT take a full SAT test. Here's how it did: <https://t.co/734sPFU3HY>, December 2022.
- [21] Alan D. Thompson. GPT-3.5 IQ testing using Raven's Progressive Matrices, 2022.

- [22] Tom Goldstein [@tomgoldsteincs]. Training PaLM takes 3.2 million kilowatt hours of power. If you powered TPUs by riding a bicycle, and you pedaled hard (nearly 400 watts), it would take you 1000 years to train PaLM, not including bathroom breaks. In that time, you'd make 320 trips around the globe!, July 2022.
- [23] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention Is All You Need, December 2017. arXiv:1706.03762 [cs].

2

Background

There are a lot of definitions to get started, but as with programming languages, preloading the work makes it easier to scale. This is the mathematical source code of string diagrams, which is only necessary if we need to show that something new is a symmetric monoidal category or if we are tinkering deeply, so it can be skipped for now. The important takeaway is that string diagrams are syntax, with morphisms in symmetric monoidal categories as semantics.

Definition 2.1.1 (Category). A *category* C consists of the following data

- A collection $\text{Ob}(C)$ of *objects*
- For every pair of objects $A, B \in \text{Ob}(C)$, a set $C(A, B)$ of *morphisms* from a to b .
- Every object $a \in \text{Ob}(C)$ has a specified morphism 1_a in $C(a, a)$ called the *identity morphism* on a .
- Every triple of objects $A, B, C \in \text{Ob}(C)$, and every pair of morphisms $f \in C(A, B)$ and $g \in C(b, c)$ has a specified morphism $(f; g) \in C(a, c)$ called the *composite* of f and g .

This data has to satisfy two coherence conditions, which are:

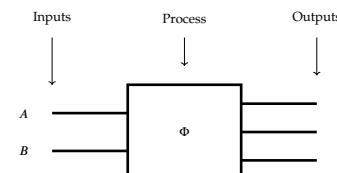
UNITALITY: For any morphism $f : a \rightarrow b$, $1_a; f = f = f; 1_b$

ASSOCIATIVITY: For any four objects A, B, C, D and three morphisms $f : A \rightarrow B$, $g : B \rightarrow C$, $h : C \rightarrow D$, $(f; g); h = f; (g; h)$

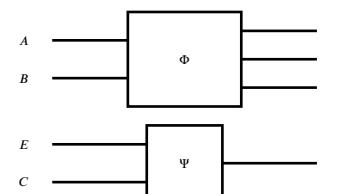
2.1 Process Theories

This section seeks to introduce process theories via string diagrams. The margin material will provide the formal mathematics of string diagrams from the bottom-up. The main body develops process theories via string diagrams by example, through which we develop towards a model of linguistic spatial relations – words like "to the left of" and "between" – which are a common ground of competence we all possess. Here we only focus on geometric relations between points in two dimensional Euclidean space equipped with the usual notions of metric and distance, providing adequate foundations to follow [talkspace], in which I demonstrate how text circuits can be obtained from sentences and how such text circuits interpreted in the category of sets and relations **Rel** provides a semantics for such sentences. This motivates the question of how to express the (arguably more primitive [piaget]) linguistic topological concepts – such as "touching" and "inside" – the mathematics of which will be in Chapter ??; the reader may skip straight to that chapter after this section if they are uninterested in syntax. We close this section with a brief note on how process theories relate to mathematical foundations and computer science.

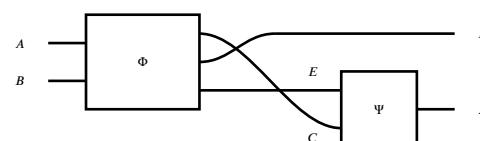
A *process* is something that transforms some number of input system types to some number of output system types. We depict systems as wires, labelled with their type, and processes as boxes. Unless otherwise specified, we read processes from left to right.



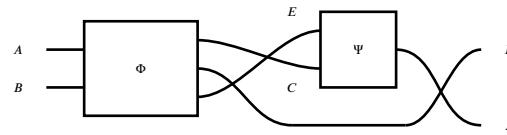
Processes may compose in parallel, depicted as placing boxes next to each other.



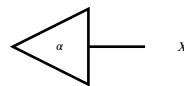
Processes may compose sequentially, depicted as connecting wires of the same type.



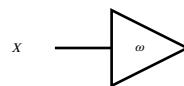
In these diagrams only input-output connectivity matters: so we may twist wires and slide boxes along wires to obtain different diagrams that still refer to the same process. So the diagram below is equal to the diagram above.



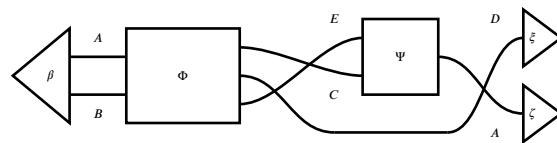
Some processes have no inputs; we call these *states*.



Some processes have no outputs; we call these *tests*.



A process with no inputs and no outputs is a *number*; the number tells us the outcome of applying tests to a composite of states modified by processes.



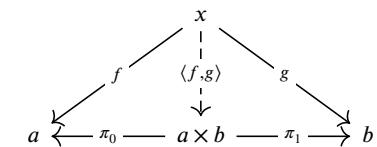
A process theory is given by the following data:

- A collection of systems
- A collection of processes along with their input and output systems
- A methodology to compose systems and processes sequentially and in parallel, and a specification of the unit of parallel composition.
- A collection of equations between composite processes

Example 2.1.13 (Linear maps with direct sum). Systems are finite-dimensional vector spaces over \mathbb{R} . Processes are linear maps, expressed as matrices with entries in \mathbb{R} .

Sequential composition is matrix multiplication. Parallel composition of systems is the direct sum of vector

Definition 2.1.2 (Categorical Product). In a category C , given two objects $a, b \in \text{Ob}(C)$, the *product* $A \times B$, if it exists, is an object with projection morphisms $\pi_0 : A \times B \rightarrow A$ and $\pi_1 : A \times B \rightarrow B$ such that for any object $x \in \text{Ob}(C)$ and any pair of morphisms $f : X \rightarrow A$ and $g : X \rightarrow b$, there exists a unique morphism $f \times g : X \rightarrow A \times B$ such that $f = (f \times g); \pi_0$ and $g = (f \times g); \pi_1$. This is a mouthful which is easier expressed as a commuting diagram as below. The dashed arrow indicates uniqueness. $A \times B$ is a product when every path through the diagram following the arrows between two objects is an equality.



The idea behind the definition of product is simple: instead of explicitly constructing the cartesian product of sets from within, let's say a *product is as a product does*. For objects, the cartesian product of sets $A \times B$ is a set of pairs, and we may destruct those pairs by extracting or projecting out the first and second elements, hence the projection maps π_0, π_1 . Another thing we would like to do with pairs is construct them; whenever we have some A -data and B -data, we can pair them in such a way that construction followed by destruction is lossless and doesn't add anything. In category-theoretic terms, we select 'arbitrary' A - and B -data by arrows $f : X \rightarrow A$ and $g : X \rightarrow B$, and we declare that $f \times g : X \rightarrow A \times B$ is the unique way to select corresponding tuples in $A \times B$. This design-pattern of "for all such-and-such there exists a unique such-and-such" is an instance of a so-called *universal property*, the purpose of which is to establish isomorphism between operationally equivalent implementations.

To understand what this style of definition gives us, let's revisit Kuratowski's and Wiener's definitions of cartesian product, which are, respectively:

$$A^K \times B := \left\{ \{\{a\}, \{a, b\}\} \mid a \in A, b \in B \right\}$$

$$A^W \times B := \left\{ \{\{a, \emptyset\}, b\} \mid a \in A, b \in B \right\}$$

Keeping overset-labels and using maplet notation, the associated projections are:

$$\pi_0^K := \{\{a\}, \{a, b\}\} \mapsto a$$

$$\pi_1^K := \{\{a\}, \{a, b\}\} \mapsto b$$

$$\pi_0^W := \{\{a, \emptyset\}, b\} \mapsto a$$

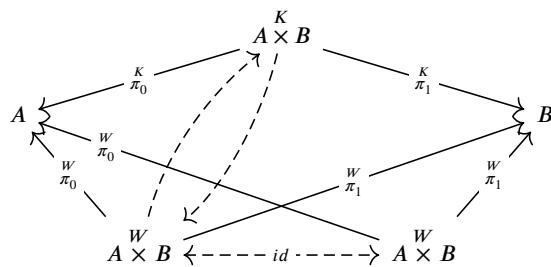
$$\pi_1^W := \{\{a, \emptyset\}, b\} \mapsto b$$

And maps f, g into A and B are tupled by the following:

$$f^K \times g := x \mapsto \{\{f(x)\}, \{f(x), g(x)\}\}$$

$$f^W \times g := x \mapsto \{\{f(x), \emptyset\}, g(x)\}$$

Both satisfy the commutative diagram defining the product. Something neat happens when we pick $A^K \times B$ to be the arbitrary X for the product definition of $A^W \times B$ and vice versa. We get to mash the commuting diagrams together:



spaces \oplus . The parallel composition of matrices \mathbf{A}, \mathbf{B} is the block-diagonal matrix

$$\begin{bmatrix} \mathbf{A} & \mathbf{0} \\ \mathbf{0} & \mathbf{B} \end{bmatrix}$$

The unit of parallel composition is the singleton 0-dimensional vector space. States are row vectors. Tests are column vectors. The numbers are \mathbb{R} . Usually the monoidal product is written with the symbol \otimes , which clashes with notation for the hadamard product for linear maps, while the process theory we have just described takes the direct sum \oplus to be the monoidal product.

Example 2.1.14 (Sets and functions with cartesian product). Systems are sets A, B . Processes are functions between sets $f : A \rightarrow B$. Sequential composition is function composition. Parallel composition of systems is the cartesian product of sets: the set of ordered pairs of two sets.

$$A \otimes B = A \times B := \{(a, b) \mid a \in A, b \in B\}$$

The parallel composition $f \otimes g : A \times C \rightarrow B \times D$ of functions $f : A \rightarrow B$ and $g : C \rightarrow D$ is defined:

$$f \otimes g := (a, c) \mapsto (f(a), g(c))$$

The unit of parallel composition is the singleton set $\{\star\}$. There are many singletons, but this presents no problem for the later formal definition because they are all equivalent up to unique isomorphism. States of a set A correspond to elements $a \in A$ – we forgo the usual categorical definition of points from the terminal object in favour of generalised points from the monoidal perspective. Every system A has only one test $a \mapsto \star$; this is since the singleton is terminal in **Set**. So there is only one number.

Example 2.1.15 (Sets and relations with cartesian product). Systems are sets A, B . Processes are relations between sets $\Phi \subseteq A \times B$, which we may write in either direction $\Phi^* : A \nrightarrow B$ or $\Phi_* : B \nrightarrow A$. Relations between sets are equivalently matrices with entries from the boolean semiring. Relation composition is matrix multiplication with the boolean semiring. Φ^*, Φ_* are the transposes of one another. Sequential composition is relation composition:

$$A \xrightarrow{\Phi} B \xrightarrow{\Psi} C := \{(a, c) \mid a \in A, c \in C, \exists b \in B : (a, b) \in \Phi \wedge (b, c) \in \Psi\}$$

Parallel composition of systems is the cartesian product of sets. The parallel composition of relations $A \otimes C \xrightarrow{\Phi \otimes \Psi} B \otimes D$ of relations $A \xrightarrow{\Phi} B$ and $C \xrightarrow{\Psi} D$ is defined:

$$\Phi \otimes \Psi := \{(a, c), (b, d) \mid (a, b) \in \Phi \wedge (c, d) \in \Psi\}$$

The unit of parallel composition is the singleton. States of a set A are subsets of A . Tests of a set A are also

subsets of A .

2.1.1 What does it mean to copy and delete?

Now we discuss how we might define the properties and behaviour of processes by positing equations between diagrams. Let's begin simply with two intuitive processes *copy* and *delete*:



Example 2.1.16 (Linear maps). Consider a vector space \mathbf{V} , which we assume includes a choice of basis. The copy map is the rectangular matrix made of two identity matrices:

$$\Delta_{\mathbf{V}} : \mathbf{V} \rightarrow \mathbf{V} \oplus \mathbf{V} := \begin{bmatrix} \mathbf{1}_{\mathbf{V}} & \mathbf{1}_{\mathbf{V}} \end{bmatrix}$$

The delete map is the column vector of 1s:

$$\epsilon_{\mathbf{V}} : \mathbf{V} \rightarrow \mathbf{0} := \begin{bmatrix} 1 \\ \vdots \\ 1 \end{bmatrix}$$

Example 2.1.17 (Sets and functions). Consider a set A . The copy function is defined:

$$\Delta_A : A \rightarrow A \times A := a \mapsto (a, a)$$

The delete function is defined:

$$\epsilon_A : A \rightarrow \{\star\} := a \mapsto \star$$

Example 2.1.18 (Sets and relations). Consider a set A . The copy relation is defined:

$$\Delta_A : A \nrightarrow A \times A := \{(a, (a, a)) \mid a \in A\}$$

The delete relation is defined:

$$\epsilon_A : A \nrightarrow \{\star\} := \{(a, \star) \mid a \in A\}$$

We may verify that, no matter the concrete interpretation of the diagram in terms of linear maps, functions

The two unique arrows between $\overset{K}{X}$ and $\overset{W}{X}$ are format-conversions, and we know by definition that the unique arrow that performs format conversion from $\overset{W}{X}$ to itself in the bottom face is the identity. In maplet notation, the conversion from $A \overset{K}{\times} B \rightarrow A \overset{W}{\times} B$ is $\{\{a\}, \{a, b\}\} \mapsto \{\{a, \emptyset\}, b\}$, and similarly for the other direction. Because these conversions are uniquely determined arrows, their composite is also uniquely determined, and we know their composite is equal to the identity. So, the nontrivial conversions witness an *isomorphism* between $A \overset{K}{\times} B$ and $A \overset{W}{\times} B$; a pair of maps $X \rightarrow Y$ and $Y \rightarrow X$ such that their loop-composites equal identities. This in a nutshell is the category-theoretic approach to overcoming the bureaucracy of syntax: use universal properties (or whatever else) encode your intents and purposes, establish isomorphisms, and then treat isomorphic things as "the same for all intents and purposes". The idea of treating isomorphic objects as the same is ingrained in category theory, so isomorphism notation \simeq is often just written as equality $=$; going forward we will use equality notation unless there are good reasons to remember that we only have isomorphisms.

Definition 2.1.3 (Functor). A functor $F : \mathcal{C} \rightarrow \mathcal{D}$ (read: with domain a category \mathcal{C} and codomain a category \mathcal{D}) consists of two suitably related functions. An object function $F_0 : \text{Ob}(\mathcal{C}) \rightarrow \text{Ob}(\mathcal{D})$ and a morphism function (equivalently viewed as a family of functions indexed by pairs of objects of \mathcal{C}) $F_1(X, Y) : \mathcal{C}(X, Y) \rightarrow \mathcal{D}(F_0 X, F_0 Y)$. F_1 must map identities to identities – i.e., be such that for all $X \in \mathcal{C}$, $F_1(1_X) = 1_{F_0 X}$ – and F_1 must map composites to composites – i.e., for all $X, Y, Z \in \text{Ob}(\mathcal{C})$ and all $f : X \rightarrow Y$ and $g : Y \rightarrow Z$, $F_1(f; g) = F_1 f; F_1 g$.

Functors in short map categories to categories, preserving the structure of identities and composition. They are the essence of "structure preserving transformation". Insofar as semantics is the science of finding structure-preserving transformations that tell us when syntactic things are equal, functors are just that. They are incredibly useful and mysterious and worth internalising in a way I am not adept enough to impress by example in this margin. For us, for now, they are just stepping stones to define transformations *between functors*.

Definition 2.1.4 (Natural Transformation). A natural transformation $\theta : F \Rightarrow G$ for (co)domain-aligned functors $F, G : \mathcal{C} \rightarrow \mathcal{D}$ is a family of morphisms in \mathcal{D} indexed by objects $X \in \mathcal{C}$ such that for all $f : X \rightarrow Y$ in \mathcal{C} , the following commuting diagram holds in \mathcal{D} :

$$\begin{array}{ccc} FX & \xrightarrow{Ff} & FY \\ | & & | \\ \theta_X & \downarrow & \downarrow \theta_Y \\ GX & \xrightarrow{Gf} & GY \end{array}$$

or relations, the following equations characterise a cocommutative comonoid internal to a monoidal category.

The diagram shows three equations labeled (2.1).
1. Coassociativity: Two diagrams of a node with three outgoing edges are shown as equal. The left diagram has edges entering from the top-left and bottom-left, and exiting to the top-right, top-left, and bottom-right. The right diagram has edges entering from the top-right and bottom-right, and exiting to the top-left, top-right, and bottom-left.
2. Cocommutativity: Two diagrams of a node with two outgoing edges are shown as equal. The left diagram has a crossing loop between the two edges. The right diagram has the edges swapped.
3. Counitality: Two diagrams of a node with one outgoing edge are shown as equal. The left diagram has a small loop at the node. The right diagram is a simple horizontal line.
(2.1)

It is worth pausing here to think about how one might characterise the process of copying in words; it is challenging to do so for such an intuitive process. The diagrammatic equations, when translated into prose, provide an answer.

Coassociativity: says there is no difference between copying copies.

Cocommutativity: says there is no difference between the outputs of a copy process.

Counitality: says that if a copy is made and one of the copies is deleted, the remaining copy is the same as the original.

Insofar as we think this is an acceptable characterisation of copying, rather than specify concretely what a copy and delete does for each system X we encounter, we can instead posit that so long as we have processes $\Delta_X : X \otimes X \rightarrow X$ and $\epsilon_X : X \rightarrow I$ that obey all the equational constraints above, Δ_X and ϵ_X are as good as a copy and delete.

Example 2.1.19 (Not all states are copyable). Call a state *copyable* when it satisfies the following diagrammatic equation:

The diagram shows two diagrams separated by an equals sign. The left diagram is a node with two incoming arrows from the left and two outgoing arrows to the right. The right diagram is a node with two incoming arrows from the left and two outgoing arrows to the right, where the two outgoing arrows are swapped relative to the incoming ones.

In the process theory of sets and functions, all states are copyable. Not all states are copyable in the process theories of sets and relations and of linear maps. For example, consider the two element set $\mathbb{B} := \{0, 1\}$, and let $T : \{\star\} \nrightarrow \mathbb{B} := \{(\star, 0), (\star, 1)\} \simeq \{0, 1\}$. Consider the composite of T with the copy relation:

$$T; \Delta_{\mathbb{B}} := \{(\star, (0, 0)), (\star, (1, 1))\} \simeq \{(0, 0), (1, 1)\}$$

This is a perfectly correlated bipartite state, and it is not equal to $\{0, 1\} \times \{0, 1\}$, so T is not copyable.

Remark 2.1.20. The copyability of states is a special case of a more general form of interaction with the copy relation:

A cyan map that satisfies this equation is said to be a homomorphism with respect to the commutative comonoid. In the process theory of relations, those relations that satisfy this equation are precisely the functions; in other words, this diagrammatic equation expresses *determinism*.

Here is an unexpected consequence. Suppose we insist that *to copy* in principle also implies the ability to copy *anything* – arbitrary states. From Example 2.1.19 and Remark 2.1.20, we know that this demand is incompatible with certain process theories. In particular, this demand would constrain a process theory of sets and relations to a subtheory of sets and functions. The moral here is that process theories are flexible enough to meet ontological needs. A classical computer scientist who works with perfectly copyable data and processes might demand universal copying along with Equations 2.1, whereas a quantum physicist who wishes to distinguish between copyable classical data and non-copyable quantum data might taxonomise copy and delete as a special case of a more generic quasi-copy and quasi-delete that only satisfies equations 2.1. In fact, quantum physicists *do* do this; see Dodo: [].

2.1.2 What is an update?

In the previous section we have seen how we can start with concrete examples of copying in distinct process theories, and obtain a generic characterisation of copying by finding diagrammatic equations copying satisfies in each concrete case. In this section, we show how to go in the opposite direction: we start by positing diagrammatic equations that characterise the operational behaviour of a particular process – such as *updating* – and it will turn out that any concrete process that satisfies the equational constraints we set out will *by our own definition* be an update.

Perhaps the most familiar setting for an update is a database. In a database, an *entry* often takes the form of pairs of *fields* and *values*. For example, where a database contains information about employees, a typical entry might look like:

```
< NAME:Jono Doe, AGE:69, JOB:CONTENT CREATOR, SALARY:$420, ... >
```

There are all kinds of reasons one might wish to update the value of a field: Jono might legally change their name, a year might pass and Jono's age must be incremented, Jono might be promoted or demoted or get a raise and so on. It was the concern of database theorists to formalise and axiomatise the notion of updating

Definition 2.1.5 (Cat). *Cat* is a (2-)category where the objects are (1-)categories as defined above, the morphisms are functors, and the (2-)morphisms are natural transformations. (2-)morphisms are morphisms between morphisms that we discuss in more detail in Section ???. There's no "set of all sets" paradox here by construction; *Cat* is slightly more than a category as we have seen so far because of the (2-)morphisms. We're introducing this just to state that the definition of product also works here so that we can consider product categories $C \times D$, whose objects are pairs of objects and morphisms pairs of morphisms.

Definition 2.1.6 (Monoidal Category). A monoidal category consists of a category \mathcal{C} , a functor $\otimes : \mathcal{C} \times \mathcal{C} \rightarrow \mathcal{C}$, a monoidal unit object $I \in \text{Ob}(\mathcal{C})$, and the following natural isomorphisms – i.e. natural transformations with inverses, where multiple bar notation indicates variable object argument positions: an associator $\alpha : ((-\otimes=)\otimes\equiv) \mapsto (-\otimes(=\otimes\equiv))$, a right unit $\rho : -\otimes I \mapsto -$, and a left unit $\lambda : I\otimes- \mapsto -$. These natural isomorphisms must in addition satisfy certain *coherence* diagrams, to be displayed shortly.

Theorem 2.1.7 (Coherence for monoidal categories). The following pentagon and triangle diagrams are conditions in the definition of a monoidal category. When they hold, all composites of associators and unitors (and their inverses) are isomorphisms. 1 denotes identities.

$$\begin{array}{ccc}
 & ((W \otimes X) \otimes (Y \otimes Z)) & \\
 \alpha \nearrow & \downarrow & \downarrow \alpha \\
 (W \otimes (X \otimes (Y \otimes Z))) & & (((W \otimes X) \otimes Y) \otimes Z) \\
 \downarrow 1 \otimes \alpha & & \uparrow \alpha \otimes 1 \\
 (W \otimes ((X \otimes Y) \otimes Z)) & & ((W \otimes (X \otimes Y)) \otimes Z) \\
 & \alpha \searrow & \\
 & (X \otimes (I \otimes Y)) & \\
 & \downarrow 1 \otimes \lambda & \nearrow \alpha \\
 & (X \otimes Y) & \xleftarrow{\rho \otimes 1}
 \end{array}$$

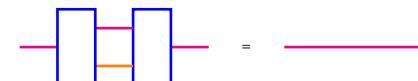
Remark 2.1.8 (Coherence). Coherence is about getting rid of syntactic bureaucracy. Addition for example is a binary associative operation, and knowing that $(x + (y + z)) = ((x + y) + z)$ is what allows us to safely drop all possible bracketings and just write $1 + 2 + 3$. We further know that addition is a monoid with unit 0 , so we can always write $x + 0 = x = 0 + x$. Now the situation is that we have replaced the associativity equation with associator natural transformations $((X \otimes Y) \otimes Z) \xrightarrow{\alpha_{XYZ}} (X \otimes (Y \otimes Z))$, and unit equations with left and right unitors $(X \otimes I) \xrightarrow{\rho_X} X \xrightarrow{\lambda_X} (I \otimes X)$. Recalling that we're happy with isomorphism in place of equality, we would like to know that every possible composite of these structural operations is an isomorphism.

the value of a field – *independently of the specific programming language implementation of a database* – so that they had reasoning tools to ensure program correctness []. The problem is reducible to axiomatising a *rewrite*: we can think of updating a value as first calculating the new value, then *putting* the new value in place of the old. Since often the new value depends in some way on the old value, we also need a procedure to *get* the current value. That was a flash-prehistory of *bidirectional transformations* []. Following the monoidal generalisation of lenses in [], a rewrite as we have described above is specified by system diagrammatic equations in the margin, each of which we introduce in prose.

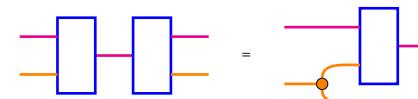
PutPut: Putting in one value and then a second is the same as deleting the first value and just putting in the second.



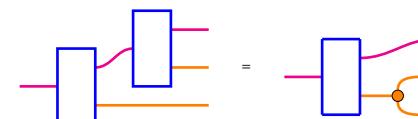
GetPut: Getting a value from a field and putting it back in is the same as not doing anything.



PutGet: Putting in a value and getting a value from the field is the same as first copying the value, putting in one copy and keeping the second.



GetGet: Getting a value from a field twice is the same as getting the value once and copying it.

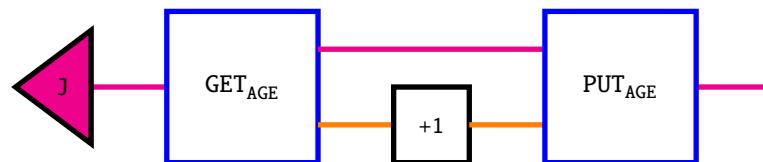


These diagrammatic equations do two things. First, they completely specify what it means to get and put values in a field in an implementation independent manner; it doesn't matter whether database entries are encoded as bitstrings, qubits, slips or paper or anything else, what matters is the interaction of get and put. Second, the diagrammatic equations give us the right to call our processes *get* and *put* in the first place: we define what it means to get and put by outlining the mutual interactions of get, put, copy, and delete. These

two points are worth condensing and rephrasing:

A kind of process is determined by patterns of interaction with other kinds of processes.

Now we can diagrammatically depict the process of updating Jono's age, by [getting](#) Jono's [age value](#) from their [entry](#), incrementing it by 1, and [putting](#) it back in.



2.1.3 Pregroup diagrams and correlations

Let's revisit the copy and delete maps for a moment. Suppose our process theory is such that for every wire X , there is a unique copy map δ_X such that every state on X is copyable and deletable. A consequence of this assumption is that every state is \otimes -separable (read *tensor separable*):

placeholder

But there are certainly process theories in which we don't want this. For example, if states are random variables and parallel composition is the product of independent random variables (as is the case in Markov categories for probability theory []) then copying a random variable gives a perfectly correlated pair of variables, which cannot be expressed as the product of a pair of independent random variables.

2.1.4 Equational Constraints and Frobenius Algebras

2.1.5 Processes, Sets, and Computers

OBJECTION: BUT WHAT ARE THE THINGS THAT THE PROCESSES OPERATE ON?

This is a common objection from philosophers who want their ontologies tidy. The claim roughly goes that you can't really reason about processes without knowing the underlying objects that participate on those processes, and since set theory is the only way we know how to spell out objects intensionally in this way, we should stick to sets. In simpler terms, if we're drawing only functions as (black)-boxes in our diagrams, how will we know what they do to the elements of the underlying sets?

The short answer is that – perhaps surprisingly – reasoning process-theoretically is mathematically equivalent to reasoning about sets and elements for all practical purposes; it is as if whatever is going on *out there* is

Definition 2.1.9 (Symmetric Monoidal Category). A symmetric monoidal category is a monoidal category with an additional natural isomorphism $\theta : - \otimes - = \leftrightarrow = \otimes -$. Coherence requires the following pair of hexagons

$$\begin{array}{ccc}
(X \otimes (Y \otimes Z)) & \xrightarrow{\theta} & (Z \otimes (X \otimes Y)) \\
\alpha^{-1} \downarrow & & \downarrow \alpha \\
(X \otimes (Z \otimes Y)) & & ((Z \otimes X) \otimes Y) \\
1 \otimes \theta \downarrow & & \downarrow \theta^{-1} \\
(X \otimes (Z \otimes Y)) & \xrightarrow{\alpha} & ((X \otimes Z) \otimes Y) \\
\\
(X \otimes (Y \otimes Z)) & \xrightarrow{\theta} & ((Y \otimes Z) \otimes X) \\
\alpha \downarrow & & \downarrow \alpha^{-1} \\
((X \otimes Y) \otimes Z) & & (Y \otimes (Z \otimes X)) \\
\theta^{-1} \downarrow & & \downarrow 1 \otimes \theta \\
((Y \otimes X) \otimes Z) & \xrightarrow{\alpha^{-1}} & (Z \otimes (X \otimes Y))
\end{array}$$

Remark 2.1.10 (Coherence for symmetric monoidal categories). Going back to addition, we can rearrange sums by commutativity of the addition monoid: $x + y = y + x$. In the monoidal setting, the natural isomorphism θ is the twisting of wires. Just as we would like to have addition equations such as $(x + (y + z)) = ((z + y) + x)$, we would like the twists θ to behave well with respect to the bundling of objects and morphisms with the associator α .

indifferent to whether we describe using a language of only nouns or only verbs.

In the case of set theory, let's suppose that instead of encoding functions as sets, we treat functions as primitive, so that we have a process theory where wires are labelled with sets, and functions are process boxes that we draw. The problem we face now is that it is not immediately clear how we would access the elements of any set using only the diagrammatic language. The solution is the observation that the elements $\{x \mid X\}$ of a set X are in bijective correspondence with the functions from a singleton into X : $\{f(\star) \mapsto x \mid \{\star\} \xrightarrow{f} X\}$. In prose, for any element x in a set X , we can find a function that behaves as a pointer to that element $\{\star\} \rightarrow X$. So the states we have been drawing, when interpreted in the category of sets and function, are precisely elements of the sets that label their output wires.

The full and formal answer will require the reader to see Section ?? which spells out the category theory underpinning process theories. The caveat here is that process theories work for all *practical* purposes, so I make no promises about how diagrams work for the kind of set theories that deals with hierarchies of infinities that set theorists do. For other issues concerning for instance the set of all functions between two sets, that requires symmetric monoidal closure, for which there exist string-diagrammatic formalisms [].

OBJECTION: BUT IF THEY'RE EXPRESSIVELY THE SAME, WHAT'S THE POINT?

The following rebuttal draws on Harold Abelson's introductory lecture to computer science [] (in which string diagrams appear to introduce programs without being explicitly named as such).

There is a distinction between declarative and imperative knowledge. Declarative knowledge is *knowing-that*, for example, 6 is the square root of 36, which we might write $6 = \sqrt{36}$. Imperative knowledge is *knowing-how*, for example, to obtain the square root of a positive number, for instance, by Heron's iterative method: to obtain the square root of Y , make a guess X , and take the average of X and $\frac{Y}{X}$ until your guess is good enough.

Computer science concerns imperative knowledge. An obstacle to the study of imperative knowledge is complexity, which computer scientists manage by black-box abstraction – suppressing irrelevant details, so that for instance once a square root procedure is defined, the reasoner outside the system does not need to know whether the procedure inside is an iterative method by Heron or Newton, only that it works and has certain properties. These black-boxes can be then composed into larger processes and procedures within human cognitive load.

Abstraction also yields generality. For example, in the case of addition, it is not only numbers we may care to add, but perhaps vectors, or the waveforms of signals. So there is an abstract notion of addition which we concretely instantiate for different domains that share a common interface; we may decide for example that all binary operations that are commutative monoids are valid candidates for what it means to be an addition operation.

In this light, string diagrams are a natural metalanguage for the study of imperative knowledge; string diagrams in fact independently evolved within computer science from flowcharts describing processes. Pro-

cess theories, which are equations or logical sentences about processes, allow us to reason declaratively about imperative knowledge. Moreover, string diagrams as syntactic objects can be interpreted in various concrete settings, so that the same diagram serves as the common interface for a process like addition, with compliant implementation details for each particular domain spelled out separately.

2.2 An introduction to weak n -categories for formal linguists

Geometrically, a set is a collection of labelled zero-dimensional points. A category is a collection of one-dimensional labelled arrows upon a set of zero-dimensional points (that satisfy certain additional conditions, such as having identity morphisms and associativity of composition.) Naturally, we might ask what happens when we generalise to more dimensions, asking for two-dimensional somethings going between one-dimensional arrows, and three-dimensional somethings going between two-dimensional somethings, and so on. This is the gist of an n -category, where n is a positive integer denoting dimension.

There is a fork in the road in generalisation. First, different choices of what n -dimensional somethings could be give different conceptions of n -category, because there are multiple mathematically well-founded choices for filling in the blank of points, lines, ??. Simplices, cubical sets, and globular sets are three common options; though the names are fancy, they correspond to triangular, cubical, and circular families of objects indexed by dimension.

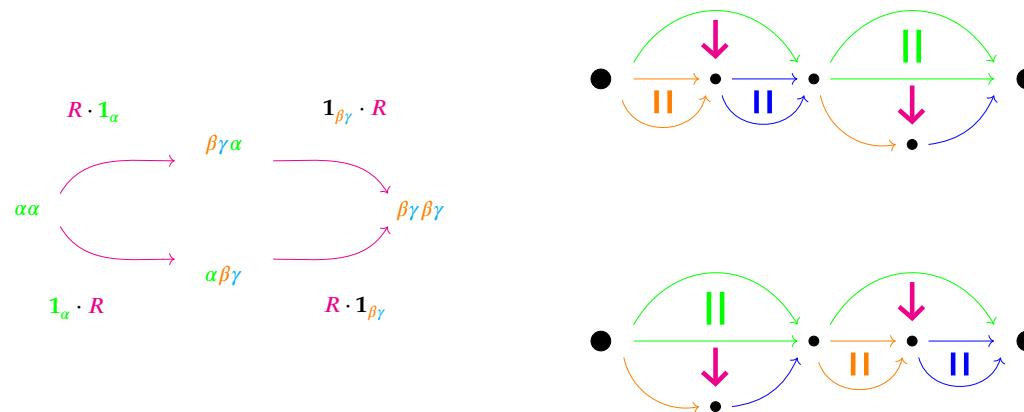
Second, there is a distinction between *strict* and *weak* n -categories. Just as in regular or 1-category theory we are interested in objects up to isomorphism rather than equality – because isomorphic objects in a category are as good as one another – lifting this philosophy to n -categories gives us *weak* n -categories. In a strict k -category, all of the j -dimensional morphisms for $j > k$ are identity-equalities. In a weak k -category, all equalities for dimensions $j > k$ are replaced by isomorphisms all the way up: which means that a k -equality $\alpha = \beta$ in the strict setting is replaced by a pair of $k + 1$ -morphisms witnessing the isomorphism of α and β , and that pair of $k + 1$ -morphisms has a pair of $k + 2$ morphisms witnessing that they are $k + 1$ -isomorphic, and so on for $k + 3$ and all the way up. Unsurprisingly, strict n -categories are easy to formalise, and weak n -categories are hard. A conjecture or guiding principle like the Church-Turing thesis holds for weak n -categories that they should all be equivalent to one another, whatever equivalence means.

Mathematicians, computer scientists, and physicists may have good reasons to work with weak n -categories CITE, but what is the value proposition for formal linguists? A philosophical draw for the formal semanticist is that insofar as semantics is synonymy – the study of when expressions are equivalent – weak n -categories are an exquisite setting to control and study sameness in terms of meta-equivalences. A practical draw for the formal syntactician is that weak n -categories provide a natural setting to model generic rewriting systems, and that is what we will focus on here. In this setting we will introduce weak n -categories in the homotopy.io formulation, along the way showing how they provide a common setting to formalise string-rewrite and tree-adjoining grammars, setting the stage for us to specify a circuit-adjoining grammar for text circuits later on.

2.2.1 String-rewrite systems as 1-object-2-categories

Say we have an alphabet $\Sigma := \{\alpha, \beta, \gamma\}$. Then the Kleene-star Σ^* consists of all strings (including the empty string ϵ) made up of Σ , and we consider formal languages on Σ to be subsets of Σ^* . Another way of viewing Σ^* is as the free monoid generated by Σ under the binary concatenation operation ($_ \cdot _$) which is associative and unital with unit ϵ , the empty string. Associativity and unitality are precisely the conditions of composition of morphisms in categories, so we have yet another way to express Σ^* as a finitely presented category; we consider a category with a single object \star , taking ϵ to be the identity morphism 1_\star on the single object, and we ask for the category obtained when we consider the closure under composition of three non-identity morphisms $\alpha, \beta, \gamma : \star \rightarrow \star$. In this category, every morphism $\star \rightarrow \star$ corresponds to a string in Σ^* . We illustrate this example in the margins. A string-rewrite system additionally consists of a finite number of string-transformation rules. Building on our example, we might have a named rule $R : \alpha \mapsto \beta \cdot \gamma$, which we illustrate in Figure 2.2.1.

WE CONSIDER REWRITES TO BE EQUIVALENT, BUT NOT EQUAL. In a string-rewrite system, rewrites are applied one at a time. This means that even for our simple example, there are two possible rewrites from $\alpha \cdot \alpha$ to obtain $\beta \cdot \gamma \cdot \beta \cdot \gamma$. Here are the two rewrites viewed in two equivalent ways, first on the left informally where strings are nodes in a graph and rewrites are labelled transitions and secondly on the right as two distinct commuting 2-diagrams.



What should we say about how these two different rewrites relate to each other? Let's say Alice is a formal linguist who is only interested in what strings are reachable from others by rewrites – this is *de rigueur* when we consider formal languages to be subsets of Σ^* . She might be happy to declare that these two rewrites are simply equal; categorically this is tantamount to her declaring that any two 2-cells in the 1-object-2-category that share the same source and target are in fact the same, or equivalently, that any n -cells for $n \geq 3$ are iden-

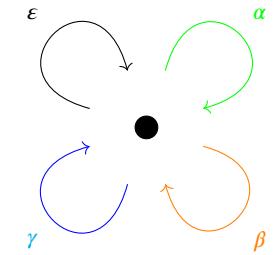


Figure 2.1: The category in question can be visualised as a commutative diagram.

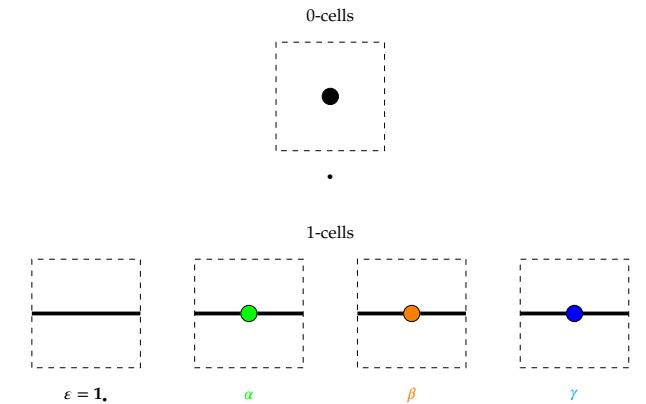


Figure 2.2: When there are too many generating morphisms, we can instead present the same data as a table of n -cells; there is a single 0-cell \star , and three non-identity 1-cells corresponding to α, β, γ , each with source and target 0-cells \star . Typically identity morphisms can be omitted from tables as they come for free. Observe that composition of identities enforces the behaviour of the empty string, so that for any string x , we have $\epsilon \cdot x = x = \epsilon \cdot x$.



Figure 2.3: For a concrete example, we can depict the string $\alpha \cdot \gamma \cdot \gamma \cdot \beta$ as a morphism in a commuting diagram.

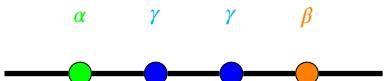


Figure 2.4: The string-diagrammatic view, where \star is treated as a wire and morphisms are treated as boxes or dots is an expression of the same data under the Poincaré dual.

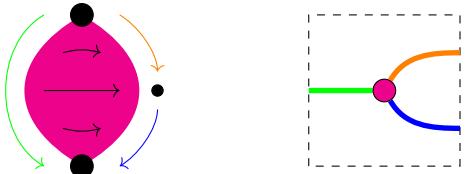
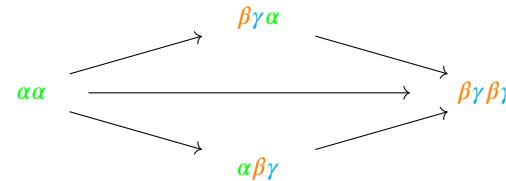
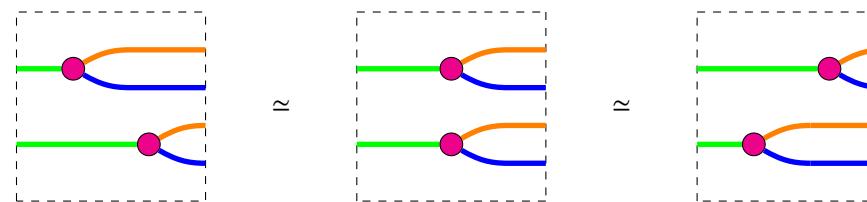


Figure 2.5: We can visualise the rule as a commutative diagram where R is a 2-cell between the source and target 1-cells. Just as 1-cells are arrows between 0-cell points in a commuting diagram, a 2-cell can also be conceptualised as a directed surface from a 1-cell to another. Taking the Poincaré dual of this view gives us a string diagram for the 2-cell R .

tities. In fact, what Alice really cares to have is a category where the objects are strings from Σ^* , and the morphisms are a reachability relation by rewrites; this category is *thin*, in that there is at most one arrow between each pair of objects, which forgets what rewrites are applied.



Let's say Bob is a different sort of formal linguist who wants to model the two rewrites as nonequal but equivalent, with some way to keep track of how different equivalent rewrites relate to one another. Bob might want this for example because he wants to show that head-first rewrite strategies are the same as tail-first, so he wants to keep the observation that the two rewrites are equivalent in that they have the same source and target, while keeping the precise order of rewrites distinct. This order-independence of disjoint or non-interfering rewrites is reflected in the interchange law for monoidal categories, which in the case of our example is depicted as:



In fact, Bob gets to express a new kind of rewrite in the middle: the kind where two non-conflicting rewrites happen *concurrently*. The important aspect of Bob's view over Alice's is that equalities have been replaced by syntactically unequal rewrites. This demotion of equalities to isomorphisms means that Bob is dealing with a *weak* 1-object-2-category; Bob does have 3-cells that relate different 2-cells with the same source and target, but all of Bob's n -cells for $n \geq 3$ are isomorphisms, rather than equalities.

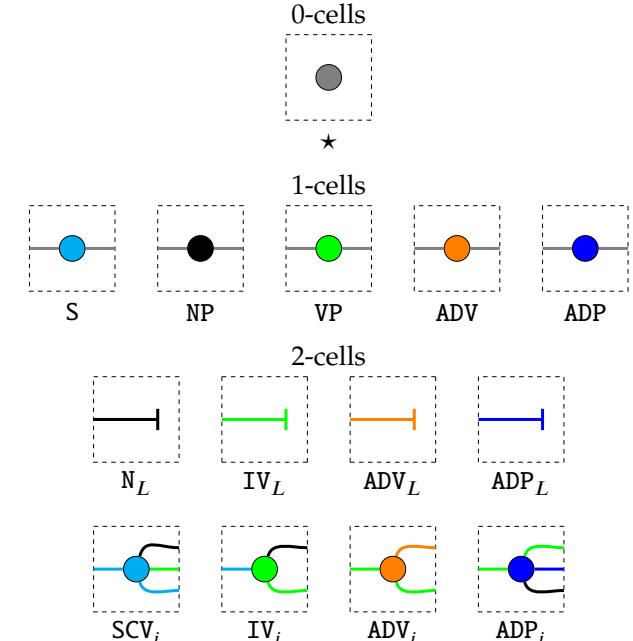
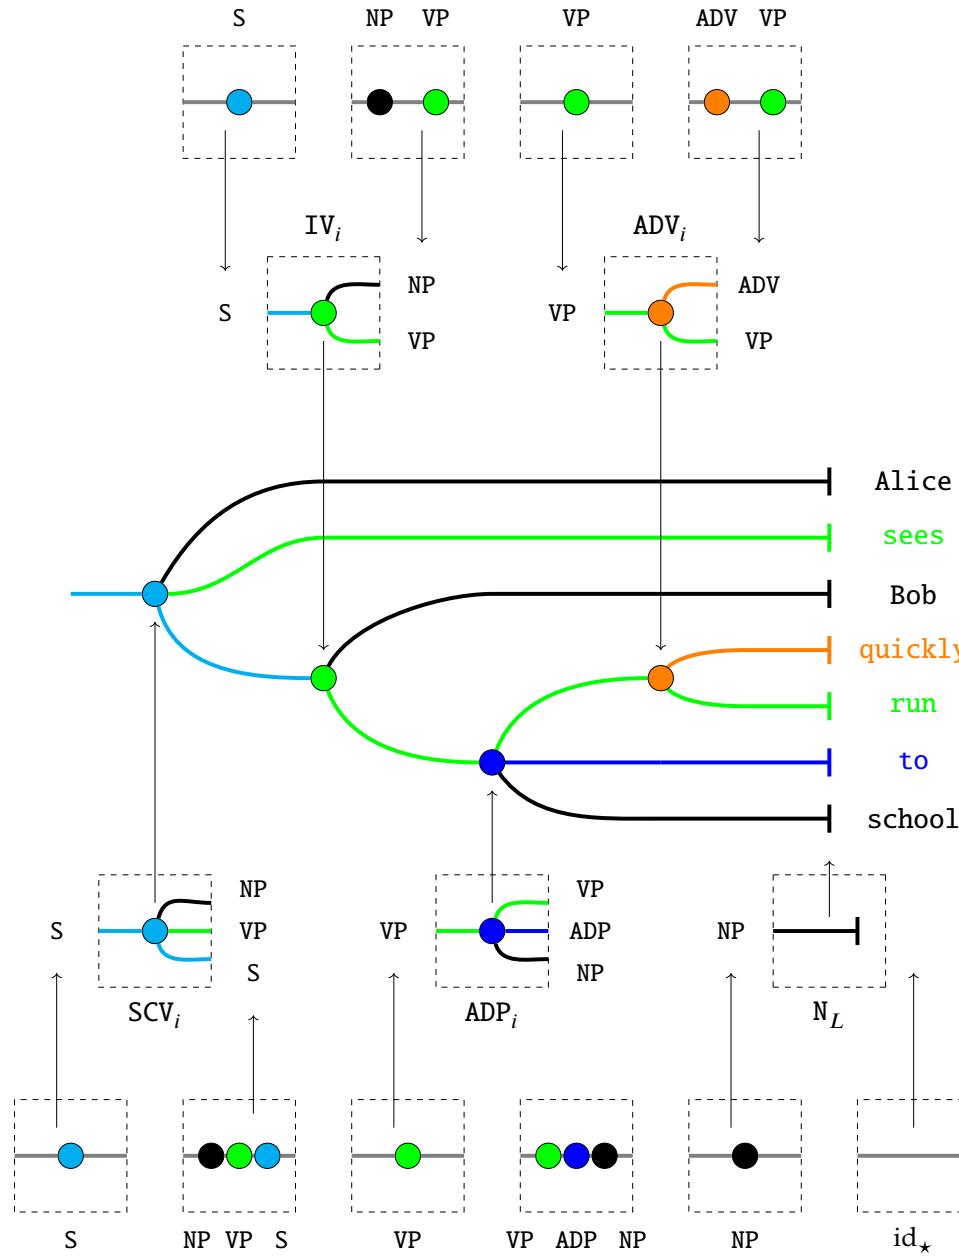


Figure 2.6: A context-free grammar for Alice sees Bob quickly run to school. We can describe a context-free grammar with the same combinatorial rewriting data that specifies planar string diagrams as we have been illustrating so far. One aspect of rewrite systems we adapt for now is the distinction between terminal and nonterminal symbols; terminal symbols are those after which no further rewrites are possible. We capture this string-diagrammatically by modelling terminal rewrites as 2-cells with target equal to the 1-cell identity of the 0-cell \star , which amounts to graphically terminating a wire. The generators subscripted L (for *label* or *leaf*) correspond to terminals of the CFG, and represent a family of generators indexed by a lexicon for the language. The generators subscripted i (for *introducing a type*) correspond to rewrites of the CFG. Reading the central diagram in the main body from left-to-right, we additionally depict the breakdown of the derivation in terms of rewrites of lower dimension from our signature.

The essence of a tree-*adjoining* grammar is as follows: whereas for a CFG one grows the tree by appending branches and leaves at the top of the tree (substitution), in a TAG one can also sprout subtrees from the middle of a branch (adjoining). Now we show that this gloss is more formal than it sounds, by the following steps. First we show that the 2-categorical data of a CFG can be transformed into 3-categorical data – which we call *Leaf-Ansatz* – which presents a rewrite system that obtains the same sentences as the CFG, by a bijective correspondence between composition of 2-cells in the CFG and constructed 3-cells in the leaf-ansatz. These 3-cells in the leaf ansatz correspond precisely to the permitted *substitutions* in a TAG. Then we show how to model *adjoining* as 3-cells. Throughout we work with a running example, the CFG grammar introduced earlier. Here is the formal but unenlightening definition of *elementary* tree adjoining grammars which we will convert to diagrams. We will deal with the extensions of links and local constraints to adjoining shortly.

Definition 2.2.1 (*Elementary Tree Adjoining Grammar: Classic Computer Science style*). An elementary TAG is a tuple

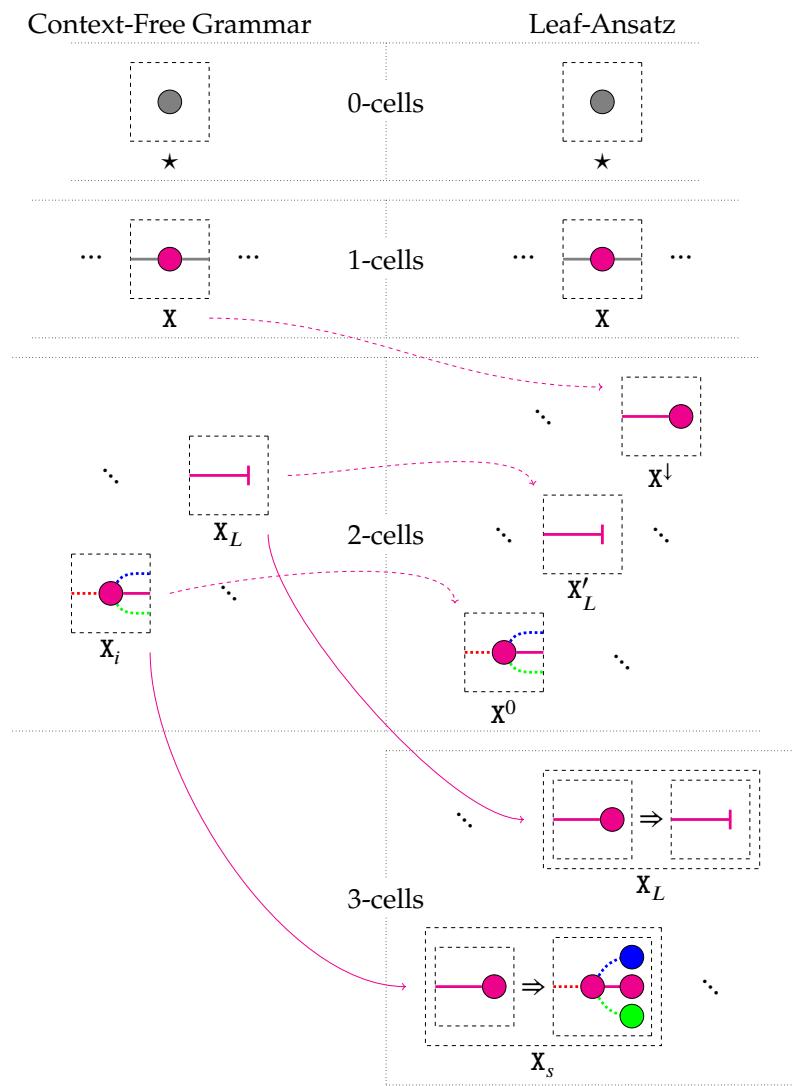
$$(\mathcal{N}, \mathcal{N}^\downarrow, \mathcal{N}^*, \Sigma, I, A)$$

The first four elements of the tuple are referred to as *non-terminals*. They are:

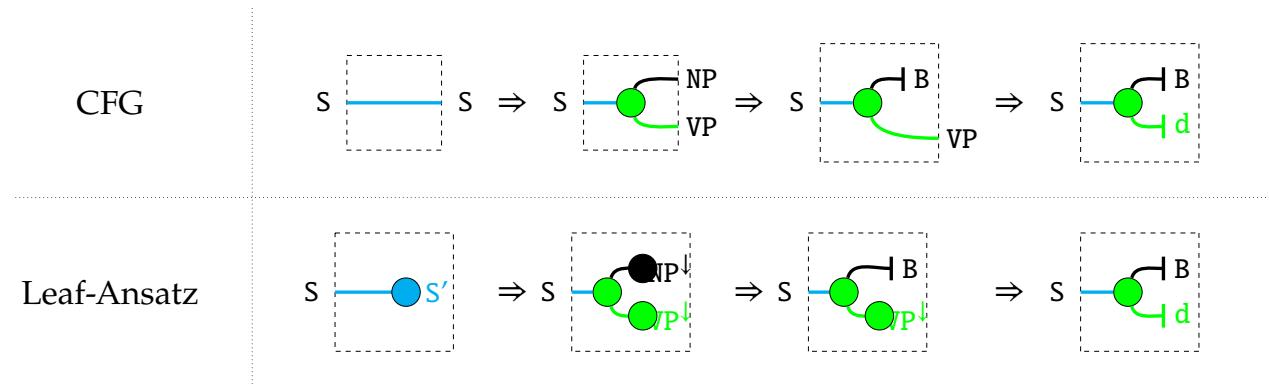
- A set of *non-terminal symbols* \mathcal{N} – these stand in for grammatical types such as NP and VP.
- A bijection $\downarrow: \mathcal{N} \rightarrow \mathcal{N}^\downarrow$ which acts as $X \mapsto X^\downarrow$. Nonterminals in \mathcal{N} are sent to marked counterparts in \mathcal{N}^\downarrow , and the inverse sends marked nonterminals to their unmarked counterparts. These markings are *substitution markers*, which are used to indicate when certain leaf nodes are valid targets for a substitution operation – discussed later.
- A bijection $*: \mathcal{N} \rightarrow \mathcal{N}^*$ – the same idea as above. This time to mark *foot nodes* on auxiliary trees, which is structural information used by the adjoining operation – discussed later.

2.2.2 Tree Adjoining Grammars

Construction 2.2.2 (Leaf-Ansatz of a CFG). Given a signature \mathfrak{G} for a CFG, we construct a new signature \mathfrak{G}' which has the same 0- and 1-cells as \mathfrak{G} . See the dashed magenta arrows in the schematic below: for each 1-cell wire type X of \mathfrak{G} , we introduce a *leaf-ansatz* 2-cell X^\downarrow . For each leaf 2-cell X_L in \mathfrak{G} , we introduce a renamed copy X'_L in \mathfrak{G}' . Now refer to the solid magenta: we construct a 3-cell in \mathfrak{G}' for each 2-cell in \mathfrak{G} , which has the effect of systematically replacing open output wires in \mathfrak{G} with leaf-ansatzes in \mathfrak{G}' .



Example 2.2.3. The leaf-ansatz construction just makes formal the following observation: there are multiple equivalent ways of modelling terminal symbols in a rewrite system considered string-diagrammatically. One way (which we have already done) is to treat non-terminals as wires and terminals as effects, so that the presence of an open wire available for composition visually indicates non-terminality. Another (which is the leaf-ansatz construction) treats all symbols in a rewrite system as leaves, where bookkeeping the distinction between (non-)terminals occurs in the signature. So for a sentence like Bob drinks, we have the following derivations that match step for step in the two ways we have considered.



Proposition 2.2.4 (Leaf-ansatzes of CFGs are precisely TAGs with only initial trees and substitution). *Proof.*

By construction. Consider a CFG given by 2-categorical signature \mathfrak{G} , with leaf-ansatz signature \mathfrak{G}' . The types X of \mathfrak{G} become substitution marked symbols X^\downarrow in \mathfrak{G}' . The trees X_i in \mathfrak{G} become initial trees X^0 in \mathfrak{G}' . The 3-cells X_s of \mathfrak{G}' are precisely substitution operations corresponding to appending the 2-cells X_i of \mathfrak{G} . \square

Σ is a set of *terminal symbols* – these stand in for the words of the natural language being modelled. I and A are sets of *elementary trees*, which are either *initial* or *auxiliary*, respectively. *initial trees* satisfy the following constraints:

- The interior nodes of an initial tree must be labelled with nonterminals from \mathcal{N}
- The leaf nodes of an initial tree must be labelled from $\Sigma \cup \mathcal{N}^\downarrow$

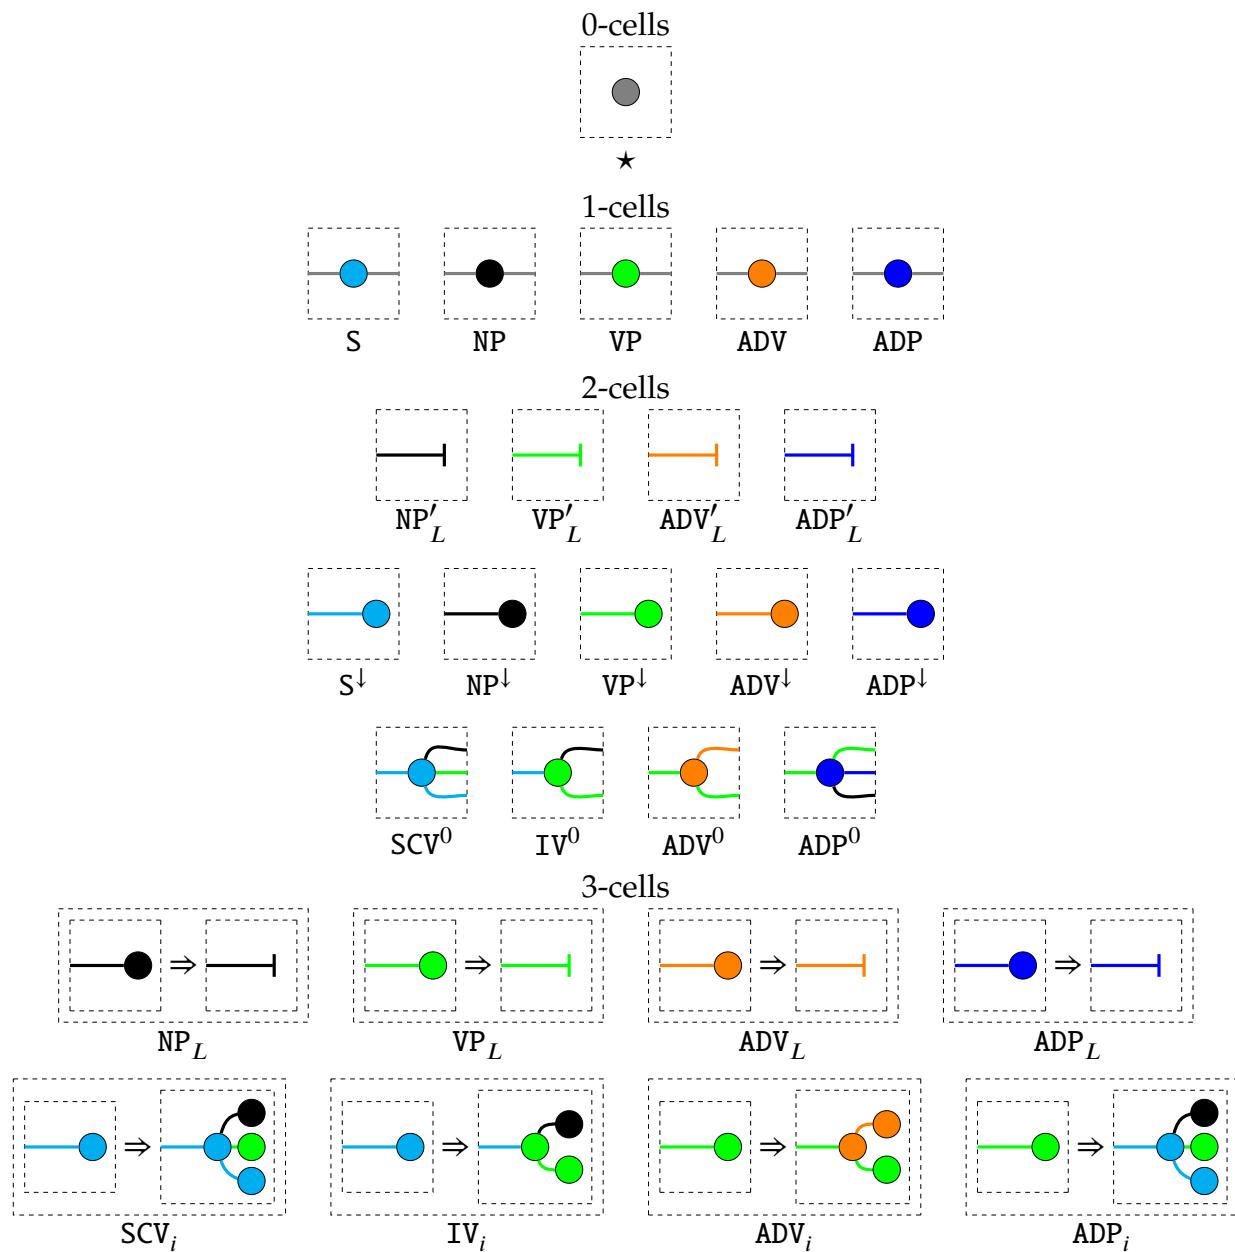
Auxiliary trees satisfy the following constraints:

- The interior nodes of an auxiliary tree must be labelled with nonterminals from \mathcal{N}
- Exactly one leaf node of an auxiliary tree must be labelled with a foot node $X^* \in \mathcal{N}^*$; moreover, this labelled foot node must be the marked counterpart of the root node label of the tree.
- All other leaf nodes of an auxiliary tree are labelled from $\Sigma \cup \mathcal{N}^\downarrow$

Further, there are two operations to build what are called *derived trees* from elementary and derived trees. These operations are called *substitution* and *adjoining*.

- *Substitution* replaces a substitution marked leaf node X^\downarrow in a tree α with another tree α' that has X as a root node.
- *Adjoining* takes auxiliary tree β with root and foot nodes X, X^* , and a derived tree γ at an interior node X of γ . Removing the X node from γ separates it into a parent tree with an X -shaped hole for one of its leaves, and possibly multiple child trees with X -shaped holes for roots. The result of adjoining is obtained by identifying the root of β with the X -context of the parent, and making all the child trees children of β 's foot node X^* .

Figure 2.7: Leaf-ansatz signature of Alice sees Bob quickly run to school CFG



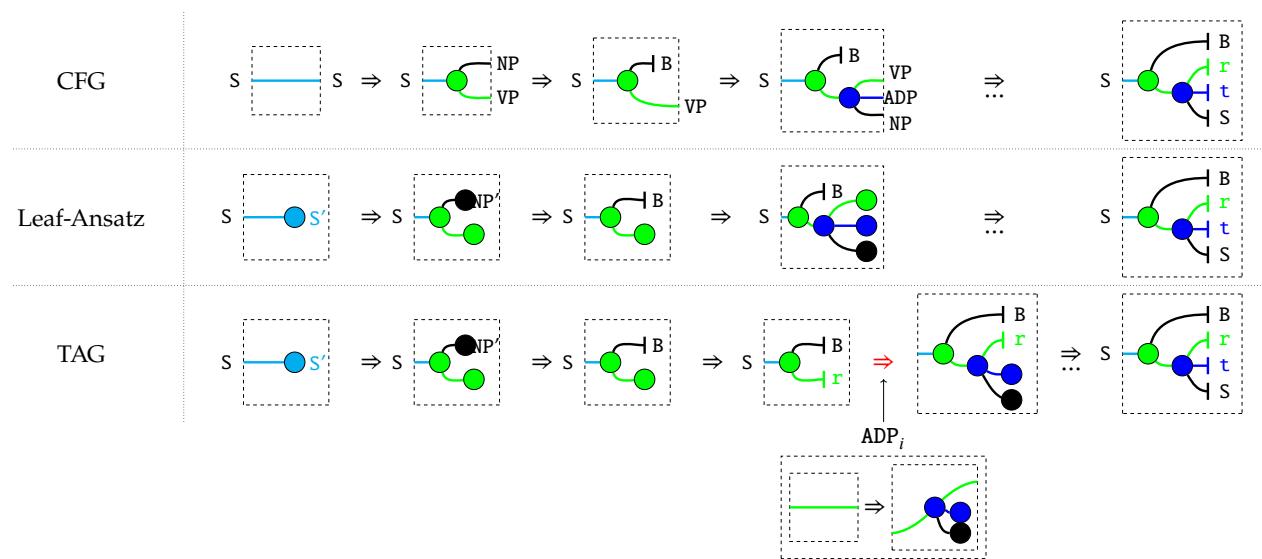


Figure 2.8: Adjoining is sprouting subtrees in the middle of branches. One way we might obtain the sentence Bob runs to school is to start from the simpler sentence Bob runs, and then refine the verb runs into runs to school. This refinement on part of an already completed sentence is not permitted in CFGs, since terminals can no longer be modified. The adjoining operation of TAGs gets around this constraint by permitting rewrites in the middle of trees.

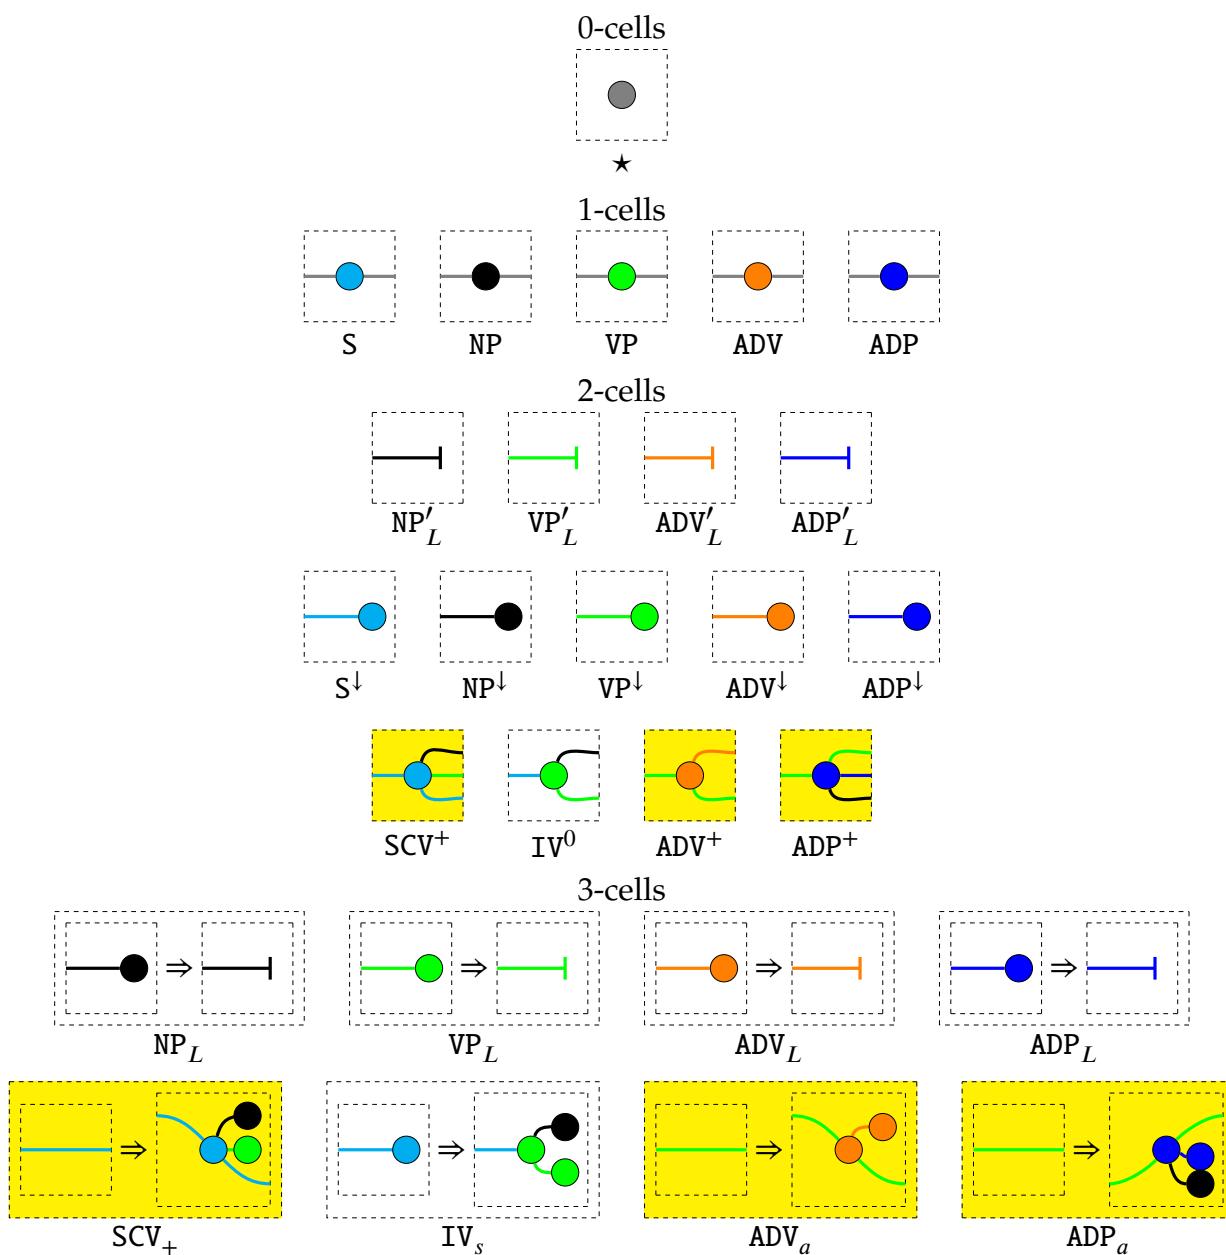


Figure 2.9: TAG signature of Alice sees Bob quickly run to school. The highlighted 2-cells are auxiliary trees that replace CFG 2-cells for verbs with sentential complement, adverbs, and adpositions. The highlighted 3-cells are the tree adjoining operations of the auxiliary trees.

The construction yields as a corollary an alternate proof of Theorem [Joshi 6.1.1]...

Corollary 2.2.5. For every context-free grammar \mathfrak{G} there exists a tree-adjoining grammar \mathfrak{G}' such that \mathfrak{G} and \mathfrak{G}' are strongly equivalent – both formalisms generate the same set of strings (weak equivalence) and the same abstract syntactic structures (in this case, trees) behind the strings (strong equivalence).

Proof. Proposition 2.2.4 provides one direction of both equivalences. For the other direction, we have to show that each auxiliary tree (a 2-cell) and its adjoining operation (a 3-cell) in \mathfrak{G}' corresponds to a single 2-cell tree of some CFG signature \mathfrak{G} , which we demonstrate by construction. See the example above; the highlighted 3-cells of \mathfrak{G}' are obtained systematically from the auxiliary 2-cells as follows: the root and foot nodes X, X^* indicate which wire-type to take as the identity in the left of the 3-cell, and the right of the 3-cell is obtained by replacing all non- X open wires Y with their leaf-ansatzes Y^\downarrow . This establishes a correspondence between any 2-cells of \mathfrak{G} considered as auxiliary trees in \mathfrak{G}' . \square

2.2.3 Tree adjoining grammars with local constraints

The usual conception of TAGs includes two extensions to the basic definition presented above. First, there may be *local constraints* on adjoining, which only allows certain trees to be adjoined at certain nodes. Second, TAGs may have links, which are extra edges between nodes obeying a c-command condition. Here we deal with local constraints; dealing with links requires the introduction of braiding.

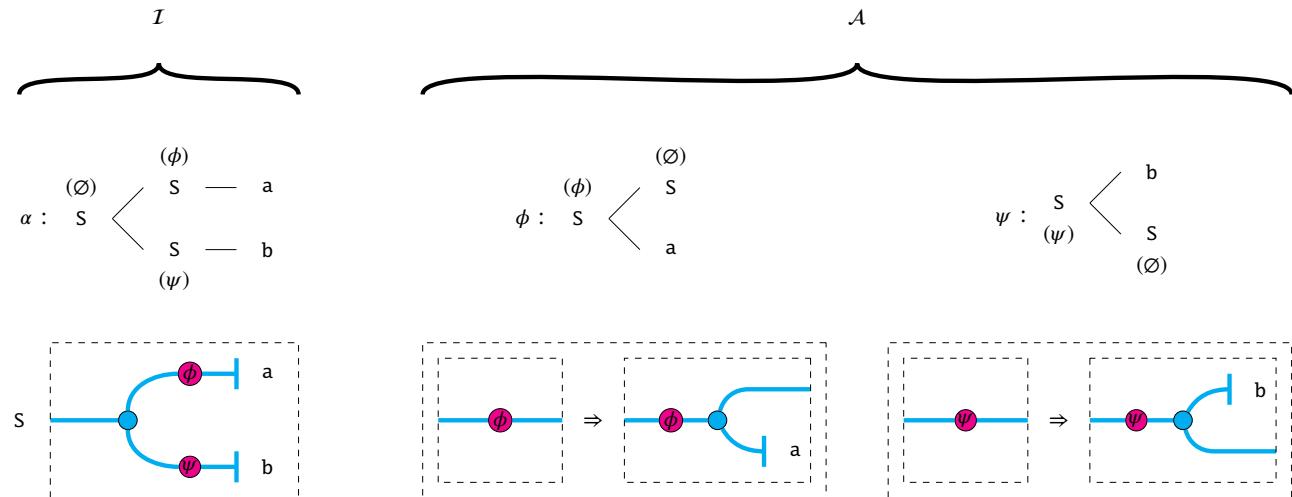
The n -categorical approach easily accommodates local constraints. For (SA), whereas before we take the source of adjoining rewrites to be identities, we can instead introduce ansatz endomorphisms that are rewritable to the desired subsets. For (NA), we assert that identities have no rewrites (NA). For (OA), we can make a distinction between finished and unfinished derivations, where we require that unfinished derivations are precisely those that still contain an obligatory-rewrite endomorphism.

Example 2.2.7 (Selective and null adjoining diagrammatically). The below is a reproduction of Example 2.5 of [Joshi] which demonstrates the usage of selective and null adjoining. The notation from [Joshi] is presented first, followed by their corresponding representations in an n -categorical signature. The initial tree is presented as a 2-cell where the (SA) rules are rewritable nodes, that serve as sources of rewrites in the 3-cell

Definition 2.2.6 (TAG with local constraints: CS-style).
[Joshi] $G = (I, A)$ is a TAG with local constraints if for each node n and each tree t , exactly one of the following constraints is specified:

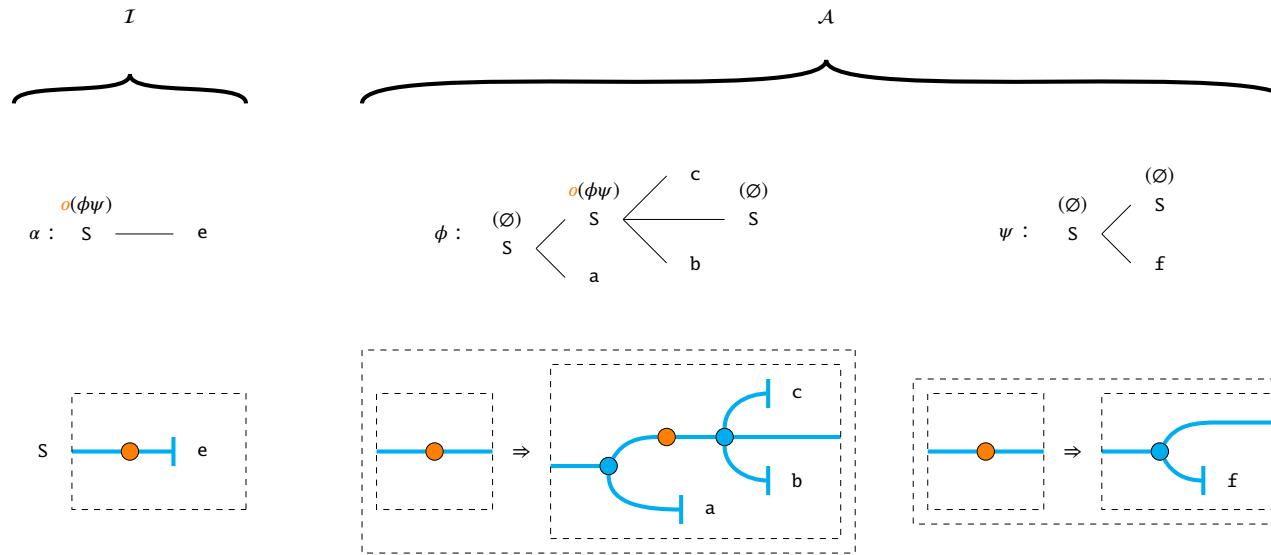
1. Selective adjoining (SA): Only a specified subset $\bar{\beta} \subseteq A$ of all auxiliary trees are adjoinable at n .
2. Null adjoining (NA): No auxiliary tree is adjoinable at the node n .
3. Obligatory Adjoining (OA): At least one out of all the auxiliary trees adjoinable at n must be adjoined at n .

presentations of the auxiliary trees.



Example 2.2.8 (Obligatory adjoining diagrammatically). The below is a reproduction of Example 2.11 of [Joshi] which demonstrates the usage of obligatory adjoining, marked orange. The notation from [Joshi] is presented first, followed by their corresponding representations in an n -categorical signature. The initial tree is presented as a 2-cell where the (OA) rule is given its own 2-cell, which is the source of rewrites in 3-cell presentations of auxiliary trees. We may capture the obligatory nature of the rewrite by asking that finished

derivations contain no instance of the orange 2-cell.



2.2.4 Braiding, symmetries, and suspension

Before we can model TAGs with links, we must introduce the concepts of braiding and symmetries, which we have seen in the diagrammatic setting already as wires twisting past one another. In our current setting of 1-object-3-categories for TAGs, the diagrams we are dealing with are all planar – i.e. wires may not cross – but this is a restriction we must overcome when we are dealing with links.

First we observe that in a 1-object-2-categorical setting, a morphism from the identity ϵ on the base object \star to itself would, in our analogy with string-rewrite systems, be a rewrite from the empty string to itself.

For example, a rewrite R may introduce a symbol from the empty string and then delete it. A rewrite S may create a pair of symbols from nothing and then annihilate them.

$$R := \epsilon \mapsto x \mapsto \epsilon \quad S := \epsilon \mapsto a \cdot b \mapsto \epsilon$$

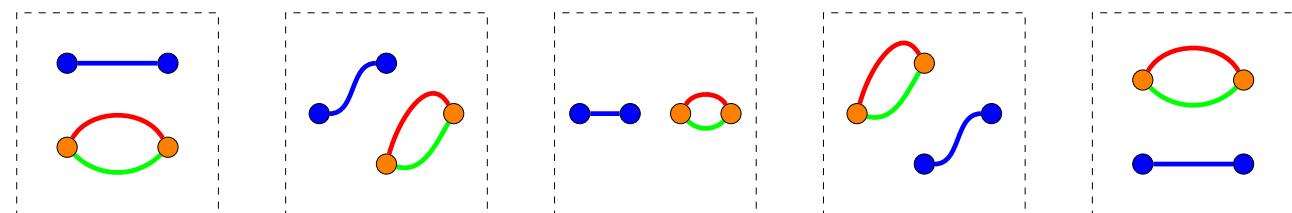
In our analogy with string rewrite systems, we might like that the following rewrites are equivalent, while respecting that they are not equal.

1. Start and finish R and S concurrently, R on the left and S on the right: $\epsilon = \epsilon \cdot \epsilon \mapsto x \cdot a \cdot b \mapsto \epsilon \cdot \epsilon = \epsilon$
2. Start R first. Start S on the right concurrently as R finishes: $\epsilon \mapsto x = x \cdot \epsilon \mapsto \epsilon \cdot a \cdot b = a \cdot b \mapsto \epsilon$
3. Start and finish R first, and then start and finish S : $\epsilon \mapsto x \mapsto \epsilon \mapsto a \cdot b \mapsto \epsilon$

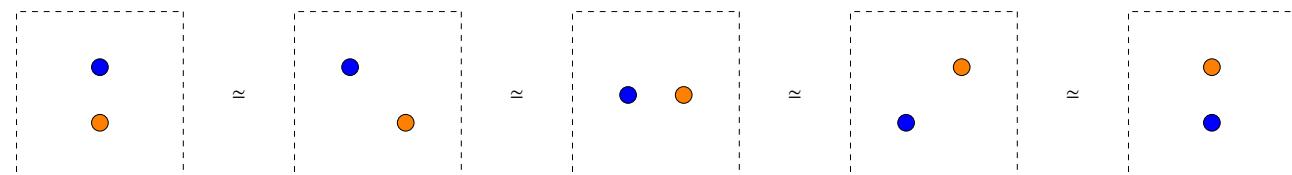
4. Start S first. Start R on the right concurrently as S finishes: $\varepsilon \mapsto a \cdot b = a \cdot b \cdot \varepsilon \mapsto \varepsilon \cdot x = x \mapsto \varepsilon$

5. Do S and R concurrently, S on the left and R on the right: $\varepsilon = \varepsilon \cdot \varepsilon \mapsto a \cdot b \cdot x \mapsto \varepsilon \cdot \varepsilon = \varepsilon$

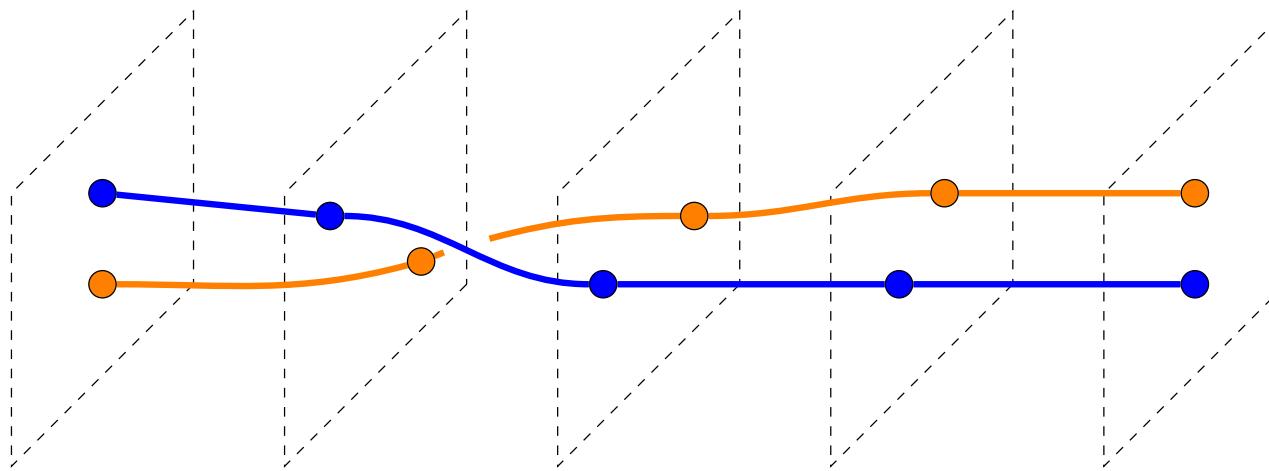
Diagrammatically, we may represent these rewrites from left to right as follows, representing x, a, b as blue, red, and green wires respectively:



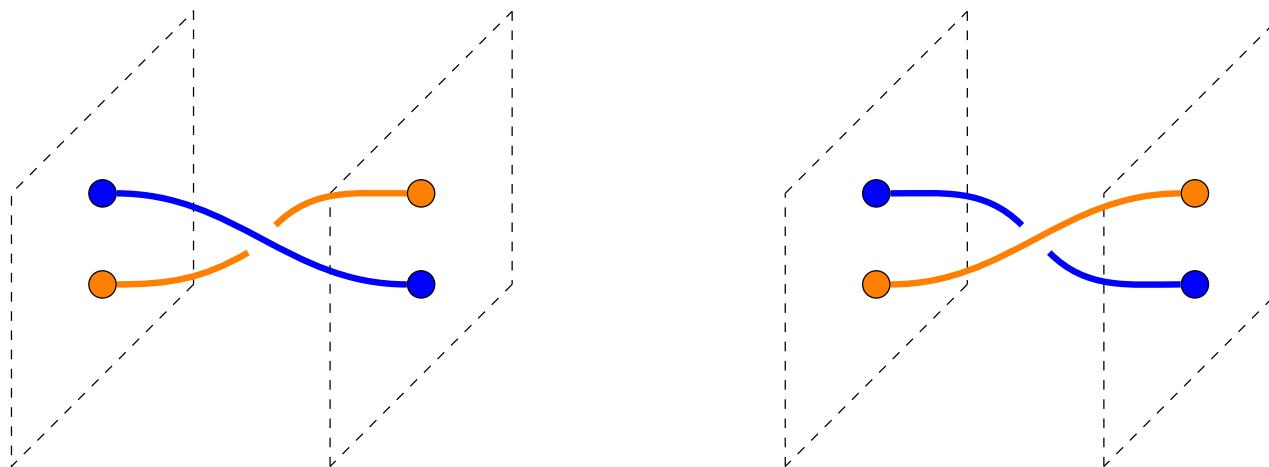
Such rewrites from the empty string to itself are more generally called *scalars* in the monoidal setting, viewed 2-categorically. We may generally represent such scalars as labelled dots. A fact about scalars in a 1-object-2-category called the Eckmann-Hilton argument [CITE](#) is that dots may circle around one another, and all of those expressions are equivalent up to homotopy. The mechanism that enables this in our setting is that the empty string is equal to copies of itself, which creates the necessary space for manoeuvring; translating into the n -categorical setting, expressions are equivalent up to introducing and contracting identities.



We may view the homotopies that get us from one rewrite to another as 3-cells, which produces a braid in a pair of wires when viewed as a vignette.

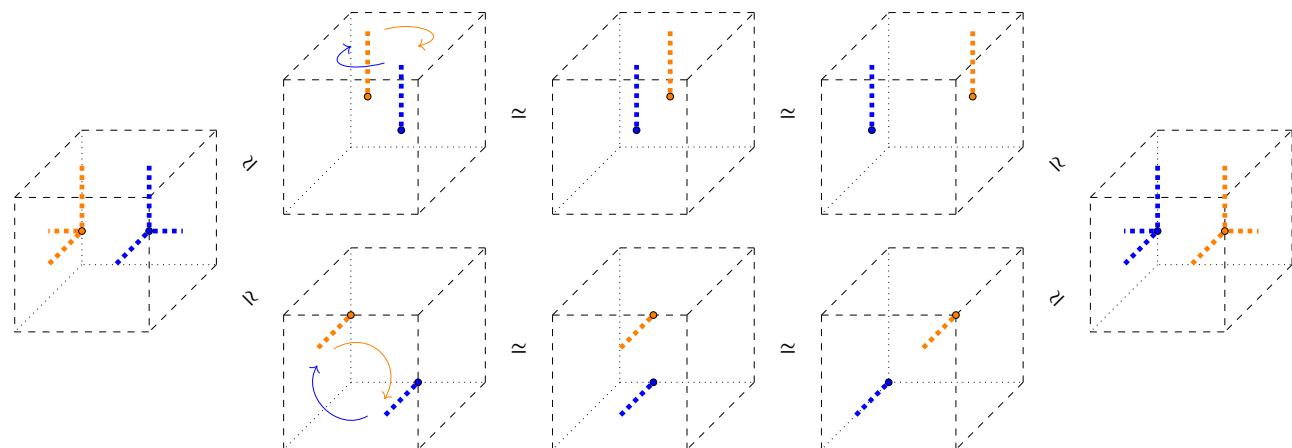


Up to processive isotopies CITE , which are continuous bijective transformations that don't let wires double back on themselves, we can identify two different braidings that are not continuously deformable to one another in the 3-dimensional space of the vignette. We may depict the two braidings as follows, where wires either go over or under one another.



Now we have a setting in which we can consider wires swapping places. However, having two different ways to swap wires presents a complication. While useful for knot-theory, in symmetric monoidal categories we don't want this distinction. Perhaps surprisingly, this distinction is eliminated if we consider swapping

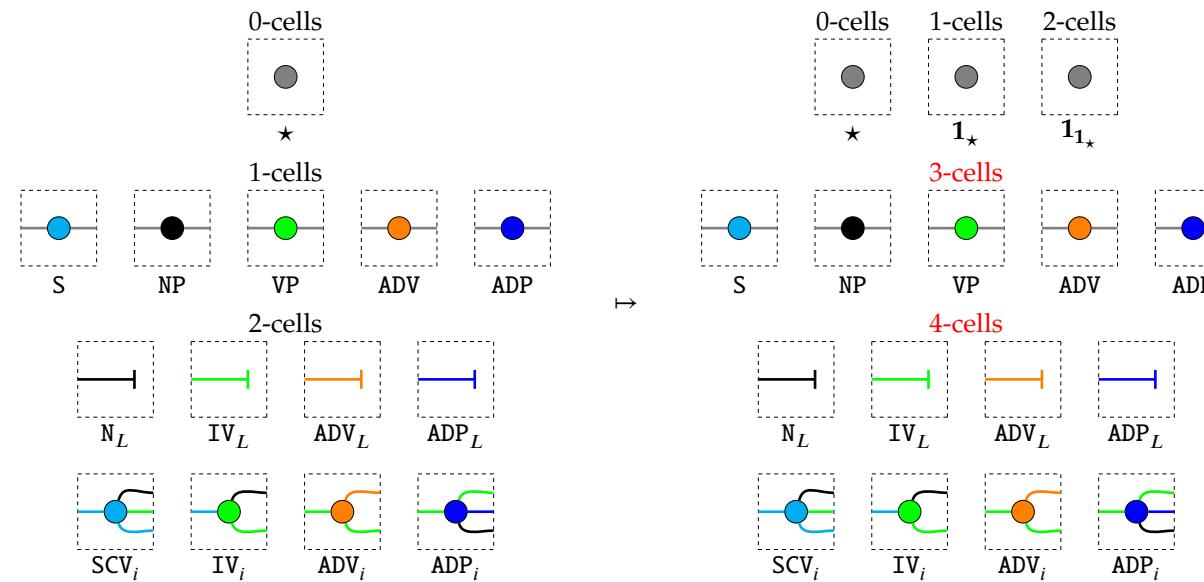
dots in the 3D volume rather than the 2D plane. Now we have to extend our analogy to reach 1-object-4-categories. Just as symbols on a 1D string were encoded as 1-cells on the 0-cell identity initially, our braidings are the behaviour of symbols on the 2D plane, encoded as 2-cells on the 1-cell identity of the 0-cell. So, to obtain symbols in a 3D volume, we want 3-cells on the 2-cell identity of the 1-cell identity of the 0-cell. That is a mouthful, so we instead clumsily denote the stacked identity as 1_{1_\star} . To obtain a dot in a 3-dimensional volume, we consider a rewrite from 1_{1_\star} to itself. These dots also enjoy a version of the Eckmann-Hilton argument in 3 dimensions (which holds for all dimensions 2 and higher). We can depict these swaps by movements in a cubic volume where each axis corresponds to a direction of composition. Whereas on the plane the dots have two ways to swap places – clockwise and counterclockwise rotation – in the volume they have two new ways to swap places – clockwise and counterclockwise in the new dimension. Shown below are two ways to swap left-to-right sequentially composed dots by clockwise rotations in the forward-backward and up-down directions of composition:



Considering these changes of dot positions in 3D as a 4D vignette gives us braidings again. But this time, they are equivalent up to processive isotopy – in other words, any two ways of swapping the dots in the volume are continuously deformable to one another in the 4-dimensional space of the vignette. The reason is intuitive: a clockwise and counterclockwise braiding along one composition axis can be mutually deformed by making use of the extra available axis of composition. So we have eliminated the differences between the two kinds of braidings.

To recap: We follow the convention that object dimensions start at 0. In a 1-object-2-category we obtain planar string diagrams for monoidal categories, which are equivalent up to processive isotopies in an ambient 2-dimensional space. In a 1-object-3-category, we obtain string diagrams for braided monoidal categories that are equivalent up to processive isotopy in an ambient 3-dimensional space – the page with depth. In a

1-object-4-category, we obtain string diagrams for symmetric monoidal categories, which are equivalent up to processive isotopies in an ambient 4-dimensional space – while it is difficult to visualise 4-dimensionally, diagrammatically this simplifies dramatically to just allowing wires to cross and only caring about whether input-output connectivity from left to right makes sense. Now observe that for any 1-object-2-category, we can obtain a 1-object-4-category by promoting all 2-cells and higher to sit on top of $1_{1\star}$, in essence turning dots on a line into dots in a volume; this procedure is called *suspension* CITE . For example, taking our CFG signature from earlier, suspension promotes 1-cells to 3-cells and 2-cells to 4-cells. The resulting signature gives us the same diagrams, now with the added ability to consider diagrams equivalent up to twisting wires, which models a string-rewrite system with free swapping of symbol order.



We call the above a 1-object-4-category since there is a single 0-cell object, and the highest dimension we consider is 4. Symmetric monoidal categories are equivalently seen as 1-object-4-categories CITE , which are in particular obtained by suspending 1-object-2-categories for planar string diagrams. To summarise, by appropriately suspending the signature, we obtain a formal way to power up our diagrams to permit twisting wires, as in symmetric monoidal categories.

Remark 2.2.9 (THE IMPORTANT TAKEAWAY!). Now have a combinatoric way to specify string diagrams that generalises PROPs for symmetric monoidal categories, which in addition permits us to reason with directed diagram rewrites. So we have a trinity of presentations of the same mathematical concept. String diagrams are the *syntax*. Symmetric monoidal categories are the *semantics*. n -categories are the *finite combi-*

natoric presentation of theories and rewrite systems upon them. By analogy, n -categories provide a *vocabulary* and *rewrites* for string-diagram *expressions*, which are interpreted with respect to the *context* of a symmetric monoidal category.

2.2.5 TAGs with links

Definition 2.2.10 (TAGs with links: CS-style). The following is a reproduction of the discussion on nodes with links from [Joshi]. Linking can be defined for any two nodes in an elementary tree. However, in the linguistic context we will require the following conditions to hold for a link in an elementary tree. If a node n_1 is linked to a node n_2 then:

1. n_2 *c-commands* n_1 , (i.e., n_2 is not an ancestor of n_1 , and there exists a node m which is the immediate parent node of n_2 , and an ancestor of n_1).
2. n_1 and n_2 have the same label.
3. n_1 is the parent only of a null string, or terminal symbols.

A TAG *with links* is a TAG in which some of the elementary trees may have links as defined above.

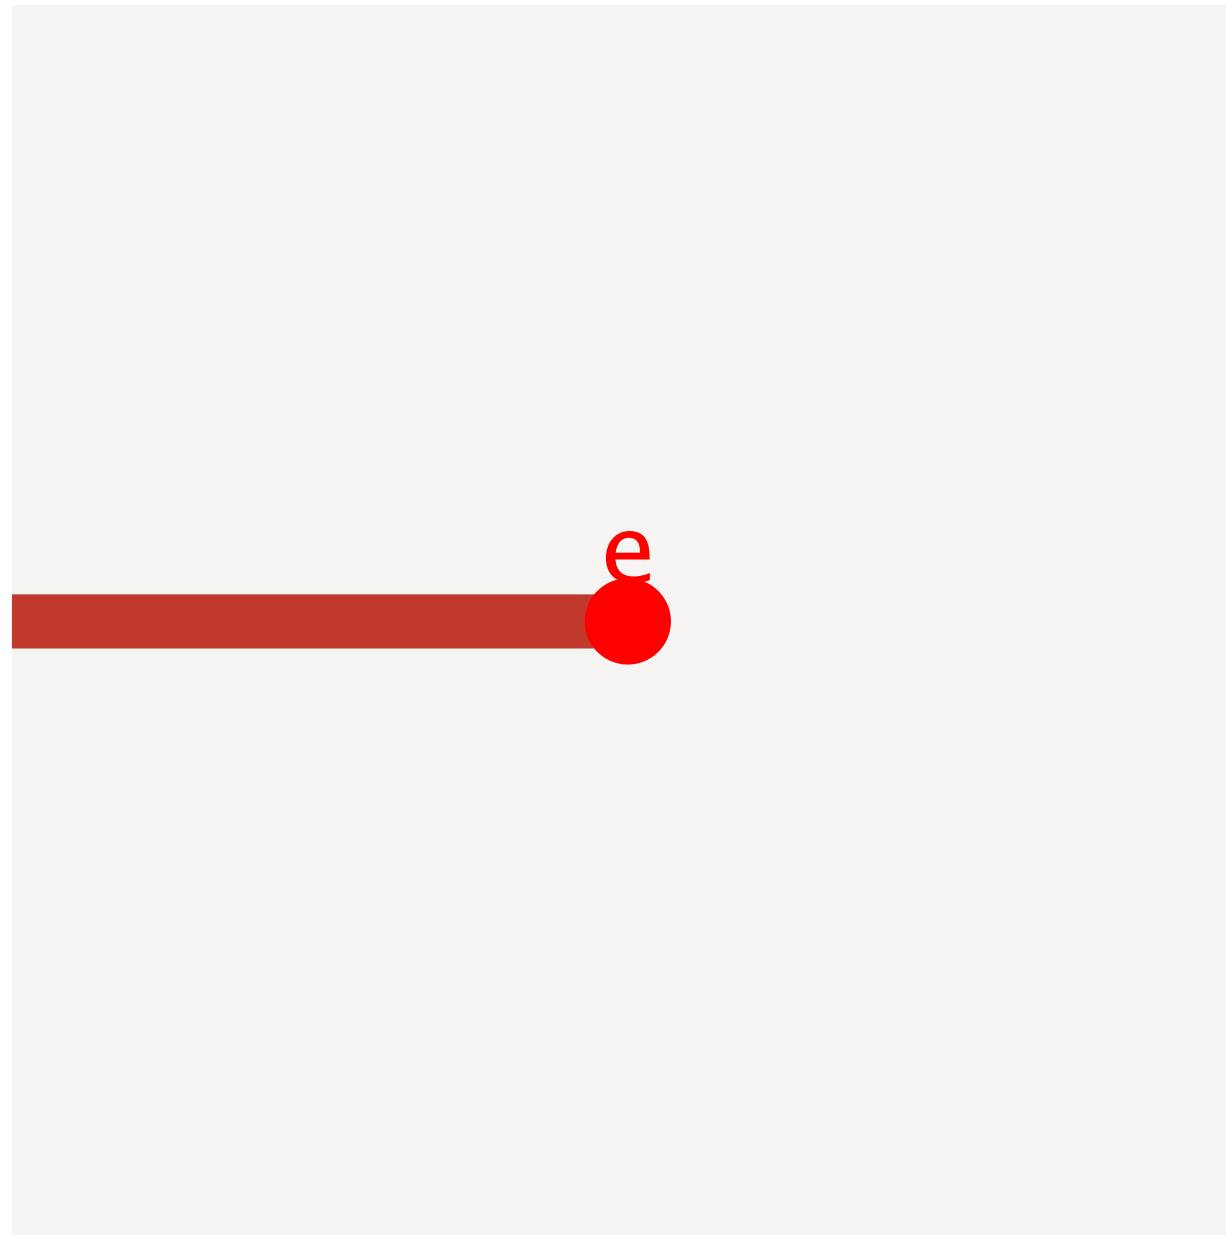
Example 2.2.11. The following is a reproduction in homotopy.io of Example 2.4 of [Joshi], illustrating TAGs with links; as a proof assistant for weak n -categories, homotopy.io automatically performs tree adjunction after the signature has been properly spelled out. The trees given in Examples 2.4 are:

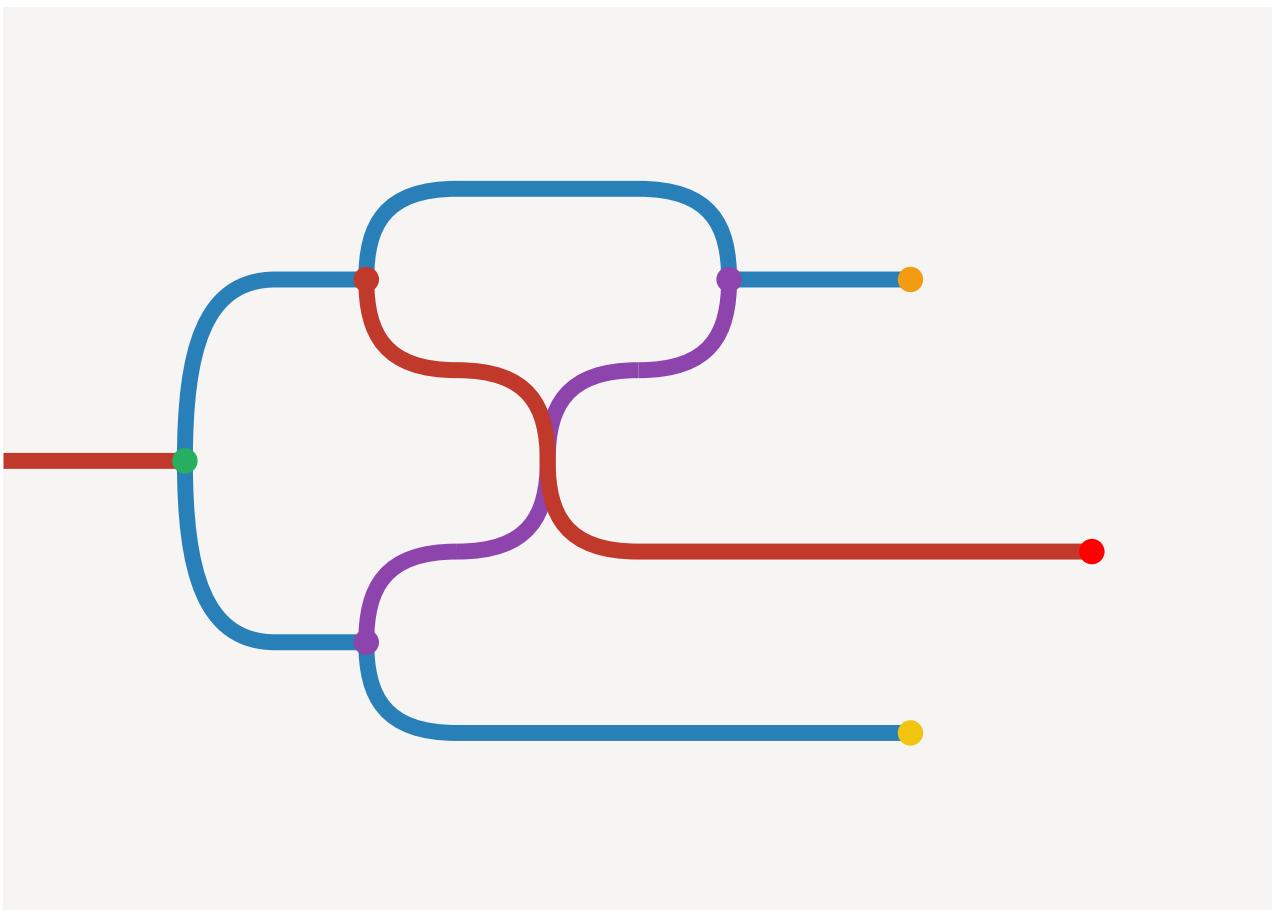
placeholder

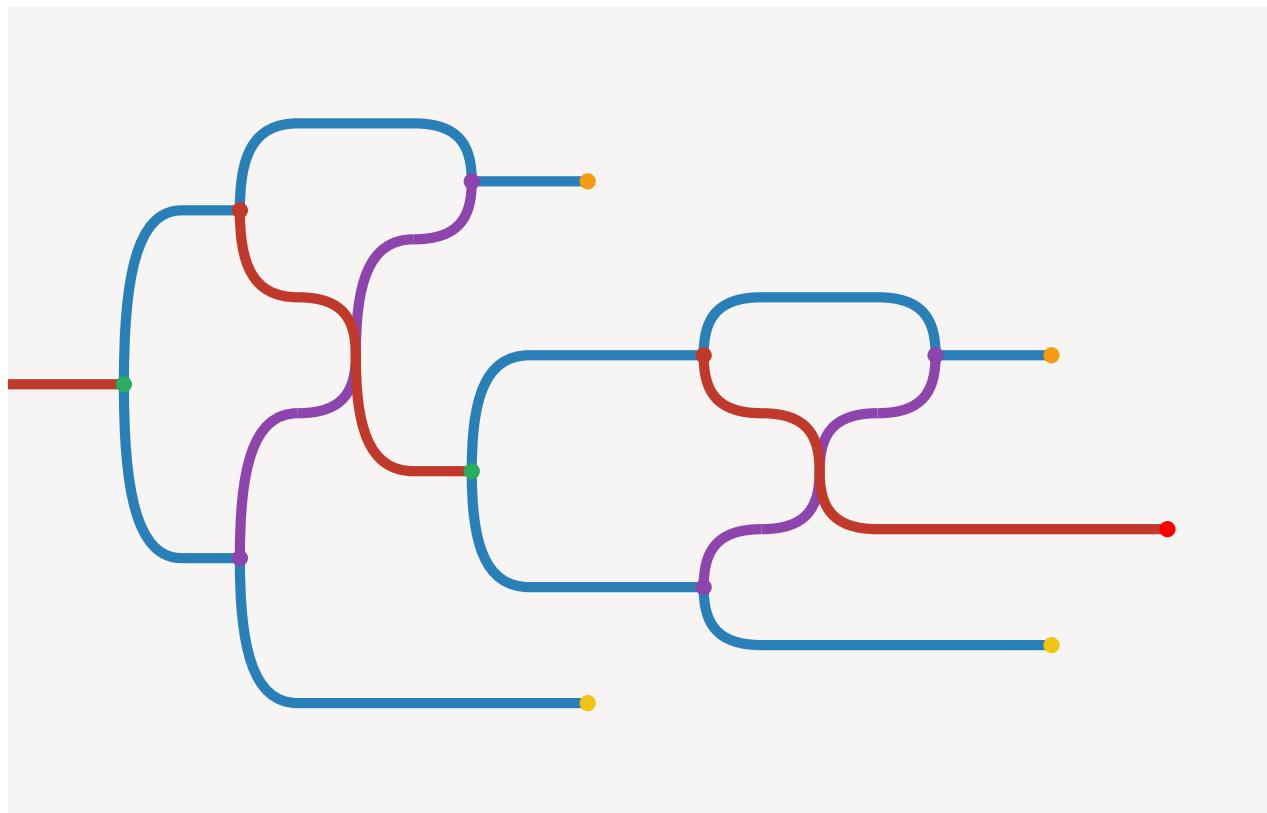
Which we interpret as expressions comprised of the following signature:

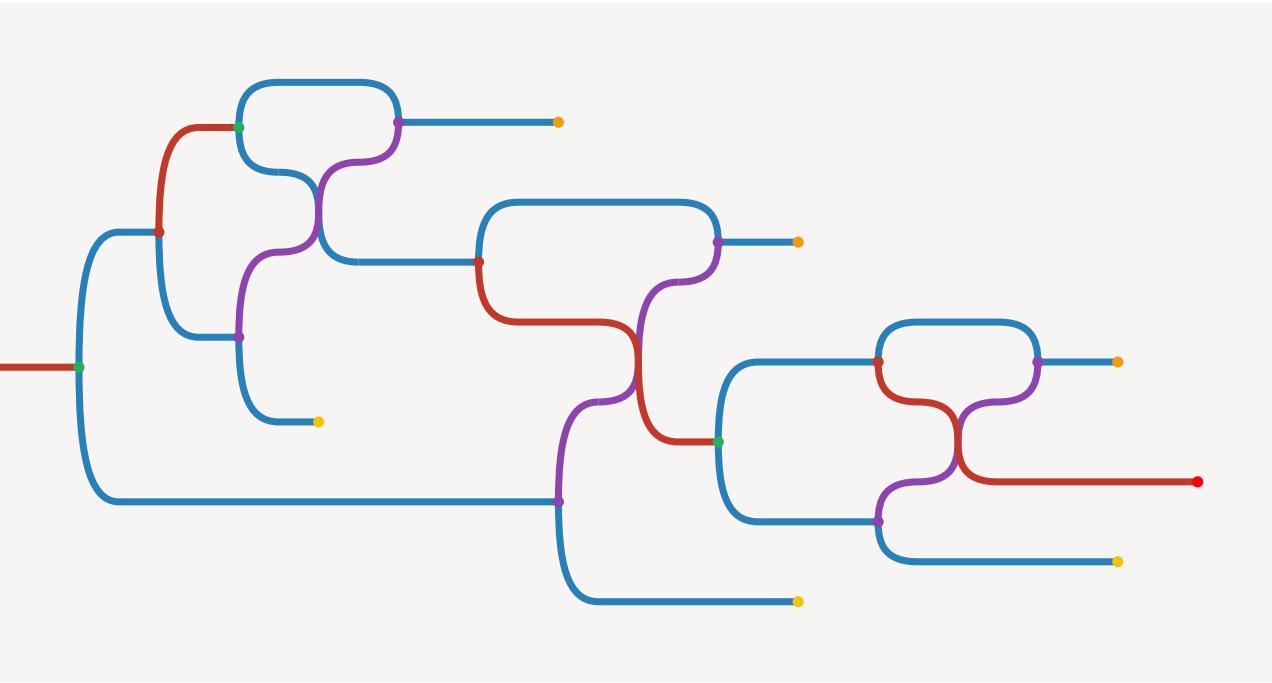
placeholder

In the process of interpretation, we introduce a link wire-type (in purple), and include directed link generation and elimination morphisms for the T wire-type (in blue). A necessary step in the process of interpretation (which for us involves taking a Poincaré dual to interpret nodes as wires) is a typing assignment of the tree-branches connected to terminal nodes, which we have opted to read as sharing a T -type for minimality, though we could just as well have introduced a separate label-type wire.









N.B. In practice when using `homotopy.io` for the symmetric monoidal setting, it is simpler to suspend symmetric monoidal signatures to begin at 4-cells rather than 3-cells. The reason for this is that under- and over-braids still exist in the symmetric monoidal setting, and while sequentially composed braids are homotopically equivalent to the pair of identities, they are not uniquely so, thus these homotopies must be input manually. By beginning at 4-cells (or higher, due to the stabilisation hypothesis CITE), braid-eliminations are unique up to homotopy and can be performed more easily in the proof assistant.

Now we have enough to spell out full TAGs with local constraints and links as an n -categorical signature. To recap briefly, we have seen that we can model the passage from CFGs to TAGs in the n -categorical setting, and then how selective and null adjoining rules by specifying endomorphisms on wires as the source of rewrites, and finally that we can model links in the symmetric monoidal setting. Maintaining planarity (which is important for a left-to-right order of terminals for spelling out a language) and modelling obligatory rewrites can be done by requiring that *finished* derivations are precisely those whose only twists are link-wires, and have no sources for obligatory adjoins.

Definition 2.2.12 (Tree Adjoining Grammars with local constraints and links in `homotopy.io`). A *Linguistic Tree Adjoining Grammar* is given by the following data:

$$(\mathcal{N}, \mathcal{N}^\downarrow, \mathcal{N}^*, \Sigma, \mathcal{I}, \mathcal{A}, \mathfrak{S} \subseteq \mathcal{P}(\mathcal{A}), \square, \lozenge, \mathfrak{L} \subseteq \mathcal{N})$$

The initial elements are the same as an elementary tree-adjoining grammar and obey the same constraints.

The modifications are:

- \mathcal{I} is a nonempty set of *initial* constrained-linked-trees.
- \mathcal{A} is a nonempty set of *auxiliary* constrained-linked-trees.
- \mathfrak{S} is a set of sets of *select* auxiliary trees.
- \square, \diamond are fresh symbols. \square marks *obligatory adjoins*, and \diamond marks *optional adjoins*.
- \mathfrak{L} is a set permissible *link types* among nonterminals or T .

A *constrained-linked-tree*, or CL-tree, is a pair consisting of:

- A tree where each internal node is an element of $\mathcal{N} \times \mathfrak{S} \times \{\square, \diamond\} \times \{*, \bar{*}\}$, and each leaf is an element of $\mathcal{N} \times \mathfrak{S} \times \{\square, \diamond\} \times \{*, \bar{*}\} \cup \Sigma$. In prose, each label is either a terminal symbol (as a leaf), or otherwise a nonterminal, along with a subset of auxiliary trees that indicates select adjoins (or null-adjoins when the subset is \emptyset), a marker indicating whether those adjoins are obligatory or optional, and a marker indicating whether the node is a foot node or not. Observe that there is no need to indicate when a node is a valid target for substitution since that function is subsumed by null adjoining rules.
- A set of ordered pairs of nodes (n_1, n_2) of the tree such that:
 1. n_2 *c-commands* n_1 , (i.e., n_2 is not an ancestor of n_1 , and there exists a node m which is the immediate parent node of n_2 , and an ancestor of n_1).
 2. n_1 and n_2 share the same type $T \in \mathcal{N}$ and $T \in \mathfrak{L}$, or both n_1, n_2 are terminals.
 3. n_1 is the parent of terminal symbols, or childless.

We spell out how this data becomes an n -categorical signature by enumerating cell dimensions:

0. A single object \star
1. None.
2. None.
3. • For each $T \in \mathcal{N}$, a cell $T : \mathbf{1}_{1_\star} \rightarrow \mathbf{1}_{1_\star}$.
 - $T : \mathbf{1}_{1_\star} \rightarrow \mathbf{1}_{1_\star}$ A wire for terminal symbols.
 - For each $L \in \mathfrak{L}$, a cell $L : \mathbf{1}_{1_\star} \rightarrow \mathbf{1}_{1_\star}$.

4. We distinguish between *atomic* and *composite* generators at this level. The trees themselves are composites of *atomic* generators. The *atomic* generators are:
 - For each node n that occurs in either \mathcal{I} or \mathcal{A} , we populate cells by a case analysis:
 - If n is a terminal $\sigma \in \Sigma$, we create a cell $\sigma : T \rightarrow \mathbf{1}_{1_{\star}}$.
 - If $n = (T, S, \dagger, \bar{*})$, we create $S_T^\square : T \rightarrow T$ if the node is marked obligatory ($\dagger = \square$), and a cell $S_T^\diamond : T \rightarrow T$ otherwise.
 - If $n = (T, S, \dagger, *)$, it is a foot note, for which we create a cell $S_T^\square : T \rightarrow \mathbf{1}_{1_{\star}}$ if the node is marked obligatory ($\dagger = \square$), and a cell $S_T^\diamond : T \rightarrow \mathbf{1}_{1_{\star}}$ otherwise.
 - For each $L \in \mathfrak{L}$ (which is also a type $T \in \mathcal{N}$) a pair of cells $T^L := T \rightarrow T \otimes L$ and $T_L := L \otimes T \rightarrow T$.
 - For each node p of type T_p in either \mathcal{I} or \mathcal{A} with a nonempty left-to-right list of children $C_p := < c_1, c_2, \dots, c_i, c \dots c_n >$ with types T_i , a branch cell $C_p : T_p \rightarrow \bigotimes_{i=1}^n T_i$.

We represent trees by composite generators, defined recursively. For a given tree \mathcal{T} in either \mathcal{I} or \mathcal{A} , we define a composite generator beginning at the root. Where the root node is $p = (T, S, \dagger)$, we begin with the cell S_T^\dagger . For branches, we compose the branch cell C_p to this cell sequentially. If p has a child c that has a link, we do a case analysis. If that child c -commands the other end of the link we generate the first half of the linking wire by composing T^L for the appropriate type T of the child node. Otherwise the child is c -commanded by a previously generated link, which we braid over and connect using T_L , again with the appropriate typing for the child. Now we may recurse the procedure for subtrees. If a node has no children, it is a leaf $l = (T, S, \dagger)$ or a terminal symbol. We append a terminal cell σ if l is a terminal symbol (thus killing the wire), and otherwise we leave an open T wire after appending S_T^\dagger . Altogether this obtains a 3-cell which we denote \mathcal{T} , overloading notation; different execution-orders of the above procedure evidently obtain 3-cells equivalent up to homotopy.

5.

3

String Diagrams for Text

3.1 Previously, on DisCoCat

DisCoCat is a research programme in applied mathematical linguistics that is **Distributional**, **Compositional** and **Categorical**. In this section I will recount a selective development of DisCoCat as relevant for this thesis.

3.1.1 Lambek's Linguistics

It's hard for me to do justice to Jim Lambek's life. I feel as if have been in intimate conversation with Jim throughout my research, despite our separation by time. Anyone can look up the Curry-Howard-Lambek correspondence and follow the rabbit hole to see Jim's broad reach and lasting impact on category theory. I know that he was a jovial man who always carried a good sense of humour and a wad of twenties. I also can't do better than Moortgat's history and exposition of typological grammar in [CITE](#), so I will borrow Moortgat's phrasing and summarise Lambek's role in the story. Typological grammar originated in two seminal papers by Lambek in 1958 and 1961 [CITE](#), where Lambek sought "to obtain an effective rule (or algorithm) for distinguishing sentences from non-sentences, which works not only for the formal languages of interest to the mathematical logician, but also for natural languages [...]"'. The method is to assign grammatical categories – parts of speech such as nouns and verbs – logical formulae. Whether a sentence is grammatical or not is obtained from deduction using these logical formulae in a Gentzen-style sequent proof.

Figure 3.1: In English, we may consider a noun to have type n , and an transitive english verb $(n/s) \setminus n$, to yield a well-formedness proof of Bob drinks beer. The type formation rules for such a grammar are intuitive. Apart from a stock of basic types \mathbb{B} that contains special final types to indicate sentences, we have two type formation operators $(-/-)$ and $(-\backslash -)$, which along with their elimination rules establish a requirement that grammatical categories require other grammatical categories to their left or right. This is the essence of Lambek's calculi NL and L. CCGs keep the same minimal type-formations, but include extra sequent rules such as type-raising and cross-composition.



THAT'S GRAMMAR!

$$\begin{array}{c}
 \text{Bob: n} \quad \text{drinks: } (n/s) \backslash n \quad \text{beer: n} \\
 \cdot \qquad \qquad \qquad \hline
 \vdots \qquad \qquad \qquad \text{drinks_beer: } n/s \qquad \qquad \qquad (<) \\
 \hline
 \text{Bob_drinks_beer: s} \qquad \qquad \qquad (>)
 \end{array}$$

Figure 3.2: We can notice an asymmetry in the above formulation when we examine the transitive verb type $(n/s) \setminus n$ again; it asks first for a noun to the right, and then a noun to the left. We could just as well have asked for the nouns in the other order with the typing $(n/s) \setminus n$ and obtained all of the same proofs.

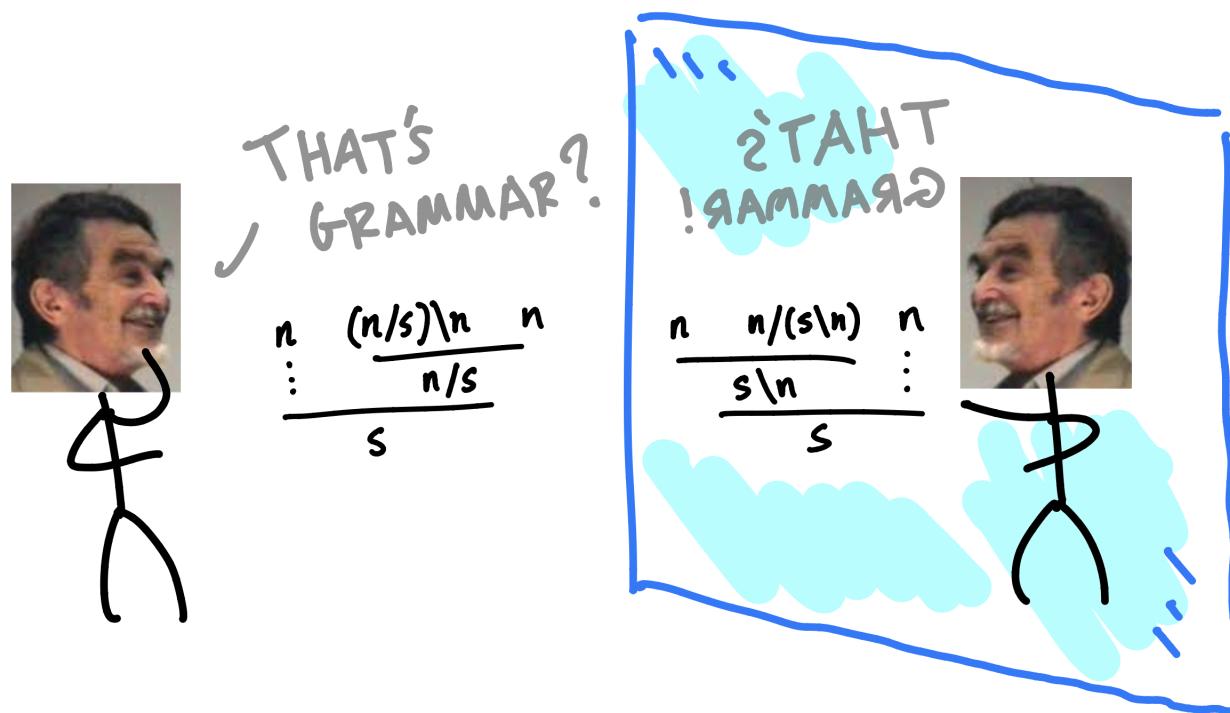
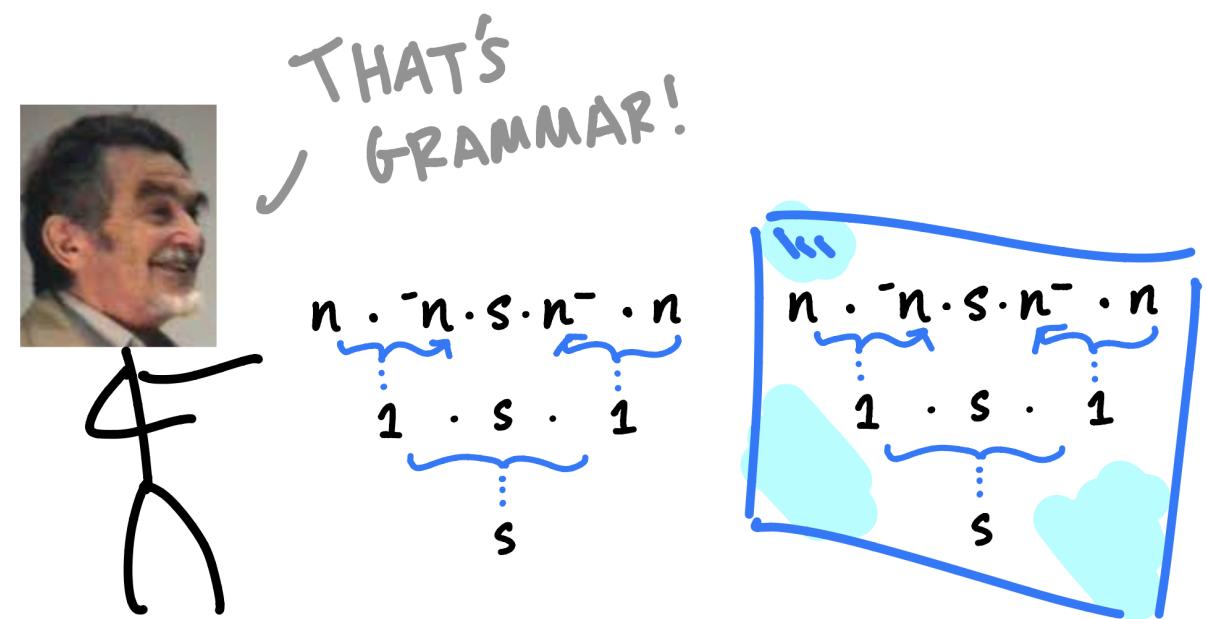


Figure 3.3: To eliminate this asymmetry, Lambek devised pregroup grammars. Whereas a group is a monoid with inverses up to left- and right-multiplication, a pregroup weakens the requirement for inverses so that all elements have distinct left- and right- inverses, denoted x^{-1} and ${}^{-1}x$ respectively. Eliminating or introducing inverses is a non-identity relation on elements of the pregroup, so we have axioms of the form e.g. $x \cdot {}^{-1}x \rightarrow 1 \rightarrow {}^1x \cdot x$. In this formulation, denoting the multiplication with a dot, both $(n/s) \setminus n$ and $(n/s) \setminus n$ become ${}^{-1}n \cdot s \cdot n^{-1}$, which just wants a noun to the left and a noun to the right in whatever order to eliminate the flanking inverses to reveal the embedded sentence type. Now we can obtain the same proof of correctness as a series of algebraic reductions.

$$\begin{array}{ll} n \cdot ({}^{-1}n \cdot s \cdot n^{-1}) \cdot n & (3.1) \\ \rightarrow & (n \cdot {}^{-1}n) \cdot s \cdot (n^{-1} \cdot n) & (3.2) \\ \rightarrow & 1 \cdot s \cdot 1 & (3.3) \\ \rightarrow & s & (3.4) \end{array}$$



3.1.2 Coecke's Composition

Figure 3.4: Meanwhile, an underground grunge vagabond moonlighting as a quantum physicist moonlighting as a computer scientist was causing a shortage of cigars and whiskey in a small English town. He noticed a funny thing about the composition of multiple non-destructive measurements of a quantum system, which was that information could be carried, or flow, between them. So he wrote a paper [CITE](#), which contained informal diagrams that looked like this.

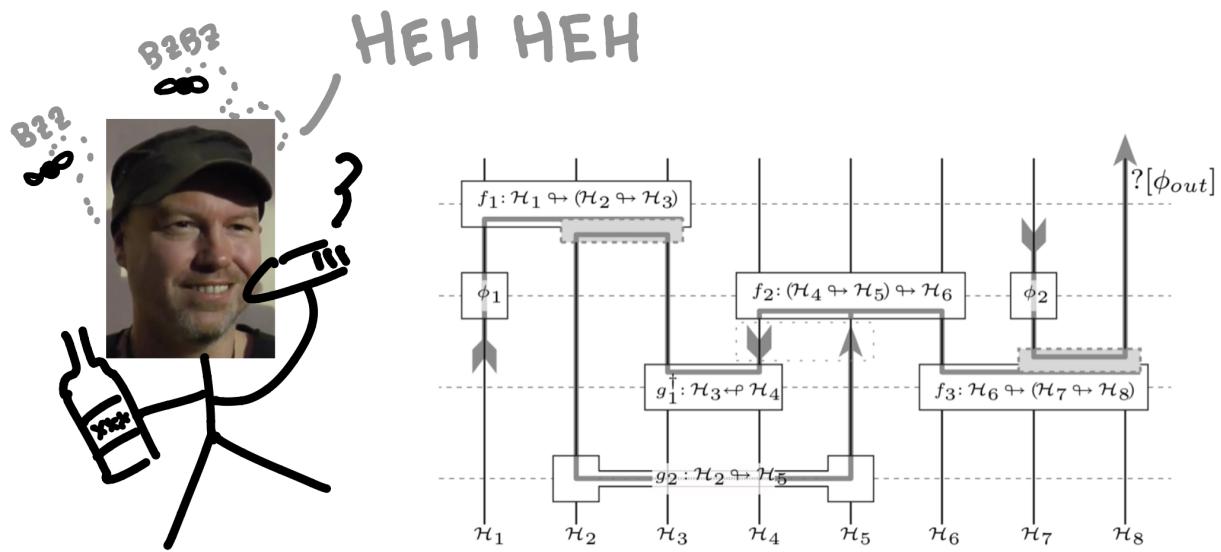
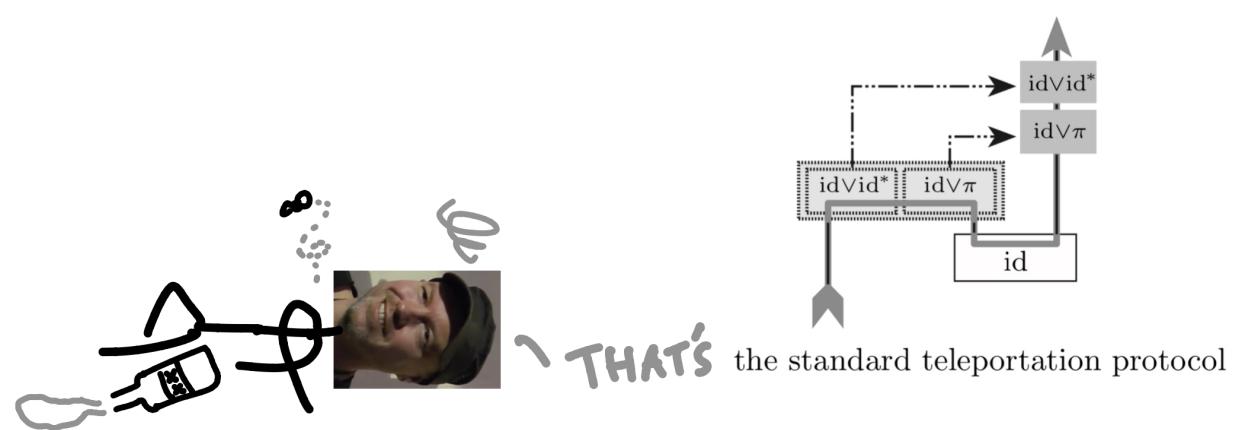


Figure 3.5: There were two impressive things about these diagrams. First, the effects such as transparencies for text boxes and curved serifs for angled arrows give a modern feel, but they were done manually in macdraw, the diagrammatic equivalent of sticks and stones. Second, though the diagrams were informal, they provided a way to visualise and reason about entanglement that was impossible by staring at the equivalent matrix formulation of the same composite operator. The most important diagram for our story was this one, which captures the information flow of quantum teleportation.



3.1.3 Categorical quantum mechanics

Figure 3.6: Category theorists and physicists such as Abramsky and Baez were excited about these diagrams, which looked like string diagrams waiting to be made formal. The graphical cups and caps in the important diagram were determined to correspond to a special form of symmetric monoidal closed category called strong compact closed.

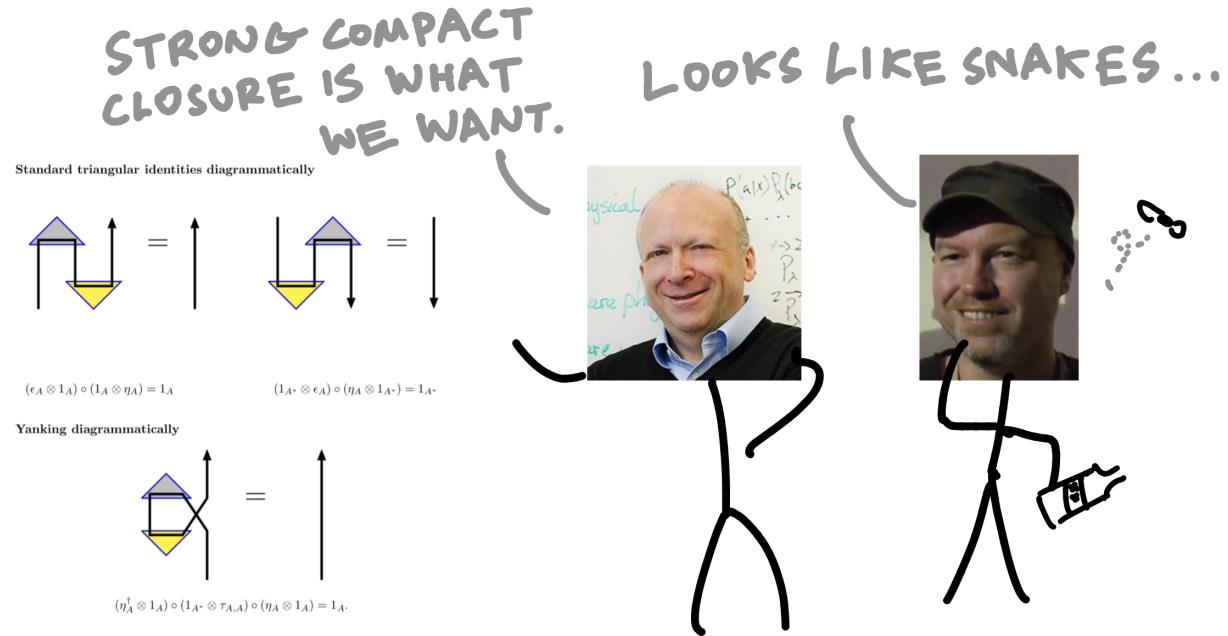


Figure 3.7: Diagrammatically, reasoning in a strongly compact closed category amounts to ignoring the usual requirement of processiveness and forgetting the distinction between inputs and outputs, so that "future" outputs could curl back and be "past" inputs. This formulation also gave insight into the structure of quantum mechanics. For example, the process-state duality of strong compact closure manifested as the Choi-Jamiołkowski isomorphism.



Figure 3.8: However, dealing with superpositions necessitated using summation operators within diagrams, which is cumbersome to write especially when dealing with even theoretically simple Bell states. An elegant diagrammatic simplification arose with the observation that special- \dagger -frobenius algebras, or spiders, correspond to choices of orthonormal bases CITE in **FdHilb**, the ambient setting of finite-dimensional hilbert spaces.

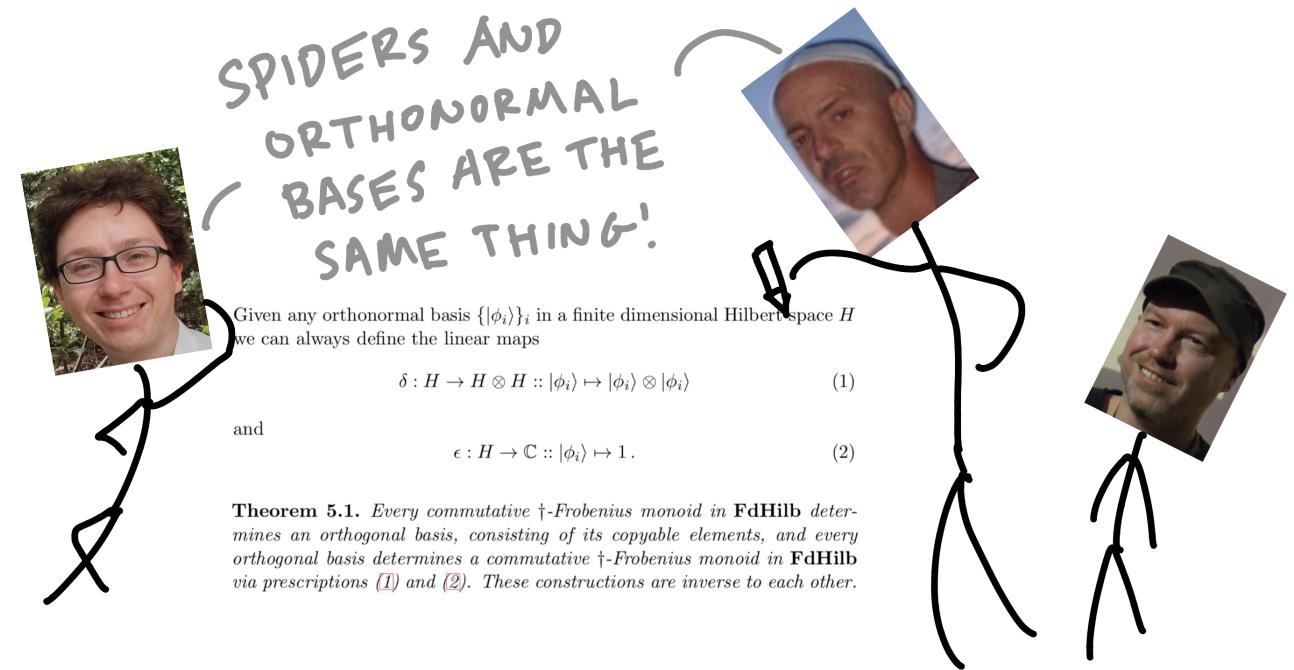


Figure 3.9: Not only did this remove the need for summation operators, it also revealed that strong compact closure was a derived, rather than fundamental structure, since spiders induce compact closed structure.

classical quantum structuralism

Figure 3.10: And so the stage was set for a purely diagrammatic treatment of ZX quantum mechanics. The story of ZX diverges away from our interest, so I will summarise what happened afterwards. In no particular order, the development of ZX went on to accommodate a third axis of measurement to yield a ZXW calculus CITE, the systems were proven to be complete CITES, there are at the time of writing two expository books CITES, and ZX-variants are becoming an industry standard for quantum circuit specification and rewriting CITE.

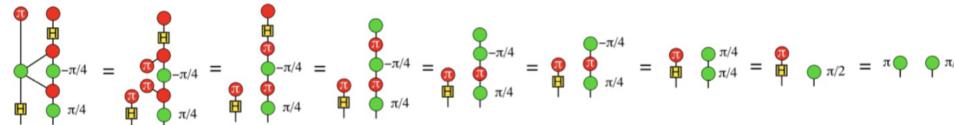


FIG. 1: Graphical simulation of quantum Fourier transform on the input $|10\rangle$.

THAT'S QUANTUM
MECHANICS!



3.1.4 Enter computational linguistics

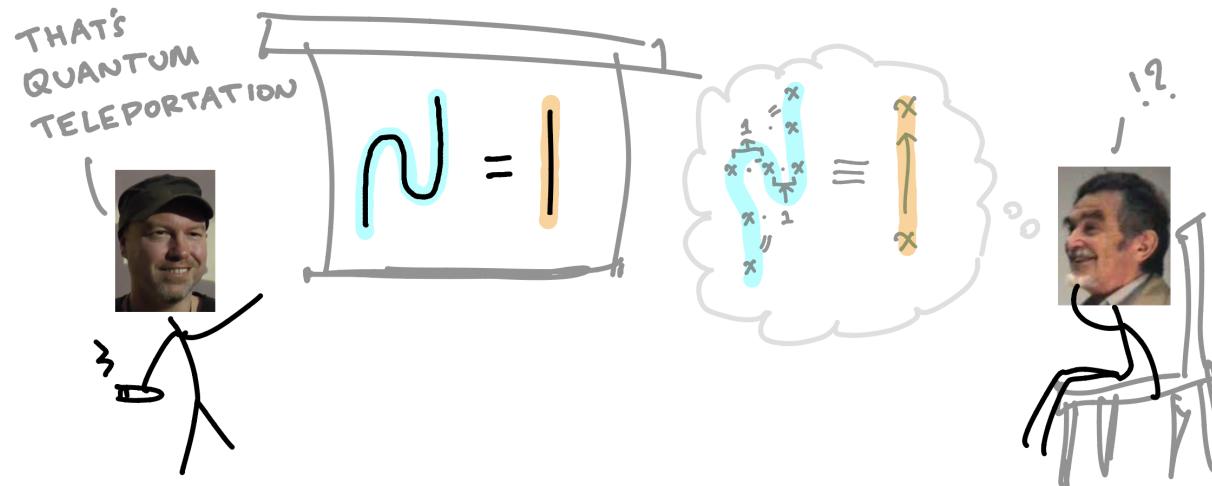


Figure 3.11: Somewhere in Canada at the turn of the millennium, Bob met Jim, who saw something familiar about the diagram for quantum teleportation. The snake equation for compact closure looked a lot like the categorified version of introducing and eliminating pregroup types.

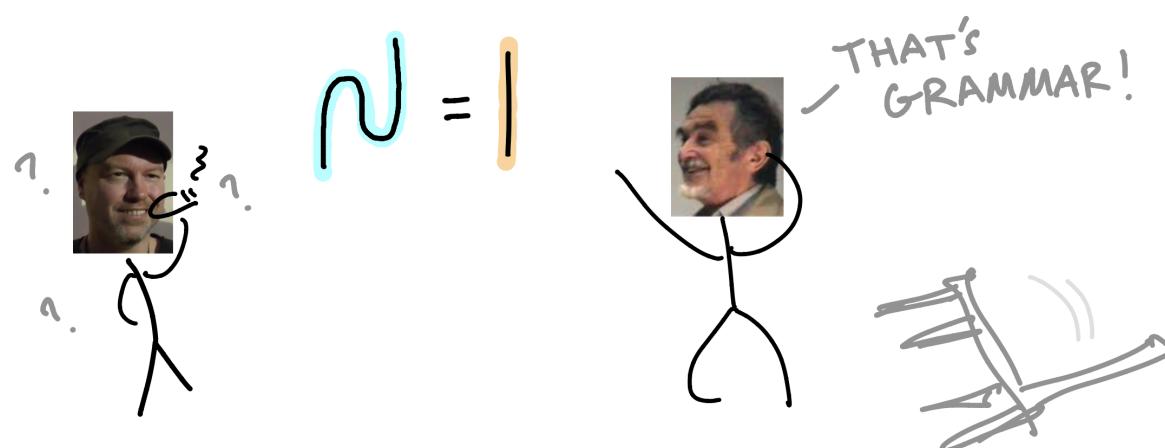


Figure 3.12: Bob and Jim's meeting put the adjectives *compositional* and *categorical* on the same table, but the cake wasn't ready. Two more actors Steve and Mehrnoosh were required to introduce *distributional*, which refers to Firth's maxim CITE "you shall know a word by the company it keeps". In its modern incarnation, this refers generally to vector-based semantics for words, where it is desirable but not necessarily so (as in the case of generic latent space embeddings by an autoencoder) that proximity of vectors models semantic closeness.



where the reversed triangles are now the corresponding Dirac-bra's, or in vector space terms, the corresponding functionals in the dual space. This simplifies the expression that we need to compute to:

$$(\langle \vec{v} | \otimes 1_S \otimes \langle \vec{v} |) |\vec{\Psi}\rangle$$

Figure 3.13: Steve Clark was a professor in the computer science department at Oxford, and he was wondering how to compose vector-based semantic representations. Steve asked Bob, who realised suddenly what Jim was talking about. Mediated by the linguistic expertise of Mehrnoosh who was a postdoctoral researcher in Oxford at the time, pregroup diagrams were born. The basic types n and s are assigned finite-dimensional vector spaces, concatenation of types the hadamard product \otimes , and by the isomorphism of dual spaces in finite dimensions there is no need to keep track of the left- and right- inverse data. Words become vectors, and pregroup reductions become bell-states, or bell-measurements, depending on whether one reads top-down or bottom-up. There was simply no other game in town for an approach to computational linguistics that combined linguistic compositionality with distributional representations.

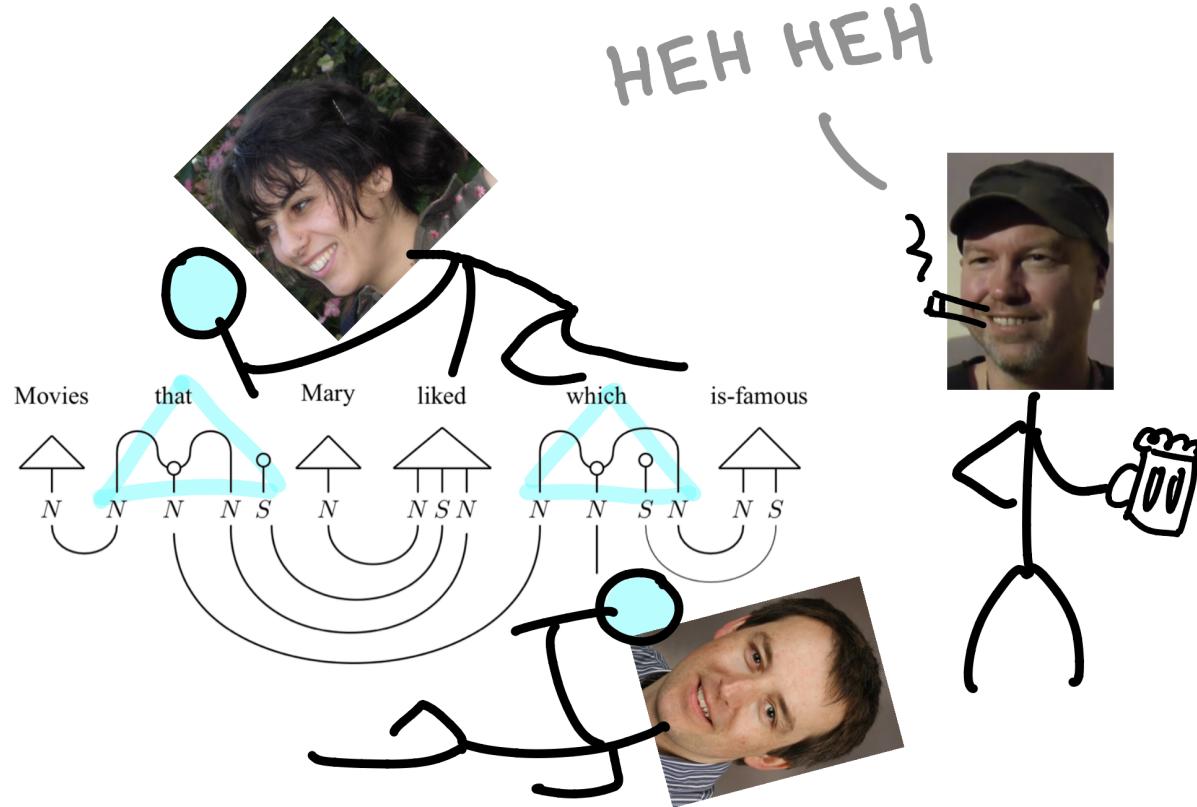


Figure 3.14: In *the frobenius anatomy of relative pronouns*[CITE](#), the trio realised that spiders could play the role of relative pronouns, which was genuinely novel linguistics. If one follows the noun-wire of "movies", one sees that by declaring the relative pronoun to be a vector made up of a particular bunch of spiders-as-multiwires, "movies" is copied to be related to the "liked" word, copied again by "which" to be related to the "is-famous" word, and a third time to act as the noun in the whole noun-phrase. This discovery clarified a value proposition: insights from quantum theory could be applied in the linguistic setting, and linguistics offered a novel use-case for quantum computers. For example, density matrices were used to model semantic ambiguity [CITE](#), and natural language experiments were performed on real quantum computers [CITE](#).

Definition 4. We define the category **ConvexRel** as having convex algebras as objects and convex relations as morphisms, with composition and identities as for ordinary binary relations.

Given a pair of convex algebras (A, α) and (B, β) we can form a new convex algebra on the cartesian product $A \times B$, denoted $(A, \alpha) \otimes (B, \beta)$, with mixing operation:

$$\sum_i p_i |(a_i, b_i)\rangle \mapsto \left(\sum_i p_i a_i, \sum_i p_i b_i \right)$$

This induces a symmetric monoidal structure on **ConvexRel**. In fact, the category **ConvexRel** has the necessary categorical structure for categorical compositional semantics:

Theorem 1. The category **ConvexRel** is a compact closed category. The symmetric monoidal structure is given by the unit and monoidal product outlined above. The caps for an object (A, α) are given by:

$$\circlearrowleft : I \rightarrow (A, \alpha) \otimes (A, \alpha) :: \{(*, (a, a)) \mid a \in A\}$$

the cups by:

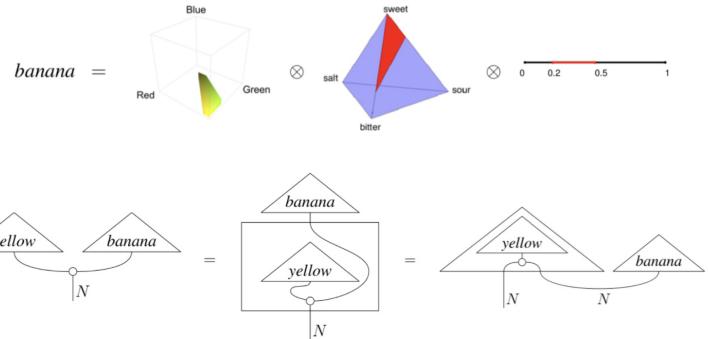
$$\circlearrowright : (A, \alpha) \otimes (A, \alpha) \rightarrow I :: \{((a, a), *) \mid a \in A\}$$

and more generally, the multi-wires by:

$$\dots : A \otimes \dots \otimes A \rightarrow A \otimes \dots \otimes A :: \{((a, \dots, a), (a, \dots, a)) \mid a \in A\}$$



BUT WHERE CAN WE FIND A SENTENCE-SPACE
BIG ENOUGH FOR SPATIAL RELATIONS ON MANY THINGS?



$$\begin{aligned} \text{yellow banana} &= (1_N \otimes \epsilon_N)(\text{yellow}_{adj} \otimes \text{banana}) \\ &= (1_N \otimes \epsilon_N)\{(\vec{x}, \vec{x}') | x_{colour} \in \text{yellow}\} \\ &\quad \otimes \{(R, G, B) | (0.9R \leq G \leq 1.5R), (R \geq 0.3), (B \leq 0.1)\} \\ &\quad \otimes \text{Conv}(\{t_{sweet}, 0.25t_{sweet} + 0.75t_{bitter}, 0.7t_{sweet} + 0.3t_{sour}\}) \otimes [0.2, 0.5]\} \\ &= \{(R, G, B) | (0.9R \leq G \leq 1.5R), (R \geq 0.7), (G \geq 0.7), (B \leq 0.1)\} \\ &\quad \otimes \text{Conv}(\{t_{sweet}, 0.25t_{sweet} + 0.75t_{bitter}, 0.7t_{sweet} + 0.3t_{sour}\}) \otimes [0.2, 0.5]\} \end{aligned}$$

Figure 3.15: Keeping the structure of the diagrams but seeking set-relational rather than vector-based semantics, a bridge was made between linguistics and cognitive science in *interacting conceptual spaces* CITE . Briefly, Gärdenfors posits that spatial representations of concepts mediate raw sense data and symbolic representations – e.g. red is a region in colourspace – and moreover that concepts ought to be spatially convex – e.g. mixing any two shades of red still gives red. This paper created a new point in the value proposition: that new mathematics would arise from investigating the linguistic-quantum bridge, e.g. generalised relations CITE . Although labelled as if it is the first in a series, the paper never saw a sequel by the same title, blocked by an apparently simple but actually tricky theoretical problem. The problem is that while this convex-relational story worked for conceptual adjectives modifying a single noun such for "sweet yellow bananas", there was difficulty in extending the story to work for multiple objects interacting in the same space, as in "cup on table in room". It couldn't be worked out what structure a sentence-wire in **ConvexRel** ought to have in order to accommodate (in principle) arbitrarily many objects and spatial relations between them.

DisCoCat then diverges from the story I want to tell. In no particular order, QNLP was done on an actual quantum computer CITE , some software packages were written CITE , and some art was made CITE .

3.1.5 I killed DisCoCat, and I would do it again.

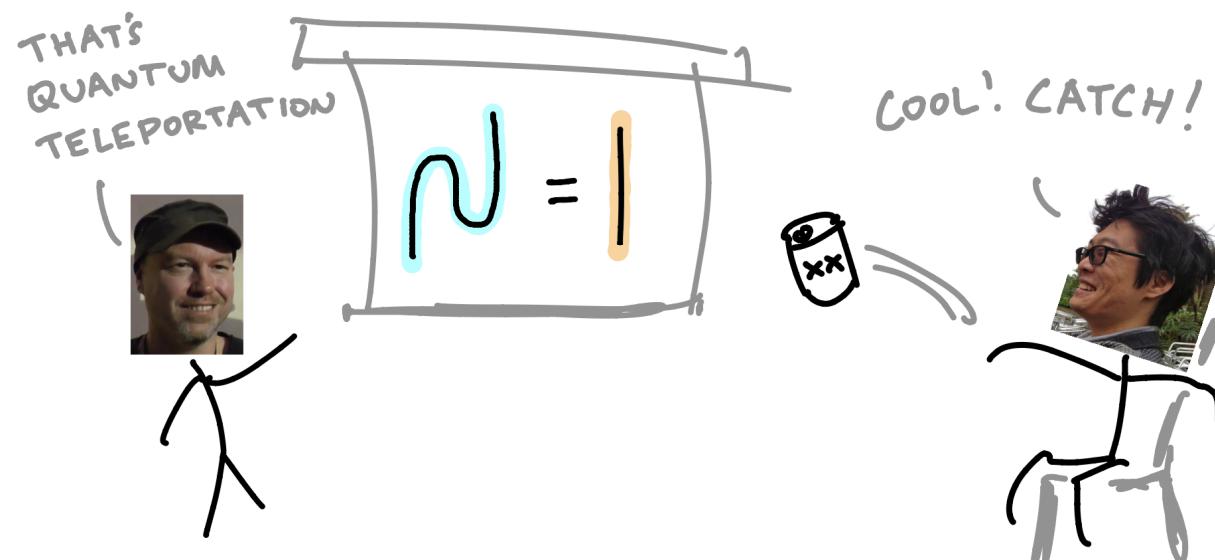
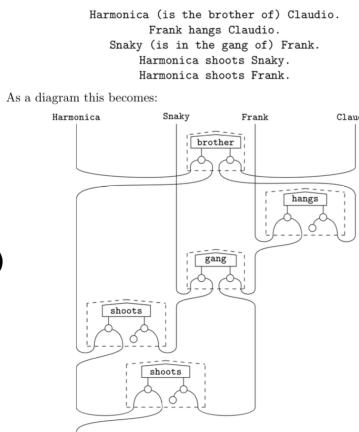


Figure 3.16: It is a common evolutionary step in linguistics that theories ‘break the sentential barrier’, moving from sentence-restricted to text- or discourse-level analysis [CITE](#). The same thing happened with DisCoCirc, due to a combination of practical constraints and theoretical ambition. On the practical side, wide tensors were (and remain) prohibitively expensive to simulate classically and actual quantum computers did not (and still do not) have many qubits, hence in practice pregroup diagrams were reduced to thinner and deeper circuits, often with the help of an additional simplifying assumption that sentence wires were pairs of noun wires in the illustrated form on the left. Theoretically, seeking dynamic epistemic logic, Bob had an epiphanous hangover (really) where he envisioned that these “Cartesian verbs” could be used in service of compositional text meanings, and he called this idea DisCoCirc [CITE](#).

Figure 3.17: I met Bob in my master’s in 2019, where he taught the picturing quantum processes course. When quantum teleportation was explained in half a minute by a diagram, I decided to pursue a DPhil in diagrammatic mathematics. In the last lecture, I threw Bob a cider, after which he seemed to like me. I did not know he was an alcoholic.

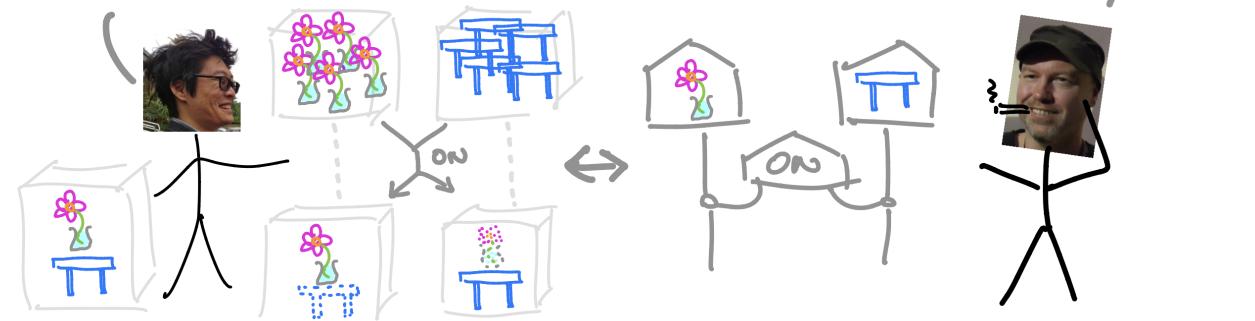
I was shanghaied into thinking about diagrams for language. I was deeply dissatisfied with the content from the standpoint my own intellectual integrity. Firstly, there seemed to me an unspoken claim that the presence of cups in pregroup diagrams (which implied a noncartesian and hence large tensor product) made it necessary to use quantum computers to effectively compute pregroup diagrams. I just could not believe that my brain required quantum computation to understand language. This implicit claim of kinship between quantum and linguistics was further entrenched by the analysis of the relative pronoun in terms of frobenius algebras, since spiders in \mathbf{Vect}^{\otimes} were the *sine qua non* of categorical quantum mechanics. The best steelman for spiders I have is that frobenius algebras (which are central to bicategories of relations CITE) just happen to be a ubiquitous mathematical structure that are well-suited to express the mathematics of connections, both in language and in quantum.

Second, representing the content of a sentence as a vector in a sentence-vector-space did not sit well with me, since this move meant that the only meaningful thing one could do with two sentences was take their inner-product as a measure of similarity. Moreover, I had the theoretical concern that language is in principle indefinitely productive, so one could construct a sentence that marshalled indefinitely many nouns, and at some point for any finite vector space s one would run out of room to encode relationships, or else they would be cramped together in a way that did not suit intuitions about the freedom of constructing meanings using language. I always believed in the existence of a simple, practical, and intuitive categorical, compositional, and distributional semantics; I just didn't believe that the role of quantum – however helpful or interesting – was *necessary*.

My first unsatisfactory attempt was in my Master's thesis CITE . It had been known for a while that a free autonomous category construction by Delpuech CITE could potentially eliminate some of the cups in pregroup diagrams, yielding what amounted to a method to transform a pregroup diagram into a monoidal string diagram in the shape of a context-free grammar tree. This trick had the limitation that freely adding directed cups and caps to a string diagrammatic signature did not turn a symmetric monoidal category into a (weakly) compact closed one, rather just into a monoidal category where the original wires had braidings, but all the new left and right dual wires did not; this presented difficulties in accounting for iterated duals for higher-order modifiers such as adverbs in grammatical types, and had nothing to say about spiders. I tried to generalise this trick to 'freely' adding arbitrary diagrammatic gadgets to string diagrams, but my assessor Samson pointed out that it was nontrivial to determine whether such constructions were faithful. In retrospect the free autonomous completion of a parameterised CITE markov category CITE is in the ballpark of dequantumfying pregroup diagrams, but I didn't learn about them until later, and that still wouldn't have addressed the issues that come with only having a sentence-wire.

Figure 3.18: Then COVID happened. During the first lockdown, I visited Bob's garden under technically legal circumstances, and I suggested a solution to the longstanding problem of representing linguistic spatial relationships. My theoretical concern was the culprit: the initial attempts at the problem failed because the approach was to find a single sentence object s in which one could paste the data of arbitrarily many distinct spatial entities. The simple solution was a change in perspective.

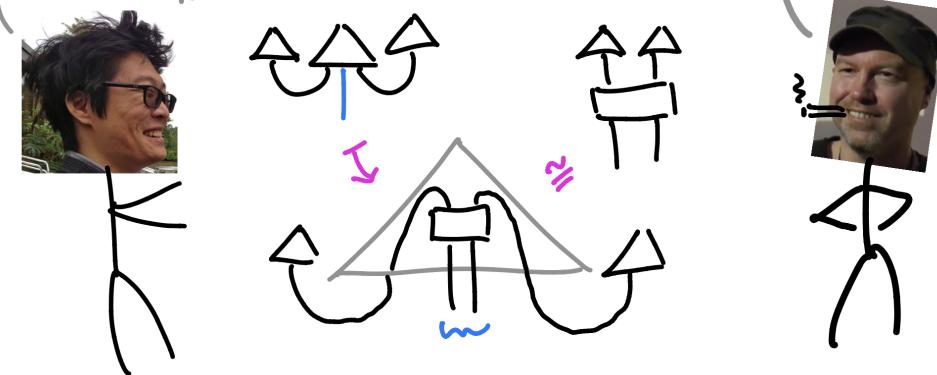
INSTEAD OF PUTTING OBJECTS IN A SHARED SPACE,
GIVE EACH OBJECT THEIR OWN COPY OF SPACE.
SPATIAL RELATIONS BECOME POSSIBILISTIC RESTRICTIONS!



THAT'S TEXT CIRCUITS!

Figure 3.19: That this move of splitting up the sentence-wire into a sentence-dependent collection of wires was sufficient to solve what had appeared to be a difficult problem prompted some re-examination of foundations. The free autonomisation trick in conjunction with sentence-wire-as-tensored-nouns seemed promising, but it became clear that right way to drown a DisCoCat thoroughly was to explain and eliminate the spiders.

THE IDEA OF INTERACTING PRIVATE SPACES GENERALISES.
IF WE PICK THE RIGHT INTERNAL WIRING,
(WE CAN GET RID OF CUPS;
NO NEED FOR QUANTUM!



WHERE ARE THE SPIDERS?

Figure 3.20: I then discovered that by interpreting spiders as the well-known "pair of pants" algebra in a compact closed monoidal setting allowed for a procedure in which the final form was purely symmetric monoidal – the absence of cups and caps meant that there was no practical necessity to interpret diagrams on quantum computers: any computer would suffice. The role of spiders for relative pronouns was illuminated in the presence of splitting the sentence wire: the pair-of-pants are the algebra of morphism composition, and splitting the sentence wire into a collection of nouns allowed relative-pronoun-spiders to pick out the participating nouns to compose relationships onto.

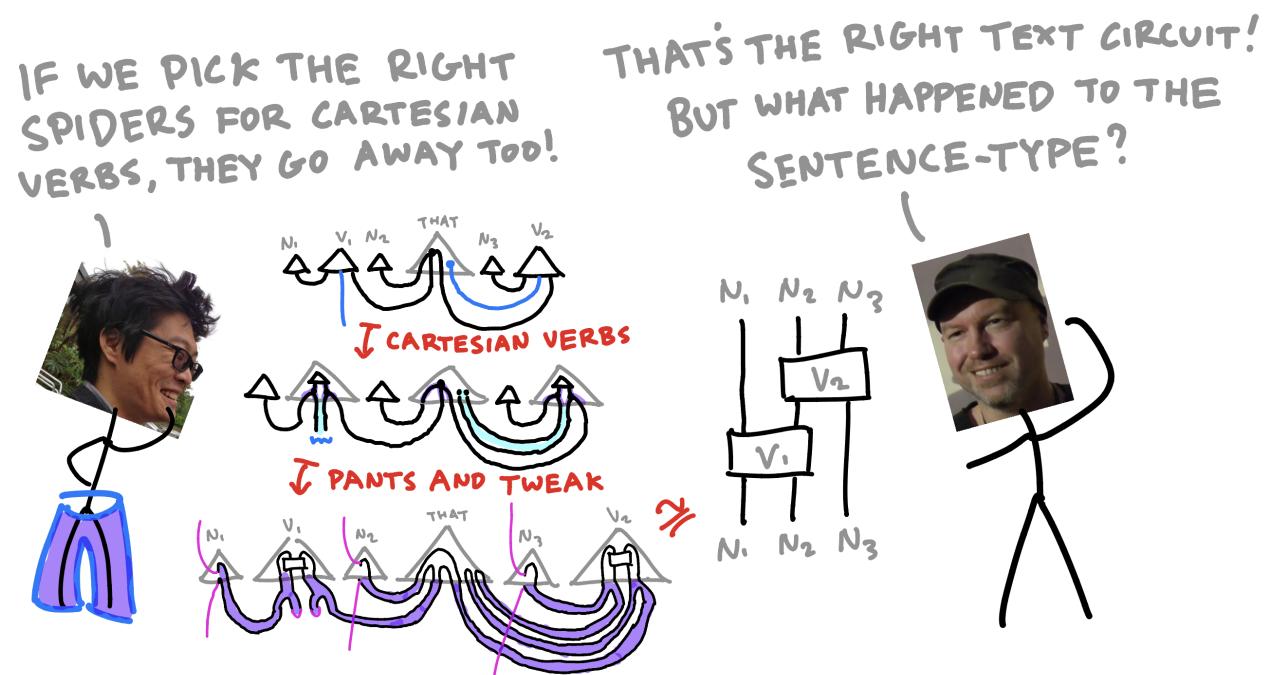
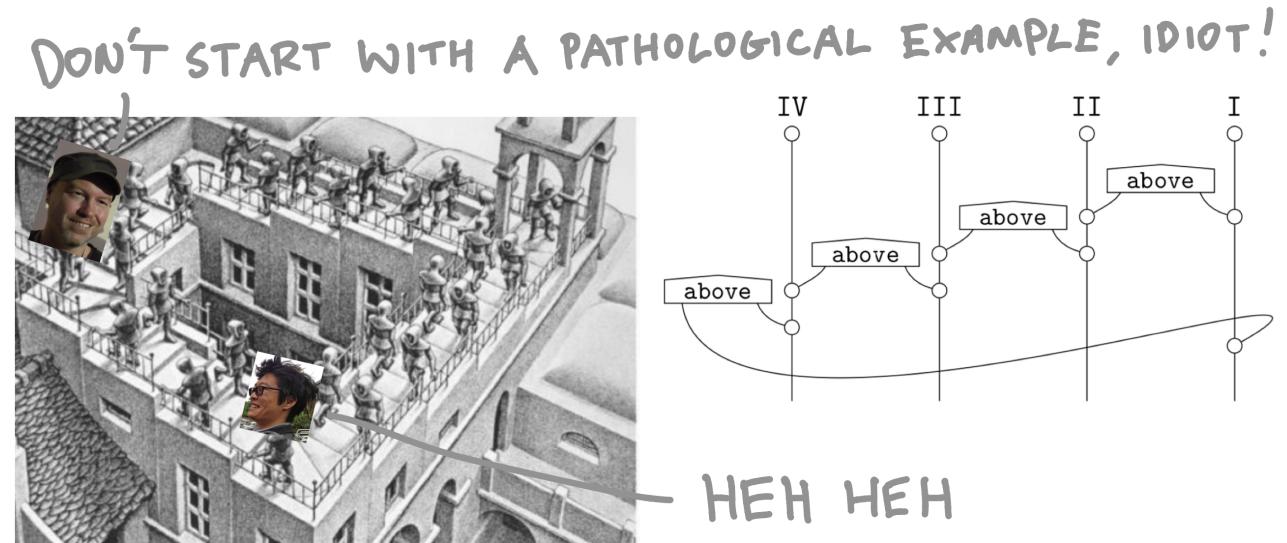


Figure 3.21: A coherent conservative generalisation of DisCoCat with less baggage had emerged, or rather, DisCoCirc was placed to formally subsume DisCoCat. It was now understood that the sentence type was a formal syntactic ansatz for the sake of grammar, which was to be interpreted in the semantic domain not as a single wire, but as a sentence-dependent collection of wires. It was further realised that the complexity of pregroup diagrams was due to grammar – the topological deformation of semantic connections to fit the one-dimensional line of language – whereas the essential connective content of language could be expressed in a simple form that distilled away the bureaucracy of syntax.



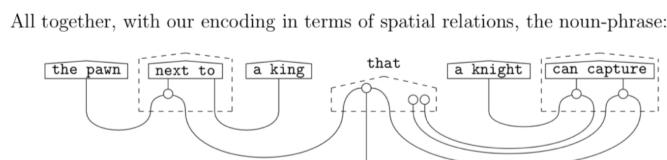
Figure 3.22: We wrote up the story about spaces in **CITE**, the spiritual successor to *interacting conceptual spaces I*. We could formally calculate the meanings of sentences that used linguistic spatial relations, all using a simple and tactile diagrammatic calculus.



BETTER?
(



NO! SHOW A CIRCUIT!



All together, with our encoding in terms of spatial relations, the noun-phrase:

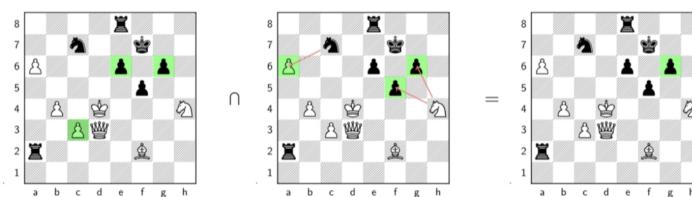
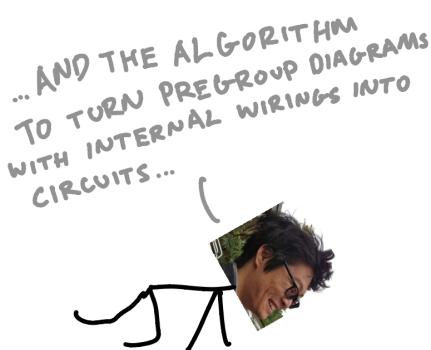


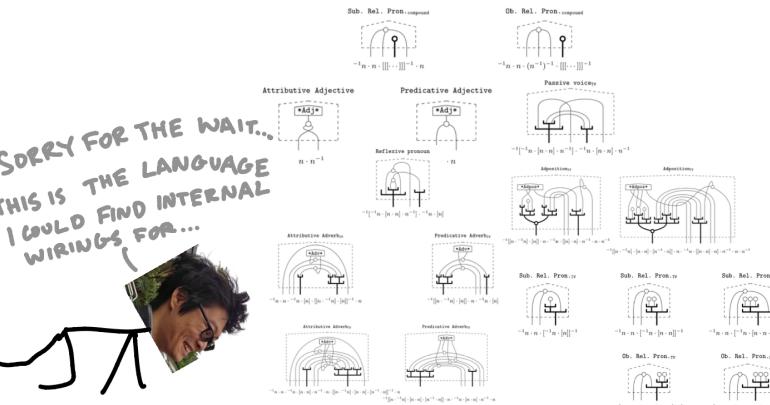
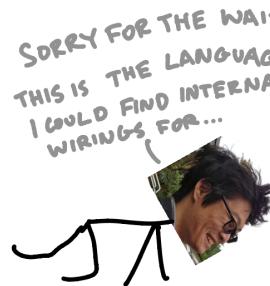
Figure 3.23: The paper on spatial relations actually came very late, because I was busy with Bob's ludicrous request to go turn "all of language" into circuits. I bitched and moaned about how I wasn't a linguist and how it was an impossible task, but I was in too deep to back out.



Figure 3.24: I suppose the nice thing about aiming for the moon is that even failure might mean you leave orbit. So I settled for what I thought was a sensible fragment of English, for which I devised internal wirings and an algorithm that transformed pregroup diagrams with the internal wirings into circuit form. Many tiring diagrams later, I presented my results in the first draft of "distilling text into circuits".



```
Algorithm 1: Grammar to Circuit Algorithm
Input: A sentence S
Result: A text circuit
Data: Parse, PrnRes, NOUNS, PRONOUNS
BASICWIRES  $\leftarrow \{w \in S \cap \text{NOUNS}\}$ ; // 'BASICWIRES' is nouns of S
DIAGRAM  $\leftarrow \text{Parse}(S)$ ; // 'DIAGRAM' is pregroup diagram with wrapping
if 'DIAGRAM' has single output wire o then
| Append appropriate deletion to o in DIAGRAM; // delete output wire
| for w  $\in$  BASICWIRES do
| | Append w's wire from DIAGRAM, creating open wire; // copy nouns
| | Attach rightmost of a copy-spider to the open wire;
| | Pull free input of copy-spider to the top of DIAGRAM;
| | Pull free output of copy-spider to the bottom of DIAGRAM;
| | Label this new wire with w;
end
while There remain wrapped wires in DIAGRAM do
| | Apply unwrapping from Definition 3; // unwrap
end
for  $w_1, w_2 \in$  BASICWIRES do
| if  $w_1 \in \text{PRONOUNS}$  and  $w_2 \notin \text{PRONOUNS}$  and  $\text{PrnRes}(w_1, w_2)$  then
| | | Attach outputs of a copy-spider to  $w_1, w_2$  inputs;
| | | Attach inputs of upside-down copy-spider to  $w_1, w_2$  outputs;
| | | relabel merged wire as  $w_2$ ; // resolve pronouns
| | end
| end
while There remain deletions, cups, or caps of non-wrapped wires
| | Use spider-fusion to remove deletions, cups, and caps;
| end
| Simplify to obtain text circuit; // clean up
| return DIAGRAM; // return text circuit
else
| | return False; // (Fail if input ungrammatical)
end
```



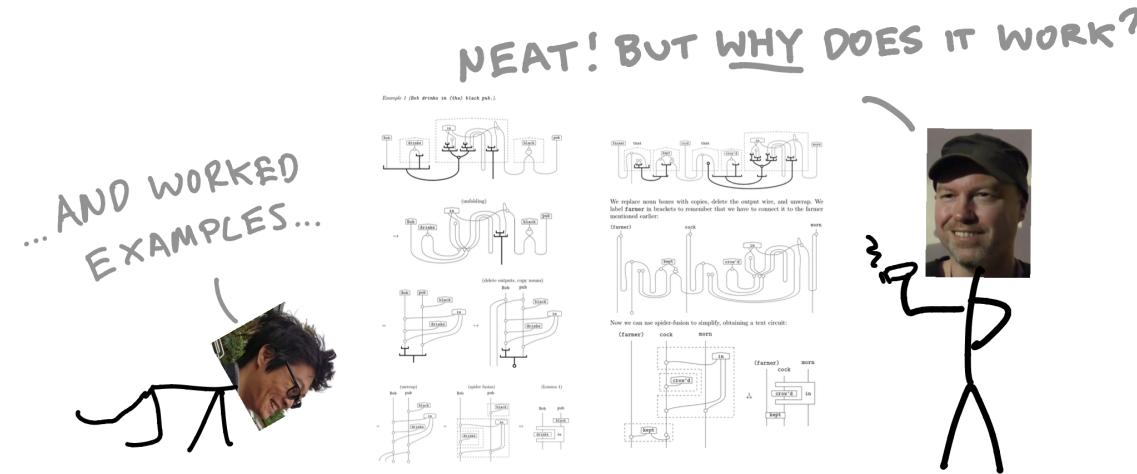


Figure 3.25: Bob had a good point. Everything worked, but we had no understanding as to why, and accordingly, whether or not it would all break. At this point in time, Jonathon Liu, who was a masters' student I taught during COVID, had committed the error of thinking diagrams were cool, and was now hanging out with me and Bob. After understanding the procedure, Jono independently devised the same arcane internal wirings as I had, but neither of us could explain how we did it. So we had evidence of an underlying governing structure that was coherent but inarticulable.



Figure 3.26: I realised that our intuitions were coming from an implicit productive grammar, rather than a parsing one, and that the path of least resistance for obtaining formal guarantees for the language-to-circuit procedure was to just handcraft a generative grammar for the fragment of language we were interested in. This meant scrapping everything in the first draft and starting again from scratch. Bob always had a word of gentle encouragement, giving me the motivation to persevere.

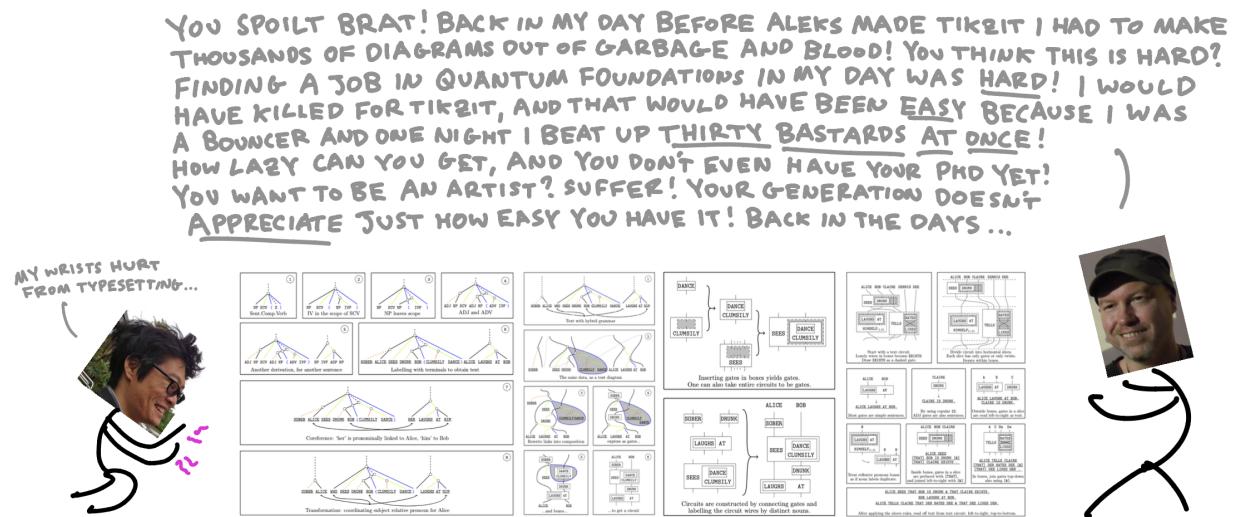
So now we had two ways to obtain text circuits. One from pregroups (which Jono had extended the technique for to CCGs in his master's thesis [CITE](#)), and one from handcrafted productive grammars. Then came time for me to write my thesis. Three salient questions arose.

Firstly, what is the relationship between these two ways of getting at text circuits?

Secondly, how do text circuits stand in relation to other generative grammars?

Thirdly, what is it that text circuits allow us to do?

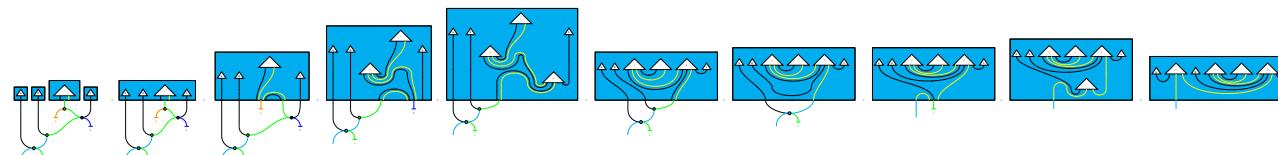
These questions are now what the rest of the thesis seeks to answer.



3.2 Internal wirings redux

3.3 How do we communicate using language?

SPEAKERS AND LISTENERS UNDERSTAND ONE ANOTHER. Obviously, natural language involves communication, which involves at minimum a speaker and a listener, or a producer and a parser. The fact that communication happens at all is an everyday miracle that any formal understanding of language must account for. The miracle remains so even if we cautiously hedge to exclude pragmatics and context and only encompass small and boring fragments of factual language like *Alice sees Bob quickly run to school*. At minimum, we should be able to model a single conversational turn, where a speaker produces a sentence from a thought and a listener parses it to recover the thought. Here is a sequence of diagram equations that demonstrates mathematically how the miracle works for two toy grammars, for the sentence *Alice sees Bob quickly run to school*. On the left we have a grammatical structure obtained from a context-free grammar, and we have equations from a discrete monoidal fibration all the way to the right, where we obtain a pregroup representation of the same sentence. Going from right to left recovers the correspondence in the other direction.



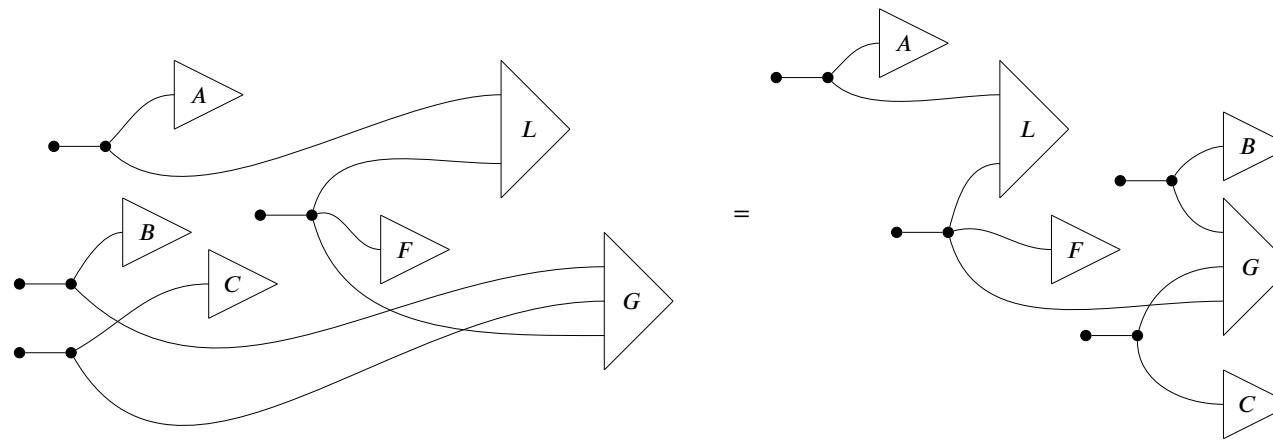
A POST-LLM THEORY OF LANGUAGE REQUIRES COHERENCE BETWEEN THEORIES OF SYNTAX – ONE FOR THE SPEAKER AND ONE FOR THE LISTENER – SUCH THAT THEIR UNDERLYING STRUCTURES ARE STRONGLY EQUIVALENT AND ARE AMENABLE IN PRINCIPLE TO DISTRIBUTIONAL OR VECTOR SEMANTICS. Anything short of this fails to provide any understanding of language beyond what can be gleaned from an LLM. If you don't have two coherent theories of syntax for speaker and listener, you cannot account for communication. If you don't have a matching theory of semantics that can be learned in principle from data and represented as vectors-and-operations-upon-them, your theory will never be of practical relevance; **N.B.** go collect an award and dance on Minsky's grave if you figure out how to learn insert-FOL-variant-here representations of dictionary words from data. The stakes are, at worst, the entirety of linguistics as an enterprise of inquiry. At best, it's all made up and nothing really matters so I will subscribe to this standard anyway just because. To the best of my knowledge there isn't any, even partial account of language that satisfies the bare minimum. Perhaps this is due to siloisation, because I don't doubt that all of the ingredients are already present in the literature: there are popular accounts of how grammar composes (albeit practically worthless) truth-theoretic semantics

[Heim and Kratzer], and there are too many equivalence proofs between different grammatical formalisms, but I can't find anything that puts the minimum all at once in the same bowl. If it does exist already, my cake is probably prettier and tastier. So in this thesis I will supply a post-LLM theory of language, and I will endeavour to do it as diagrammatically as possible. I only care to get the absolute basics right by my own standard, and I will accordingly treat everything else about language like pragmatics and gesture and whatever else as an afterthought, *as it should be*: consider air resistance only after you know how motion works in a vacuum.

3.3.1 Grammars of speakers and listeners

There is a distinction between *grammars of the speaker* – which produce sentences – and *grammars of the listener* – which deduce sentences. Viewed as mathematical processes, the two kinds of grammars go in opposite directions; speaking grammars (e.g. string-rewrite systems) start with some grammatical structure, and require informational input (e.g. which rule comes next) to produce lists of words – sentences. Conversely, listening grammars (e.g. typological grammars) start with a sentence, and require informational input (e.g. grammatical typing and which proof rule to try next) to deduce or parse a grammatical structure. Since we can understand each other, these two types of grammar must enjoy a systematic correspondence, and if one believes that semantics is compositional according to syntax, then the correspondence must further explain how both speaking and listening grammars manipulate the same underlying semantic expression, whatever that may be.

Here are some naïve observations on the nature of speaking and listening. Let's suppose that a speaker, Preube, wants to communicate a thought to Fondo. Just as a running example that does not affect the point, let's say we can gloss the thought in first order logic as $\exists a \exists b \exists c \exists f : A(a) \wedge B(b) \wedge C(c) \wedge F(f) \wedge L(a, f) \wedge G(b, c, f)$. In diagrammatic first order logic [], this is equivalently presented as the following diagrams (and any other diagram that agrees up to connectivity)

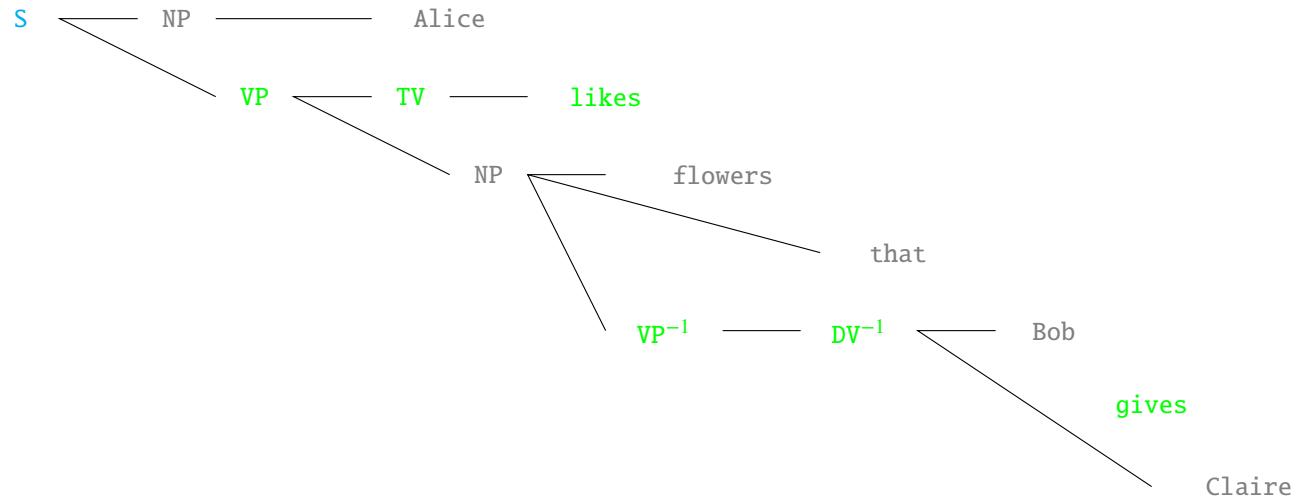


Preube and Fondo cooperate to achieve the miracle; Preube encodes his thoughts – a structure that isn't a one-dimensional string of symbols – into a one-dimensional string of symbols. And then Fondo does the reverse, turning a one-dimensional string of symbols into a thought-structure like that of Preube's. What I mean by "like that of Preube's" is that both Preube and Fondo agree on the structure of entities and relations up to the words for those entities and relations. For example, Preube could ask Fondo comprehension questions such as WHO GAVE WHAT? TO WHO?, and if Fondo can always correctly answer – e.g. BOB GAVE FLOWERS. TO CLAIRE. – then both Preube and Fondo agree on the relational structure of the communicated thought to the extent permitted by language. It may still be that Preube and Fondo have radically different internal conceptions of what FLOWERS or GIVING or BEETLES IN BOXES are, but that is alright: we only care that the *interacting structure* of the thought-relations in each person's head are the same, not their specific representations.

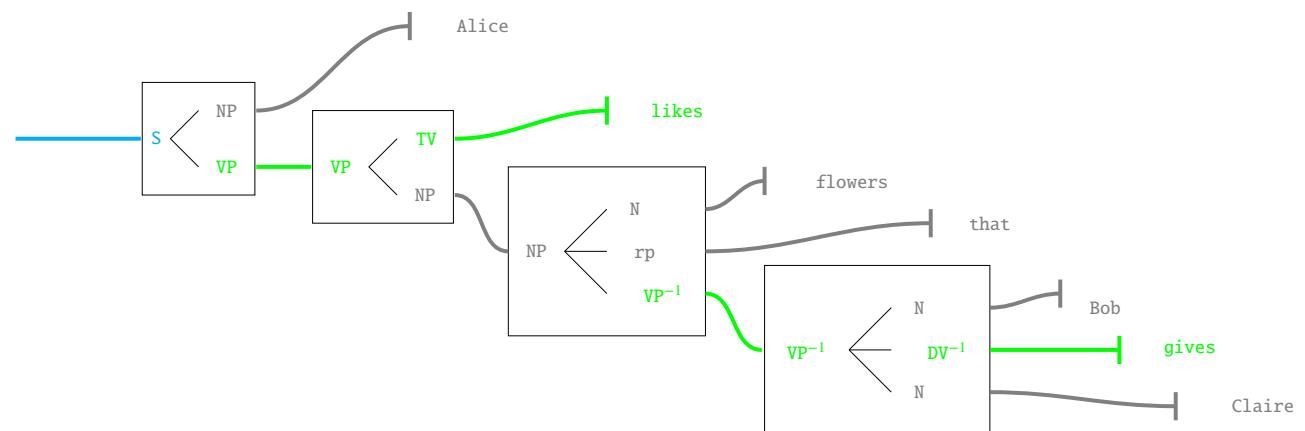
We assume Preube and Fondo speak the same language, so both know how words in their language correspond to putative building blocks of thoughts, and how the order of words in sentences and special grammatical words affect the (de-/re-)construction procedures. Now we have to explain how it is that the two can do this for infinitely many thoughts, and new thoughts never encountered before. Using string diagrams, this is surprisingly easy, because string diagrams are algebraic expressions that are invariant under certain topological manipulations that make it easy to convert between different shapes of language. Here is a sketch example that we will soon have the mathematics to express formally:

Example 3.3.1 (Alice likes flowers that Bob gives Claire.). Let's say Preube is using a context-free

grammar to produce sentences, and Fondo a pregroup grammar. The rule of the game is that Preube and Fondo can agree on a string-diagrammatic encoding strategy before having to communicate with each other. Here is one such strategy: Preube might generate the example sentence like so:

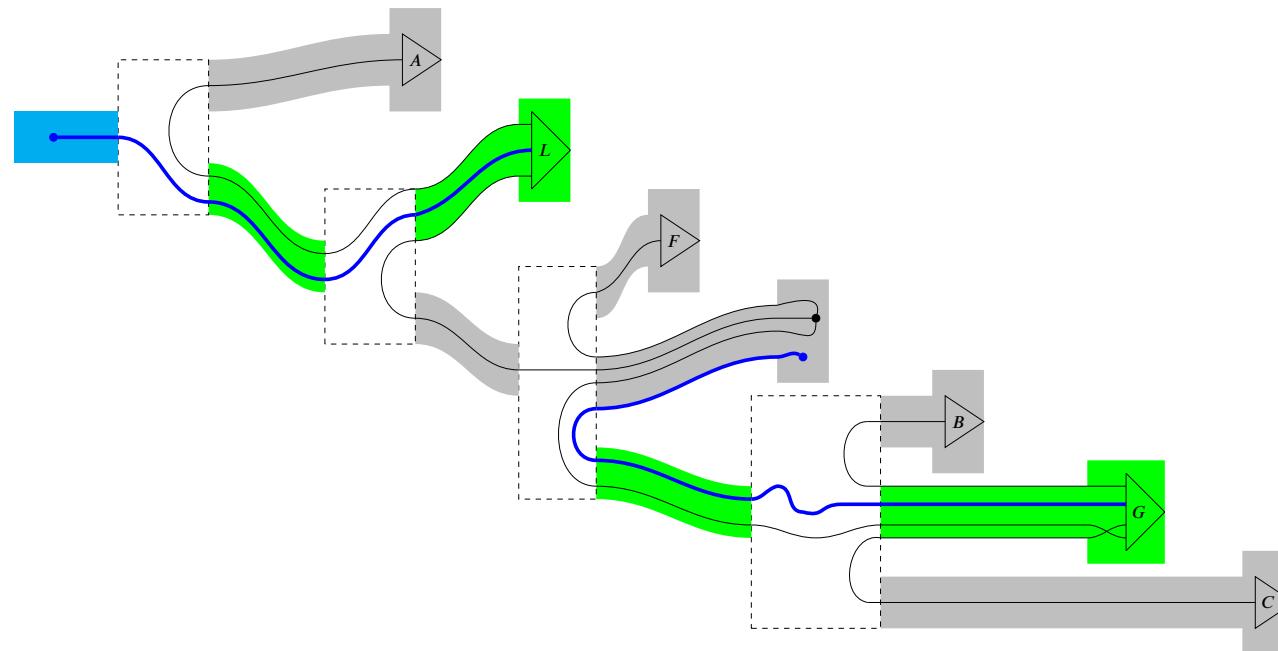


Mathematically, it makes no difference if we take the Poincaré dual of the tree, so that zero-dimensional nodes become one-dimensional wires, and branchings become zero-dimensional points linking wires – but we can just as well depict those points as boxes to label them more clearly.

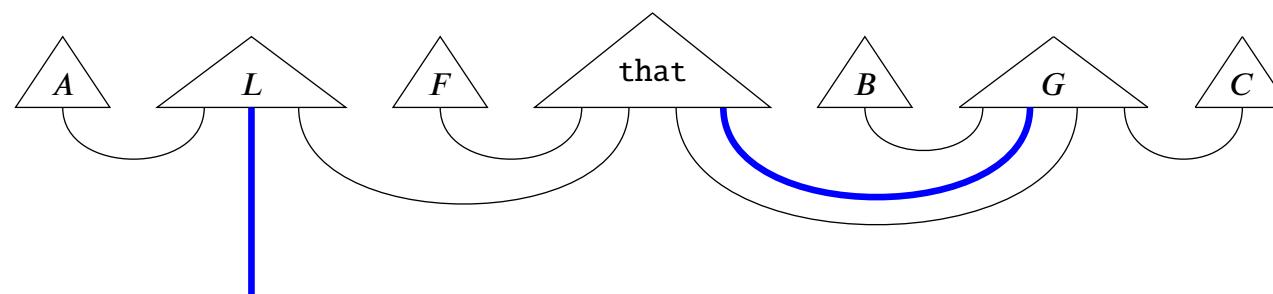


Now that Preube can express their grammatical structure string-diagrammatically, they can try to deform their first-order-logic diagram – representing what they mean to communicate – subject to the constraint that every one of their branchings (the structure of the CFG) is something recoverable by Fondo using just

pregroup reductions. To do so, Preube introduces a formal blue wire to mimic Fondo's sentence-type, and stuffs some complexity inside the labels in the form of internal wirings: a multiwire configuration for *that*, and a twist for *gives*. Those internal wirings are the content of Preube and Fondo's shared strategy.

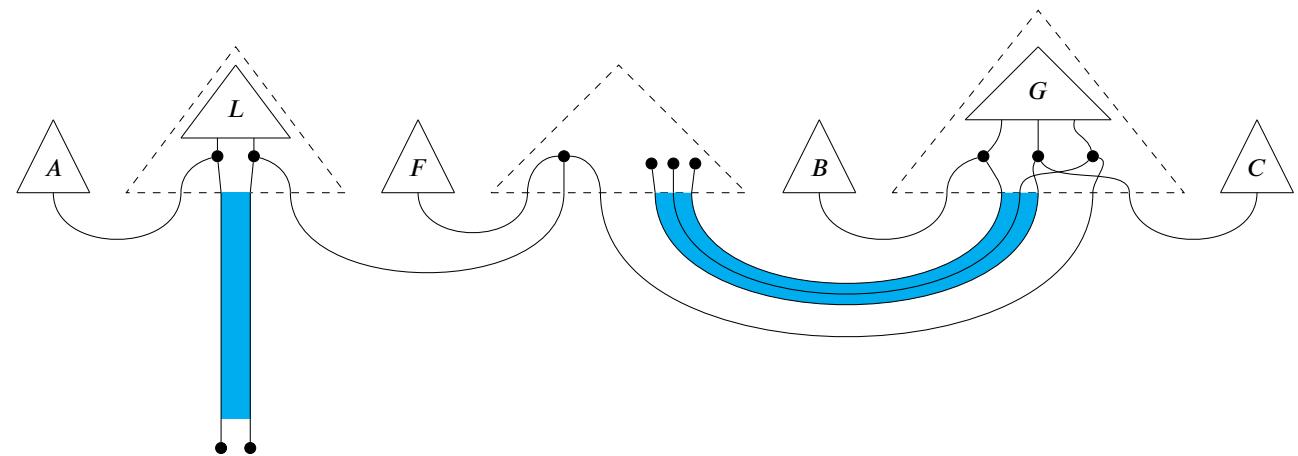


So, when Fondo receives the sentence, Fondo's pregroup derivation yields a pregroup diagram that is connectively equivalent to what Preube stuffed inside the context-free grammar structure. So now the two have strong equivalence between their grammars in the sense that every one of Preube's branches is resolved by one of Fondo's reductions.

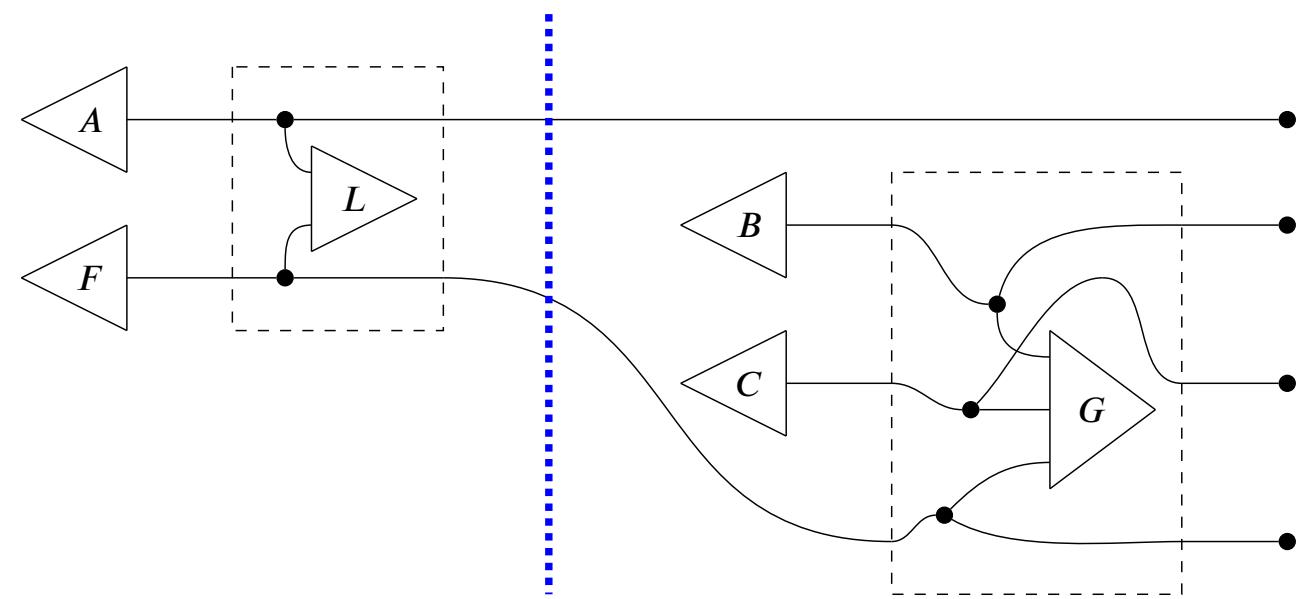


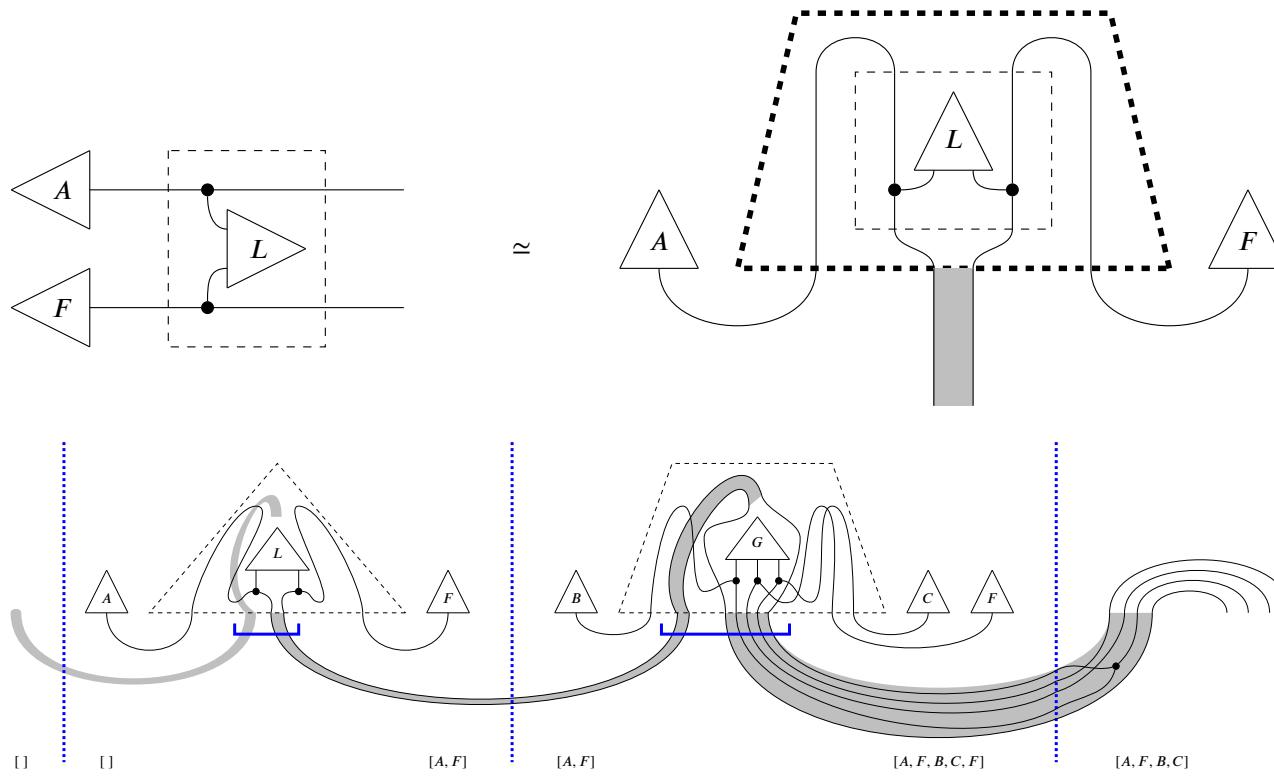
Now to fully recover Preube's intended FOL-diagram, Fondo refers to the internal wirings that form their

shared strategy, and fills those in.



Example 3.3.2 (Bob gives Claire flowers. Alice likes flowers.).





The nature of their challenge can be summarised as an asymmetry of information. The speaker knows the structure of a thought and has to supply information or computation in the form of choices to turn that thought into text. The listener knows only the text, and must supply information or computation to deduce the thought behind it. By this perspective, language is a shared and cooperative strategy to solve this (de/en)coding interaction. I will now outline the constraints imposed by this interaction, illustrating them simply using string-rewrite grammars for the speaker and pregroup grammars for the listener.

SPEAKERS CHOOSE. The speaker Preube must supply decisions about phrasing a thought in the process of speaking it. At some point at the beginning of an utterance, Preube has a thought but has not yet decided how to say it. Finding a particular phrasing requires choices to be made, because there are many ways to express even simple relational thoughts. For example, the relational content of our running example might be expressed in at least two ways (glossing over determiners):

Alice likes flowers that Bob gives Claire.

Bob gives Claire flowers. Alice likes (those) flowers.

Whether those decisions are made by committee or coinflips, they represent information that must be supplied to Preube in the process of producing language. For this reason, we consider context-free-grammars (and more generally, other string-rewrite systems) to be *grammars of the speaker*, or *productive grammars*. The start symbol S is incrementally expanded and determined by rule-applications that are selected by the speaker. The important aspect here is that the speaker has an initial state of information S that requires more information as input in order to arrive at the final sentence. Note that the concept of productive grammars are not exhausted by string-rewrite systems, merely that string-rewrite systems are a prototype that illustrate the concept well.

LISTENERS DEDUCE. The listener Fondo must supply decisions about which words are grammatically related, and how. Like right-of-way in driving, sometimes these decisions are settled by convention, for example, subject-verb-object order in English. Sometimes sophisticated decisions need to be made that override or are orthogonal to conventions, as will be illustrated in the closing discussions and limitations section of this chapter. Since Fondo has to supply information in the form of choices in the process of converting text into meaning, we consider *parsing grammars* – such as all typological grammars, including pregroups and CCGs – to be *grammars of the listener*.

THE SPEAKER'S CHOICES AND THE LISTENER'S DEDUCTIONS MUST BE RELATED. The way the speaker decomposes the thought into words in text in the speaker's grammar must allow the listener to reconstruct the thought in the listener's grammar. Even in simple cases where both parties are aiming for unambiguous communication, the listener still must make choices. This is best illustrated by introducing two toy grammars – we pick a context-free grammar for the speaker and a pregroup grammar for the listener, because they are simple, planar, and known to be weakly equivalent.

3.3.2 Discrete Monoidal Fibrations

We introduce the concept of a discrete monoidal fibration: a mathematical bookkeeping tool that relates kinds of choices speakers and listeners make when generating and parsing text respectively. We proceed diagrammatically. The first concept is that of a *monoidal functor box*.

Scholium 3.3.3. Functor boxes are from CITE . Expressing the coherence conditions of monoidal functors using equations involving functor boxes as below is not new CITE . The idea of a functor being simultaneously monoidal and a fibration is not new CITE . What is new is minor: the express requirement that the lifts of the fibration satisfy interchange, which is in general not guaranteed by just having a functor be monoidal and a (even discrete) fibration [fosco].

Figure 3.27: Suppose we have a functor between monoidal categories $\mathbf{F} : \mathcal{C} \rightarrow \mathcal{D}$. Then we have this diagrammatic representation of a morphism $\mathbf{F}A \xrightarrow{\mathbf{F}f} \mathbf{F}B$ in \mathcal{D} .

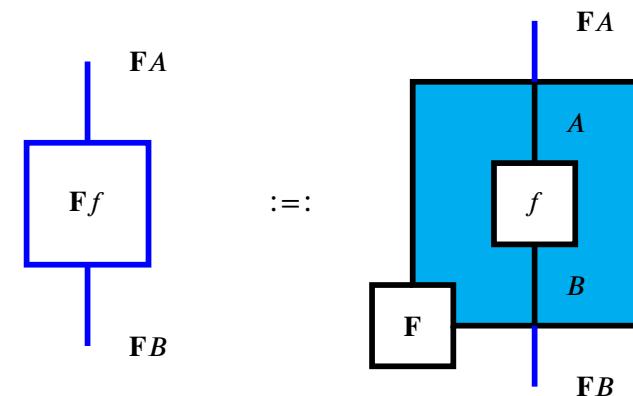


Figure 3.28: The use of a functor box is like a window from the target category \mathcal{D} into the source category \mathcal{C} ; when we know that a morphism in \mathcal{D} is the image under \mathbf{F} of some morphism in \mathcal{C} , the functor box notation is just a way of presenting all of that data at once. Since \mathbf{F} is a functor, we must have that $\mathbf{F}f; \mathbf{F}g = \mathbf{F}(f; g)$. Diagrammatically this equation is represented by freely splitting and merging functor boxes vertically.

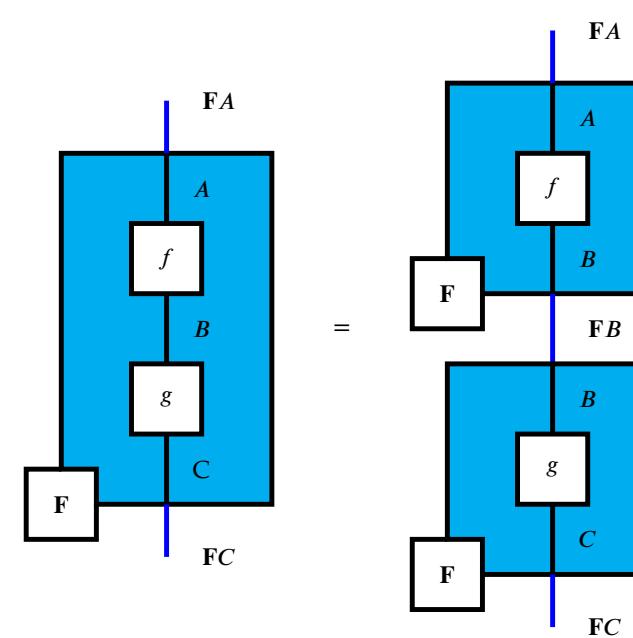


Figure 3.29: Assume that \mathbf{F} is strict monoidal; without loss of generality by the strictification theorem CITE, this lets us gloss over the associators and unitors. For \mathbf{F} to be strict monoidal, it has to preserve monoidal units and tensor products on the nose: i.e. $\mathbf{F}I_C = I_D$ and $\mathbf{F}A \otimes_D \mathbf{F}B = \mathbf{F}(A \otimes_C B)$. Diagrammatically these structural constraints amount to these equations.

The diagram consists of two parts separated by an equals sign (=).
The left part shows a white square labeled 'F' positioned above a blue rectangle divided into three horizontal sections labeled 'A' and 'B'. Below this is another blue rectangle labeled 'F(A ⊗ B)'. A vertical blue line connects the top 'F' to the middle 'A' section, and another vertical blue line connects the bottom 'F(A ⊗ B)' to the bottom 'B' section.
The right part shows a dashed rectangle divided into three horizontal sections labeled 'FA', 'A', and 'FB'. Below this is another dashed rectangle labeled 'F(A ⊗ B)'. A vertical blue line connects the top 'FA' to the middle 'A' section, and another vertical blue line connects the bottom 'FB' to the bottom 'B' section. This indicates that $\mathbf{F}(A \otimes_C B) = FA \otimes_D FB$.

Figure 3.30: What remains is the monoidality of \mathbf{F} , which is the requirement $\mathbf{F}f \otimes \mathbf{F}g = \mathbf{F}(f \otimes g)$. Diagrammatically, this equation is represented by freely splitting and merging functor boxes horizontally; analogously to how splitting vertically is the functor-boxes' way of respecting sequential composition, splitting horizontally is how they respect parallel composition.

The diagram consists of two parts separated by an equals sign (=).
The left part shows a large blue rectangle divided into four horizontal sections labeled 'A', 'C', 'B', and 'D'. Within this rectangle are two smaller white squares labeled 'f' and 'g'. Vertical blue lines connect 'f' to 'A' and 'g' to 'C'. Below this is a blue rectangle labeled 'F(A ⊗ B ⊗ C ⊗ D)'.
The right part shows two separate blue rectangles. The first is labeled 'FA' and 'A' at the top, with a 'f' square and a vertical blue line connecting it to 'A'. The second is labeled 'FC' and 'C' at the top, with a 'g' square and a vertical blue line connecting it to 'C'. Below these is a blue rectangle labeled 'F(A ⊗ B ⊗ C ⊗ D)'. This indicates that $\mathbf{F}(f \otimes g) = \mathbf{F}f \otimes \mathbf{F}g$.

Figure 3.31: And for when we want \mathbf{F} to be a (strict) symmetric monoidal functor, we are just asking that boxes and twists do not get stuck on one another.

The diagram consists of three parts separated by equals signs (=).
The first part shows a blue rectangle divided into two horizontal sections labeled 'A' and 'B'. A white square labeled 'F' is at the top left. Two blue lines, labeled 'FA' and 'FB', connect the 'F' to the 'A' and 'B' sections respectively. A blue twist arrow (an 'x' shape) connects the 'FA' line to the 'FB' line.
The second part shows the same blue rectangle with 'A' and 'B' sections. The 'F' square is now at the bottom left. The 'FA' line is at the top and the 'FB' line is at the bottom. The blue twist arrow connects the 'FA' line to the 'FB' line.
The third part shows the same blue rectangle with 'A' and 'B' sections. The 'F' square is at the bottom left. The 'FA' line is at the top and the 'FB' line is at the bottom. The blue twist arrow is now positioned such that it does not 'stick' on the 'F' square, instead curving around it to connect the 'FA' and 'FB' lines.

Figure 3.32: To motivate fibrations, first observe that by the diagrammatic equations of monoidal categories and functor boxes we have so far, we can always "slide out" the contents of a functor box out of the bottom. When can we do the reverse? That is, take a morphism in \mathcal{D} and *slide it into* a functor box? We know that in general this is not possible, because not all morphisms in \mathcal{D} may be in the image of \mathbf{F} . So instead we ask "under what circumstances" can we do this for a functor \mathbf{F} ? The answer is when \mathbf{F} is a discrete fibration.

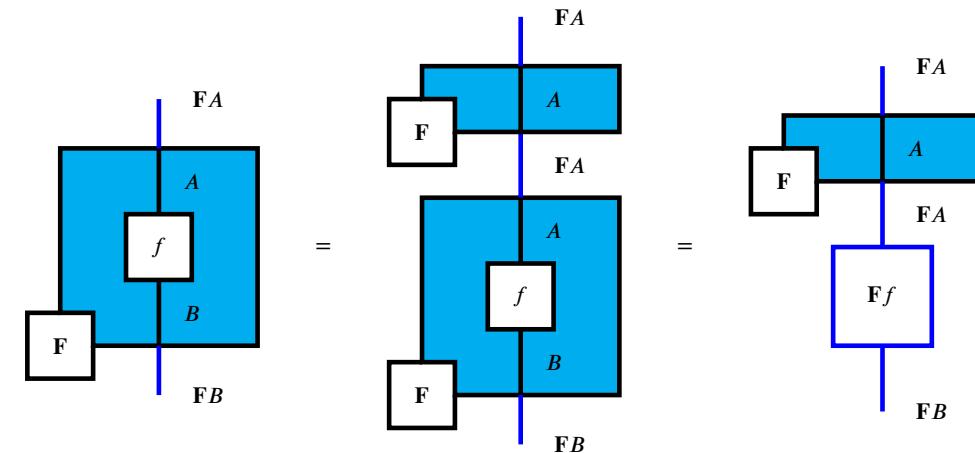


Figure 3.33:

Definition 3.3.4 (Discrete opfibration). $\mathbf{F} : \mathcal{C} \rightarrow \mathcal{D}$ is a *discrete fibration* when: for all morphisms $f : \mathbf{F}A \rightarrow B$ in \mathcal{D} with domain in the image of \mathbf{F} , there exists a unique object Φ_f^A and a unique morphism $\phi_f : A \rightarrow \Phi_f^A$ in \mathcal{C} , such that $f = \mathbf{F}\phi_f$. Diagrammatically, we can present all of the above as an equation reminiscent of sliding a morphism *into* a functor box from below.

$$\forall f : \mathbf{F}A \rightarrow B \in \mathcal{D}$$

$$\exists! \varphi_f : A \rightarrow \Phi_f^A \in \mathcal{C}$$

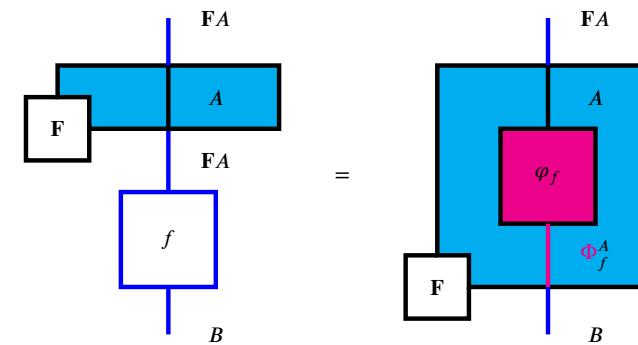


Figure 3.34:

Definition 3.3.5 (Monoidal discrete opfibration). We consider \mathbf{F} to be a (*strict, symmetric*) monoidal discrete opfibration when it is a (strict, symmetric) monoidal functor, a discrete opfibration, and the depicted equations relating lifts to interchange hold. The diagrammatic motivation for the additional coherence equations is that – if we view the lifts of opfibrations as sliding morphisms into functor boxes – we do not want the order in which sliding occurs to affect the final result. In this way, lifts behave as ‘graphical primitives’ in the same manner as interchange isotopies and symmetry twists.

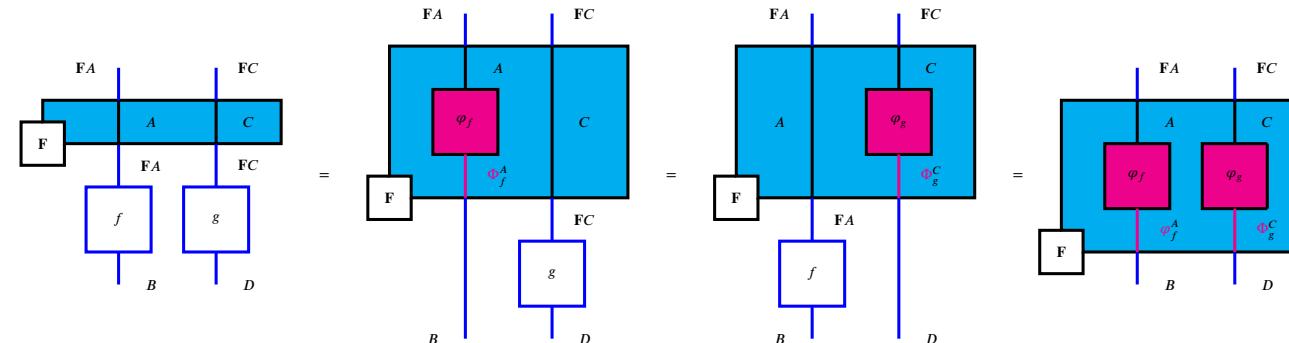


Figure 3.35: We aim to be able to use discrete monoidal functor-boxes like so. In the leftmost diagram, we would like to graphically introduce pregroup states. In the first equation (isomorphism), we would like to use the monoidal condition of the functor to horizontally merge functor boxes. In the second equation, we would like to use the discrete fibration condition of the functor to expand the box downwards, converting a string diagram obtained from a context-free grammar into a pregroup diagram. Observe that the adverb quickly has its label vertically flipped, alongside the adposition to and the sentential-complement verb sees. This is by design for all grammatical categories where pregroup typings are contextually dependent, as will be illustrated in Figure 3.4.

Figure 3.36: quickly could find itself modifying an intransitive (single noun) or transitive (two noun) verb. Suppose that it is the job of some process q' to handle intransitive verbs, and similarly q'' to handle transitive ones. We use the functor for bookkeeping, by asking it to send both q' and q'' to the dependent label \bar{q} . Treating the label as a test rather than a state allows the fibration-box to choose the right version based on the domain wires as it expands top-down.

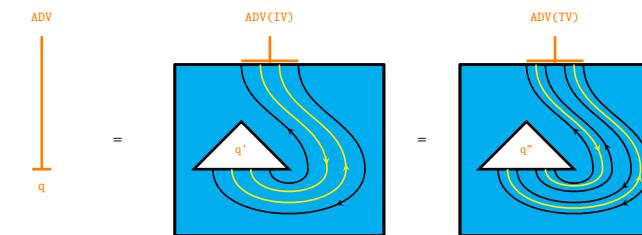
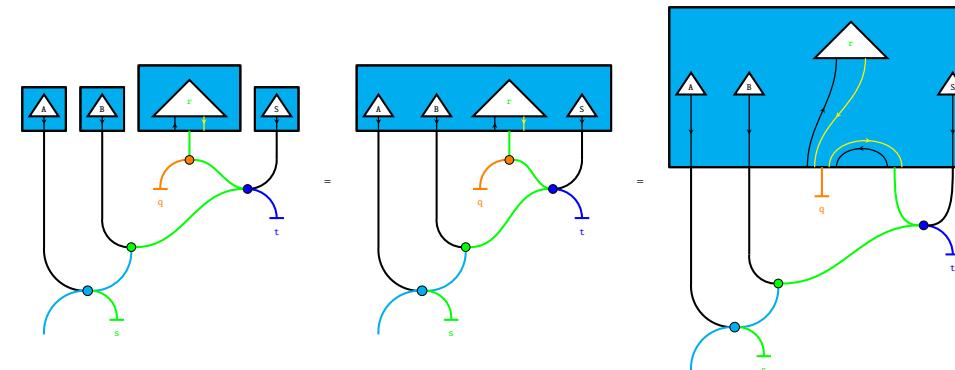
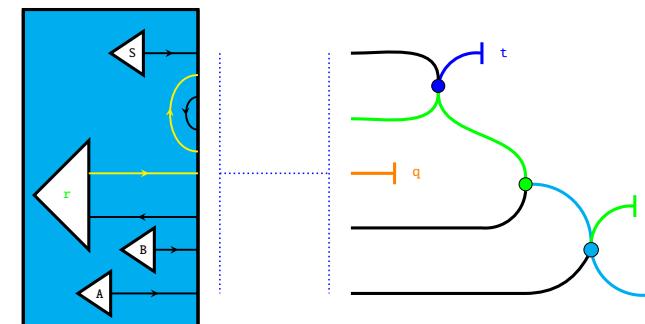
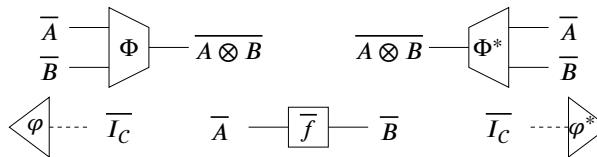


Figure 3.37: However, this procedure as described is at risk of being ill-defined. Observe that in the third diagram of Figure 3.4, the assignment of wires in the domain of the functor to wires in the codomain of the functor is only declared by diagrammatic grouping; if we consider the algebraic data available in the third diagram, really all we have is the data in the figure. How do we know which wires in the domain of the functor correspond to which wires in the codomain? Resolving this issue is the purpose of the next section.



Definition 3.3.6 (Strictified string diagrams). (Presentation taken from CITE) Fix an arbitrary (non-strict) monoidal category \mathcal{C} . The strictification $(\bar{\mathcal{C}}, \bullet)$ is defined as follows (where strictness of $\bar{\mathcal{C}}$ entitles use of string-diagrammatic notation):

- (1) Objects \bar{A} for each $A \in \mathcal{C}$
- (2) The following generators, with $\bar{f} : \bar{A} \rightarrow \bar{B}$ for each $f \in \mathcal{C}(A, B)$, where we adopt the convention of notating the monoidal unit with a dashed line:



- (3) The following functoriality equations:

$$\begin{array}{ccc} \bar{A} & \xrightarrow{\boxed{1_A}} & \bar{A} \\ \bar{A} & = & \bar{A} \end{array} \quad \begin{array}{ccc} \bar{A} & \xrightarrow{\bar{f}} & \bar{A} \\ \bar{A} & = & \bar{A} \end{array}$$

$$\begin{array}{ccc} \bar{f} & \bar{g} & = \\ \bar{f} ; \bar{g} & = & \bar{f} ; \bar{g} \end{array}$$

- (4) The following adapter equations:

$$\begin{array}{ccc} \Phi^* & \xrightarrow{\bar{f}} & \Phi \\ \Phi^* & \xrightarrow{\bar{g}} & \Phi \end{array} \quad \begin{array}{ccc} \Phi^* & \xrightarrow{\bar{f} \otimes g} & \Phi \\ \Phi^* & = & \Phi \end{array}$$

$$\begin{array}{ccc} \Phi & \xrightarrow{f \otimes g} & \Phi^* \\ \Phi & = & \Phi^* \end{array}$$

$$\begin{array}{ccc} \varphi & \dashrightarrow & \varphi^* \\ \varphi & = & \square \end{array}$$

$$\begin{array}{ccc} \varphi^* & \dashrightarrow & \varphi \\ \varphi^* & = & \dots \end{array}$$

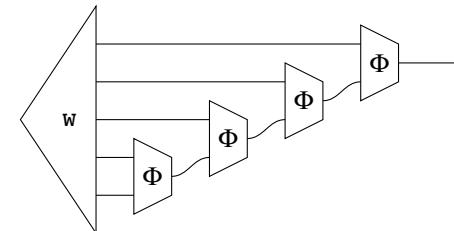
3.3.3 Strictified diagrams for monoidal categories

The crux of the issue sketched in Figure 3.4 is that while pregroup proofs – viewed as sequent trees – syntactically distinguish the roots of subtrees, interpretation as pregroup diagrams in a monoidal category forgets the subtree structure of the specific proof the diagram arises from. But it is precisely this forgotten structure that contains the algebraic data we require to keep track of (co)domain data diagrammatically. So a solution would be to force the diagrams in the blue domain recording pregroup data to hold onto this proof structure. For this purpose we use strictified diagrams for monoidal categories, defined in the margins.

We are seeking some way to algebraically group or bracket together pregroup types that arise from a single word, in a distinguished way from concatenation-as-tensor. In this way we can preserve the structure of pregroup-sequent proofs: grouping indicates a node in the proof-tree, while tensor indicates parallel composition of proof trees. With strictified diagrams, we can model grouping (ab)using biased tensor structure, i.e. treating $\cdots ((A \otimes B) \otimes C) \cdots \otimes \cdots Z$ as a bracketed expression $[A \otimes B \cdots \otimes Z]$.

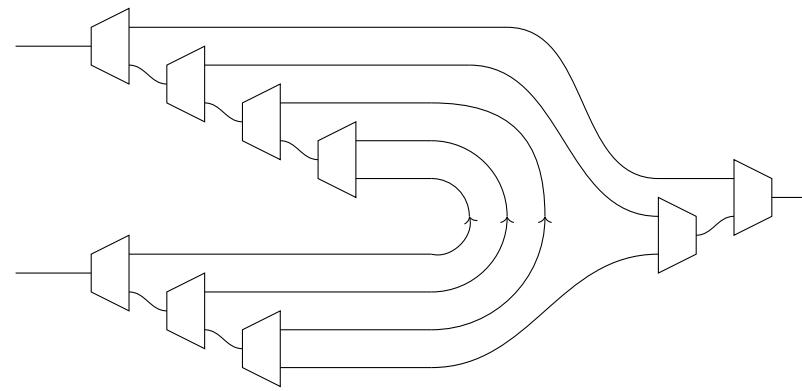
Construction 3.3.8 (Pregroups with bracketing). Where \mathcal{P} is a monoidal category generated by pregroup states and (directed) cups, we define pregroups-with-bracketing as subcategory of the free strictification $\bar{\mathcal{P}}$, generated by:

- For each pregroup state $w : I \rightarrow \bigotimes_i X_i$, a strictified state generator $w : I \rightarrow (\cdots ((X_1 \otimes X_2) \otimes \cdots X_i) \cdots)$ with left-nested syntactic tensors. To illustrate, a state with 5 wires would correspond to a generator as follows:



- Let $[A \cdot B \cdots Z]$ denote the left-nested tensoring $((A \otimes B) \cdots \otimes Z)$, and let \mathbf{X} denote $(\bigotimes_i X_i)$. For each directed cap $\mathbf{X} \otimes \mathbf{X}^{-1} \rightarrow I$ (and symmetrically for caps of the other direction and cups), and for each pair of bracketed types $[A \cdot \mathbf{X}]$ and $[\mathbf{X}^{-1} \cdot B]$, we ask for a generator that fully detensors, applies the directed cup, and then retensors. Diagrammatically, this amounts to asking for generators that look like the following,

that mimick a single proof step.



Construction 3.3.8 solves the (co)domain assignment problem by organising the data as a monoidal co-functor.

- (3) The following representations of the natural isomorphisms in the definition of a monoidal category:

$$\bar{\alpha} := \begin{array}{c} \Phi^* \\ \swarrow \quad \searrow \\ \Phi^* & \Phi \\ \downarrow & \downarrow \\ \Phi & \Phi \end{array}$$

$$\overline{\alpha^{-1}} := \begin{array}{c} \Phi^* \\ \swarrow \quad \searrow \\ \Phi^* & \Phi \\ \downarrow & \downarrow \\ \Phi & \Phi \end{array}$$

$$\bar{\lambda} := \begin{array}{c} \Phi^* \\ \swarrow \quad \searrow \\ \Phi^* & \varphi^* \\ \downarrow & \downarrow \end{array} \quad \bar{\rho} := \begin{array}{c} \Phi^* \\ \swarrow \quad \searrow \\ \varphi^* & \Phi \\ \downarrow & \downarrow \end{array}$$

$$\overline{\lambda^{-1}} := \begin{array}{c} \varphi \\ \swarrow \quad \searrow \\ \Phi \end{array} \quad \overline{\rho^{-1}} := \begin{array}{c} \varphi \\ \swarrow \quad \searrow \\ \Phi \end{array}$$

Proposition 3.3.7 ($\bar{\mathcal{M}}$ and \mathcal{M} are monoidally equivalent). *Proof.*
[CITE](#) □

Cofunctors are arguably necessary even for vanilla DisCoCat. Pregroups by themselves are thin rigid autonomous categories, meaning there is at most one morphism between any pair of objects and that there are directed cups and caps in addition to plain monoidal structure. Contra the aims of DisCoCat, this implies there is no nontrivial strong monoidal functor from pregroups into FinVect^\otimes , where strength for a monoidal functor means that $Fx \otimes Fy \simeq F(x \otimes y)$ in the codomain. This is because for any object x in the pregroup, any strong monoidal F into FinVect^\otimes that interprets duals in the usual way must identify the following diagrams up to isomorphism of finite vector spaces with their duals:

placeholder

The only vector space Fx for which this holds is the base field of Vect ; the monoidal unit. The same argument applies for functors into other strong compact closed categories, such as variants of Rel – another common family of semantic domains – where the only valid target is the singleton as monoidal unit.

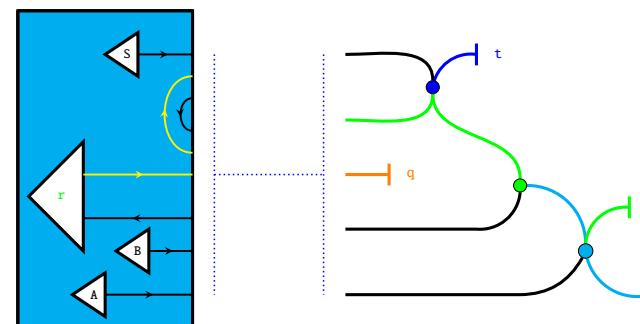
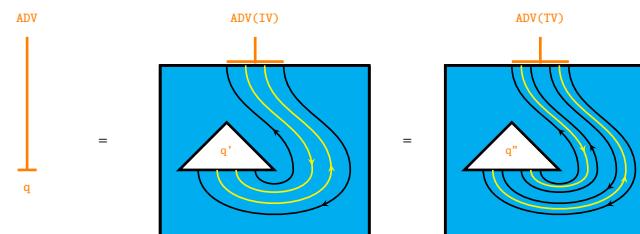
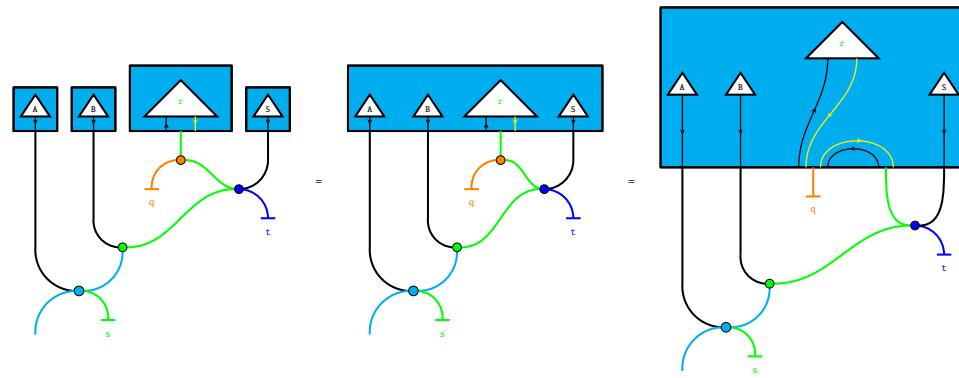
3.3.4 Monoidal Cofunctors

Asking for opfibrations to be discrete is an unnecessarily strict condition. Often rather than asking for uniqueness, we are happy to supply a choice of lift...

Definition 3.3.9 (Bijective-on-objects functor). A functor $F : \mathcal{C} \rightarrow \mathcal{D}$ is *bijective-on-objects* if for all $d \in \mathcal{D}$, there exists a $c \in \mathcal{C}$ such that $Fc = d$.

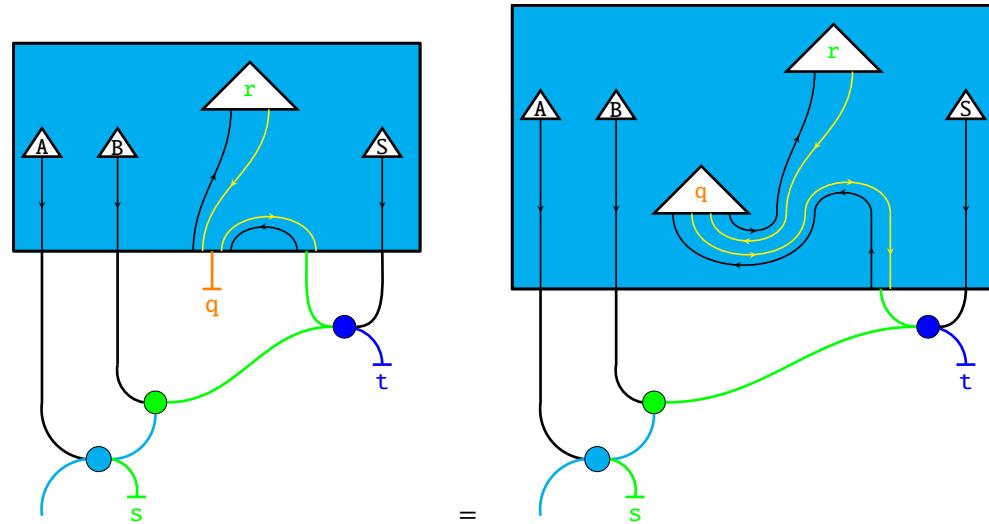
Proposition 3.3.10.

3.4 Communicative constraints as a cofunctor from productive to parsing grammar

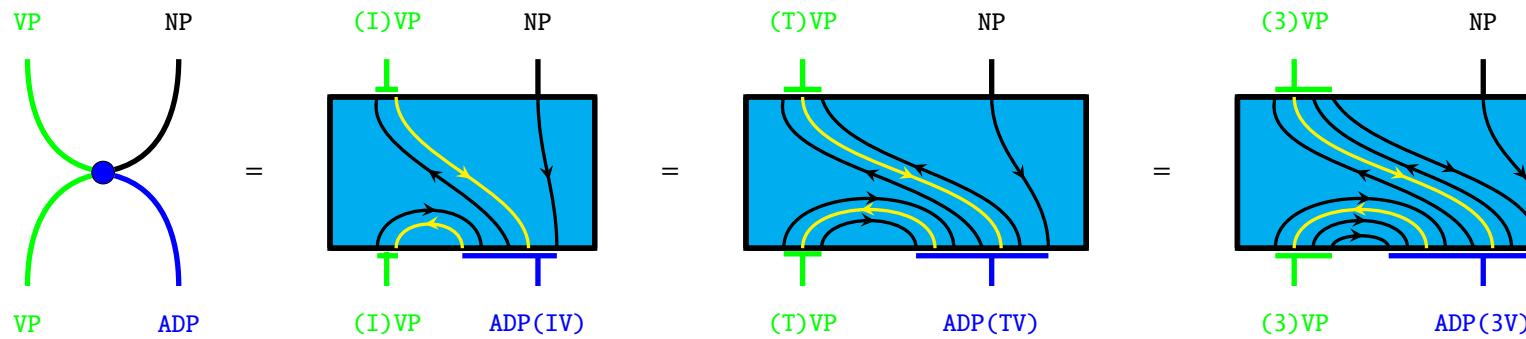


Since the functor is a monoidal discrete fibration, it introduces the appropriate choice of quickly when we

pull the functor-box down, while leaving everything else in parallel alone.

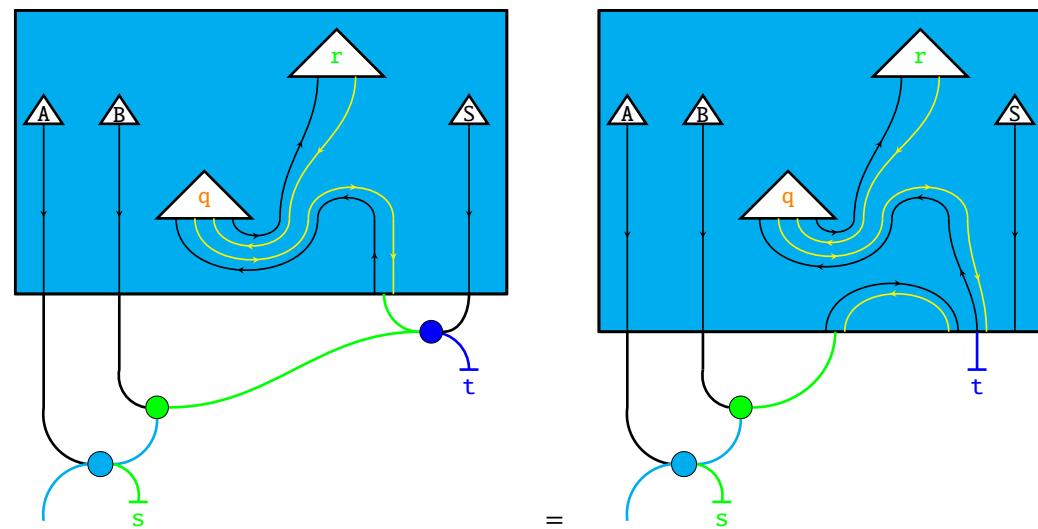


Adpositions can apply for verbs of any noun-arity. We again use the fiber of the functor for bookkeeping by asking it to send all of the following partial pregroup diagrams to the adposition generator. We consider the pregroup typing of a verb of noun-arity $k \geq 1$ to be ${}^{-1}n \cdot s \cdot n^{-1} \cdots n^{-1}$.

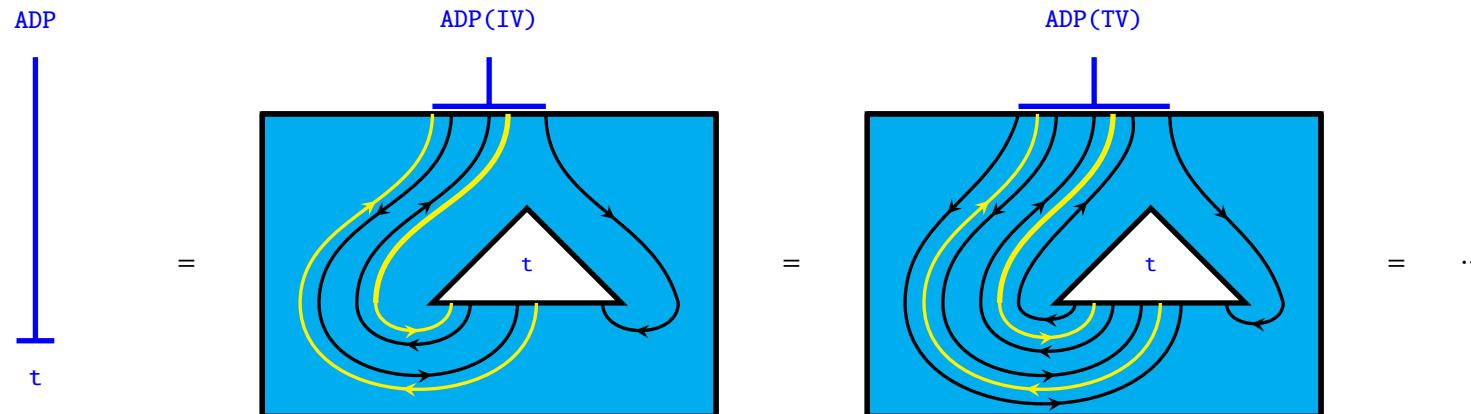


When we pull down the functor-box, the discrete fibration introduces the appropriate choice of diagram from

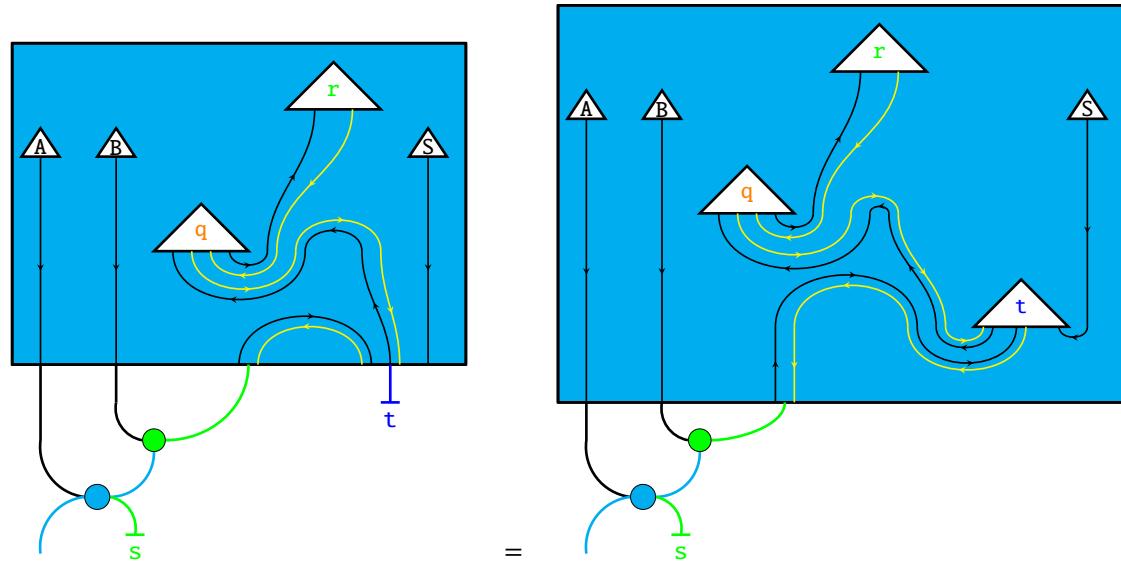
above, corresponding to the intransitive verb case.



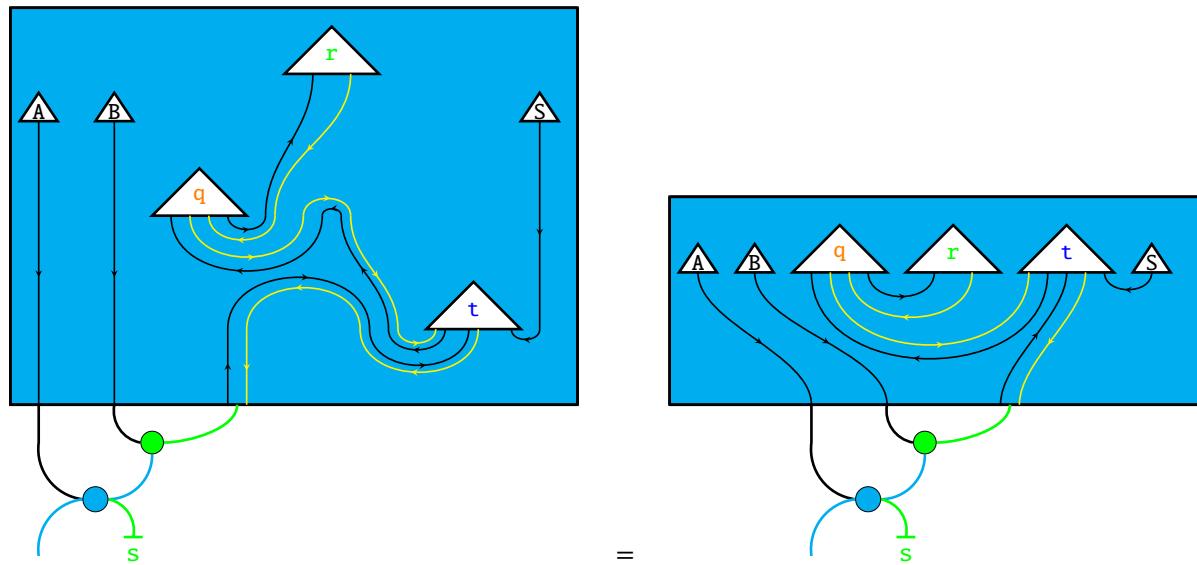
Similarly to quickly, we suppose we have a family of processes for the word *to*, one for each noun-arity of verb.



Again the discrete fibration introduces the appropriate choice of α when we pull the functor box down.

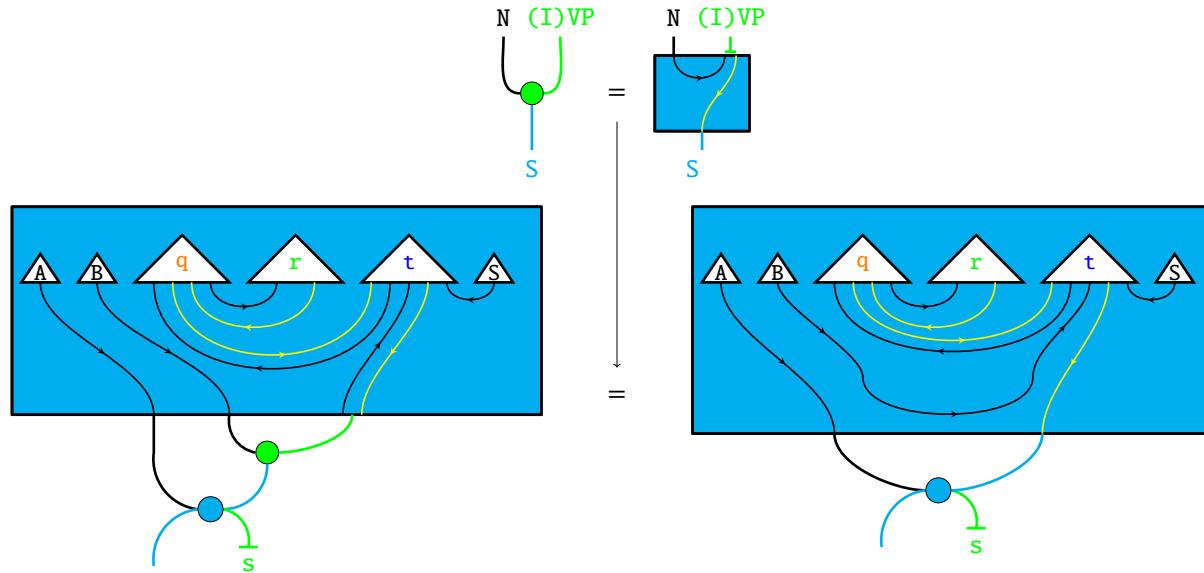


Now we visually simplify the inside of the functor-box by applying yanking equations.

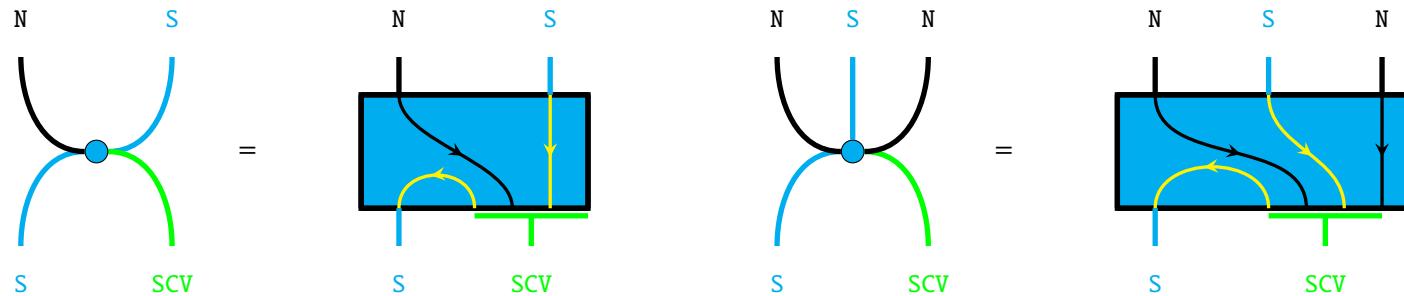


Similarly as before, we can pull the functor-box past the intransitive verb node. There is only one pregroup

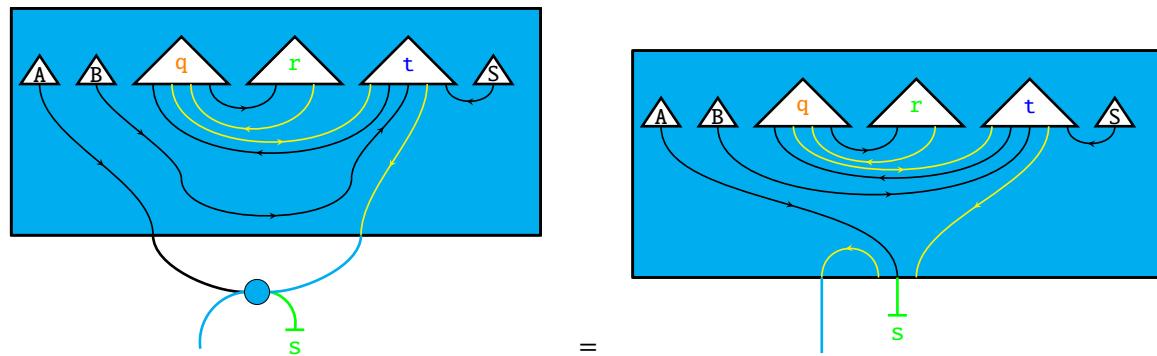
type ${}^{-1}n \cdot s$ that corresponds to the grammatical category **(I)VP**.



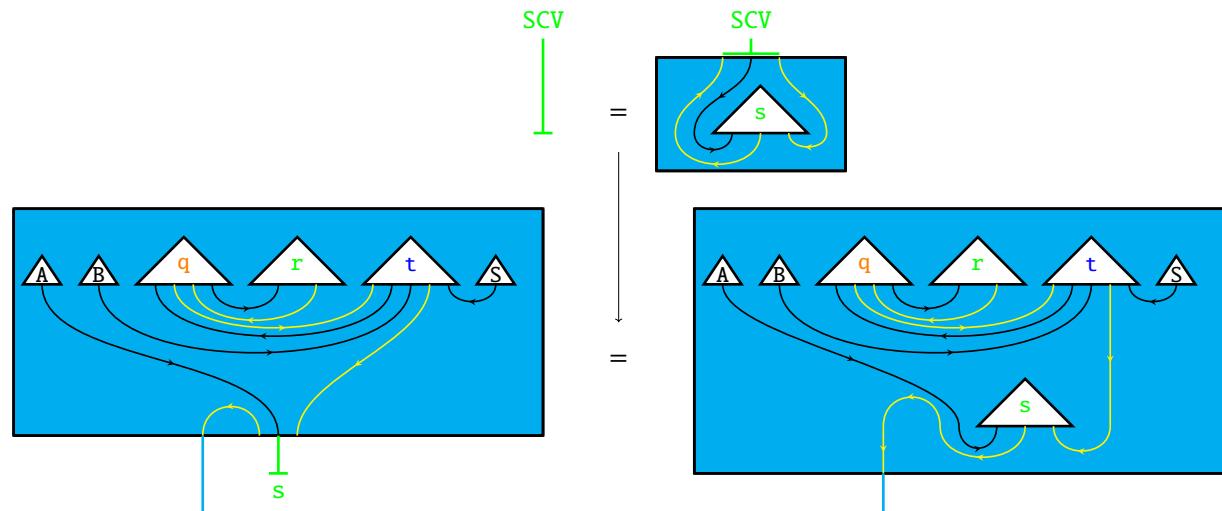
Proceeding similarly, we can pull the functor-box past the sentential-complement-verb node. There are multiple possible pregroup types for **SCV**, depending on how many noun-phrases are taken as arguments in addition to the sentence. For example, in **Alice sees [sentence]**, **sees** returns a sentence after taking a noun to the left and a sentence to the right, so it has pregroup typing ${}^{-1}n \cdot s \cdot s^{-1}$. On the other hand, for something like **Alice tells Bob [sentence]**, **tells** returns a sentence after taking a noun (the teller) to the left, a noun (the tellee) to the left, and a sentence (the story) to the left, so it has a pregroup typing ${}^{-1}n \cdot s \cdot n^{-1} \cdot s^{-1}$. These two instances of sentential-complement-verbs are introduced by different nodes. We can record both of these pregroup typings in the functor by asking for the following:



Pulling down the functor box:

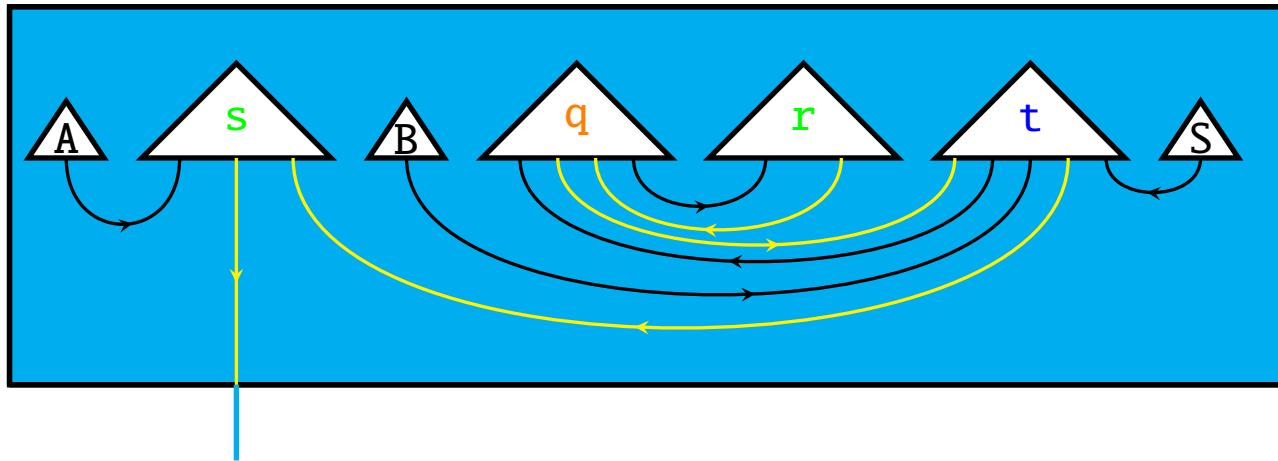


As before, we can ask the functor to send an appropriate partial pregroup diagram to the dependent label **sēe**.



Now again we can visually simplify using the yanking equation and isotopies, which obtains a pregroup

diagram.



The pregroup diagram corresponds to a particular pregroup proof of the syntactic correctness of the sentence Alice sees Bob run quickly to school.

$$\begin{array}{c}
 \frac{\textcolor{brown}{q} : (-^1 n \cdot s) \cdot (-^1 n \cdot s)^{-1} \quad \textcolor{green}{r} : ^{-1} n \cdot s}{\textcolor{brown}{q} \sqcup \textcolor{green}{r} : ^{-1} n \cdot s} \quad \frac{}{\textcolor{blue}{t} : ^{-1} (-^1 n \cdot s) \cdot (-^1 n \cdot s) \cdot n^{-1}} \\
 \frac{\textcolor{brown}{q} \sqcup \textcolor{green}{r} \sqcup \textcolor{blue}{t} : (-^1 n \cdot s) \cdot n^{-1}}{\textcolor{brown}{q} \sqcup \textcolor{green}{r} \sqcup \textcolor{blue}{t} \sqcup \textcolor{black}{S} : ^{-1} n \cdot s} \quad \textcolor{black}{S} : n
 \end{array}$$

$$\frac{A : n \quad \textcolor{green}{s} : ^{-1} n \cdot s \cdot s^{-1}}{A \sqcup \textcolor{green}{s} : s \cdot s^{-1}} \quad \frac{B : n}{B \sqcup \textcolor{brown}{q} \sqcup \textcolor{green}{r} \sqcup \textcolor{blue}{t} \sqcup \textcolor{black}{S} : s}$$

$$\textcolor{brown}{A} \sqcup \textcolor{green}{s} \sqcup \textcolor{black}{B} \sqcup \textcolor{brown}{q} \sqcup \textcolor{green}{r} \sqcup \textcolor{blue}{t} \sqcup \textcolor{black}{S} : s$$

3.5 Relating circuits and CFGs fibrationally

3.6 Where internal wirings come from

3.6.1 Discrete monoidal fibrations for grammatical functions

3.7 Discussion and Limitations

WHAT ARE THE ASSUMPTIONS AND LIMITATIONS? In my view, the preceding analysis is fair if one entertains the following three commitments.

1. At some level, semantics is compositional, and syntax directs this composition.
2. Speakers produce sentences, and listeners parse sentences.
3. Speakers and listeners understand each other, insofar as the compositional structure of their semantic representations are isomorphic.

Insofar as compositionality entails that infinite ends can be achieved by finite combinatorial means, discrete monoidal fibrations are bookkeeping for an idealised structural correspondence between the components of productive and parsing grammars, and internal wirings arise as balancing terms in the bookkeeping.

The first assumption establishes an idealised view of communication and compositionality where there are no extraneous rules in the language, i.e. that a particular phrase of five or sixty-seven words is to be parsed exceptionally. This is not the case in natural languages, where everyday idioms may be considered semantically atomic despite being compositionally decomposable. For example, in Mandarin, 马上 is the concatenation of "horse" and "up", and would be "on horseback" if interpreted literally, but is treated as an adverbial "as soon as possible" in imperative contexts. I use the hedging phrase "at some level" in the first assumption to describe the compositionality of semantics just to indicate an assumption that we are not dealing with exceptional rules all the way up.

The second assumption commits to an idealisation that speakers and listeners communicate for the purposes of exchanging propositional information as well-formed and disambiguated sentences, which is clearly not all that language is for. I can promise nothing regarding questions, imperatives, speech acts, and so on.

The third assumption asks that one entertain string diagrams as representative of what the content of language is, and even so, it still requires some elaboration on what is meant by "understanding", as it is obviously untrue that everyone understands one another. I do not mean understanding in the strong sense as a form of telepathy of mental states, c.f. Wittgenstein's beetle-in-a-box thought experiment. I mean that insofar as the speaker and listener both have their own ideas about cats, sitting, space, and mats, their respective mental models of the cat sat on the mat are indistinguishable as far as meaningfully equivalent syntactic re-presentations and probings go; for instance, both speakers ought to agree that the mat is beneath the cat, and both speakers ought to agree despite the concrete images in their minds that there is insufficient information to know the colour of the cat from the sentence alone, and so on. This is a shallow form of understanding; consider the case where one communicator is a human with mental models encoded in meat and another is an LLM with tokens encoding who-knows-what – they may be in perfect agreement about

rephrasings of texts for an arbitrary finite amount of communication, even if the representations of the latter are not compositional. It would be nice to ask that "mutual understanding" requires structurally equivalent (as opposed to extensionally indistinguishable) meaning-representation mechanisms between language agents c.f. Chomsky's universal grammar, but our means of achieving mutual understanding in practice seems to align with the shallow view: we pose comprehension challenges and ask clarifying questions all at the level of language, without taking a scalpel to the other's head. Depending on one's view of what understanding language entails, it may be that humans and LLMs both understand language in their own way, but mutual understanding between the two kinds is an illusion.

Despite these limitations, I believe that this formal approach to grounding relationships between productive and parsing grammars in mathematical considerations surrounding communication has some merits.

THEORIES OF GRAMMAR BY THEMSELVES ARE INSUFFICIENT TO ACCOUNT FOR COMMUNICATION. At minimum, for every grammar that produces sentences, one must also provide a corresponding parsing grammar. A theory of grammar that only produces correct sentences or correct parses is a 'theory' of language outperformed in every respect by an LLM. So we must distinguish between grammars of the speaker and listener, and then investigate how they cohere.

COHERENCE OF THEORIES OF GRAMMAR IS INEXTRICABLE FROM SEMANTICS. We are interested in the ideal of communication, the end result of a single turn of which is that both speaker and listener have the same semantic information, whether that be a logical expression or something else. A consequence of this criterion is that in order to obtain an adequate account of communication, we must seek a relation between grammars and semantics beyond weak and strong equivalence of pairs of theories of grammar.

Firstly, "weak equivalence" between grammar formalisms in terms of possible sets of generated sentences is insufficient. Weak equivalence proofs are mathematical busywork that have nothing to do with a unified account of syntax and semantics. For example, merely demonstrating that, e.g. pregroup grammars and context-free grammars can generate the same sentences [] only admits the possibility that a speaker using a context-free grammar and a listener using a pregroup grammar *could* understand each other, without providing any explanation *how*. But we already know that users of language *do* understand one another more-or-less, so the exercise is more-or-less pointless.

Secondly, "strong equivalence" that seeks equivalence at a structural level between theories of syntax often helps, but is not always necessary. I will explain by analogy. Theories of syntax are like file formats, e.g. .png or .jpeg for images. A model for a particular language is a particular file or photograph. The task here is to show that two photographs in different file formats that both purport to model the same language are really photographs of the same thing from different perspectives. It is overkill to demonstrate that all .pngs and .jpeg are structurally bijective, just as it is overkill to show that, say, context-free grammars are strongly equivalent to pregroup grammars, because there are context-free and pregroup grammars that generate sets

of strings that have nothing to do with natural language. It could just as well be that there is a pair of productive and parsing grammar-formats that are not strongly equivalent, but happen to coincide for a particular natural language – in this sense, asking for a discrete monoidal fibration is a way to check a weaker condition than strong equivalence that achieves the more specific aim of determining whether a pair of productive and parsing grammars for a language are plausible models.

A systematic analysis of communication requires intimacy with specific grammars and a specific semantics. Specific grammars – and not formats of grammar, such as "all CFGs" – that model natural languages, even poorly, are the only relevant objects of study for any form of language intended to communicate information. Once you have a specific grammar that produces sentences in natural language, then to explain communication, you must supply a specific partnered parsing grammar such that on the produced sentences, both grammars yield the same semantic objects by a Montagovian approach, broadly construed as a homomorphism from syntax to semantics. On this account, syntax does not hold a dictatorship over semantics, but we can find duarchies, and in these duumvirates the two syntaxes and the semantics mutually constrain one another.

IT IS WORTH NOTING THAT IN PRACTICE, NEITHER GRAMMAR NOR MEANING STRICTLY DETERMINES THE OTHER.

Clearly there are cases where grammar supercedes: when Fondo hears `man bites dog`, despite his prior prejudices and associations about which animal is more likely to be biting, he knows that the `man` is doing the biting and the `dog` is getting bitten. Going the other way, there are many cases in which the meaning of a subphrase affects grammatical acceptability and structure.

Example 3.7.1 (Exclamations: how meaning affects grammar). The following examples from [Lakofflecture] illustrate how whether a phrase is an *exclamation* affects what kinds of grammatical constructions are acceptable. By this argument, to know whether something is an exclamation in context is an aspect of meaning, so we have cases where meaning determines grammar. Observe first that the following three phrases are all grammatically acceptable and mean the same thing.

`nobody knows` how many beers Bob drinks

`who knows` how many beers Bob drinks

`God knows` how many beers Bob drinks

The latter two are distinguished when `God knows` and `who knows` are exclamations. First, the modularity of grammar and meaning may not match when an exclamation is involved. For example, negating the blue text, we obtain:

`somebody knows` how many beers Bob drinks

`who doesn't know` how many beers Bob drinks

God doesn't know how many beers Bob drinks

The first two are acceptable, but mean different things; the latter means to say that everyone knows how many beers Bob drinks, which is stronger than the former. The last sentence is awkward: unlike in the first two cases, the quantified variable in the (gloss) $\dots \neg \exists x_{Person} \dots$ of God knows is lost, and what is left is a literal reading $\dots \neg \text{knows}(\text{God}, \dots) \dots$. Second, whether a sentence is grammatically acceptable may depend on whether an exclamation is involved. God knows and who knows can be shuffled into the sentence to behave as an intensifier as in:

Bob drinks God knows how many beers

Bob drinks who knows how many beers

But it is awkward to have:

Bob drank nobody knows how many beers

And it is not acceptable to have:

Bob drank Alice knows how many beers

Example 3.7.2 (Garden path sentences). So-called "garden path" sentences illustrate that listeners have to make choices to resolve lexical ambiguities. One such garden-path sentence is The old man the boat, where typically readers take The old man as a noun-phrase and the boat as another noun-phrase. We can sketch how the readers might think with a (failed) pregroup grammar derivation:

$$\begin{array}{c} \text{the : } n \cdot n^{-1} \quad \text{old : } n \cdot n^{-1} \\ \hline \text{the_old : } n \cdot n^{-1} \quad \text{man : } n \quad \text{the : } n \cdot n^{-1} \quad \text{boat : } n \\ \hline \text{the_old_man : } n \quad \text{the_boat : } n \\ \text{Not a sentence!} \end{array}$$

So the reader has to backtrack, taking The old as a noun-phrase and man as the transitive verb. This yields a sentence as follows:

$$\begin{array}{c} \text{the : } n \cdot n^{-1} \quad \text{old : } n \\ \hline \text{the_old : } n \quad \text{man : } n^{-1} \cdot n \cdot s \cdot n^{-1} \quad \text{the : } n \cdot n^{-1} \quad \text{boat : } n \\ \hline \text{the_old_man : } s \cdot n^{-1} \quad \text{the_old : } n \\ \text{the_old_man_the_boat_ : } s \end{array}$$

Garden-path sentences illustrate that listeners must make choices about what grammatical roles to assign words. We make these kinds of contextual decisions all the time with lexically ambiguous words or highly homophonic languages like Mandarin; garden-path sentences are special in that they trick the default strategy badly enough that the mental effort for correction is noticeable.

Example 3.7.3 (Ambiguous scoping). Consider the following sentence:

Everyone loves someone

The sentence is secretly (at least) two, corresponding to two possible parses. The usual reading is (glossed) $\forall x \exists y : \text{loves}(x, y)$. The odd reading is $\exists y \forall x : \text{loves}(x, y)$: a situation where there is a single person loved by everyone. We can sketch this difference formally using a simple combinatory categorial grammar.

$$\begin{array}{c}
 \frac{\text{everyone} : (n \multimap s) \multimap s \quad \text{loves} : n \multimap s \multimap n}{\text{everyone_loves} : s \multimap n} \qquad \text{someone} : (s \multimap n) \multimap s \\
 \hline
 \text{everyone_loves_someone} : s
 \end{array}$$

$$\begin{array}{c}
 \frac{\text{loves} : n \multimap s \multimap n \quad \text{someone} : (s \multimap n) \multimap s}{\text{loves_someone} : n \multimap s} \\
 \hline
 \text{everyone_loves_someone} : s
 \end{array}$$

CCGs have functorial semantics in any cartesian closed category, such as one where morphisms are terms in the lambda calculus and composition is substitution. So we might specify a semantics as follows:

$$\llbracket \text{everyone} \rrbracket = \lambda(\lambda x. V(x)). \forall x : V(x) \tag{3.5}$$

$$\llbracket \text{loves} \rrbracket = \lambda x \lambda y. \text{loves}(x, y) \tag{3.6}$$

$$\llbracket \text{someone} \rrbracket = \lambda(\lambda y. V(y)). \exists y : V(y) \tag{3.7}$$

Now we can plug-in these interpretations to obtain the two different meanings. We decorate with corners just to visually distinguish which bits are partial first-order logic.

$$\begin{array}{c}
 \frac{\lambda(\lambda x. V(x)). \Gamma \forall x : V(x)^\sqcap : (n \multimap s) \multimap s \quad \lambda x \lambda y. \Gamma \text{loves}(x, y)^\sqcap : n \multimap s \multimap n}{\lambda y. \Gamma \forall x : \text{loves}(x, y)^\sqcap : s \multimap n} \qquad \frac{\lambda(\lambda y. V(y)). \Gamma \exists y : V(y)^\sqcap : (s \multimap n) \multimap s}{\Gamma \exists y \forall x : \text{loves}(x, y)^\sqcap : s} \\
 \hline
 \Gamma \exists y \forall x : \text{loves}(x, y)^\sqcap : s
 \end{array}$$

$$\begin{array}{c}
 \frac{\lambda(\lambda x. V(x)). \Gamma \forall x : V(x)^\sqcap : (n \multimap s) \multimap s \quad \lambda x \lambda y. \Gamma \text{loves}(x, y)^\sqcap : n \multimap s \multimap n \quad \lambda(\lambda y. V(y)). \Gamma \exists y : V(y)^\sqcap : (s \multimap n) \multimap s}{\lambda x. \Gamma \exists y : \text{loves}(x, y)^\sqcap : n \multimap s} \\
 \hline
 \Gamma \forall x \exists y : \text{loves}(x, y)^\sqcap : s
 \end{array}$$

Example 3.7.4. (grammar pieces – as simple as it gets) Bob runs. Bob quickly runs. Bob drinks beer.
Bob quickly drinks beer.

3.8 A generative grammar for text circuits

3.8.1 Text Circuits

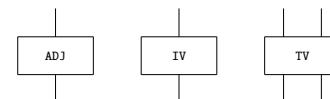
Our *text circuits* are made up of three ingredients:

- wires
- boxes, or gates
- boxes with holes that fit a box, or 2nd order gates

Firstly, nouns are represented by wires, each ‘distinct’ noun having its own wire:

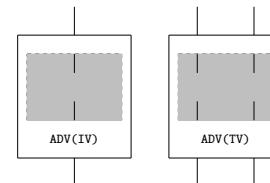


We represent adjectives, intransitive verbs, and transitive verbs by gates acting on noun-wires:

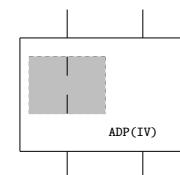


Since a transitive verb has both a subject and an object noun, that will then be two noun-wires, while adjectives and intransitive verbs only have one.

Adverbs, which modify verbs, we represent as boxes with holes in them, with a number of dangling wires in the hole indicating the shape of gate expected, and these should match the input- and output-wires of the box with the whole:

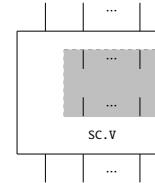


Similarly, adpositions also modify verbs, by moreover adding another noun-wire to the right:

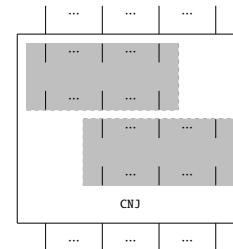


¹ The ‘dot dot dot’ notation within boxes is graphically formal . Interpretations of such boxes were earlier formalised in .
;;; and

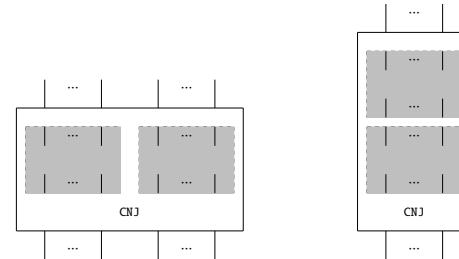
For verbs that take sentential complements and conjunctions, we have families of boxes to accommodate input circuits of all sizes. They add another noun-wire to the left of a circuit:¹



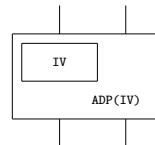
while conjunctions are boxes that take two circuits which might share labels on some wires:



As special cases, when the noun-wires of two circuits are disjoint (left), or coincide (right), conjunctions are depicted as follows:

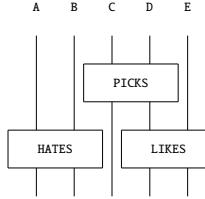


Of course filled up boxes are just gates

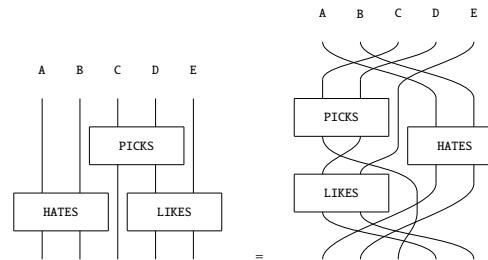


and we will now discuss how those compose. Gates compose sequentially by matching labels on some of their noun-wires and in parallel when they share no noun-wires, to give text circuits, which by convention we

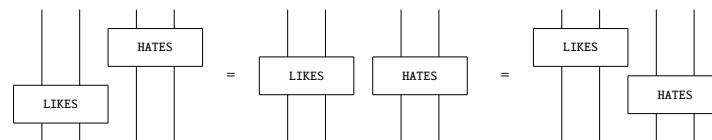
read from top-to-bottom:



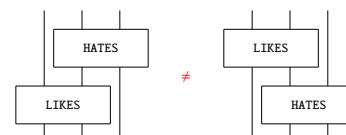
Convention 3.8.1. Sometimes we allow wires to twist past each other, and we consider two circuits the same if their gate-connectivity is the same:



Since only gate-connectivity matters, we consider circuits the same if all that differs is the horizontal positioning of gates composed in parallel:

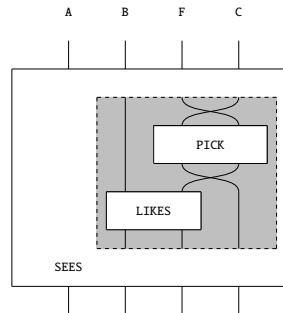


We do care about output-to-input connectivity, so in particular, we **do not** consider circuits to be equal up to sequentially composed gates commuting past each other:



Example 3.8.2. The sentence ALICE SEES BOB LIKES FLOWERS THAT CLAIRE PICKS can intuitively be given

the following text circuit:



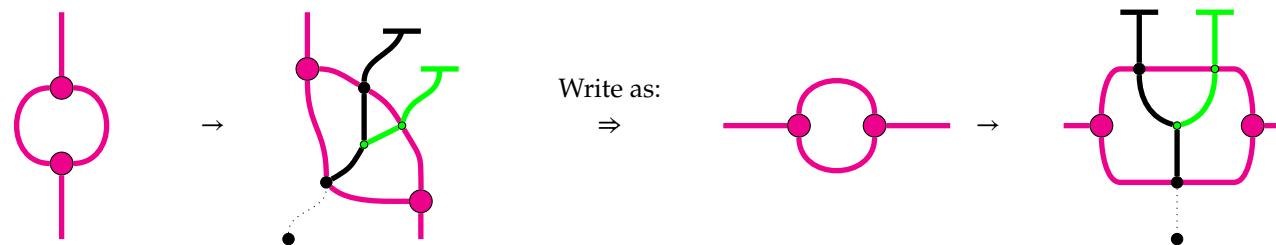
Example 3.8.3. Figure 3.38 illustrates how to directly generate text circuits, by providing an example analogous to the text generated in Figure ?? where we used our hybrid grammar.

3.8.2 *A circuit-growing grammar*

There are many different ways to write an n -categorical signature that generates circuits. Mostly as an illustration of expressive capacity, I will provide a signature where the terms "surface" and "deep" structure are taken literally as metaphors; the generative grammar will grow a line of words in (roughly) syntactic order, and like mushrooms on soil, the circuits will appear as the mycelium underneath words.

SIMPLIFICATIONS: Propositions only, no determiners, only one tense, no morphological agreement between nouns and their verbs and referring pronouns, and we assume that adverbs, adverbial adjunctions, and adjectives stack indefinitely and without further order requirements; e.g. *Yesterday yesterday yesterday Alice happily secretly finds red big toy shiny car that he gives to Bob.* we consider grammatical enough. For now, we consider only the case where adjectives and adverbs appear before their respective noun or verb. Note that all of these limitations can principle be overcome by the techniques we developed in Section 2.2.5 for restricted tree-adjoining and links.

HOW TO READ THE FOLLOWING DIAGRAMS: We work in a dimension where wires behave symmetric monoidally by homotopy, and the signature still works if interpreted in a compact closed setting. We start each derivation with a pink sentence bubble, which we depict for aesthetic purposes as horizontal, which amounts to picking slightly different axes by which to read diagrams. The insides of the sentence bubble will fill up with a circuit, and labelled words will appear on the top surface bubble like stylised mushrooms, to be read off left-to-right. We will only express the rewrite rules; the generators of lower dimension are implicit.

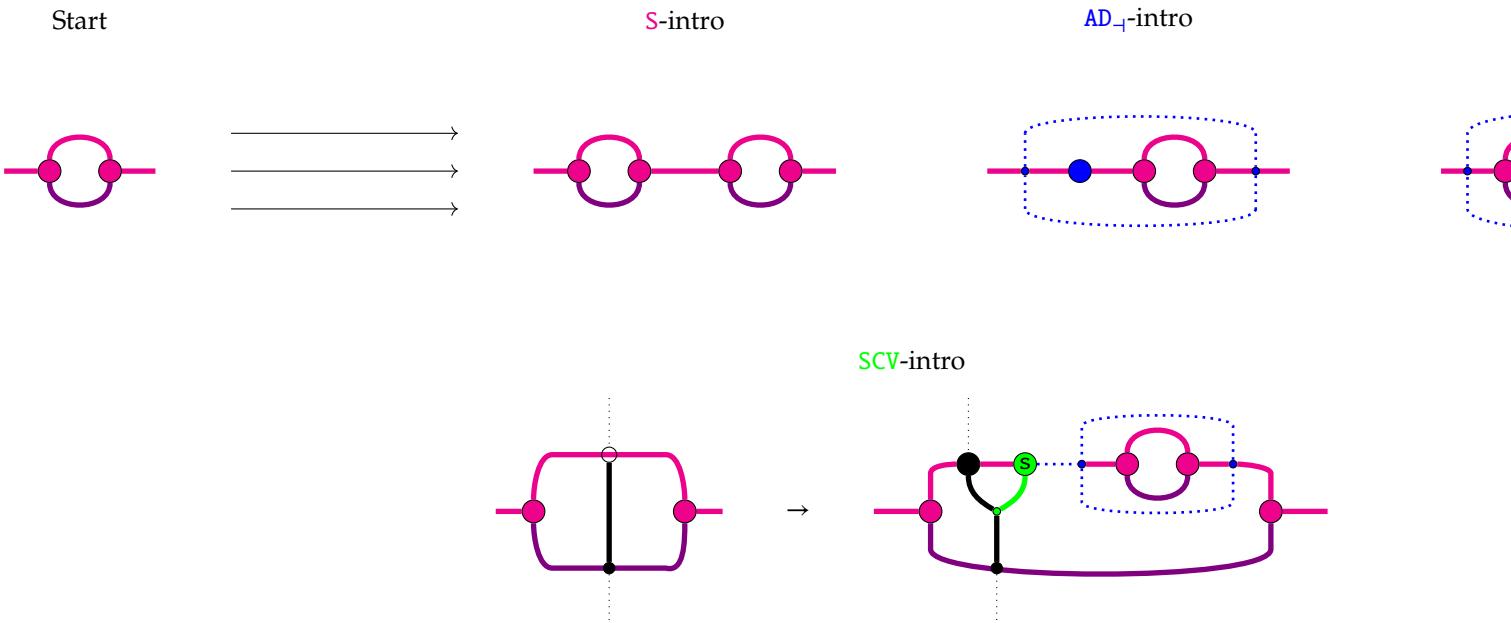


THE LEXICON:

3.8.3 *Sentences*

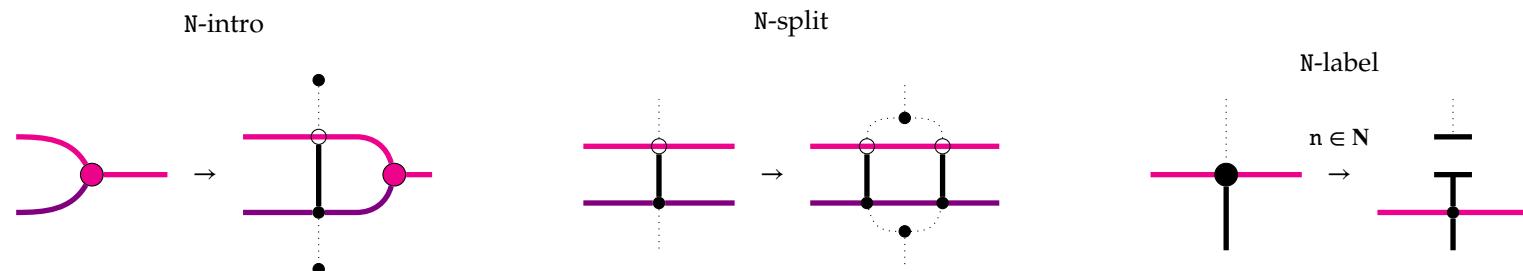
Before the contents of a sentence are even decided, we may decide to (top row, left to right) get another sentence ready, introduce an adverbial adjunction (such as *yesterday*), or introduce a conjunction of two sentences (such as *because*). Introducing new words yields dots that carry information of the grammatical category of the word, and there are separate rewrites that allow dots to be labelled according to the lexicon. In the bottom, we have a rewrite that allows a sentence with one unlabelled noun to be the subject of a "sentential complement verb", abbreviated SCV: these are verbs that take sentences as objects rather than other nouns, and they are typically verbs of cognition, perception, and causation, such as *Alice suspects Bob drinks*.

The blue-dotted lines are just syntactic guardrails that correspond to the holes in boxes of circuits later.

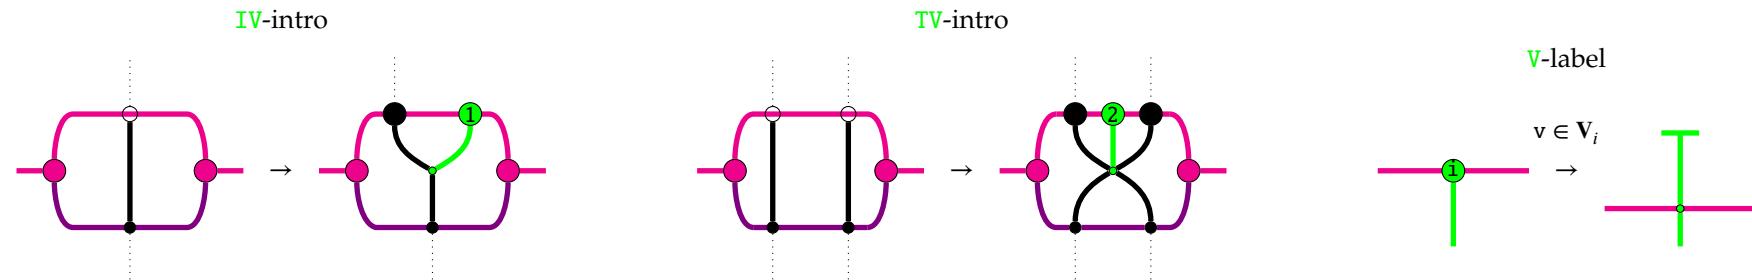


3.8.4 Simple sentences

Within each sentence bubble, each derivation starts life as a "simple sentence", which only involves nouns and a single verb, which is either an intransitive verb that takes a single noun argument, or a transitive verb that takes two. You just can't have a (propositional) sentence without at least a noun and a verb. Within each sentence, we may start introducing nouns. From left-to-right; we may introduce a new noun (which comes with tendrils that extrude outside the sentence bubble for later use to resolve pronominal reference); split a noun so that the same noun-label is used multiply; and label a noun *if it is saturated* – depicted by a solid black circle which includes a copy of the label for bookkeeping purposes when resolving references.



Nouns require verbs in order to be saturated. From left-to-right; if there is precisely one unlabelled noun, we may introduce an unlabelled intransitive verb and saturate the noun so that it is now ready to grow a label; or if there are two unlabelled nouns, we may introduce an unlabelled transitive verb on the surface and saturate the two nouns that will be subject and object; and verbs may be labelled.

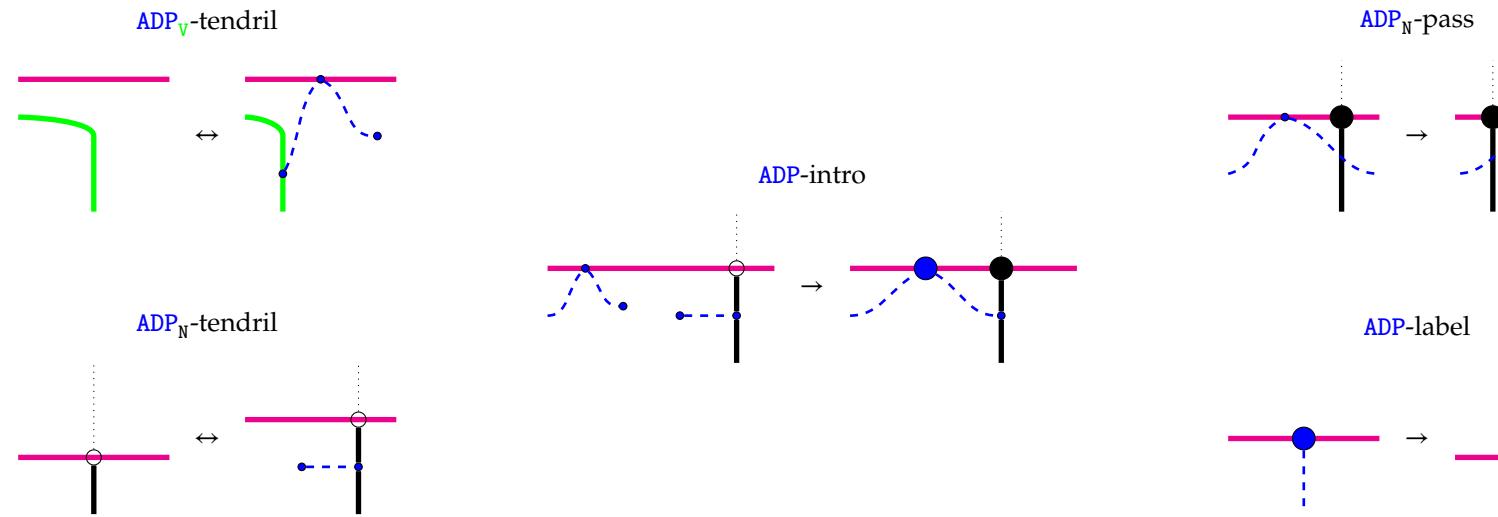


3.8.5 Modifiers

Modifiers are optional parts of sentences that modify (and hence depend on there being) nouns and verbs. We consider adjectives, adverbs, and adpositions. From left to right; we allow adjectives to sprout immediately before a saturated noun, and we allow adverbs to sprout immediately before any verb.



Adpositions modify verbs by tying in an additional noun argument; e.g. while *runs* is intransitive, *runs towards* behaves as a transitive verb. Some more advanced technology is required to place adpositions and their thematic nouns in the correct linear order on the surface. In the left column; an adposition tendril can sprout from a verb via an unsaturated adposition, seeking an unsaturated noun to the right; an unsaturated noun can sprout an tendril seeking a verb to connect to on the left. Both of these rewrites are bidirectional, as tendrils might attempt connection but fail, and so be retracted. In the centre, when an unsaturated adposition and its tendril find an unsaturated noun, they may connect, saturating the adposition so that it is ready to label. In the right column; an unsaturated adposition may move past a saturated noun in the same sentence, which allows multiple adpositions for the same verb; finally, a saturated adposition can be labelled.



3.8.6 Rewriting to circuit-form

RESOLVING REFERENCES

CONNECTING CIRCUITS

Example 3.8.4.

3.8.7 Extensions I: relative and reflexive pronouns

SUBJECT RELATIVE PRONOUNS

Example 3.8.5.

OBJECT RELATIVE PRONOUNS

Example 3.8.6.

REFLEXIVE PRONOUNS

Example 3.8.7.

3.8.8 *Extensions II: grammar equations*

ATTRIBUTIVE VS. PREDICATIVE MODIFIERS

Example 3.8.8.

COPULAS

Example 3.8.9.

POSSESSIVE PRONOUNS

Example 3.8.10.

3.8.9 *Extensions III: higher-order modifiers*

INTENSIFIERS

Example 3.8.11.

COMPARATIVES

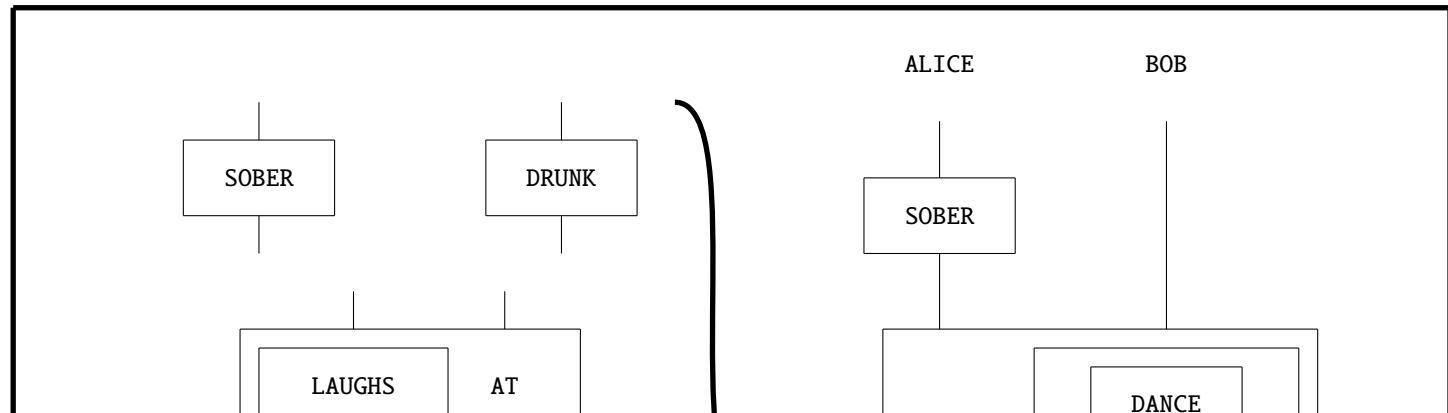
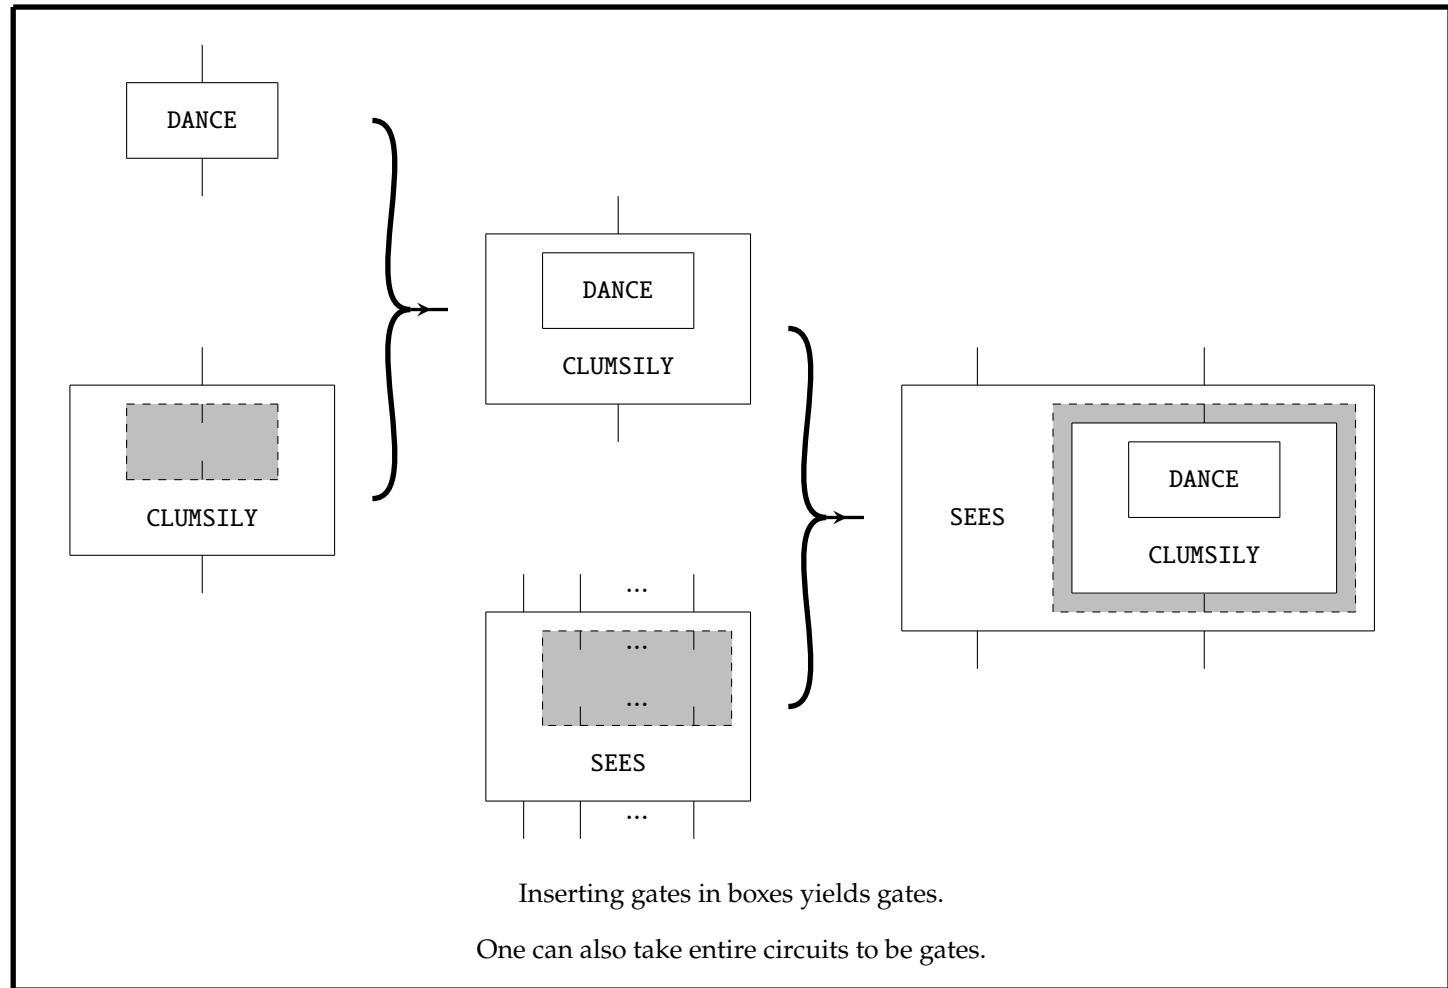
Example 3.8.12.

3.8.10 *Equivalence to internal wirings*

3.8.11 *Text circuit theorem*

3.8.12 *Related work*

Figure 3.38: Generating text circuits directly.



4

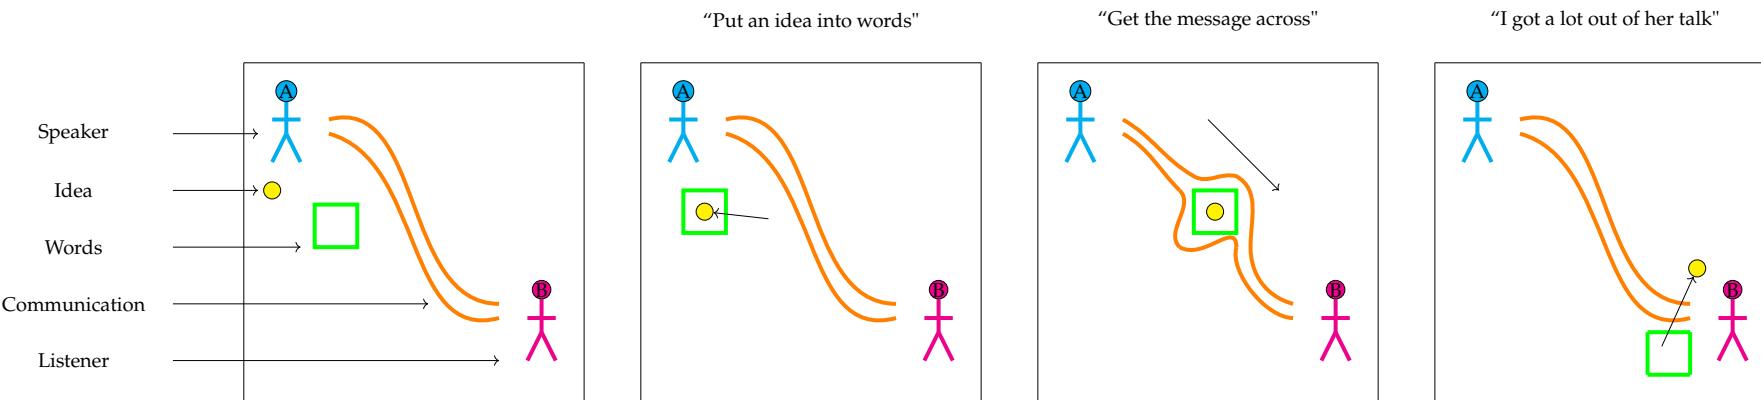
Continuous relations: a palette for toy models

4.1 Continuous Relations: A concept-compliant setting for text circuits

In this chapter, we introduce *continuous relations*, which are a naive extension of the category **Top** of topological spaces and continuous functions towards continuous relations. We choose this category (as opposed to plain **Rel**, the category of sets and relations) because it satisfies several requirements arrived at by introspection of some of the demands of modelling language. These justifications involve basic reflections upon language use, and the consequent mathematical constraints those affordances impose on any interpretation of text circuits in a symmetric monoidal category; *A priori* it could well be that there is no non-trivial process theory that satisfies these constraints, so the onus is on us to show that there exists such a process theory. I outline these justifications – mostly intended for readers interested in how any of this relates to the cognitive aspect of text – after this subsection. The justifications can also be skipped and optionally revisited after the reader is more familiar with what **ContRel** is.

Second, we introduce **ContRel** diagrammatically. In the appendix for this chapter we do the bookwork demonstrating that it is a symmetric monoidal category, and we relate it to the well studied categories **Rel** and **Loc**. To the best of my knowledge, the study of this category is a novel contribution, for reasons I list prefacing the bookwork.

Third, once we have defined a stage to perform calculations in **ContRel**, our aim is to introduce some actors. We seek to make formal the kinds of informal schemata we might doodle on paper to animate various processes occurring in space. One good reason for doing this is to establish formal foundations for the semantics of metaphor, some of the most commonly used of which involve spatial processes in a way that is fundamentally topological []. For example, in the *conduit metaphor* [], words are considered *containers* for ideas, and communication is considered a *conduit* along which those containers are sent.



If you are already happy to treat such doodles as formal, then I think you're alright, you can skip the rest of this chapter. If you are a category theorist or just curious, hello, how are you, please do write me to share your thoughts if you care to, and please skip the rest of this paragraph. If you are still reading, I assume you are some kind of smelly epistemic-paranoiac Bourbaki-thrall sets-and-lambdas math-phallus-worshipping truth-condition-blinded symbol-pusher who takes things too seriously. I hope you choke on the math. I will begin the intimidation immediately.

We provide a generalisation (Definition 4.5.5) of special commutative frobenius algebras in **ContRel** that cohere with idempotents in the category. The relation of (\dagger)-special commutative algebras in **FdHilb** to model observables in quantum mechanics is well-studied [], as is the role of idempotents in generalisations of quantum logic to arbitrary categories [], therefore this generalisation may be viewed as a unification of these ideas to define doodles in **ContRel**. N.B. we are not quite taking the Karoubi envelope of **ContRel**, as we are restricting the idempotents we wish to consider to only those that also behave appropriately as observables. The reason

.8 which provides a diagrammatic characterisation result that allows us to precisely identify when any idempotent in the category. The discrete topology thus acts as a set of labels, where the (pre)images of each element under section and retract behave as expected. As an interlude, we demonstrate how we may construct families of such idempotent-coherent special commutative frobenius generalisation (c.f. the monoidal computer framework in []) of **FinRel** equipped with a Turing object [], thus satisfying Justification configuration spaces of collections of shapes up to rigid displacement, and we develop a relational analogue of homotopy to model choice, along with appropriate extensions to accommodate nonrigid phenomena. If you are still reading and not already a category theorist, then jargon now, because you are probably either questioning or deeply committed to your false ideals, both of which I respect, but I do everything diagrammatically as much as possible.

out there?

SYNTACTIC MONOIDAL CATEGORY. We want to work process-theoretically as much as possible, because, as we have seen, this potentially allows us to set the stage for setting to generalise to other process theories. Also, we want an excuse to build up and play with a symmetric monoidal category and see what work is involved in doing so.

SYNTACTIC CONCEPTS. Second, we have some cognitive considerations: how do we move between concepts – however they are represented and manipulated? Here we sidestep the debate around what concepts actually *are*, aligning ourselves close to Gärdenfors: we assume that concepts are entities that organise concepts of similar domain – such as colour, taste, motion – and that regions of these spaces correspond to concepts. This view is empirically supported by empirical data [], but even if he is wrong, he is at least interestingly so for our purposes. As an example, the classic example of a colour space has been studied and implemented: there are many different embeddings of the space of visible colours in Euclidean space []. In this setting, we can imagine the action of categorising a particular point in colour space as blue by checking to see whether that point falls within the region in the space that is associated with. So we find that this view is conducive to modelling concepts as spatial entities – a very permissive and expressive way to calculate interesting things – whilst also having the ability to handle them using symbolic labels. The question remains: how do we encode concepts in space and sets of labels? This leads us to the following consideration: In a category for concept-compliant text-circuits \mathcal{T} , any split idempotent must possess a split idempotent through a set of concept-labels, thus encoding the association between concepts-quaspace and concepts-

Since we have a stock of symbols referring to entities in space, we can start talking about pairs of entities (e.g. red or blue), sub-colours, arbitrary relations between the set of entities in one space and entities of another (e.g. how the colour of a banana relates to the taste). Given enough time and patience, we can linguistically construct – at least – any finite relation between finite sets. As we challenge the reader to consider that every time we define a new concept in this way, we can again treat it as a symbolic concept, that is, label it with symbols that travel along wires in our framework, we are asking that any putative process theory that satisfies these considerations must be able to encode a concept, some wire Ξ for nouns, such that all finite relations fit into it. This leads us to the following consideration: A category for text-circuits must be able to have an object Ξ such that all of **FinRel** can be encoded within Ξ somehow. But we already know how to encode using split idempotents. We can combine the two considerations of the last two paragraphs into one requirement. In prose; the first item gives us a method within the category to encode concepts-quaspace; the second item gives us a noun-wire in \mathcal{T} that permits us to encode and manipulate concepts however we like.

we would like to treat finite relations.

Requirement 4.1.1. We call a text circuit category \mathcal{T} *concept-compliant* if:

1. Any wire Γ used to model a conceptual space possesses a split idempotent through a set of concept-symbols \mathfrak{L}
2. Anything one can do with sets of concept-symbols in **FinRel** must be doable in \mathcal{T} ; in particular, there exists a noun-wire Ξ such that **FinRel** embeds as a category into the subcategory of \mathcal{T} generated by the split idempotents on Ξ

JUSTIFICATION 3: WE THINK TOPOLOGY IS A GOOD SETTING TO MODEL CONCEPTS SPATIALLY.

Where we differ from Gärdenfors is that we only ask for topological spaces, rather than his stronger requirements for metric spaces and convex concepts, so we are following the spirit but not the letter. There are several reasons for this choice. First, topology is a primitive mathematical framework for space; all metric spaces are topological spaces, but not vice versa, so this is a conservative generalisation of Gärdenfors. Second, there are technical reasons that **ContRel** is desirable. For example, we can process-theoretically characterise continuous maps from the unit interval fairly easily to model things going on in space and time; without topology around, for instance in the setting of just **Rel**, I do not know if it is even possible to pick out the continuous maps from all the others purely process-theoretically. Third and perhaps most interestingly, there are also good reasons to think that this is the right way to think about conceptual spaces. We can view topology as a framework for conceptual spaces where we consider the open sets of a topology to be the concepts. Éscardo provides a the following correspondence [], which we extend with "Point" and "Subset", and an additional column for interpretation as conceptual space:

Topology	Type Theory	Conceptual Spaces
Space	Type	Conceptual space
Point	Element of set	Copyable instance
Subset	Subset	Instance
Open set	Semi-decidable set	Concept
Closed set	Set with semi-decidable complement	-
Clopen set	Decidable set	A concept the negation of which is also a concept
Discrete topology	Type with decidable equality	A conceptual space where any collection of instances forms a concept
Hausdorff topology	Type with semi-decidable inequality of elements	A conceptual space where any pair of distinct instances can be described as belonging to two disjoint
Compact set	Exhaustively searchable set, in a finite number of steps	A conceptual space \mathfrak{C} such that for any joint concept R on $\mathfrak{C} \times \mathfrak{D}$, $\forall c \in \mathfrak{C} R(c, -)$ is a concept in

ContRel reflects the above correspondences diagrammatically. Open sets are precisely tests, and it is always possible to construct finite intersections of opens using copy-relations. Modelling concepts as open sets or tests aligns them with semi-decidability. Given any state-instance, we can test whether it overlaps with a concept graphically: success returns a unit scalar, failure returns the zero scalar. Moreover, another independent diagrammatic calculus for concepts by Tull [] agrees with ours; concepts are effects, copy maps take intersections of concepts, and so on. Tull's diagrams are interpreted in **Stoch**, the category of stochastic processes, and as a whole his design philosophy closely adheres to Gärdenfors. All this is to suggest that if you take Gärdenfors seriously, then there is something worth taking seriously about modelling concepts as effects in a monoidal category with copy-maps. However, taking diagrammatic conceptual spaces seriously also yields a no-go result for topological spaces – which includes Gärdenfors' metric spaces with convex concepts: you can only do first-order logic with equality on your concepts if your base space is discrete and finite. We explain this below.

The correspondence between spaces and type-theory extends to conceptual spaces as follows: "niceness conditions on your conceptual space correspond to the ability to form new concepts using logical operations." For example, this means that if we denote colourspace with \mathfrak{C} , we can only construct a concept different

g a Hausdorff space, such as Euclidean space. If we want to model not X as the complement of a colour X in colourspace, ask-
 res \mathbb{C} be locally indiscrete – i.e. every open set is also closed; Euclidean space is not locally indiscrete, so we cannot use set-
 use something else, like the interior of the complement. If we want to have access to universal quantifiers in colourspace, so that
 ch as the taste that apples of all colours have in common, then we require \mathbb{C} to be compact, which Euclidean space is
 Here then is the conflict: if we take modelling concepts with spaces seriously, and we also care to do logic with concepts, there
 make equality as a concept same colour on $\mathbb{C} \times \mathbb{C}$ requires that \mathbb{C} have discrete topology, but discrete topologies on infinite sets are
 universal quantification at the same time. Conversely, if you want universal quantification on colourspace, then you can at best
 s as a concept, never exact. Now we can summarise this tension. All topological spaces, including Gärdenfors conceptual spaces
 start off with regular logic – \exists, \wedge, T – for free []. In order to obtain first-order logic with equality on the points of the space, the un-
 support universal quantification) and discrete (to support equality of points), and the latter condition implies locally indiscrete
 s with the discrete topology are compact; you can only do first-order logic with equality on your concepts if your base space is

perspective of the ancient observation that logical thought really does seem symbolic. It is also important to note that the essential
 went this no-go; in our language, asking for split idempotents through (finite) discrete topologies is precisely what allows logical
 esentations of concepts, when those split idempotents exist. What this no-go does mean is that nobody has to waste their time
 ceptual spaces and first-order logic-with-equality where every logical operator is viewed as a space-preserving transformation
 ebehaves properly on the space of points; there is only "good enough".

nt a symmetric monoidal category to keep the prospect of general, process-theoretic reasoning. We want the objects of this cate-
 we want to model conceptual spaces and calculate interesting things. While the category **Top** of topological spaces and contin-
 monoidal with respect to categorical product, for linguistic and concept-related considerations we want finite relations to come
 use **Top** is cartesian monoidal it doesn't work for our purposes. So we construct **ContRel**.

4.2 Continuous Relations

Reminder 4.2.1 (Topological Space). A *topological space* is a pair (X, τ) , where X is a set, and $\tau \subset \mathcal{P}(X)$ are the *open sets* of X , such that:

"nothing" and "everything" are open

$$\emptyset, X \in \tau$$

Arbitrary unions of opens are open

$$\{U_i : i \in I\} \subseteq \tau \Rightarrow \bigcup_{i \in I} U_i \in \tau$$

Finite intersections of opens are open $n \in \mathbb{N}$:

$$U_1, \dots, U_n \in \tau \Rightarrow \bigcap_{1 \dots, i, \dots, n} U_i \in \tau$$

Reminder 4.2.2 (Relational Converse). Recall that a relation $R : S \rightarrow T$ is a subset $R \subseteq S \times T$.

$$R^\dagger : T \rightarrow S := \{(t, s) : (s, t) \in R\}$$

Reminder 4.2.3 (Continuous function). A function between sets $f : X \rightarrow Y$ is a continuous function between topologies $f : (X, \tau) \rightarrow (Y, \sigma)$ if

$$U \in \sigma \Rightarrow f^{-1}(U) \in \tau$$

where f^{-1} denotes the inverse image.

Reminder 4.2.6 (Product Topology). We denote the product topology of X^τ and Y^σ as $(X \times Y)^{(\tau \times \sigma)}$. $\tau \times \sigma$ is the topology on $X \times Y$ generated by the basis $\{t \times s : t \in \mathfrak{b}_\tau, s \in \mathfrak{b}_\sigma\}$, where \mathfrak{b}_τ and \mathfrak{b}_σ are bases for τ and σ respectively.

Reminder 4.2.7 (Product of relations). For relations between sets $R : X \rightarrow Y, S : A \rightarrow B$, the product relation $R \times S : X \times A \rightarrow Y \times B$ is defined to be

$$\{((x, a), (y, b)) : (x, y) \in R, (a, b) \in S\}$$

TO THE BEST OF MY KNOWLEDGE, THE STUDY OF **ContRel** IS A NOVEL CONTRIBUTION. I VENTURE TWO POTENTIAL REASONS.

FIRST, IT IS BECAUSE AND NOT DESPITE OF THE NAÏVITY OF THE CONSTRUCTION. Usually, the relationship between **Rel** and **Set** is often understood in sophisticated general methods which are inappropriate in different ways. I have tried applying Kriesli machinery which generalises to "relationification" of arbitrary categories via appropriate analogs of the powerset monad to relate **Top** and **ContRel**, but it is not evident to me whether there is such a monad. The view of relations as spans of maps in the base category should work, since **Top** has pullbacks, but this makes calculation difficult and especially cumbersome when monoidal structure is involved. The naïve approach I take is to observe that the preimages of functions are precisely relational converses when functions are viewed as relations, so the preimage-preserves-opens condition that defines continuous functions directly translates to the relational case.

SECOND, THE RELATIONAL NATURE OF **ContRel** MEANS THAT THE CATEGORY HAS POOR EXACTNESS PROPERTIES. Even if the sophisticated machinery mentioned in the first reason do manage to work, relational variants of **Top** are poor candidates for any kind of serious mathematics because they lack many limits and colimits. Since we take an entirely "monoidal" approach – a relative newcomer in terms of mathematical technique – we are able to find and make use of the rich structure of **ContRel** with a different toolkit.

In the end, we want to formalise doodles, so perhaps there is some virtue in proceeding by elementary means.

Recall that functions are relations, and the inverse image used in the definition of continuous maps is equivalent to the relational converse when functions are viewed as relations. So we can naively extend the notion of continuous maps to continuous relations between topological spaces.

Definition 4.2.4 (Continuous Relation). A continuous relation $R : (X, \tau) \rightarrow (Y, \sigma)$ is a relation $R : X \rightarrow Y$ such that

$$U \in \sigma \Rightarrow R^\dagger(U) \in \tau$$

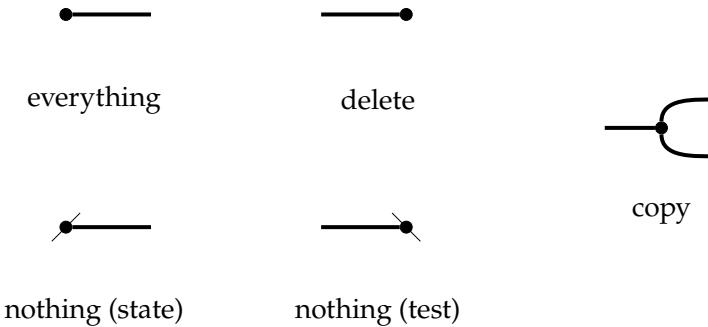
where \dagger denotes the relational converse.

Notation 4.2.5. For shorthand, we denote the topology (X, τ) as X^τ . As special cases, we denote the discrete topology on X as X^* , and the indiscrete topology X° .

The symmetric monoidal structure is that of product topologies on objects, and products of relations on morphisms.

inuous

NS FOR ANY X^τ :

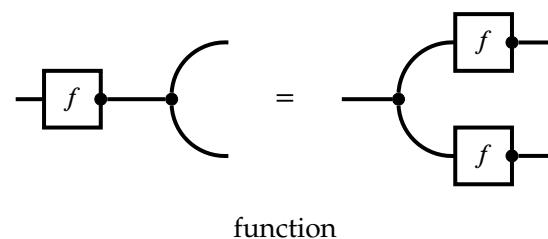
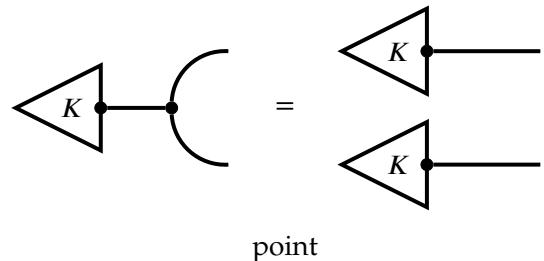


NG EQUALITIES:

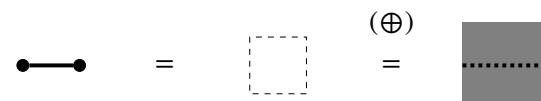
The diagram shows two equality relations:

- cocommutativity**: A string diagram showing a horizontal line with a dot at its left end, which splits into two vertical lines that then rejoin at a crossing point before continuing as a single horizontal line with a dot at its right end.
- counitality**: A string diagram showing a horizontal line with a dot at its left end, which splits into two vertical lines that then rejoin at a crossing point before continuing as a single horizontal line with a dot at its right end, followed by a second horizontal line with no dots.

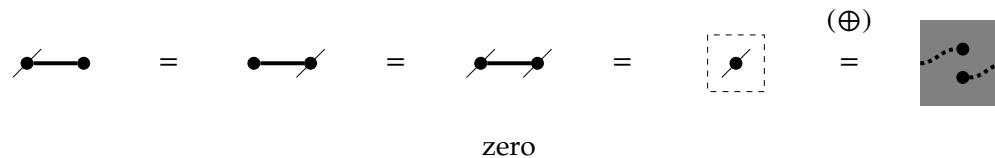
ivity



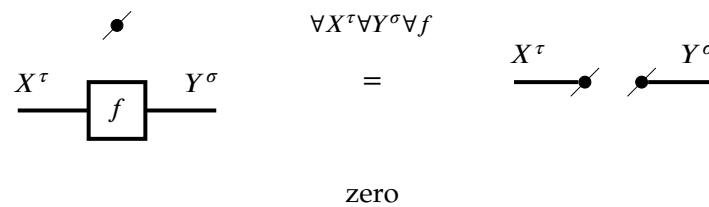
EVERYTHING, DELETE, NOTHING-STATES AND NOTHING-TESTS COMBINE TO GIVE TWO NUMBERS, ONE AND ZERO. There are extra expressions in grey squares above: they anticipate the tape-diagrams we will later use to graphically express another monoidal product of **ContRel**, the direct sum \oplus .



(b)one



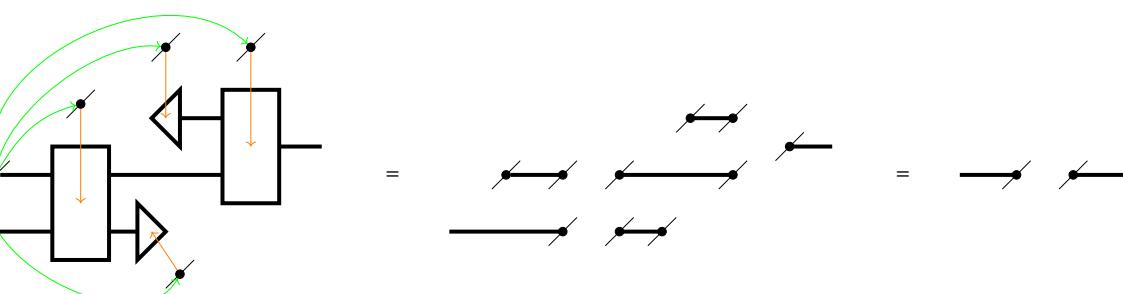
ZERO SCALARS TURN ENTIRE DIAGRAMS INTO ZERO MORPHISMS. There is a zero-morphism for every input-output pair of objects in **ContRel**.



case in **Rel**, **Vect**, and any category with biproducts and zero-objects []) is that the zero scalar is an absorbing element of multiplication: multiplying by a zero scalar sends it to its corresponding zero-morphism.

$$\begin{array}{ccccccc}
 \text{Diagram 1:} & \text{Diagram 2:} & \text{Diagram 3:} & \text{Diagram 4:} & \text{Diagram 5:} & \text{Diagram 6:} & \text{Diagram 7:} \\
 \text{Diagram 8:} & = & \text{Diagram 9:} & \text{Diagram 10:} & = & \text{Diagram 11:} & \text{Diagram 12:} \\
 \text{Diagram 13:} & = & \text{Diagram 14:} & \text{Diagram 15:} & = & \text{Diagram 16:} & \text{Diagram 17:}
 \end{array}$$

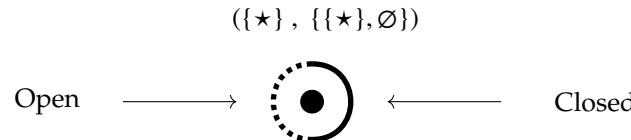
ers in a diagram, it spawns zero scalars which infect all other processes, turning them all into zero-processes. The same holds for *Take* which takes copies of itself to infect all other processes.



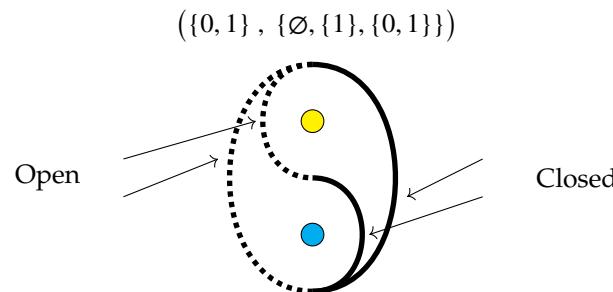
4.4 Continuous Relations by examples

Let's consider three topological spaces and examine the continuous relations between them. This way we can build up intuitions, and prove some tool results in the process.

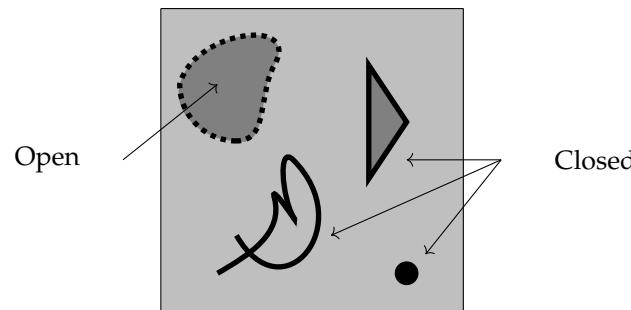
The **singleton space** consists of a single point which is both open and closed. We denote this space \bullet . Concretely, the underlying set and topology is



The **Sierpiński space** consists of two points, one of which (in yellow) is open, and the other (in cyan) is closed. We denote this space \mathcal{S} . Concretely, the underlying set and topology is:



The **unit square** has $[0, 1] \times [0, 1]$ as its underlying set. Open sets are "blobs" painted with open balls. Points, lines, and bounded shapes are closed. We denote this space \blacksquare .



- → •: There are two relations from the singleton to the singleton; the identity relation $\{(\bullet, \bullet)\}$, and the empty relation \emptyset . Both are topological.
- → \mathcal{S} : There are four relations from the singleton to the Sierpiński space, corresponding to the subsets of \mathcal{S} . All of them are topological.

$S \rightarrow \bullet$: There are four candidate relations from the Sierpiński space to the singleton, but as we see in Example 4.4.1, not all of them are topological.

NOW WE NEED SOME ABSTRACTION. We cannot study the continuous relations between the singleton and the unit square case by case. We discover that continuous relations out of the singleton indicate arbitrary subsets, and that continuous relations into the singleton indicate arbitrary opens.

$\bullet \rightarrow \blacksquare$: Proposition 4.4.3 tells us that there are as many continuous relations from the singleton to the unit square as there are subsets of the unit square.

$\blacksquare \rightarrow \bullet$: Proposition 4.4.4 tells us that there are as many continuous relations from the unit square to the singleton as there are open sets of the unit square.

THERE ARE 16 CANDIDATE RELATIONS $S \rightarrow S$ TO CHECK. A case-by-case approach won't scale, so we could instead identify the building blocks of continuous relations with the same source and target space.

WHICH RELATIONS $X^\tau \rightarrow Y^\sigma$ ARE ALWAYS CONTINUOUS?

THE EMPTY RELATION IS ALWAYS CONTINUOUS.

Proposition 4.4.6. *Proof.* The preimage of the empty relation is always \emptyset , which is open by definition. \square

FULL RELATIONS ARE ALWAYS CONTINUOUS

Proposition 4.4.8. *Proof.* The preimage of any subset of Y – open or not – under the full relation is the whole of X , which is open by definition. \square

FULL RELATIONS RESTRICTED TO OPEN SETS IN THE DOMAIN ARE CONTINUOUS.

Proposition 4.4.9. Given an open $U \subseteq X^\tau$, and an arbitrary subset $K \subset Y^\sigma$, the relation $U \times K \subseteq X \times Y$ is open.

Proof. Consider an arbitrary open set $V \in \sigma$. Either V and K are disjoint, or they overlap. If they are disjoint, the preimage of V is \emptyset , which is open. If they overlap, the preimage of V is U , which is open. \square

CONTINUOUS FUNCTIONS ARE ALWAYS CONTINUOUS.

Proposition 4.4.10. If $f : X^\tau \rightarrow Y^\sigma$ is a continuous function, then it is also a continuous relation.

Example 4.4.1 (A noncontinuous relation). The relation $\{(0, \bullet)\} \subset S \times \bullet$ is not a continuous relation: the preimage of the open set $\{\bullet\}$ under this relation is the non-open set $\{0\}$.

Terminology 4.4.2. Call a continuous relation $\bullet \rightarrow X^\tau$ a **state** of X^τ , and a continuous relation $X^\tau \rightarrow \bullet$ a **test** of X^τ .

Proposition 4.4.3. States $R : \bullet \rightarrow X^\tau$ correspond with subsets of X .

Proof. The preimage $R^\dagger(U)$ of a (non- \emptyset) open $U \in \tau$ is \star if $R(\star) \cap U$ is nonempty, and \emptyset otherwise. Both \star and \emptyset are open in $\{\star\}^*$. $R(\star)$ is free to specify any non- \emptyset subset of X . The empty relation handles \emptyset as an open of X^τ . \square

Proposition 4.4.4. Tests $R : X^\tau \rightarrow \bullet$ correspond with open sets $U \in \tau$.

Proof. The preimage $R^\dagger(\star)$ of \star must be an open set of X^τ by definition ???. $R^\dagger(\star)$ is free to specify any open set of X^τ . \square

Reminder 4.4.5 (Empty relation). The **empty relation** $X \rightarrow Y$ relates nothing. It is defined:

$$\emptyset \subset X \times Y$$

Reminder 4.4.7 (Full relation). The **full relation** $X \rightarrow Y$ relates everything to everything. It is all of $X \times Y$.

Proof. Functions are special cases of relations. The relational converse of a function viewed in this way is the same thing as the preimage. \square

Reminder 4.4.11 (Identity relation). The **identity relation** $X \rightarrow X$ relates anything to itself. It is defined:

$$\{(x, x) : x \in X\} \subseteq X \times X$$

Reminder 4.4.13 (Union, intersection, and ordering of relations). Recall that relations $X \rightarrow Y$ can be viewed as subsets of $X \times Y$. So it makes sense to speak of the union and intersection of relations, and of partially ordering them by inclusion.

Reminder 4.4.16 (Topological Basis). $\mathfrak{b} \subseteq \tau$ is a basis of the topology τ if every $U \in \tau$ is expressible as a union of elements of \mathfrak{b} . Every topology has a basis (itself).

Minimal bases are not necessarily unique.

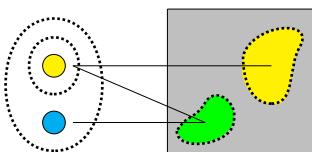


Figure 4.1: Regions of ■ in the image of the yellow point alone will be coloured yellow, and regions in the image of both yellow and cyan will be coloured green:

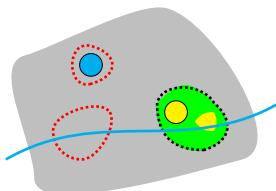


Figure 4.2: Regions in the image of the cyan point alone cannot be open sets by continuity, so they are either points or lines. Points and lines in cyan must be surrounded by an open region in either yellow or green, or else we violate continuity (open sets in red).

THE IDENTITY RELATION IS ALWAYS CONTINUOUS. The identity relation is also the "trivial" continuous map from a space to itself, so this also follows from Proposition 4.4.10.

Proposition 4.4.12. *Proof.* The preimage of any open set under the identity relation is itself, which is open by assumption. \square

GIVEN TWO CONTINUOUS RELATIONS $R, S : X^\tau \rightarrow Y^\sigma$, HOW CAN WE COMBINE THEM?

Proposition 4.4.14. If $R, S : X^\tau \rightarrow Y^\sigma$ are continuous relations, so are $R \cap S$ and $R \cup S$.

Proof. Replace \square with either \cup or \cap . For any non- \emptyset open $U \in \sigma$:

$$(R \square S)^\dagger(U) = R^\dagger(U) \square S^\dagger(U)$$

As R, S are continuous relations, $R^\dagger(U), S^\dagger(U) \in \tau$, so $R^\dagger(U) \square S^\dagger(U) = (R \square S)^\dagger(U) \in \tau$. Thus $R \square S$ is also a continuous relation. \square

Corollary 4.4.15. Continuous relations $X^\tau \rightarrow Y^\sigma$ are closed under arbitrary union and finite intersection. Hence, continuous relations $X^\tau \rightarrow Y^\sigma$ form a topological space where each continuous relation is an open set on the base space $X \times Y$, where the full relation $X \rightarrow Y$ is "everything", and the empty relation is "nothing".

A TOPOLOGICAL BASIS FOR SPACES OF CONTINUOUS RELATIONS

Definition 4.4.17 (Partial Functions). A **partial function** $X \rightarrow Y$ is a relation for which each $x \in X$ has at most a single element in its image. In particular, all functions are special cases of partial functions, as is the empty relation.

Lemma 4.4.18 (Partial functions are a \cap -ideal). The intersection $f \cap R$ of a partial function $f : X \rightarrow Y$ with any other relation $R : X \rightarrow Y$ is again a partial function.

Proof. Consider an arbitrary $x \in X$. $R(x) \cap f(x) \subseteq f(x)$, so the image of x under $f \cap R$ contains at most one element, since $f(x)$ contains at most one element. \square

Lemma 4.4.19 (Any single edge can be extended to a continuous partial function). Given any $(x, y) \in X \times Y$, there exists a continuous partial function $X^\tau \rightarrow Y^\sigma$ that contains (x, y) .

Proof. Let $\mathcal{N}(x)$ denote some open neighbourhood of x with respect to the topology τ . Then $\{(z, y) : z \in \mathcal{N}(x)\}$ is a continuous partial function that contains (x, y) . \square

Proposition 4.4.20. Continuous partial functions form a topological basis for the space $(X \times Y)^{(\tau \rightsquigarrow \sigma)}$, where the opens are continuous relations $X^\tau \rightarrow Y^\sigma$.

Proof. We will show that every continuous relation $R : X^\tau \rightarrow Y^\sigma$ arises as a union of partial functions. Denote the set of continuous partial functions \mathfrak{f} . We claim that:

$$R = \bigcup_{F \in \mathfrak{f}} (R \cap F)$$

The \supseteq direction is evident, while the \subseteq direction follows from Lemma 4.4.19. By Lemma 4.4.18, every $R \cap F$ term is a partial function, and by Corollary 4.4.15, continuous. \square

$S \rightarrow S$: We can use Proposition 4.4.20 to write out the topological basis of continuous partial functions, from which we can take unions to find all the continuous relations, which we depict in Figure 4.5.

$S \rightarrow \blacksquare$: Now we use the colour convention of the points in S to "paint" continuous relations on the unit square "canvas", as in Figures 4.1 and 4.2. So each continuous relation is a painting, and we can characterise the paintings that correspond to continuous relations $S \rightarrow \blacksquare$ in words as follows: Cyan only in points and lines, and either contained in or at the boundary of yellow or green. Have as much yellow and green as you like.

$\blacksquare \rightarrow S$: The preimage of all of S must be an open set. So the painting cannot have stray lines or points outside of blobs. The preimage of yellow must be open, so the union of yellow and green in the painting cannot have stray lines or points outside of blobs. Point or line gaps within blobs are ok. Each connected blob can contain any colours in any shapes, subject to the constraint that if cyan appears anywhere, then either yellow or green must occur somewhere. Open blobs with no lines or points outside. Yellow and green considered alone is a painting made of blobs with no stray lines or points. If cyan appears anywhere, then either yellow or green have to appear somewhere.

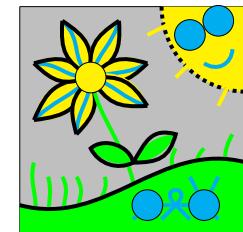


Figure 4.3: A continuous relation $S \rightarrow \blacksquare$: "Flower and critter in a sunny field".

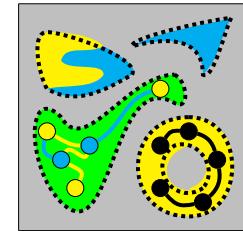


Figure 4.4: A continuous relation $\blacksquare \rightarrow S$: "still math?". Black lines and dots indicate gaps.

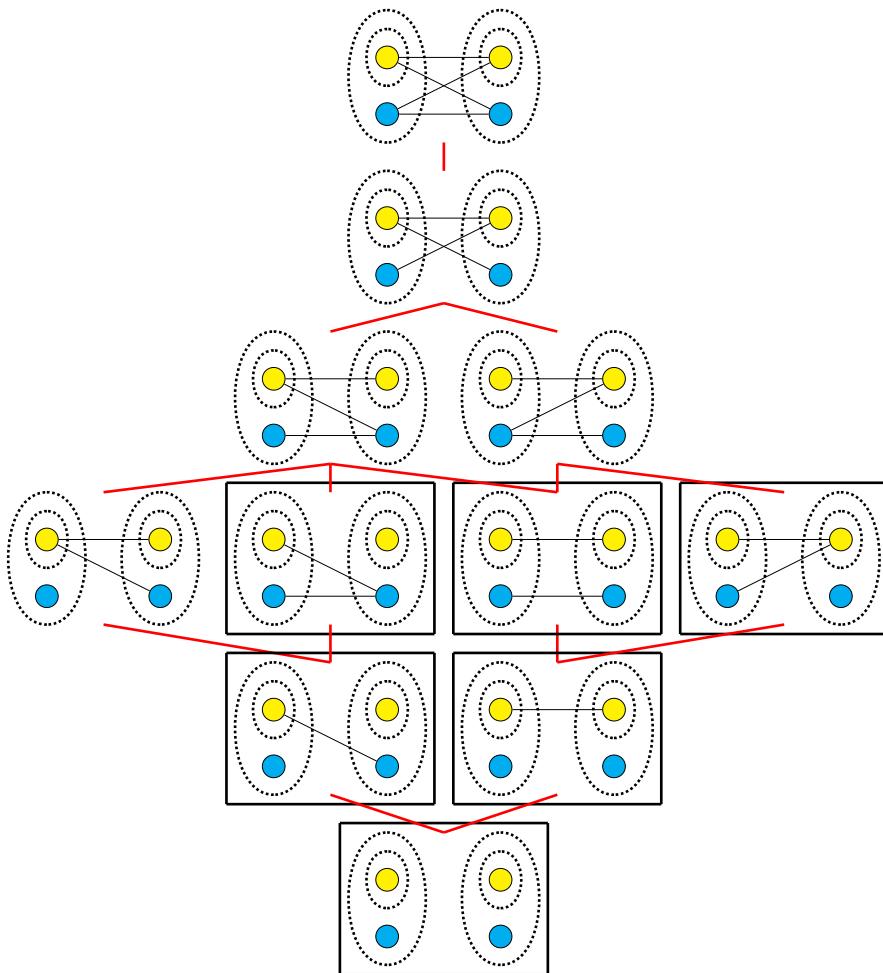


Figure 4.5: Hasse diagram of all continuous relations from the Sierpiński space to itself. Each relation is depicted left to right, and inclusion order is bottom-to-top. Relations that form the topological basis are boxed.

ONE MORE EXAMPLE FOR FUN: $[0, 1] \rightarrow \blacksquare$: We know how continuous functions from the unit line into the unit square look.

THEN WHAT ARE THE PARTIAL CONTINUOUS FUNCTIONS? If we understand these, we can obtain all continuous relations by arbitrary unions of the basis. Observe that the restriction of any continuous function to an open set in the source is a continuous partial function. The open sets of $[0, 1]$ are collections of open intervals, each of which is homeomorphic to $(0, 1)$, which is close enough to $[0, 1]$.

ANY PAINTING IS A CONTINUOUS RELATION $[0, 1] \rightarrow \blacksquare$. By colour-coding $[0, 1]$ and controlling brushstrokes, we can do quite a lot. Now we would like to develop the abstract machinery required to *formally* paint pictures with words.



Figure 4.6: continuous functions $[0, 1] \rightarrow \blacksquare$ follow the naïve notion of continuity: a line one can draw on paper without lifting the pen off the page.

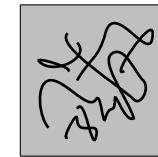


Figure 4.7: So a continuous partial function is "(countably) many (open-ended) lines, each of which one can draw on paper without lifting the pen off the page."

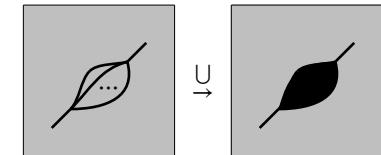


Figure 4.8: We can control the thickness of the brush-stroke, by taking the union of (uncountably) many lines.



Figure 4.9: Assign the visible spectrum of light to $[0, 1]$. Colour open sets according to perceptual addition of light, computing brightness by normalising the measure of the open set.

4.5 Populating space with shapes using sticky spiders



Figure 4.10: Like it or not, a continuous relation $[0, 1] \rightarrow \blacksquare$: "The Starry Night", by Vincent van Gogh.

Reminder 4.5.1 (copy-compare spiders of **Rel**). For a set X , the *copy* map $X \rightarrow X \times X$ is defined:

$$\{(x, (x, x)) : x \in X\}$$

the *compare* map $X \times X \rightarrow X$ is defined:

$$\{((x, x), x) : x \in X\}$$

These two maps are the (co)multiplications of special frobenius algebras. The (co)units are *delete*:

$$\{(x, \star) : x \in X\}$$

and *everything*:

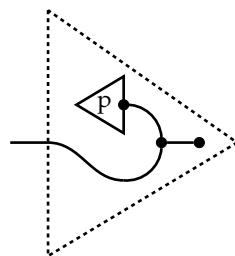
$$\{(\star, x) : x \in X\}$$

4.5.1 When does an object have a spider (or something close to one)?

Example 4.5.2 (The copy-compare spiders of **Rel** are not always continuous). The compare map for the Sierpiński space is not continuous, because the preimage of $\{0, 1\}$ is $\{(0, 0), (1, 1)\}$, which is not open in the product space of S with itself.

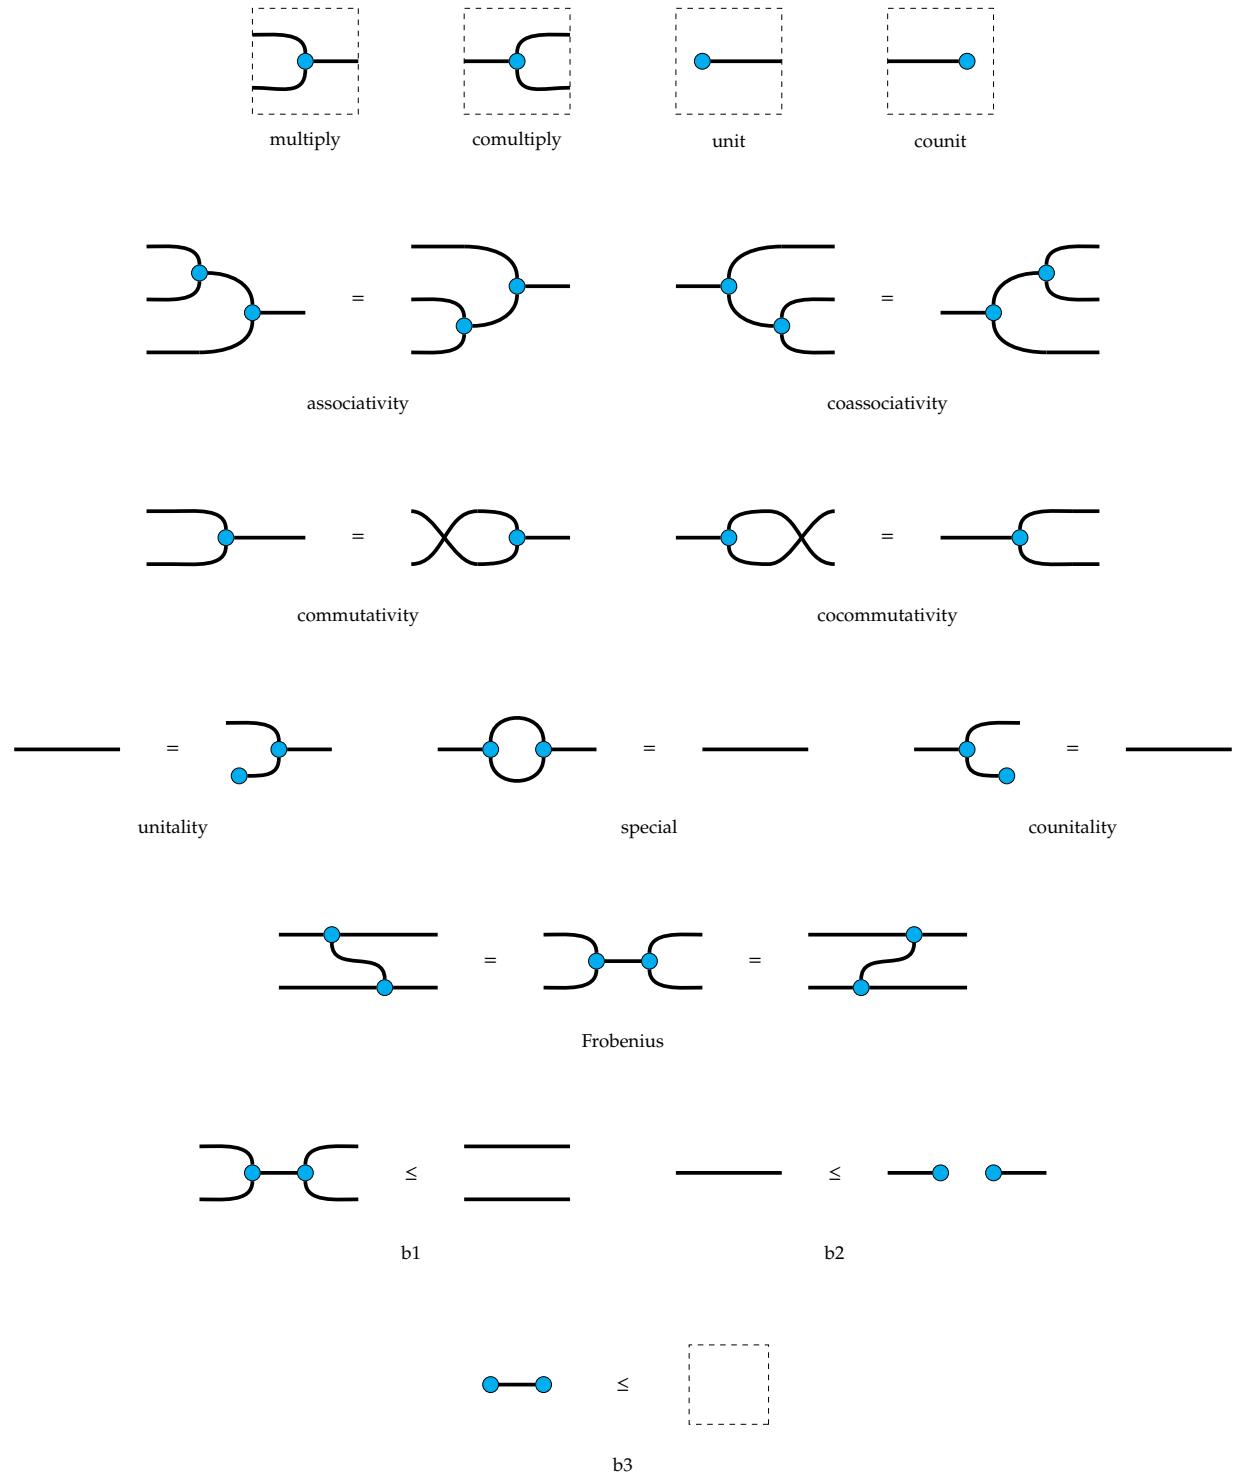
Proposition 4.5.3. The copy map is a spider iff the topology is discrete.

Proof. Discrete topologies inherit the usual copy-compare spiders from **Rel**, so we have to show that when the copy map is a spider, the underlying wire must have a discrete topology. Suppose that some wire has a spider, and construct the following open set using an arbitrary point p :

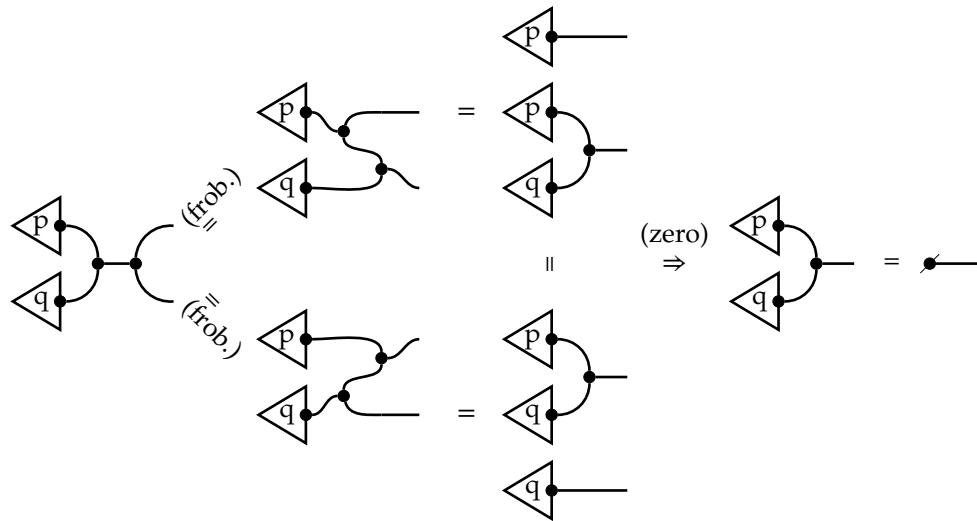


It will suffice to show that this open set is the singleton $\{p\}$ – when all singletons are open, the topology is discrete. As a lemma, using frobenius rules and the property of zero morphisms, we can show that compar-

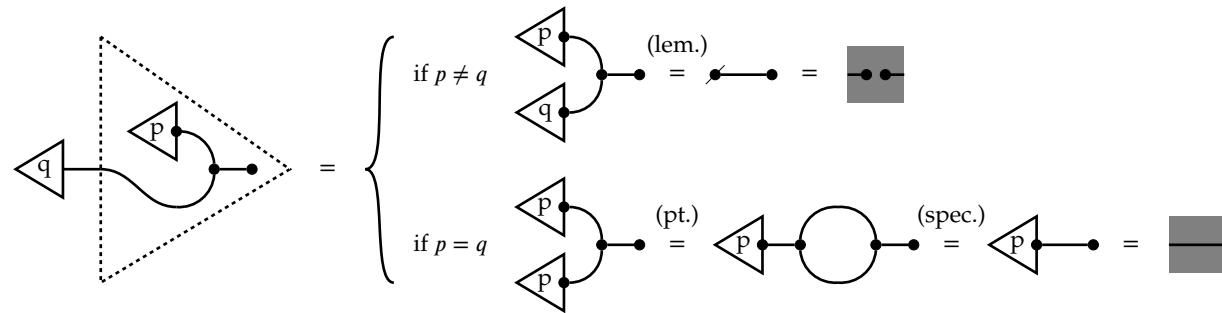
Figure 4.11: The generators (in dashed boxes) and relations that make a spider. When the spider satisfies in addition the three inequalities b1-3, we call it a **relation-spider**.



ing distinct points $p \neq q$ yields the \emptyset state.



The following case analysis shows that our open set only contains the point p .



□

Reminder 4.5.4 (Split idempotents). An **idempotent** in a category is a map $e : A \rightarrow A$ such that

$$A \xrightarrow{e} A \xrightarrow{e} A = A \xrightarrow{e} A$$

A **split idempotent** is an idempotent $e : A \rightarrow A$ along with a **retract** $r : A \rightarrow B$ and a **section** $s : B \rightarrow A$ such that:

$$A \xrightarrow{e} A = A \xrightarrow{r} B \xrightarrow{s} A$$

We can use split idempotents to transform copy-spiders from discrete topologies to almost-spiders on other spaces. We can graphically express the behaviour of a split idempotent e as follows, where the semicircles for the section and retract r, s form a visual pun.

It will be more aesthetic going forward to colour processes and treat the colours as variables instead of labelling them.

The top row shows two equivalent ways to factor a spider s into an identity wire and an idempotent e . The left diagram shows a blue semi-circular spider s with inputs X^* and Y^σ , and output Y^σ . This is followed by an equals sign. The middle diagram shows a blue circle with a dot labeled e and two wires labeled Y^σ meeting at it. This is followed by another equals sign. The right diagram shows two blue circles with dots labeled e connected by a wire, with two wires labeled Y^σ meeting at each dot.

The bottom row shows a sticky spider r with input Y^σ and output X^* . This is followed by an equals sign. The right side of the equation is a single horizontal line, representing the identity wire X^* .

A **sticky spider** (or just an e -spider, if we know that e is a split idempotent), is a spider *except* every identity wire on any side of an the idempotent e .

This diagram illustrates the **e-(co)unitality** property. It shows three equivalent configurations of a loop with a dot:

- Left: A loop with two incoming/outgoing wires, one from the top-left and one from the bottom-left, both meeting at a dot.
- Middle: A single horizontal wire with a dot in the center.
- Right: A loop with two incoming/outgoing wires, one from the top-right and one from the bottom-right, both meeting at a dot.

e-(co)unitality

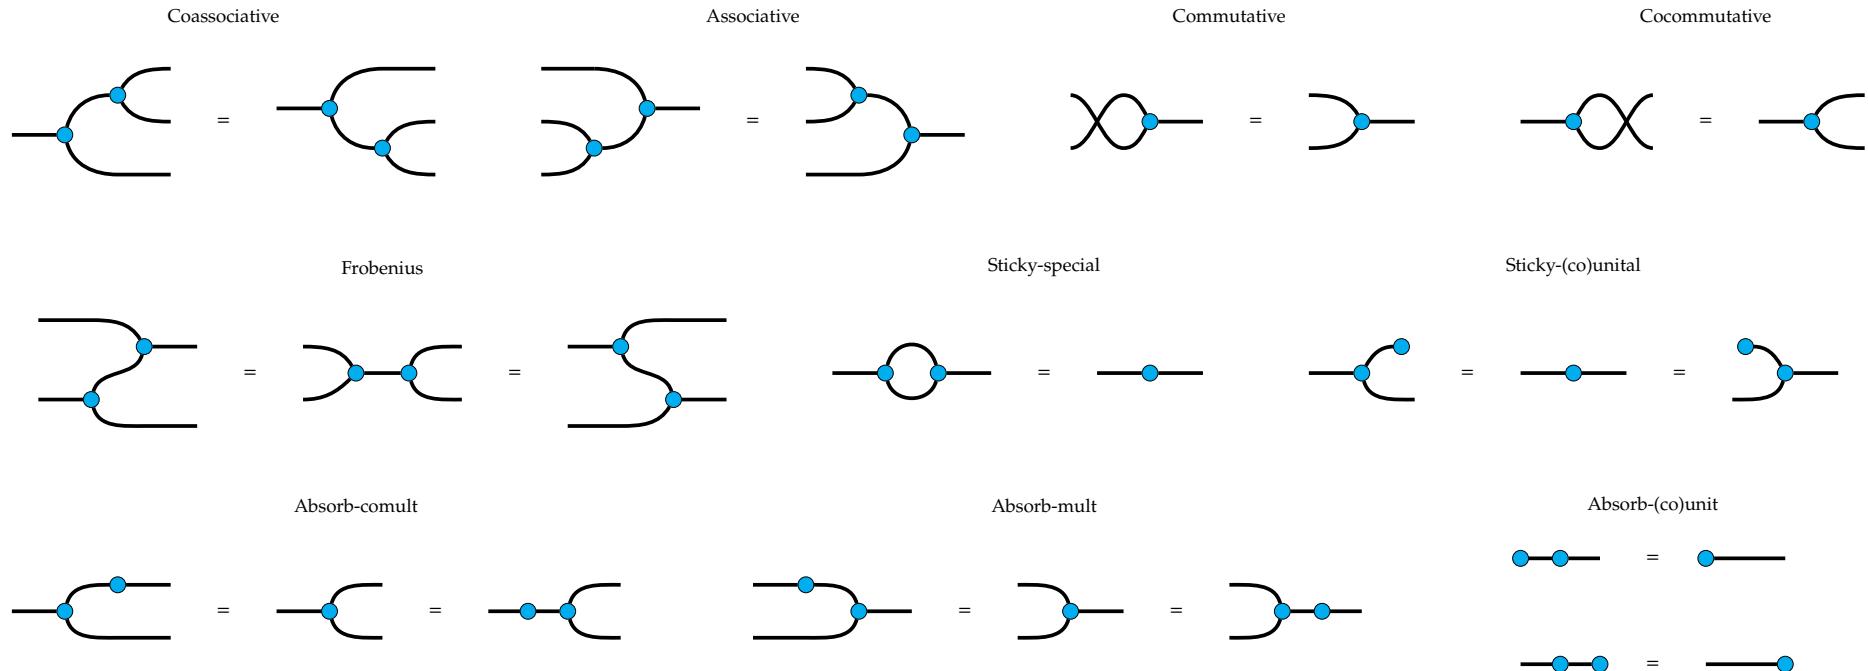
This diagram illustrates the **e-special** property. It shows two equivalent configurations of a loop with two dots:

- Left: A loop with two dots at opposite ends, one at the top-left and one at the bottom-left.
- Middle: A single horizontal wire with two dots at its ends, one at the top and one at the bottom.

e-special

A sticky spider is that one can still coalesce all connected spider-bodies together, but the 1-1 spider "sticks around" rather than achieved by the following rules that cohere the idempotent e with the (co)unit and (co)multiplications; they are the same as the frobenius algebra with two exceptions. First, where an identity wire appears in an equation, we replace it with an idempotent. components freely emit and absorb idempotents. By these rules, the usual proof [] for the normal form of spiders follows, except

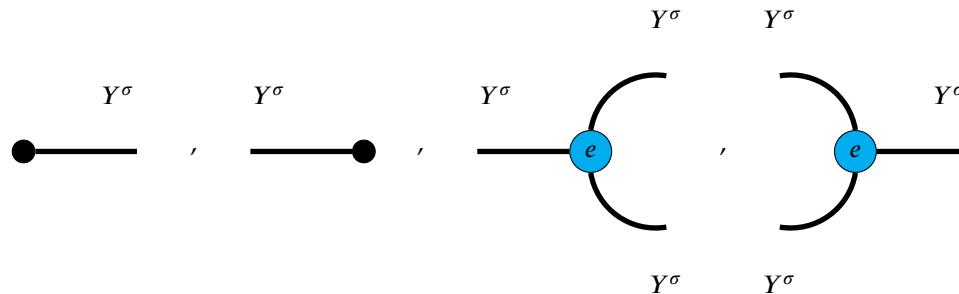
the idempotent becomes an explicit 1-1 spider, rather than the identity.



Construction 4.5.6 (Sticky spiders from split idempotents). Given an idempotent $e : Y^\sigma \rightarrow Y^\sigma$ that splits through a discrete topology X^* , we construct a new (co)multiplication as follows:



that splits through a discrete topology gives a sticky spider).



is a sticky spider

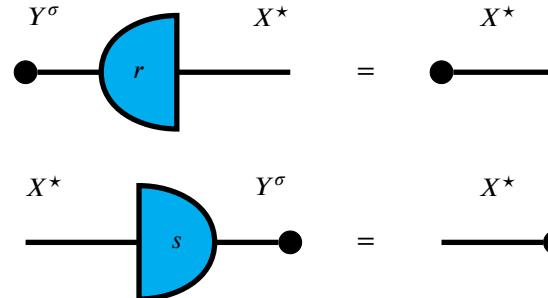
ction satisfies the frobenius rules as follows. We only present one equality; the rest follow the same idea.

(idem.)	$=$	(frob.)
$=$ 	$=$ 	$=$
(split)	$=$	(defn.)

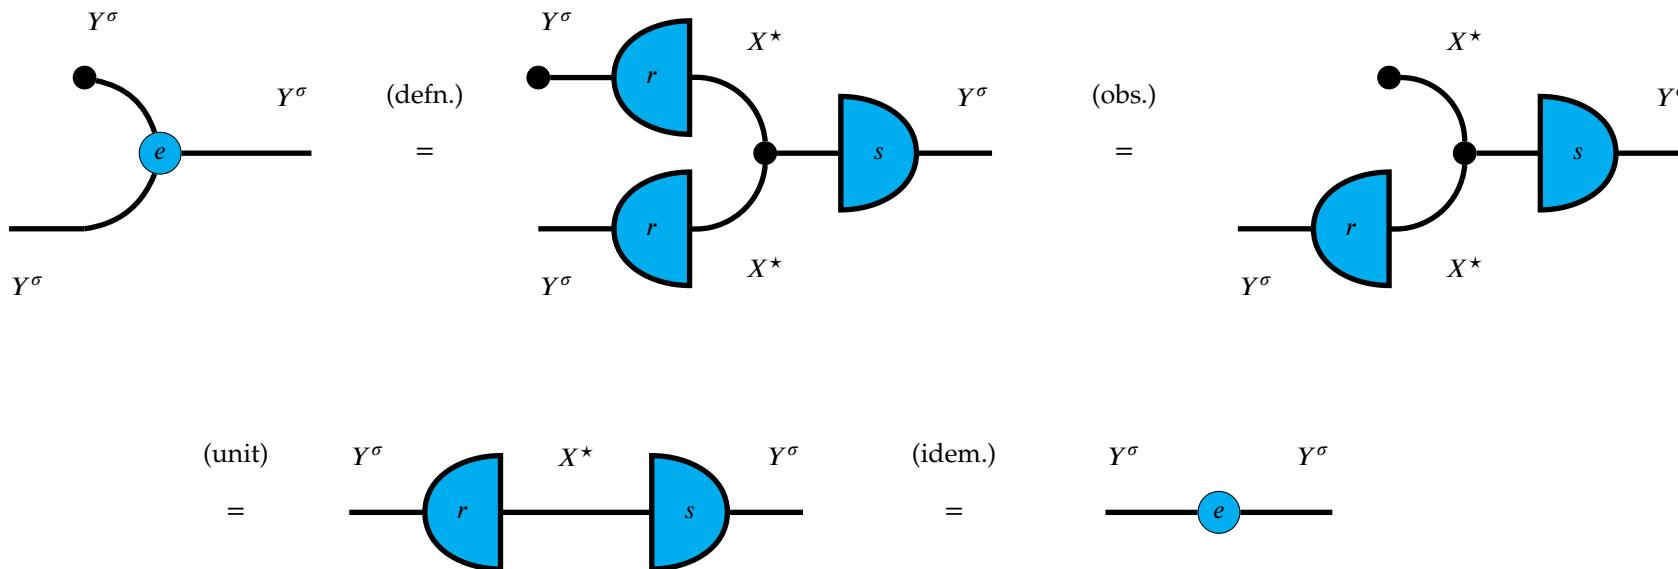
rst observe that since

$$X^* \xrightarrow{s} Y^\sigma \xrightarrow{r} X^* = X^* \xrightarrow{\text{id}} X^*$$

r must have all of X^* in its image, and s must have all of X^* in its preimage, so we have the following:



Now we show that e-unitality holds:

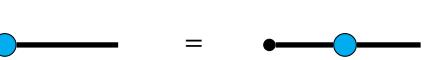


The proofs of e-countability, and e-speciality follow similarly.

We can prove a partial converse of Proposition 4.5.7: we can identify two diagrammatic equations that tell us precisely when a sticky spider has an idempotent that splits though some discrete topology.

Theorem 4.5.8. A sticky spider has an idempotent that splits through a discrete topology if and only if in addition to the sticky spider equalities, the following relation holds:

Unit/everything



Comult/copy



provide a map below of the various lemmas and propositions that will yield the claim.

basis decomposition : comultiplication

4.5.9

count/delete

basis decomposition : idem.

Lemma 4.5.18

basis decomposition : multiplication

Lemma 4.5.17

Prop. 4.5.21

basis decomposition : counit

Prop. 4.5.14

Lemma 4.5.15

copiable-basis

Proposition 4.5.9 (comult/copy implies counit/delete).

$$\text{Diagram showing } \text{comult} = \text{copy} \Rightarrow \text{counit} = \text{delete}$$

Proof.

(comult/copy)

(del)

$$\text{Diagram showing } \text{comult} = \text{copy} \subseteq \text{counit}$$

\Leftarrow
(del)

$$\text{Diagram showing } \text{counit} \Leftarrow \text{del}$$

\Leftarrow
(e-unit)

\Leftarrow
(copy-del)

So:

$$\text{Diagram showing } \text{counit} = \text{copy}$$

So:

$$\text{Diagram showing } \text{counit} = \text{copy} = \text{del}$$

□

Consider the set $e(\{x\})$ obtained by applying the idempotent e to a singleton $\{x\}$, and take an arbitrary element $y \in e(x)$ of this set. Diagrammatically:

$$\text{---} \cdot \text{---} \in \text{---} \cdot \text{---} \quad \Rightarrow \quad \text{---} \cdot \text{---} = \text{---} \cdot \text{---} \quad \text{Or} \quad \text{---} \cdot \text{---}$$

$$\text{---} \cdot \text{---} \neq \text{---} \cdot \text{---}$$

We seek: $\text{---} \cdot \text{---} = \text{---} \cdot \text{---}$

Using inclusion:

$$\text{---} \cdot \text{---} \supseteq \text{---} \cdot \text{---} \quad \text{(prem.)} \quad \text{---} \cdot \text{---} = \text{---} \cdot \text{---} \quad \text{(copy)} \quad \text{---} \cdot \text{---} = \text{---} \cdot \text{---} \quad \text{(comult/copy)} \quad \text{---} \cdot \text{---} = \text{---} \cdot \text{---} \quad \text{(e-special)}$$

So: $\text{---} \cdot \text{---} = \text{---} \cdot \text{---} \neq \text{---} \cdot \text{---}$

So we have the following **equality**:

$$\begin{array}{c}
 \text{Diagram 1} \quad \text{Diagram 2} \quad \text{Diagram 3} \quad \text{Diagram 4} \quad \text{Diagram 5} \\
 \text{(frob.)} \quad \text{(comult/copy)} \quad \text{(e-spider)} \quad \text{(e-copy)} \quad \text{II}
 \end{array}$$

$$\begin{array}{c}
 \text{Diagram 1} \quad \text{Diagram 2} \quad \text{Diagram 3} \quad \text{Diagram 4} \quad \text{Diagram 5} \\
 \text{II} \quad \text{II} \quad \text{II} \quad \text{II} \quad \text{II}
 \end{array}$$

Which implies:

$$\begin{array}{c}
 \text{Diagram 1} \quad \text{Diagram 2} \quad \text{Diagram 3} \quad \text{Diagram 4} \\
 \text{II} \quad \text{II} \quad \text{II} \quad \text{II}
 \end{array}$$

and symmetrically,

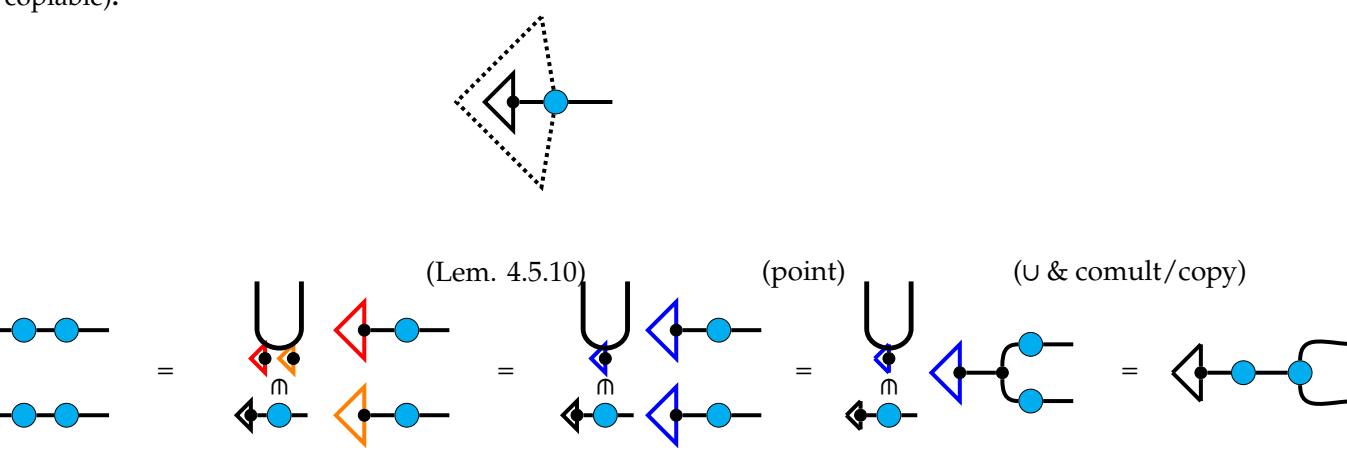
$$\begin{array}{c}
 \text{Diagram 1} \quad \text{Diagram 2} \\
 \text{II} \quad \text{II}
 \end{array}$$

So we have the claim:

$$\begin{array}{c}
 \text{Diagram 1} \quad \text{Diagram 2} \\
 \text{II} \quad \text{II}
 \end{array}$$

□

-copyable).



□

Proposition 4.5.12 (The unit is the union of all e -copiables).

$$\text{---} = \begin{array}{c} \text{U} \\ \triangleleft \\ \triangleleft \end{array}$$

Proof.

The union of *all* e -copiaables is a subset of the unit.

$$\begin{array}{ccccc}
 \text{Diagram A} & = & \text{Diagram B} & \subseteq & \text{Diagram C} \\
 \text{(Lem. 4.5.18)} & & \text{(evr.)} & & \text{(unit/evr.)}
 \end{array}$$

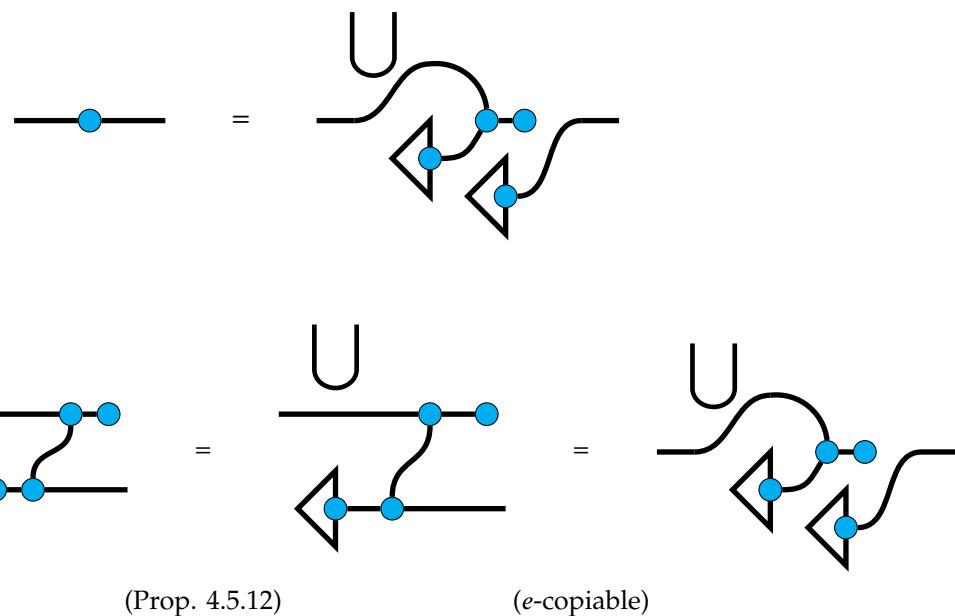
The unit is *some* union of e -copiales.

So the containment must be an equality.

$$\text{---} = \begin{array}{c} \text{U} \\ \diagdown \quad \diagup \\ \text{A} \end{array}$$

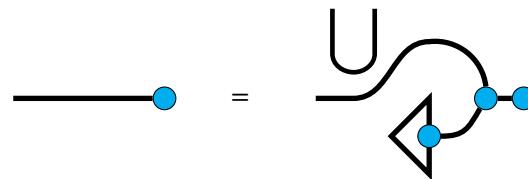
1

position of e).

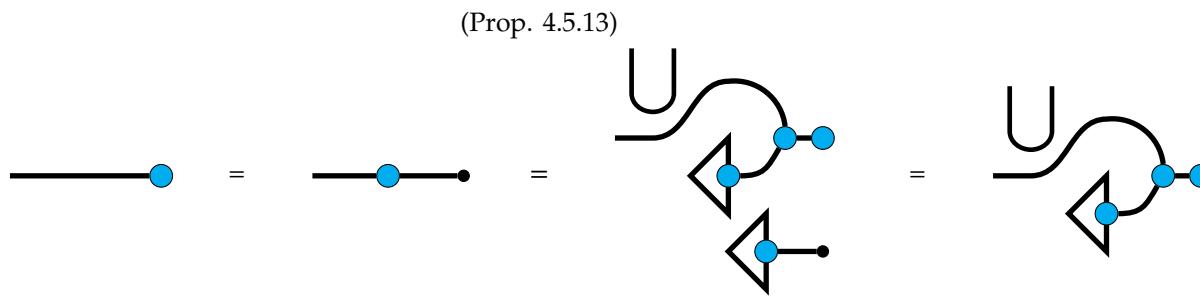


□

Proposition 4.5.14 (e -copyable decomposition of counit).



Proof.



□

HAVE LIKE AN ORTHONORMAL BASIS, AS THE FOLLOWING LEMMAS SHOW.
 (onal under multiplication).

$$\text{Diagram: } \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array} = \left\{ \begin{array}{ll} \text{Black dot} & \text{if } \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array} \neq \begin{array}{c} \text{Blue loop} \\ \text{Red loop} \end{array} \\ \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array} & \text{if } \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array} = \begin{array}{c} \text{Blue loop} \\ \text{Red loop} \end{array} \end{array} \right.$$

$$\text{Diagram: } \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array} = \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array}$$

$$\text{Diagram: } \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array} = \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array}$$

$$\text{Diagram: } \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array} = \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array} = \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array}$$

$$\text{Diagram: } \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array} = \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array}$$

$$\text{Diagram: } \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array} = \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array}$$

$$\text{Diagram: } \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array} = \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array}$$

So:

$$\text{Diagram: } \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array} \neq \text{Black dot} \Rightarrow \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array} = \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array} = \begin{array}{c} \text{Red loop} \\ \text{Blue loop} \end{array}$$

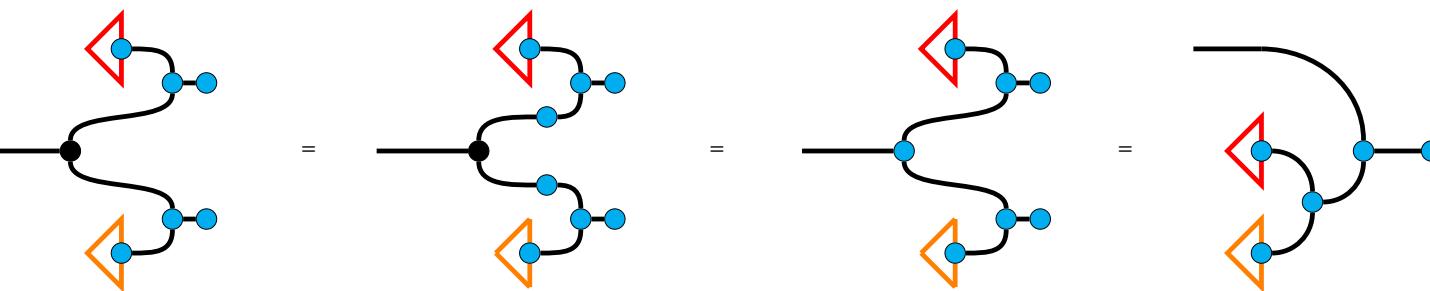
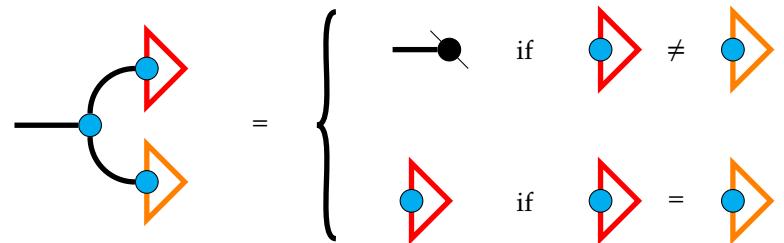
□

Convention 4.5.16 (Shorthand for the open set associated with an e -copyable). We introduce the following diagrammatic shorthand.

$$\text{Diagram: } \begin{array}{c} \text{A horizontal black line ending in an orange diamond-shaped node with a blue dot.} \\ := \\ \text{A horizontal black line with a wavy segment that loops back to form a small loop containing a blue dot, which is then connected to another blue dot.} \end{array}$$

Including the coloured dot is justified, because these open sets are co-copyable with respect to the multiplication of the sticky spider.

$$\text{Diagram sequence: } \begin{array}{ccccccc} \text{Diagram: } & \text{Diagram: } & = & \text{Diagram: } & = & \text{Diagram: } & = \\ \text{A horizontal black line with two blue dots, followed by an orange diamond-shaped node with a blue dot.} & \text{A horizontal black line with two blue dots, followed by a wavy line that loops back to form a small loop containing a blue dot, which is then connected to another blue dot.} & & \text{A horizontal black line with two blue dots, followed by an orange diamond-shaped node with a blue dot, followed by a wavy line that loops back to form a small loop containing a blue dot, which is then connected to another blue dot.} & & \text{A horizontal black line with two blue dots, followed by an orange diamond-shaped node with a blue dot, followed by a wavy line that loops back to form a small loop containing a blue dot, which is then connected to another blue dot.} & \text{A horizontal black line ending in an orange diamond-shaped node with a blue dot.} \\ := & & = & & = & & = \\ \text{Diagram: } & & & \text{Diagram: } & & \text{Diagram: } & \\ \text{A horizontal black line with two blue dots, followed by an orange diamond-shaped node with a blue dot, followed by a wavy line that loops back to form a small loop containing a blue dot, which is then connected to another blue dot.} & & & \text{A horizontal black line with two blue dots, followed by an orange diamond-shaped node with a blue dot, followed by a wavy line that loops back to form a small loop containing a blue dot, which is then connected to another blue dot.} & & \text{A horizontal black line with two blue dots, followed by an orange diamond-shaped node with a blue dot, followed by a wavy line that loops back to form a small loop containing a blue dot, which is then connected to another blue dot.} & \text{A horizontal black line ending in an orange diamond-shaped node with a blue dot.} \\ & & & & & & \end{array}$$



emma 4.5.15 to the final diagram.

1

Lemma 4.5.18 (e-copiables are e-fixpoints).

$$\begin{array}{c} \text{---} \\ \swarrow \quad \searrow \\ \text{---} \end{array} = \begin{array}{c} \text{---} \\ \swarrow \quad \circ \\ \text{---} \end{array}$$

Proof.

$$\begin{array}{ccccc} \begin{array}{c} \text{---} \\ \swarrow \quad \searrow \\ \text{---} \end{array} & = & \begin{array}{c} \text{---} \\ \swarrow \quad \circ \\ \text{---} \end{array} & = & \begin{array}{c} \text{---} \\ \swarrow \quad \circ \\ \text{---} \end{array} \\ \text{(e-counit)} & & \text{(coun/del)} & & \text{(e-copy)} \\ \begin{array}{c} \text{---} \\ \swarrow \quad \circ \\ \text{---} \end{array} & = & \begin{array}{c} \text{---} \\ \swarrow \quad \circ \\ \text{---} \end{array} & = & \begin{array}{c} \text{---} \\ \swarrow \quad \circ \\ \text{---} \end{array} \\ & & \text{---} & & \text{---} \end{array}$$

Observe that the final equation of the proof also holds when the initial e-copiable is the empty set. □

1).



(coun/del)

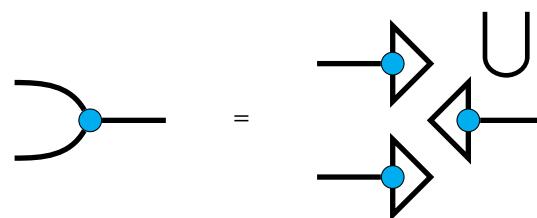
(Lem. 4.5.18)

(Prem.)

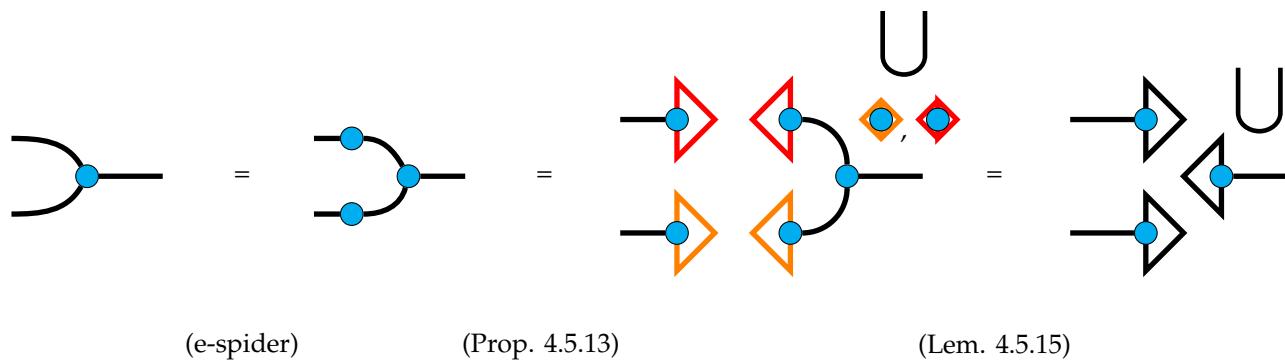


1

Proposition 4.5.20 (e -copyable decomposition of multiplication).

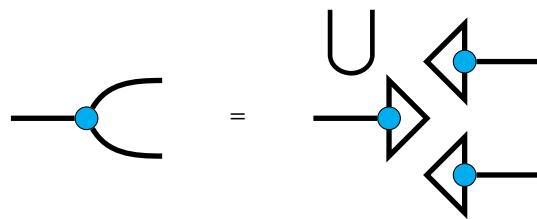


Proof.

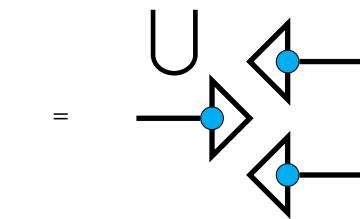


□

position of comultiplication).



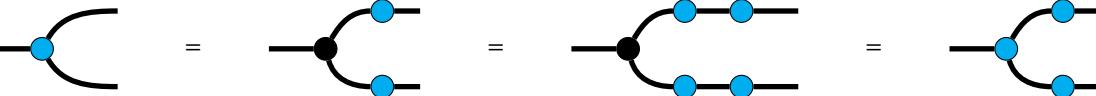
(comult/copy)



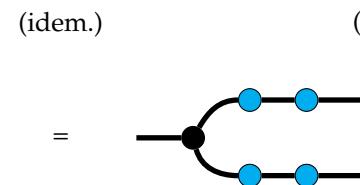
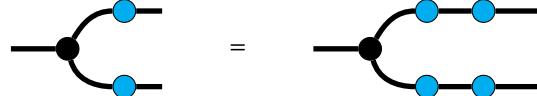
(idem.)



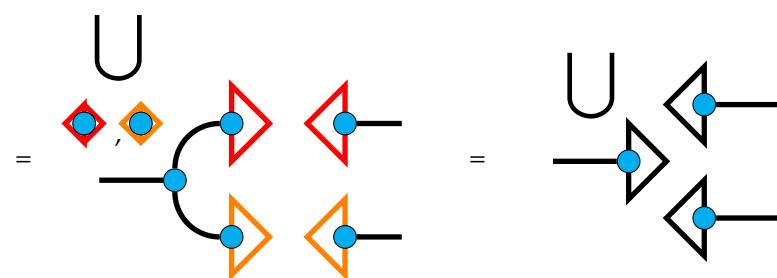
(comult/copy)



(comult/copy)



(idem.)



(Prop. 4.5.13)

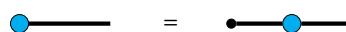
(Lem. 4.5.17)

□

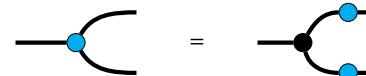
NOW WE CAN PROVE THEOREM 4.5.8.

Proof. First a reminder of the claim; we want to show that when given a sticky spider, the following relations hold if and only if the idempotent splits through a discrete topology.

Unit/everything



Comult/copy



The crucial observation is that the e -copyable decomposition of the idempotent given by Proposition 4.5.13 is equivalent to a split idempotent though the set of e -copyables equipped with discrete topology.

$$\text{---} \bullet = \text{---} \bullet \begin{array}{c} \cup \\ i \in \textcolor{blue}{I} \end{array} \text{---} \bullet = \text{---} \bullet \textcolor{blue}{I}^* \text{---}$$

(Prop. X)

$$\text{---} \bullet \textcolor{blue}{I} = \{(x, i) \mid i \in \textcolor{blue}{I}, x \in |i\rangle\}$$

$$\text{---} \textcolor{blue}{I} = \{(i, x) \mid i \in \textcolor{blue}{I}, x \in <i|\}$$

By copiable basis Proposition 4.5.12 and the decompositions Propositions 4.5.14, 4.5.20, 4.5.21, we obtain the only-if direction.

(unit/evr.)	(Prop. 4.5.12)	(Prop. 4.5.9)	(Prop. 4.5.14)
(Prop. 4.5.21)		(Prop. 4.5.20)	

the unit/everything relation, we have:

$$\begin{array}{ccccc}
 (\text{split}) & & (\text{Prop. 4.5.7}) & & (\text{split}) \\
 \bullet - \bullet & = & \bullet - \text{C} - \text{C} & = & \bullet - \text{C} - \bullet = \bullet - \bullet
 \end{array}$$

Observe that for any split idempotent, the retract must be a partial function. To see this, suppose the split idempotent $e = r; s$ is on $. .$. Seeking contradiction, if the retract is not a partial function, then there is some point $x \in X$ such that $x \in e(x)$, and the image of this point, which we denote and discriminate $a, b \in r(x) \subseteq Y$ and $a \neq b$. Because the composite $s; r = 1_Y$ of the section and retract of the section s must be total – i.e. the image $s(X) = Y$. So $x \in s(a) \cap s(b)$. Now we have that $(a, x), (b, x) \in s$, and $(x, a), (x, b) \in r$, therefore this contradicts that $s; r$ is the identity relation 1_Y .

$$\begin{array}{ccccc}
 (\text{split}) & & (\text{pfn.}) & & (\text{split}) \\
 \bullet - \bullet & = & \bullet - \text{C} - \text{C} & = & \bullet - \text{C} - \bullet = \bullet - \bullet
 \end{array}$$

□

4.6 Mathematician's endnotes

This section has two aims. First is to formally demonstrate that **ContRel** is indeed a symmetric monoidal category. Second is to investigate the relationship between **ContRel** and what seem like they should be close cousins: **Rel**, **Top**, and **Loc**. We demonstrate that **ContRel** enjoys a free-forgetful adjunction with **Rel** as expected, but **ContRel** has no forgetful functor to **Loc**. We verify that **ContRel** cannot be viewed as "powering up topology with relations" in the usual ways. Specifically, **ContRel** does not arise as conservative generalisation of the Kleisli category of the powerset monad on **Set** to **Top**, nor is it equivalent to **Span(Top)**. We provide a sketch involving display categories to attempt to explain where the topology is coming from. The failure of these (relatively sophisticated and general) techniques to modify **Top** to accommodate relations may explain why **ContRel** has no footprint in the literature, and suggests that the study of this category may be a novel contribution.

4.6.1 The category **ContRel**

Proposition 4.6.1 (**ContRel** is a category). continuous relations form a category **ContRel**.

Proof. IDENTITIES: Identity relations, which are always continuous since the preimage of an open U is itself.

COMPOSITION: The normal composition of relations. We verify that the composite $X^\tau \xrightarrow{R} Y^\sigma \xrightarrow{S} Z^\theta$ of continuous relations is again continuous as follows:

$$U \in \theta \implies S^\dagger(U) \in \sigma \implies R^\dagger \circ S^\dagger(U) = (S \circ R)^\dagger \in \tau$$

ASSOCIATIVITY OF COMPOSITION: Inherited from **Rel**. □

4.6.2 Symmetric Monoidal structure

Proposition 4.6.2. (**ContRel**, \bullet , $X^\tau \otimes Y^\sigma := (X \times Y)^{(\tau \times \sigma)}$) is a symmetric monoidal closed category.

Proof. TENSOR UNIT: The one-point space \bullet . Explicitly, $\{\star\}$ with topology $\{\emptyset, \{\star\}\}$.

TENSOR PRODUCT: For objects, $X^\tau \otimes Y^\sigma$ has base set $X \times Y$ equipped with the product topology $\tau \times \sigma$. For morphisms, $R \otimes S$ the product of relations. We show that the tensor of continuous relations is again a continuous relation. Take continuous relations $R : X^\tau \rightarrow Y^\sigma$, $S : A^\alpha \rightarrow B^\beta$, and let U be open in the product topology $(\sigma \times \beta)$. We need to prove that $(R \times S)^\dagger(U) \in (\tau \times \alpha)$. We may express U as $\bigcup_{i \in I} y_i \times b_i$, where the y_i and b_i are in the bases \mathfrak{b}_σ and \mathfrak{b}_β respectively. Since for any relations we have that $R(A \cup B) = R(A) \cup R(B)$ and $(R \times S)^\dagger = R^\dagger \times S^\dagger$:

$$\begin{aligned} (R \times S)^\dagger(\bigcup_{i \in I} y_i \times b_i) &= \bigcup_{i \in I} (R \times S)^\dagger(y_i \times b_i) \\ &= \bigcup_{i \in I} (R^\dagger \times S^\dagger)(y_i \times b_i) \end{aligned}$$

ous, $R^\dagger(y_i) \in \tau$. Symmetrically, $S^\dagger(b_i) \in \alpha$. So each $(R^\dagger \times S^\dagger)(y_i \times b_i) \in (\tau \times \alpha)$. Topologies are closed under arbitrary union, so we

HERITED FROM Rel. We will be explicit with the unitor, but for the rest, we will only check that the usual isomorphisms from **Rel** bracket-glut, we will vertically stack some tensored expressions.

as the relations $\lambda_{X^\tau} : \bullet \times X^\tau \rightarrow X^\tau := \{(\begin{pmatrix} \star \\ x \end{pmatrix}, x) \mid x \in X\}$, and we reverse the pairs to obtain the inverse $\lambda_{X^\tau}^{-1}$. These relations are y of τ with the singleton is homeomorphic to τ : $U \in \tau \iff (\bullet, U) \in (\bullet \times \tau)$. These relations are evidently inverses that compose to metric for the right unitors ρ_{X^τ} .

$\rho : ((X \times Y) \times Z)^{((\tau \times \sigma) \times \rho)} \rightarrow (X \times (Y \times Z))^{(\tau \times (\sigma \times \rho))}$ are inherited from **Rel**. They are:

$$\alpha_{X^\tau Y^\sigma Z^\rho} := \{(\left(\begin{pmatrix} x \\ y \end{pmatrix}, z\right), (x, \left(\begin{pmatrix} y \\ z \end{pmatrix}\right)) \mid x \in X, y \in Y, z \in Z\}$$

tor, observe that product topologies are isomorphic in **Top** up to bracketing, and these isomorphisms are inherited by **ContRel**. pairs of the relation reversed and is evidently an inverse that composes to the identity.

$)^{\tau \times \sigma} \rightarrow (Y \times X)^{\sigma \times \tau}$ are defined:

$$\{(\left(\begin{pmatrix} x \\ y \end{pmatrix}, \left(\begin{pmatrix} y \\ x \end{pmatrix}\right)\right) \mid x \in X, y \in Y\}$$

the isomorphisms between $X^\tau \times Y^\sigma$ and $Y^\sigma \times X^\tau$ in **Top**. They inherit everything else from **Rel**

all of the natural isomorphisms are continuous, it suffices to say that the coherences [] are inherited from the symmetric monoidal s with topologies. \square

luator.

placeholder

objects). A *biproduct* is simultaneously a categorical product and coproduct. A *zero object* is both an initial and a terminal object. sets equipped with reversible injections) and a zero object (the empty set).
object.

Proof. As in **Rel**, there is a unique relation from every object to and from the empty set with the empty topology. \square

Proposition 4.6.5. **ContRel** has biproducts.

Proof. The biproduct of topologies X^τ and Y^σ is their direct sum topology $(X \sqcup Y)^{(\tau+\sigma)}$ – the coarsest topology that contains the disjoint union $\tau \sqcup \sigma$. As in **Rel**, the (in/pro)jections are partial identities, which are continuous by construction. To verify that it is a coproduct, given continuous relations $R : X^\tau \rightarrow Z^\rho$ and $S : Y^\sigma \rightarrow Z^\rho$, where the disjoint union $X \sqcup Y$ of sets is $\{x_1 \mid x \in X\} \cup \{y_2 \mid y \in Y\}$, we observe that $R + S := \{(x_1, z) \mid (x, z) \in R\} \cup \{(y_2, z) \mid y \in S\}$ is continuous and commutes with the injections as required. The argument that it is a product is symmetric. \square

Remark 4.6.6. Biproducts yield another symmetric monoidal structure which the \times monoidal product distributes over appropriately to yield a rig category. Throughout the chapter we have been using \sqcup , but we could have also "diagrammatised" \sqcup by treating it as a monoid internal to **ContRel** viewed as a symmetric monoidal category with respect to the biproduct. There are two diagrammatic formalisms for rig categories that we could have used, $[]$ and $[\cdot]$. Neither case is perfectly suitable due to the fact that we sometimes took unions over arbitrary indexing sets, which is alright in topology but not depictable as a finite diagram in the \oplus -structure. A neat fact that follows is that a topological space is compact precisely when any arbitrarily indexed \sqcup of tests in the \times -structure is *depictable* in the \oplus -structure of either diagrammatic calculus for rig categories. **FdHilb** also has a monoidal product notated \otimes that distributes over the monoidal structure given by biproducts \oplus . In contrast, we have used \times – the cartesian product notation – for the monoidal product of **ContRel** since that is closer to what is familiar for sets.

4.6.4 **ContRel** and **Rel** are related by a free-forgetful adjunction

We provide free-forgetful adjunctions relating **ContRel** to **Rel** by "forgetting topology" and sending sets to "free" discrete topologies.

WE EXHIBIT A FREE-FORGETFUL ADJUNCTION BETWEEN **REL** AND **CONTREL**.

Lemma 4.6.7 (Any relation R between discrete topologies is continuous). *Proof.* All subsets in a discrete topologies are open. \square

Definition 4.6.8 ($L : \mathbf{Rel} \rightarrow \mathbf{ContRel}$). We define the action of the functor L :

On objects $L(X) := X^*$, (X with the discrete topology)

On morphisms $L(X \xrightarrow{R} Y) := X^* \xrightarrow{R} Y^*$, the existence of which in **ContRel** is provided by Lemma 4.6.7.

Evidently identities and associativity of composition are preserved.

Definition 4.6.9 ($R : \mathbf{ContRel} \rightarrow \mathbf{Rel}$). We define the action of the functor R as forgetting the topological structure.

On objects $R(X^\tau) := X$

On morphisms $R(X^\tau \xrightarrow{S} Y^\sigma) := X \xrightarrow{S} Y$

Evidently identities and associativity of composition are preserved.

Lemma 4.6.10 ($RL = 1_{\mathbf{Rel}}$). The composite RL (first L , then R) is precisely equal to the identity functor on **Rel**.

Proof. On objects, $FU(X) = F(X^*) = X$. On morphisms, $FU(X \xrightarrow{R} Y) = F(X^* \xrightarrow{R} Y^*) = X \xrightarrow{R} Y$ \square

Reminder 4.6.11 (Coarser and finer). Given a set of points X with two topologies X^τ and X^σ , if $\tau \subset \sigma$, we say that τ is *coarser than* σ , or σ is *finer than* τ .

uous relation). Let X^σ be coarser than X^τ . The identity relation on underlying points $X^\tau \xrightarrow{1_X} X^\sigma$ is then a continuous relation.

If any open set $U \in \sigma, U \subseteq X$ is again U . By definition of coarseness, $U \in \tau$. \square

We verify the triangular identities governing the unit and counit of the adjunction, which we first provide. By Lemma 4.6.10, we

$\Rightarrow RL$ we take to be the identity morphism:

$$\eta_X := 1_X$$

$LR \Rightarrow 1_{\mathbf{ContRel}}$ we define to be a coarsening, the existence of which in **ContRel** is granted by Lemma 4.6.12.

$$\epsilon_{X^\tau} : X^* \rightarrow X^\tau := \{(x, x) : x \in X\}$$

at an arbitrary object (set) $X \in \mathbf{Rel}$. $L(X) = X^* = LRL(X)$, where the latter equality holds because LR is precisely the identity

the left, $L(\eta_X) = L(1_X) = X^* \xrightarrow{1_X} X^* = 1_{X^*}$. For the second, $\epsilon_{L(X)} = \epsilon_{X^*} = X^* \xrightarrow{1_X} X^* = 1_{X^*}$. So we have that $L\eta; \epsilon L = L$ as

at an arbitrary object (topological space) $X^\tau \in \mathbf{ContRel}$. $R(X^\tau) = X = RLR(X^\tau)$, where the latter equality again holds because

left, $\eta_{R(X^\tau)} = \eta_X = 1_X$. For the second, $R(\epsilon_{X^\tau}) = R(X^* \xrightarrow{1_X} X^\tau) = X \xrightarrow{1_X} X = 1_X$. So $\eta R; R\epsilon = R$, as required. \square

1.1 **CONTREL** TO **LOC** HAS NO LEFT ADJOINT

ContRel to **Rel** "forgets topology while keeping the points", we might consider a forgetful functor to **Loc** that "forgets points while" that there is no such functor that forms a free-forgetful adjunction.

[] A *frame* is a poset with all joins and finite meets satisfying the infinite distributive law:

$$x \wedge (\bigvee_i y_i) = \bigvee_i (x \wedge y_i)$$

function between frames that preserves finite meets and arbitrary joins, i.e.:

$$\phi(x \wedge_A y) = \phi(x) \wedge_B \phi(y) \quad \phi(x \vee_A y) = \phi(x) \vee_B \phi(y)$$

ts and frame homomorphisms as morphisms. The category **Loc** is defined to be **Frm**^{op}.

uiations to ease the definition. The lattice of open sets of a given topology ordered by inclusion forms a frame – observe the analysis "finite intersections" : "finite meets". Closure under arbitrary joins guarantees a maximal element corresponding to the open set a setting to speak of topological structure alone, without referring to a set of underlying points, hence, pointless topology. Obvious functions, open sets in the *codomain* must correspond (uniquely) to open sets in the *domain* – so every continuous function going in the opposite direction that the function does between spaces, hence, to obtain the category **Loc** such that directions align, we that continuous relations induce frame homomorphisms in the same way. These observations give us insight into how to con-

Definition 4.6.16 ($U : \mathbf{ContRel} \rightarrow \mathbf{Loc}$). On objects, U sends a topology X^τ to the frame of opens in τ , which we denote $\hat{\tau}$.

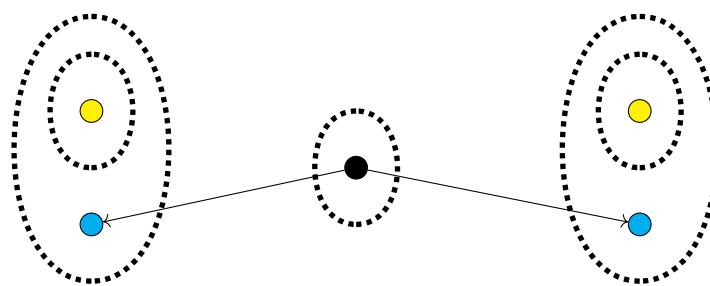
On morphisms $R : X^\tau \rightarrow Y^\sigma$, the corresponding partial frame morphism $\hat{\tau} \leftarrow \hat{\sigma}$ (notice the direction reversal for \mathbf{Loc}), we define to be $\{(U_{\in\sigma}, R^\dagger(U)_{\in\tau}) \mid U \in \sigma\}$. We ascertain that this is (1) a function that is (2) a frame homomorphism. For (1), since the relational converse picks out precisely one subset given any subset as input, these pairs do define a function. For (2), we observe that the relational converse (as all relations) preserve arbitrary unions and intersections, i.e. $R^\dagger(\bigcap_i U_i) = \bigcap_i R^\dagger(U_i)$ and $R^\dagger(\bigcup_i U_i) = \bigcup_i R^\dagger(U_i)$, so we do have a frame homomorphism. Associativity follows easily.

Proposition 4.6.17 (U has no left adjoint). *Proof.* Seeking contradiction, if U were a right adjoint, it would preserve limits. The terminal object in \mathbf{Loc} is the two-element lattice $\perp < \top$, where the unique frame homomorphism to any \mathcal{L} sends \top to the top element of \mathcal{L} and \perp to the bottom element. In $\mathbf{ContRel}$, the empty topology $\mathbf{0} = (\emptyset, \{\emptyset\})$ is terminal (and initial). However, $U\mathbf{0}$ is the singleton lattice, not $\perp < \top$ (which is the image under U of the singleton topology). \square

This is a rather frustrating result, because U does turn continuous relations into backwards frame homomorphisms on lattices of opens; see Proposition 4.4.14, and note that in the frame of opens associated with a topology, the empty set becomes the bottom element. The obstacle is the fact that the empty topology is both initial and terminal in $\mathbf{ContRel}$. We may be tempted to try treating U as a right adjoint going to \mathbf{Frm} instead, but then the monad induced by the injunction on \mathbf{Loc} would trivialise: left adjoints preserve colimits, so any putative left adjoint F must send $\perp < \top$ (initial in \mathbf{Frm} by duality) to the empty topology, and the empty topology as terminal object must be sent to the terminal singleton frame, which implies that the monad UF on \mathbf{Frm} sends everything to the singleton lattice.

4.6.5 Why not $\text{Span}(\mathbf{Top})$?

One common generalisation of relations is to take spans of monics in the base category $[]$. This actually produces a different category than the one we have defined. Below is an example of a span of monic continuous functions from \mathbf{Top} that corresponds to a relation that doesn't live in $\mathbf{ContRel}$. It is the span with the singleton as apex, with maps from the singleton to the closed points of a two Sierpiński spaces.



4.6.6 Why not a Kleisli construction on \mathbf{Top} ?

Another way to view the category \mathbf{Rel} is as the Kleisli category K_P of the powerset monad on \mathbf{Set} ; that is, every relation $A \rightarrow B$ can be viewed as a function $A \rightarrow P B$, and composition works by exploiting the monad multiplication: $A \xrightarrow{f} PB \xrightarrow{Pg} PPC \xrightarrow{\mu_{PC}} PC$. So it is reasonable to investigate whether there is a monad T on \mathbf{Top} such that K_T is equivalent to $\mathbf{ContRel}$. We observe that the usual free-forgetful adjunction between \mathbf{Set} and \mathbf{Top} sends the former to a full subcategory (of continuous functions between discrete topologies) of the latter, so a reasonable coherence condition we might ask for the putative monad T to satisfy is that it is related to P via the free-forgetful adjunction. This amounts to asking for the following commutative diagram (in addition to the usual ones stipulating that T and P are monadic):

$$\begin{array}{ccc} \mathbf{Top} & \xrightarrow{T} & \mathbf{Top} \\ \uparrow \dashv & & \downarrow \dashv \\ \mathbf{Set} & \xrightarrow{p} & \mathbf{Set} \end{array}$$

we because it witnesses K_P as precisely K_T restricted to the discrete topologies, so that T really behaves as a conservative general-commodate topologies. As a consequence of this condition, we may observe that discrete topologies X^* must be sent to discrete particular, this means the singleton topology is sent to the the discrete topology on a two-element set; $T\bullet = 2$. This sinks us. We continuous relations $X^\tau \rightarrow \bullet$ are precisely the open sets of τ , which correspond to continuous functions into Sierpiński space

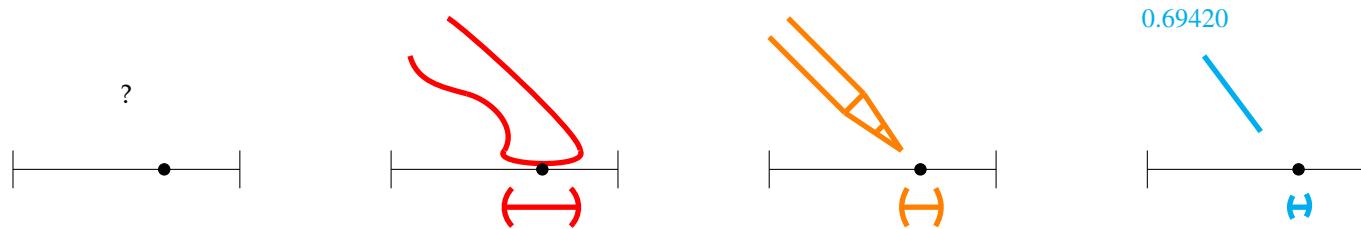
from?

ask whether **ContRel** is "giving topology to relations" or "powering up topologies with relations", but we have explored those that. It is possible that the failure of these regular avenues may explain why I had such difficulty finding any trace of **ContRel** a free-forgetful adjunction between **ContRel** and **Rel**, and if we focus on this, it is possible to crack the nut of where topology theory; here is one such sketch. Observe that the forgetful functor looks like it could be a kind of fibration, where the elements of correspond to all possible topologies on A . Moreover, these topologies may be partially ordered by coarseness-fineness to form a frame (ill suffice.) The fibre over a relation $R : A \rightarrow B$ is all pairs of topologies τ, σ such that R is continuous between A^τ and B^σ . The inuous between τ and σ , then R will be continuous for any finer topology in the domain, $\tau \leq \tau'$, and any coarser topology in over R displays a boolean-valued profunctor between preorders. So **ContRel** can be viewed as the display category induced category with preorders for objects and boolean-enriched profunctors as morphisms, and the functor encodes topological data by possible topologies, and relations to profunctors. I have deliberately left this a sketch because it doesn't seem worth it to view ex way.

worth the trouble?

duction of this section. In short, because the opens of topological spaces crudely model how we talk about concepts, and the model instances of concepts. Why this is so is best demonstrated by an illustrated example.

L FICTION. Useful, but a fiction. Suppose we have a point on a unit interval. Consider how we might tell someone else about it with a pudgy appendage, or the tip of a pencil, or give some finite decimal approximation.



But in each case we are only speaking of a vicinity, a neighbourhood, an *open set in the borel basis of the reals* that contains the point. Identifying a true point on a real line requires an infinite intersection of open balls of decreasing radius; an infinite process of pointing again and again, which nobody has the time to do. In the same way, most language outside of mathematics is only capable of offering successively finer, finite approximations to whatever it is that occurs in the mind or in reality.

MAYBE THAT EXPLAINS THE ASYMMETRY OF WHY TESTS ARE OPEN SETS, BUT WHY ARE STATES ALLOWED TO BE ARBITRARY SUBSETS? Because states in this model represent what is conceived or perceived. Suppose we have an analog photograph whether in hand or in mind, and we want to remark on a particular shade of red in some uniform patch of the photograph. As in the case of pointing out a point on the real interval, we have successively finer approximations with a vocabulary of concepts: "red", "burgundy", "hex code #800021"... but never the point in colourspace itself. If someone takes our linguistic description of the colour and tries to reproduce it, they will be off in a manner that we can in principle detect, cognize, and correct: "make it a little darker" or "add a little blue to it". That is to say, there are, in principle, differences in mind that we cannot distinguish by boundedly finite language; we would have to continue the process of "even darker" and "add a bit less blue than last time" forever. All this is just the mathematical formulation of a very common observation: sometimes you cannot do an experience justice with words, and you eventually give up with "I guess you just had to be there". Yet the experience is there and we can perform linguistic operations on it, and the states accommodate this.

TOP IS SYMMETRIC MONOIDAL CLOSED WITH RESPECT TO PRODUCT, WHY DIDN'T YOU JUST WORK THERE FROM THE START? Because **Top** is cartesian monoidal, which in particular means that there is only one test (the map into the terminal singleton topology), and worse, all states are tensor-separable. The latter fact means that we cannot reason natively in diagrams about correlated states, which are extremely useful representing entangled quantum states [dodo], and for reasoning about spatial relations [talkspace]. I'll briefly explain the gist of the analogy in prose because it is already presented formally in the cited works and elaborated in [bobcomp]. The Fregean notion of compositionality is roughly that to know a composite system is equivalent to knowing all of its parts, and diagrammatically this amounts to tensor-separability, which arises as a consequence of cartesian monoidality. Schrödinger suggests an alternative of compositionality via a lesson from entangled states in quantum mechanics: *perfect knowledge of the whole does not entail perfect knowledge of the parts*. Let's say we have information about a composite system if we can restrict the range of possible outcomes; this is the case for the bell-state, where we know that there is an even chance both qubits measure up or both measure down, and we can rule out mismatched measurements. However, discarding one entangled qubit from a bell-state means we only know that the remaining qubit has a 50/50 of measuring up or down, which is the minimal state of information we can have about a qubit. So we have a case where we can know things about the whole, but nothing about its parts. A more familiar example from everyday life is if I ask you to imagine a cup on a table in a room. There are many ways to envision or realise this scenario in your mind's eye, all drawn from a restricted set of permissible positions of the cup and the table in some room. The spatial locations of the cup and table are entangled, in that you can only consider the positions of both together. If you discard either the cup or the table from your memory, there are no restrictions about where the other object could be in the room; that is, the meaning of the utterance is not localised in any of the parts, it resides in the entangled whole.

5

Sketches of the shape of language

5.1 Topological concepts in flatland via **ContRel**

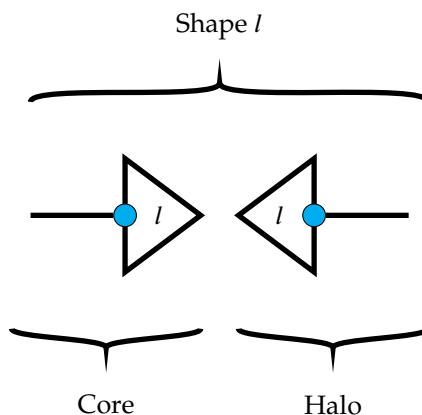
The goal of this section is to demonstrate the use of sticky spiders as formal semantics for the kinds of schematic doodles or cartoons we would like to draw. Throughout we consider sticky spiders on \mathbb{R}^2 . In Section 5.1.1, we introduce how sticky spiders may be viewed as labelled collections of shapes. In service of defining *configuration spaces* of shapes up to rigid displacement, we diagrammatically characterise the topological subgroup of isometries of \mathbb{R}^2 by building up in Sections 5.1.2 and 5.1.3 the diagrammatic presentations of the unit interval, metrics, and topological groups. To further isolate rigid displacements that arise from continuous sliding motion of shapes in the plane (thus excluding displacements that result in mirror-images), in Sections 5.1.4 and 5.1.5 we diagrammatically characterise an analogue of homotopy in the relational setting. Finally, in Sections 5.1.6 and 5.1.7 we build up a stock of topological concepts and study by examples how implementing these concepts within text circuits explains some idiosyncrasies of the theory: namely why noun wires are labelled by their noun, why adjective gates ought to commute, and why verb gates do not.

5.1.1 *Shapes and places*

Definition 5.1.1 (Labels, shapes, cores, halos). Recall by Proposition 4.5.13 that we can express the idempotent as a union of continuous relations formed of a state and test, for some indexing set of labels \mathcal{L} .

$$\text{---} \bullet = \text{---} \bullet l \cup \text{---} \bullet l$$

A *shape* is a component of this union corresponding to some arbitrary $l \in L$. So we refer to a sticky spider as a labelled collection of shapes. The state of a shape is the *halo* of the shape. The halos are precisely the copiables of the sticky spider. The test of a shape is the *core*. The cores are precisely the cocopiables of the sticky spider.

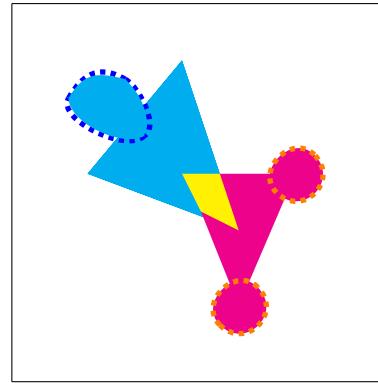


stinct cores cannot overlap). *Proof.* A direct consequence of Lemma 4.5.17. \square

n: Each core only overlaps with its corresponding halo). *Proof.* Seeking contradiction, if a core overlapped with multiple halos,

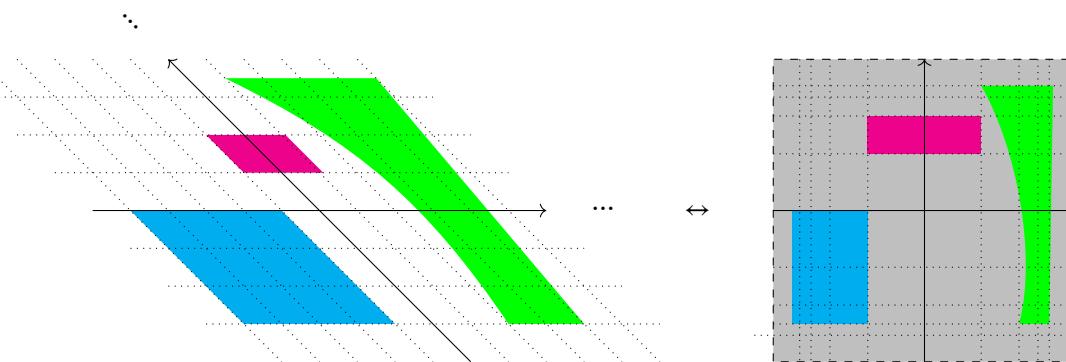
\square

n: halos may overlap). *Proof.* By example:

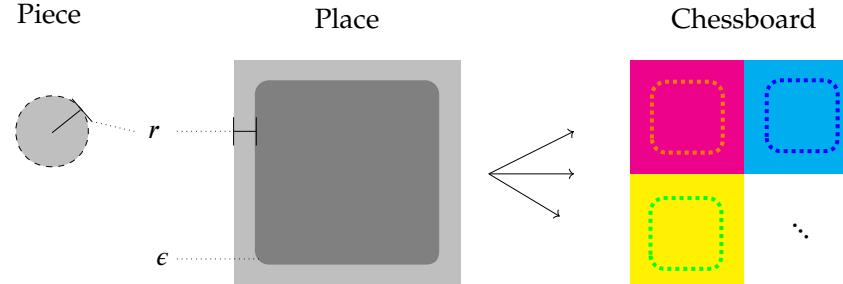


n and magenta. The halos are two triangles which overlap at a yellow region, and partially overlap with their blobby cores. The orange respectively. Observe that cores and halos do not have to be simply connected; in this example the core of the magenta s. Viewing these sticky spiders as a process, any shape that overlaps with the magenta core will be deleted and replaced by the the cyan cores and triangle. Any shape that overlaps with both the magenta and cyan cores will be deleted and replaced by the t overlaps with neither core will be deleted and not replaced. \square

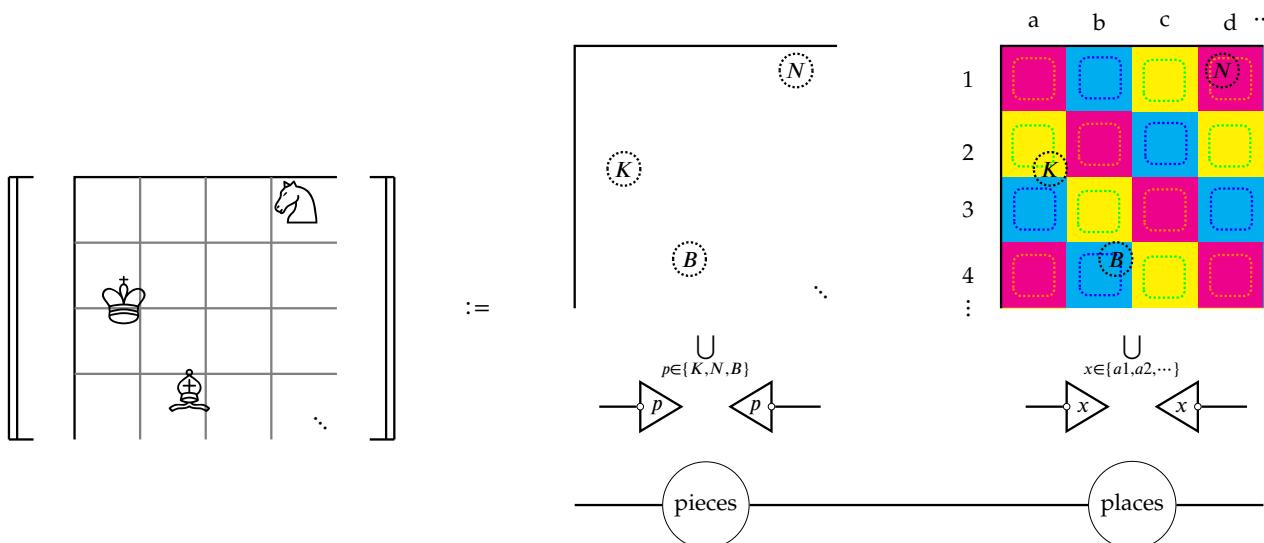
nite canvas representing all of euclidean space, properly there should be a fishbowl effect that relatively magnifies shapes close periphery, but that is only an artefact of representing all of euclidean space on a finite canvas. Since all the usual metrics are still nore this fishbowl effect and just doodle shapes as we see fit.



Example 5.1.6 (Where is a piece on a chessboard?). How is it that we quotient away the continuous structure of positions on a chessboard to locate pieces among a discrete set of squares? Evidently shifting a piece a little off the centre of a square doesn't change the state of the game, and this resistance to small perturbations suggests that a topological model is appropriate. We construct two spiders, one for pieces, and one for places on the chessboard. For the spider that represents the position of pieces, we open balls of some radius r , and we consider the places spider to consist of square halos (which tile the chessboard), containing a core inset by the same radius r ; in this way, any piece can only overlap at most one square. As a technical aside, to keep the core of the tiles open, we can choose an arbitrarily sharp curvature ϵ at the corners.



Now we observe that the calculation of positions corresponds to composing sticky spiders. We take the initial state to be the sticky spider that assigns a ball of radius r on the board for each piece. We can then obtain the set of positions of each piece by composing with the places spider. The composite (pieces;places) will send the king to $a2$, the bishop to $b4$, and the knight to $d1$, i.e. $\langle K \rangle \mapsto \langle a2 \rangle$, $\langle B \rangle \mapsto \langle b4 \rangle$ and $\langle N \rangle \mapsto \langle d1 \rangle$. In other words, we have obtained a process that models how we pass from continuous states-of-affairs on a physical chessboard to an abstract and discrete game-state.



concepts, we first need to extend our topological tools. If we have the unit interval, we can begin to define what it would mean for lines between points in those spaces), and we can also move towards defining motion as movement along a line.

is homeomorphic to the real line. How do we know when we have one of them? The following theorem provides an answer:

(X, \prec) be a topological space with a total order. If there exists a continuous map $f : X \times X \rightarrow X$ such that $\forall a, b \in X : a \prec f(a, b) \prec b$,

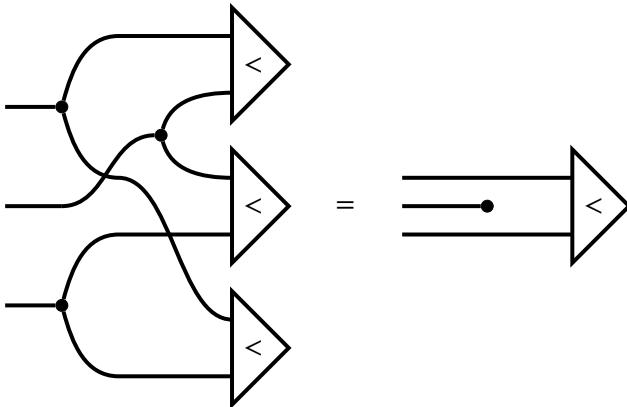
using diagrammatic equations.

is an open set – i.e. a test – on $X \times X$ that obeys the usual axiomatic rules:

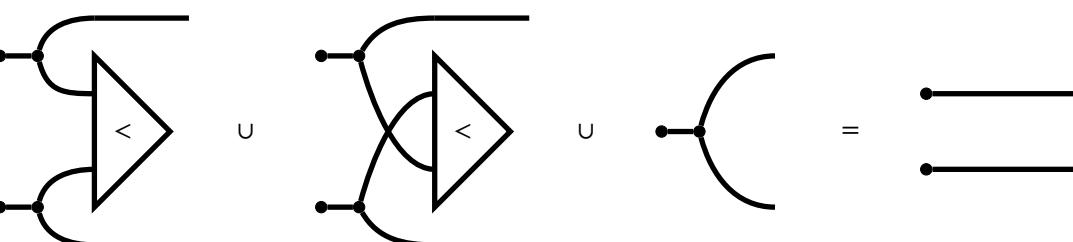
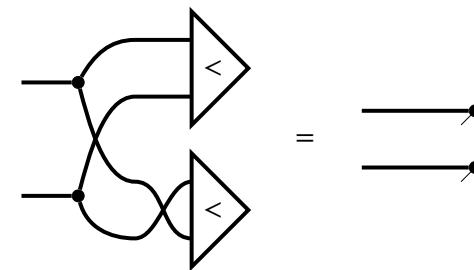
$x \not\prec x$

Transitive ($\forall xyz : x \prec y \& y \prec z \Rightarrow x \prec z$)

Antireflexive $\forall xy \neg(x \prec y \& y \prec x)$



Trichotomy ($\forall xy : x \prec y \vee y \prec x \vee x = y$)

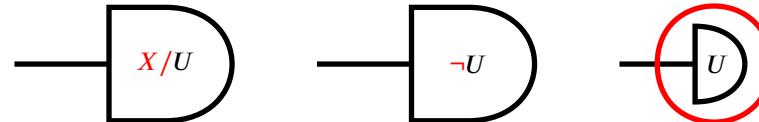


points for open intervals directly by asking for the space X to have points that are less than or greater than all other points. An-

other method, which we will use here for primarily aesthetic reasons, is to use endo combinators to define suprema. Endo combinators are like functional expressions applied to diagrams. For a motivating example, consider the case when we have a locally indiscrete topology:

Definition 5.1.8 (Locally indiscrete topology). (X, τ) is *locally indiscrete* when every open set is also closed.

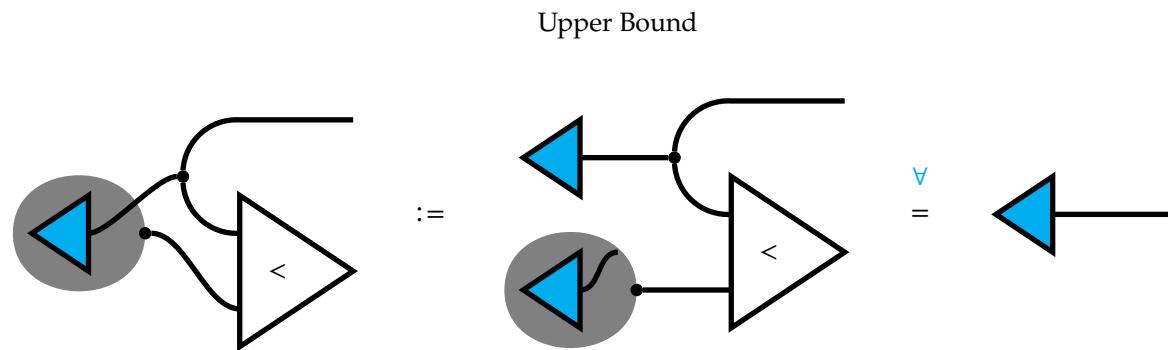
If we know that a topology is locally indiscrete and we are given an open U , we would like to notate the complement X/U – which we know to be open – as any of the following, which only differ up to notation.



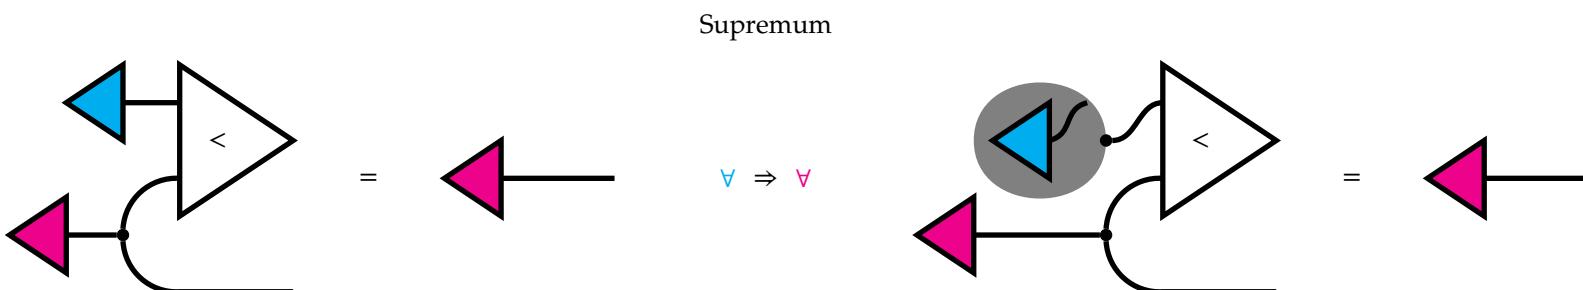
Unfortunately, the complementation operation $X/-$ is not in general a continuous relation, hence in the lattermost expression above we resort to using bubbles as a syntactic sugar. Formally, these bubbles are *endo combinators*, the semantics and notation for which we borrow and modify from [].

Definition 5.1.9 (Partial endo combinator). In a category \mathcal{C} , a *partial endo combinator* on a homset $(\mathcal{C})(A, B)$ is a function $(\mathcal{C})(A, B) \rightarrow (\mathcal{C})(A, B)$

Using this technology, we can define:



And we can add in further equations governing the upper bound endo combinator to turn it into a supremum, where the lower endpoint is obtained as the supremum of the empty set, and the upper endpoint is the supremum of the whole set.



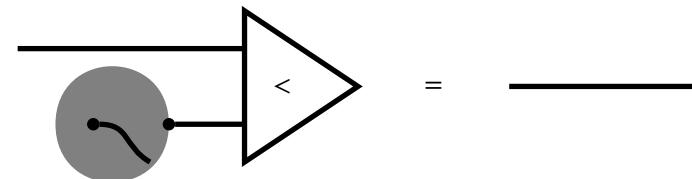
ely graphically:

lower endpoint

=



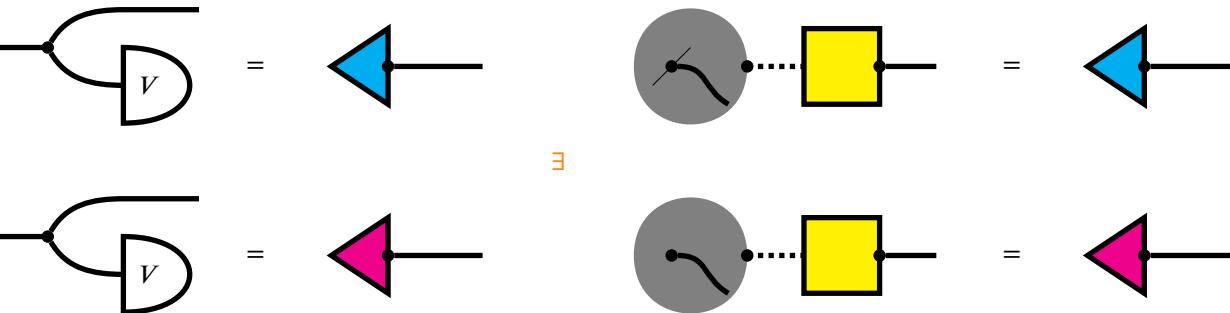
Upper endpoint



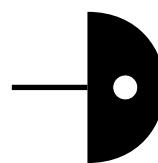
unit interval using a thick dotted wire.

an define the usual topological notion of a simply connected space: one where any two points can be connected by a continuous

V is *simply connected* when:



we will notate simply connected open sets as follows, where the hole is a reminder that simply connected spaces might still

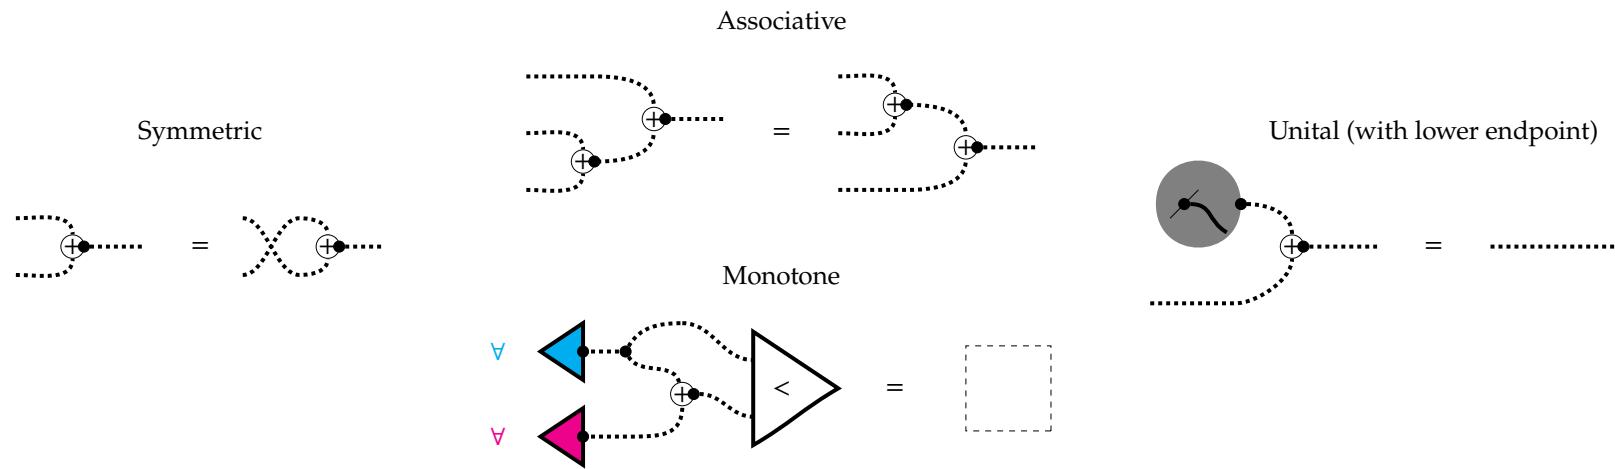


5.1.3 Displacing shapes

Static shapes in space are nice, but moving them around would be nicer. So we have to define a stock of concepts to express rigid motion. Rigidity however is a difficult concept to express in topological spaces up to homeomorphism – everyone is well aware of the popular gloss of topology in terms of coffee cups being homeomorphic to donuts. To obtain rigid transformations as we have in Euclidean space, we need to define metrics, and in order to do that, we need addition.

ADDITION

More precisely, we only need an additive monoid structure on the unit interval. We do not care about obtaining precise values from our metric, and we will not need to subtract distances from each other. All we need to know is that the lower endpoint stands in for "zero distance" – as the unit of the monoid – and that adding positive distances together will give you a larger positive distance deterministically.



METRICS

A metric on a space is a continuous map $X \rightarrow \mathbf{R}^+$ to the positive reals that satisfies the following axioms.

$$d(\textcolor{teal}{x}, \textcolor{brown}{z}) \leq d(\textcolor{teal}{x}, \textcolor{magenta}{y}) + d(\textcolor{magenta}{y}, \textcolor{brown}{z})$$

The string diagram illustrates the triangle inequality. It shows three inputs (blue, magenta, and orange) entering three parallel vertical boxes. The outputs of these boxes enter a fourth vertical box. A dotted line from the output of the fourth box goes to a diamond-shaped node labeled ' \leq '. Another dotted line from this node goes to a dashed square representing the result.

$$d(\textcolor{teal}{x}, \textcolor{magenta}{y}) = 0 \iff \textcolor{teal}{x} = \textcolor{magenta}{y}$$

The string diagram illustrates that the distance between two points is zero if and only if the points are equal. It shows two inputs (blue and magenta) entering a vertical box. The output of this box is compared against a blue input using a diamond node labeled ' \leq '. The result is then compared against a magenta input using another diamond node labeled ' \leq '. The final result is a grey circle, which is shown to be equivalent to a single magenta input.

the usual topological notion of open balls. Open balls will come in handy later, and a side-effect which we note but do not prove is that every metric space is complete. This means that for any metric space, so in the future whenever we construct spaces that come with natural metrics, we can speak of their topology.

Open ball of radius ϵ at a point \mathbf{x}

The string diagram shows the definition of an open ball. On the left, a black circle labeled ϵ, \mathbf{x} is equated to a blue input. In the middle, a double-headed arrow indicates equivalence. On the right, a blue input is compared against a magenta input using a diamond node labeled ' $<$ '. The result is compared against a white input labeled ϵ using another diamond node labeled ' $<$ '. The final result is a dashed square representing the result.

of invertible transformations of spaces $X \rightarrow X$:

A topological group G

$$\left\{ \begin{array}{c} \forall \gamma \in G \exists \gamma^- \in G \\ \gamma \circ \gamma^- = \text{id} \end{array} \right. = \text{id}$$

ISOMETRY

But recall that the collections of invertible transformations we are really interested in are the *rigid* ones, the ones that move objects in space without deforming them. We can identify when a transformation is rigid by the following criterion:

γ is an *isometry*

$$\exists \left\{ \begin{array}{c} \text{a curved path} \\ \text{from } \gamma \text{ to } \text{id} \end{array} \right. = \text{id} \rightarrow \text{blue arrowhead}$$

RIGID DISPLACEMENTS

Now we return to our sticky spiders. From now we consider sticky spiders on the open unit square, so that we can speak of shapes on a canvas. Now we will try to displace the shapes of a sticky spider. We know the planar isometries of Euclidean space can be expressed as a translation, rotation, and a bit to indicate the chirality of the shape – as mirror reflections are also an isometry.

Isometries of \mathbf{R}^2

$$\text{id} = (\mathbf{x}, \theta, c)$$

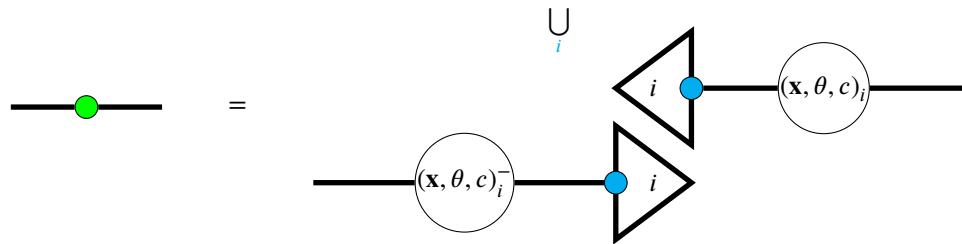
$\mathbf{x} \in \mathbf{R}^2$

$\theta \in S^1 \simeq [0, 2\pi)$

$c \in \{-1, 1\}$

With this in mind, we have the following condition relating different spiders, telling us when one is the same as the other up to rigidly displacing shapes.

Rigid displacement



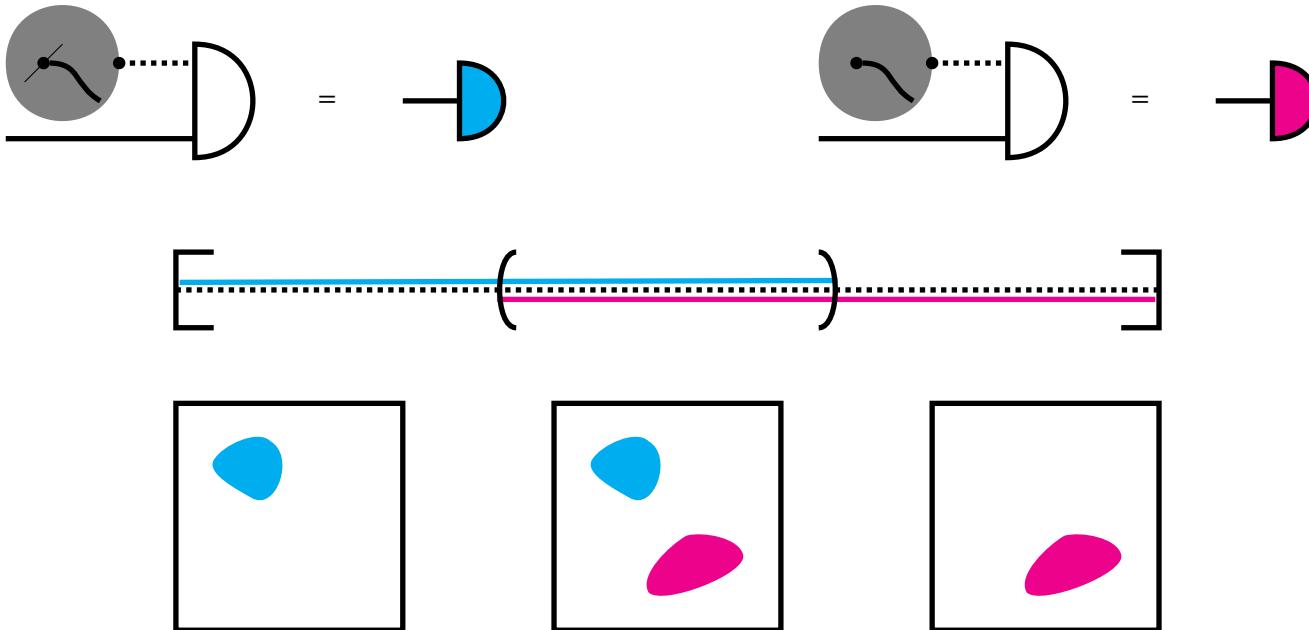
in flatland, we do not expect shapes to suddenly flip over. We would like to express just those rigid transformations that leave us really we want to only be able to slide the shapes around the canvas, not leave the canvas to flip over. So we go on to define

as in the plane from the configuration of shapes in one spider to end at the configuration of shapes in another, we ought to define continuous deformation of one map to another. However, we will have to massage the definition a little to work in our setting of

speaking of topological spaces and continuous functions, a homotopy is defined:

where f and g are continuous maps $A \rightarrow B$, a homotopy $\eta : f \Rightarrow g$ is a continuous function $\eta : [0, 1] \times A \rightarrow B$ such that

a short film where at the beginning there is an f , which continuously deforms to end the film being a g . Directly replacing "function does not quite do what we want, because we would be able to define the following "homotopy" between open sets.



What is happening in the above film is that we have our starting open set, which stays constant for a while. Then suddenly the ending open set appears, the starting open disappears, and we are left with our ending; while *technically* there was no discontinuous jump, this isn't the notion of sliding we want. The exemplified issue is that we can patch together (by union of continuous relations) vignettes of continuous relations that are not individually total on $[0, 1]$. We can patch this problem by asking for homotopies in **ContRel** to satisfy the additional condition that they are expressible as a union of continuous partial maps that are total on the unit interval.

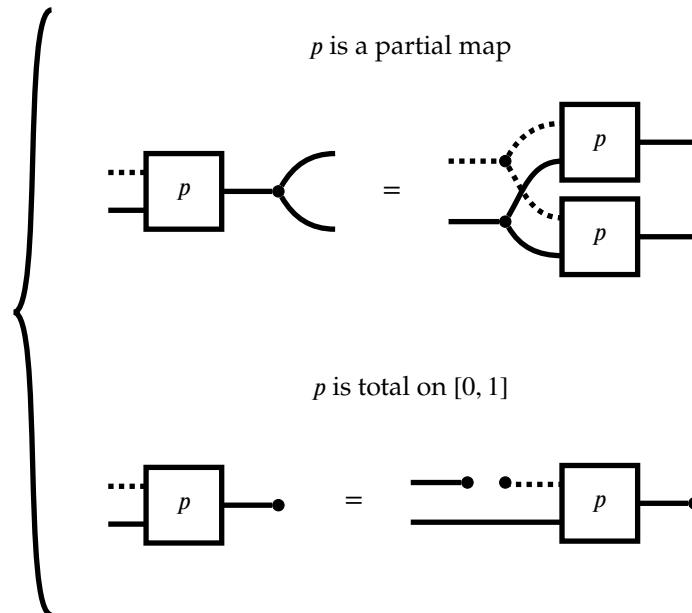
$$\eta(0, -) = \textcolor{blue}{f}(-)$$

$$\eta(1, -) = \textcolor{red}{g}(-)$$

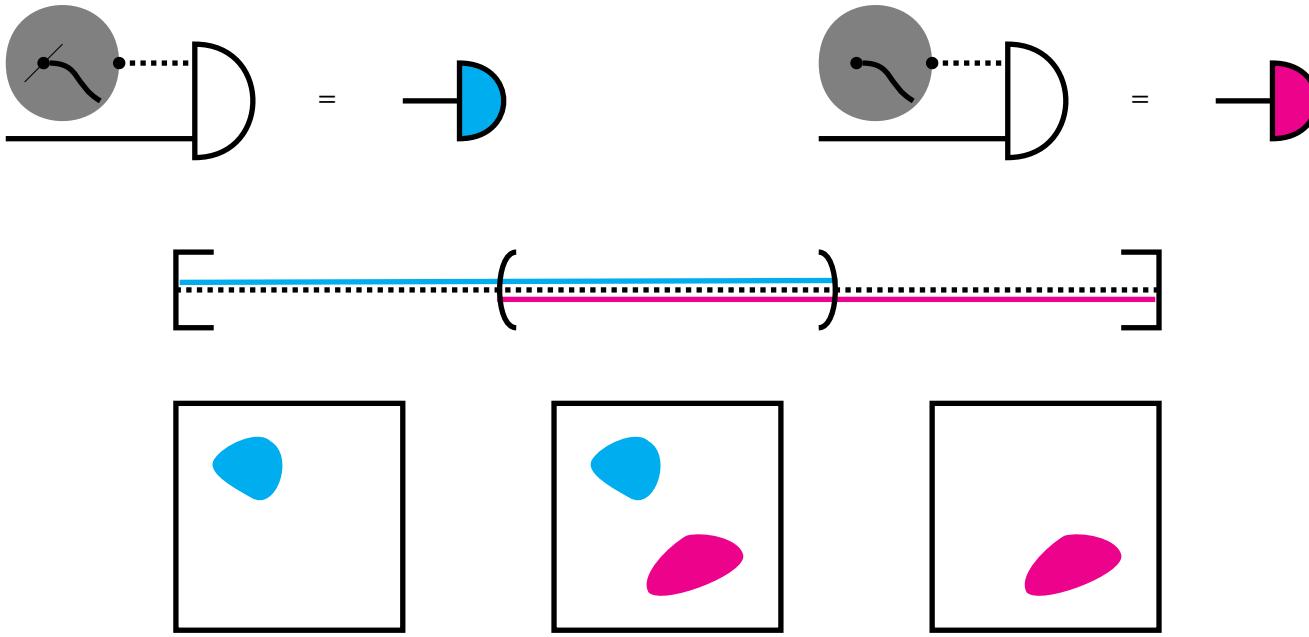


union of homotopies of partial cts. maps

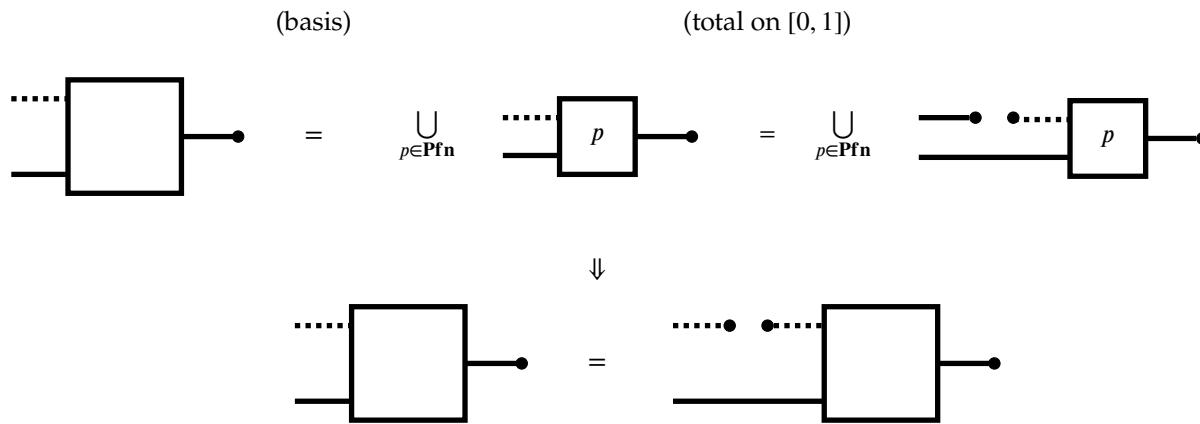
$$= \bigcup_{p \in \mathbf{Pfn}} \dots \square_p \dots$$



asking for decomposition in terms of partial comes for free by Proposition 4.4.20; the constraint of the definition is provided by
a condition than just asking that the original continuous relation be total on I :



This definition is "natural" in light of Proposition 4.4.20, that the partial continuous functions $A \rightarrow B$ form a basis for $\text{ContRel}(A, B)$: we are just asking that homotopies between partial continuous functions – which can be viewed as regular homotopies with domain restricted to the subspace topology induced by an open set – form a basis for homotopies between continuous relations.

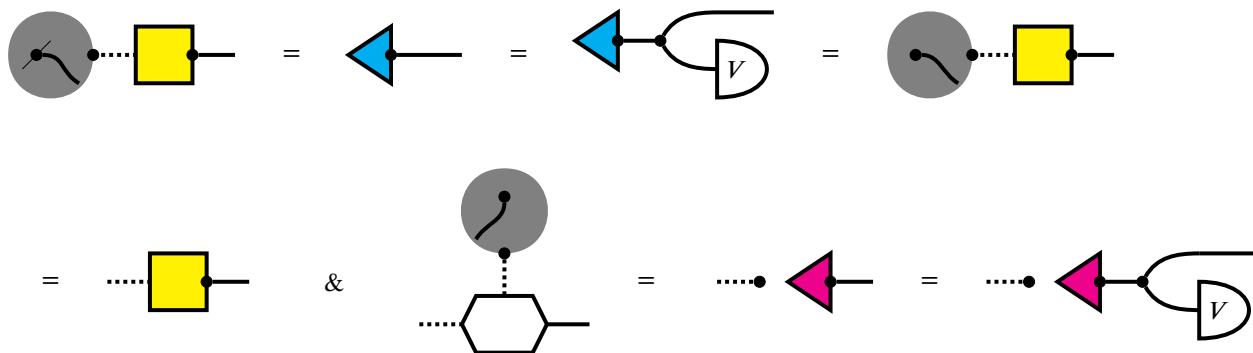


CONTRACTIBLE SPACES

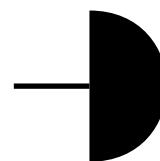
With homotopies in hand, we can define a stronger notion of connected shapes with no holes, which are usually called *contractible*. The reason for the terminology

guarantee a shape in flatland has no holes: when any loop in the shape is *contractible* to a point.

V is *contractible* when:

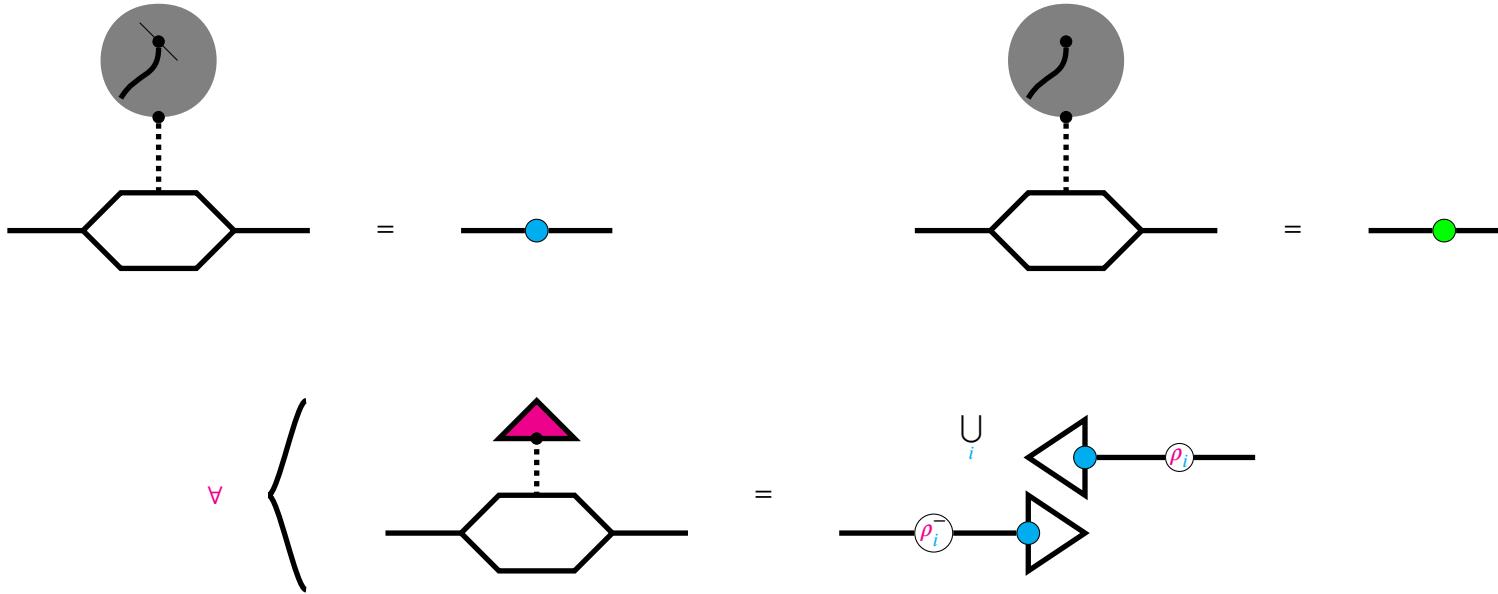


their own notation too; a solid black effect, this time with no hole.



pes. What we mean by two sticky spiders being relatable by sliding shapes is that we have a homotopy that begins at one and ends at another, and in between is itself a sticky spider related to the first by rigid displacement.

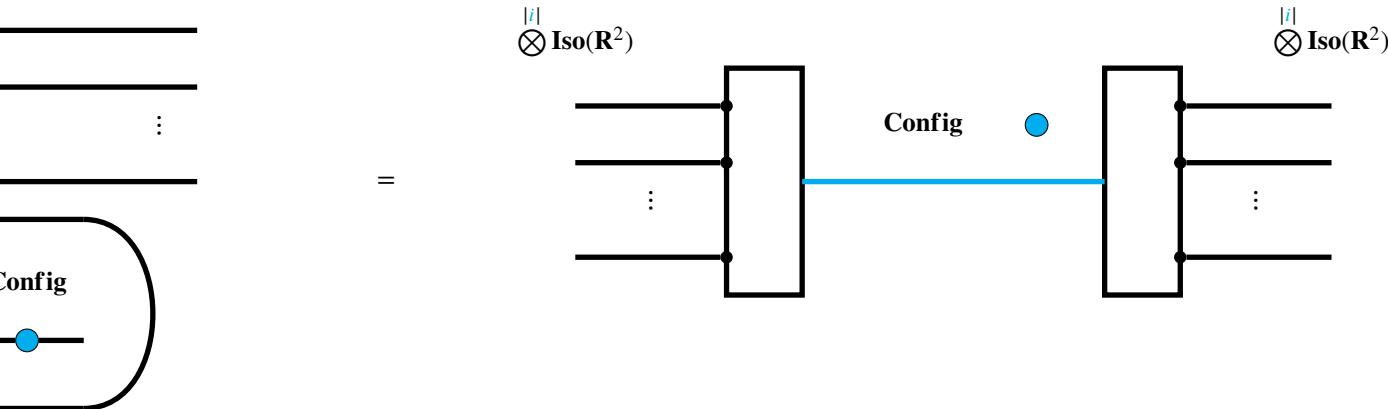
Rigid motion



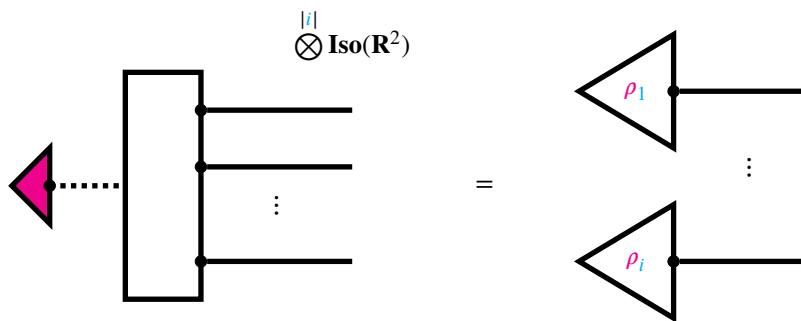
CONFIGURATION SPACES

We can depict the *configuration space* of shapes that are obtainable by displacing the shapes of a given spider by a split idempotent through the n-fold tensor of rigid transformations – a restriction to the subspace of the largest open set contained in the subset of all valid (with correct chirality) combinations of displacements that yield another spider.

Configuration space of a sticky spider

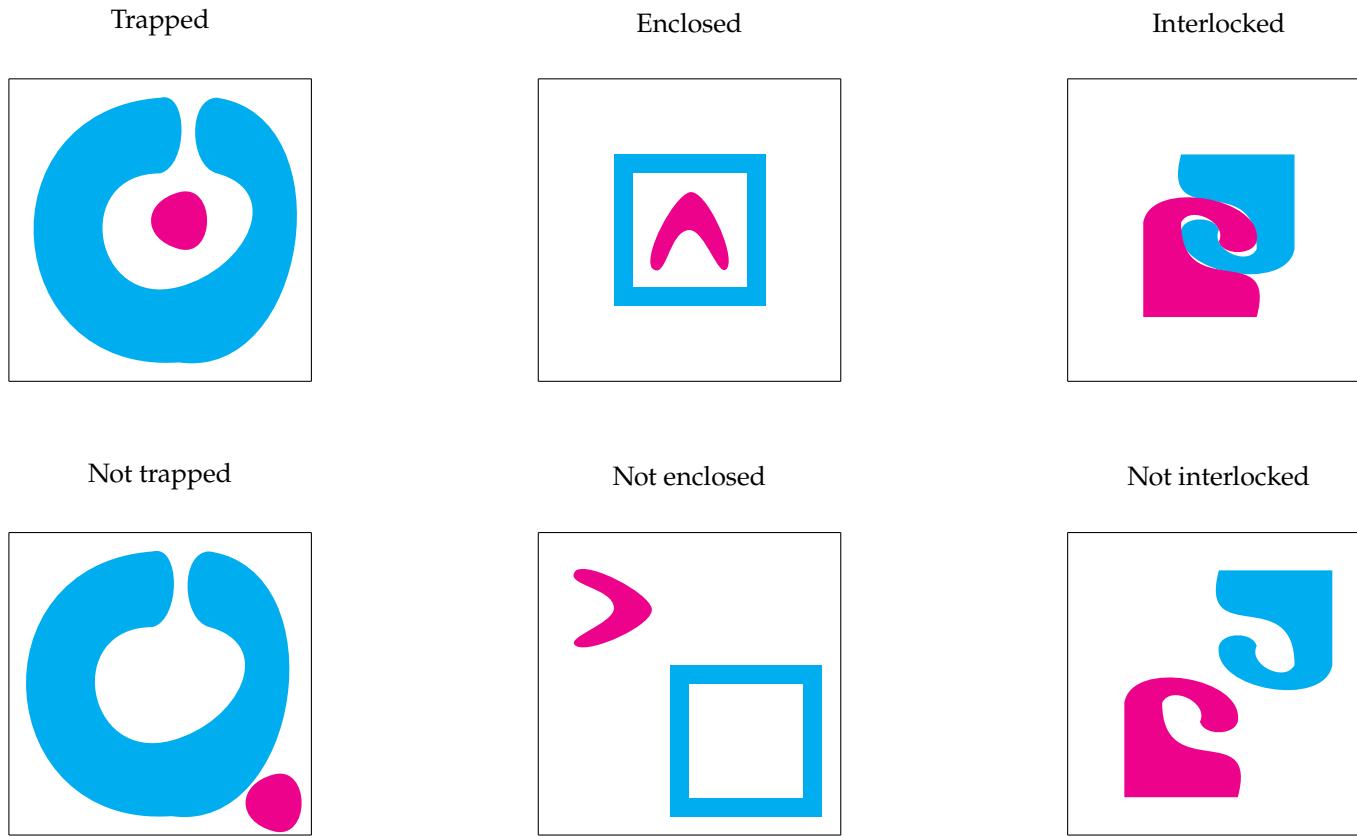


on a sticky spider as we have defined above can be captured as a continuous map from the unit interval to rigid transformations. This is precisely a continuous path in configuration space.



ts of configuration space? Evidently, there are pairs of spiders that are both valid displacements, but not mutually reachable by

ght enclose or trap other shapes, or shapes might be *interlocked*. Depicted below are some pairs of configurations that are mutually



Now we have the conceptual toolkit to begin modelling these concepts in the configuration space of a sticky spider.

5.1.6 Modelling linguistic topological concepts

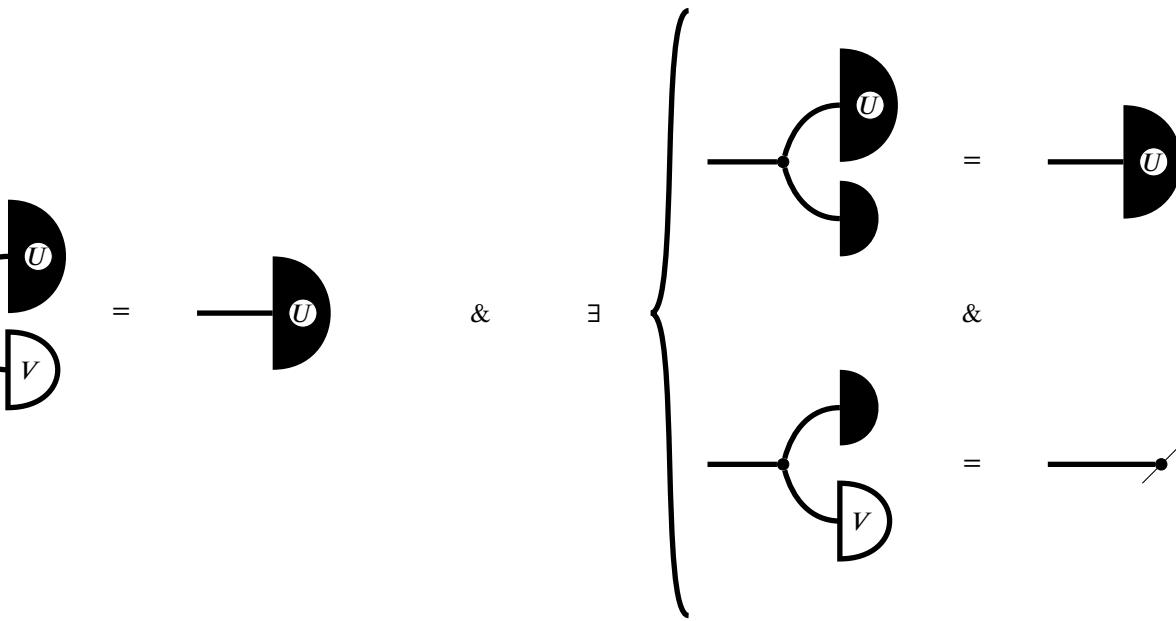
By "linguistic", I mean to refer to the kinds of concepts we use in everyday language. These are concepts that even young children have an intuitive grasp of [], but their formal definitions are difficult to pin down. One such relation modelled here – touching – is in fact a *semantic prime* []: a word that is present in essentially all natural languages that is conceptually primitive, in the sense that it resists definition in simpler terms. It is among the ranks of concepts like *wanting* or *living*, words that are understood by the experience of being human, rather than by school. As such, I make no claim that these definitions are "correct" or "canonical", just that they are good enough to build upon moving forward.

PARTHOOD

Let's say that a "part" refers to an entire simply connected component. Simply connected is already a concept in our toolkit. A shape U is disjoint from another shape V intuitively when we can cover U in a blob with no holes such that the blob has no overlap with V . So, U is a part of V when it is simply connected, wholly con-

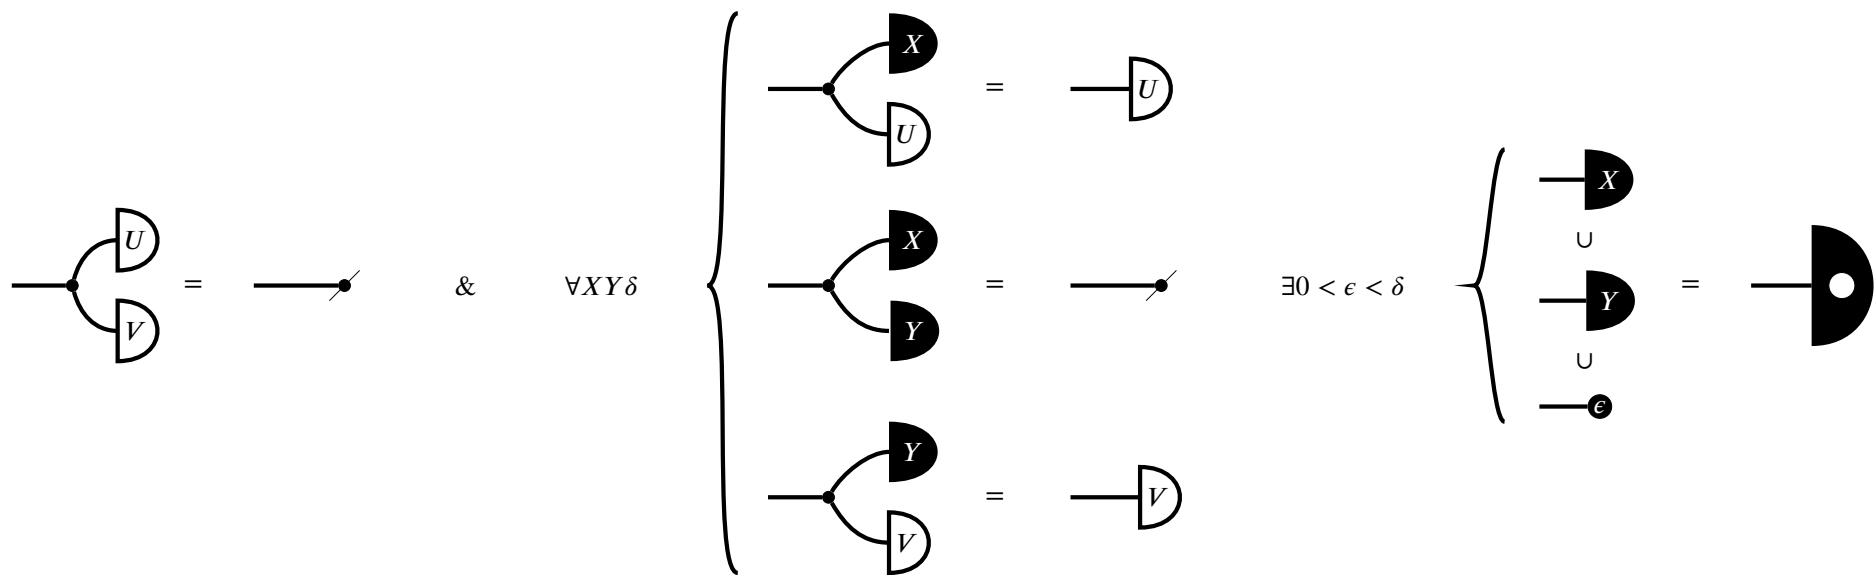
table open that is disjoint from V that covers U . Diagrammatically, this is:

U is a part of V



overlap. Two shapes are "touching" intuitively when they are as close as they can be to each other, somewhere; any closer and they can restrict our attention to the parts of the shape that are touching, and that we can fill in the holes of these parts. At the point gap – just as when we touch things in meatspace, there is a very small gap between us and the object due to the repulsive electric field – when we deal with infinitesimals we borrow the $\epsilon - \delta$ trick from mathematical analysis; for any arbitrarily small δ , we can pick an even smaller ϵ so that if we stick the ball in the gap, the ball forms a bridge that overlaps the two filled-in shapes, which allows us to draw a continuous line through them. This is:

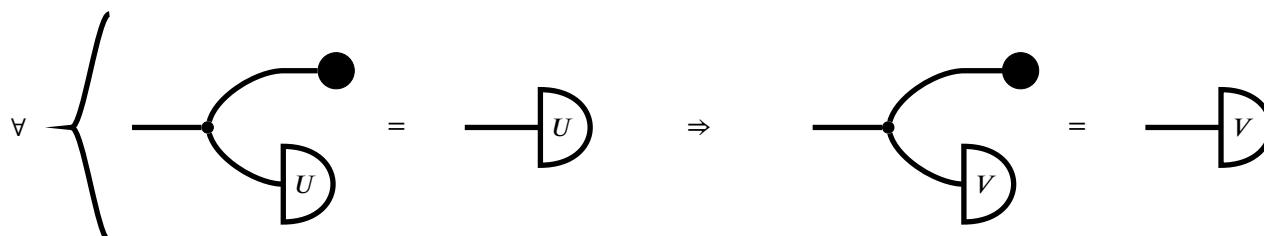
U and V are touching



WITHIN

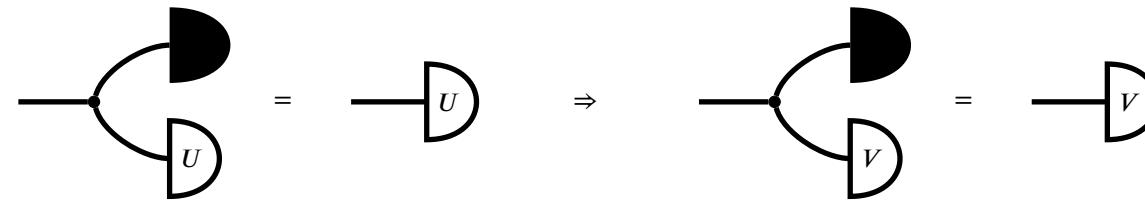
If U surrounds V , or equivalently, if V is within U , then we are saying that leaving V in almost any direction, we will see some of U before we go off to infinity. We can once again use open balls for this purpose, which correspond to possible places you can get to from a starting point x within a distance ϵ . In prose, we are asking that any open ball that contains all of U must also contain all of V .

V is within U , or U surrounds V

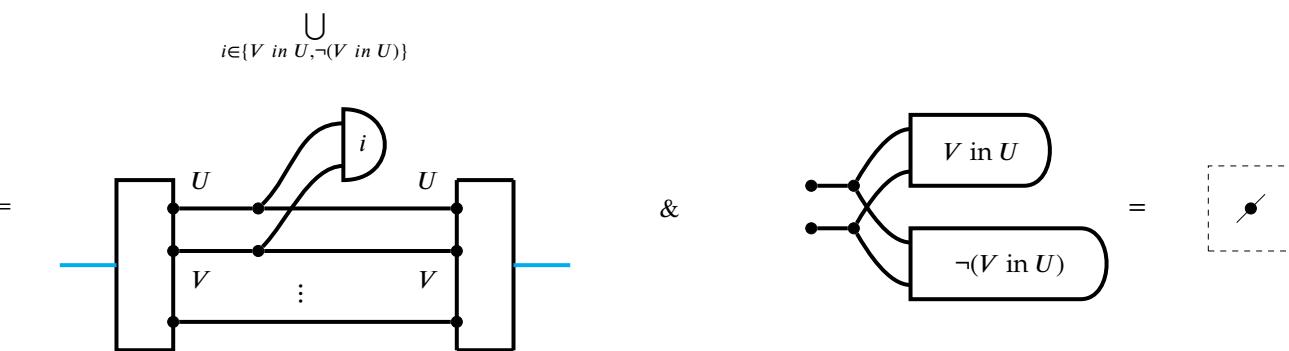


ness, which we will call enclosure. As in when we are in elevators and the door is shut, nothing gets in or out of the container. inner surrounded on all sides, and the contained shape lives within the hole. To give a real-world example, honey lives within a ther the honey is enclosed in the cell depends on whether it is sealed off from air with beeswax. So in prose we are asking that tainer with a blob, that blob must cover the contained shape. Diagrammatically, this amounts to levelling up from open balls in e sets:

U encloses V

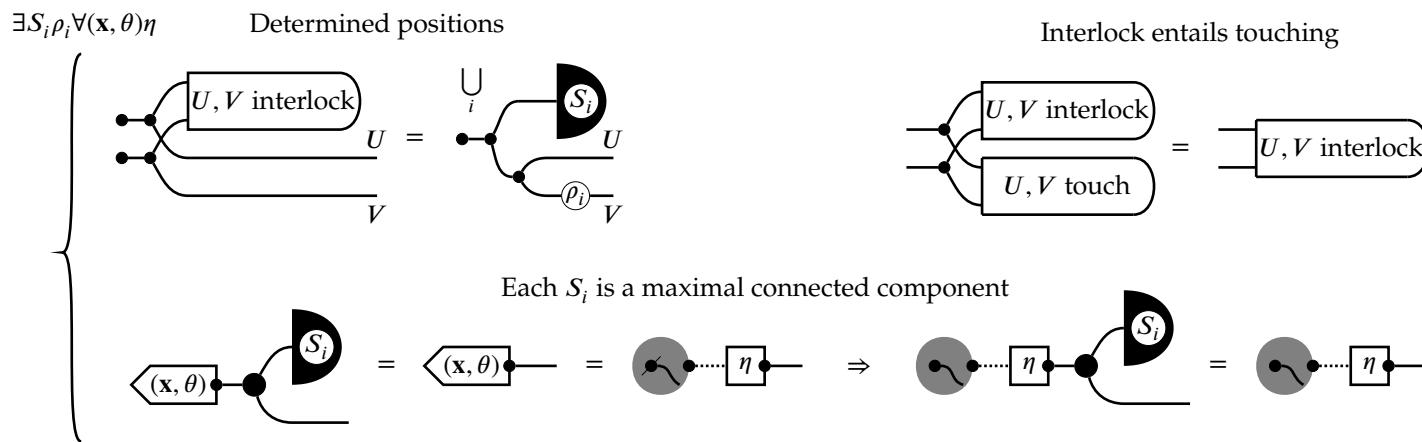


between within-ness and enclosure; for instance, standing in the stonehenge you are surrounded by the pillars, but you can always very close, such as the bars of a jail cell, a human would not be able to leave the trap while still being able to see the outside. es come into play: small animals would still consider it a case of mere within-ness, because they can still walk away between no matter how the pair of objects move rigidly, being trapped means that the trapped V stays within U . In other words, that in all other shapes, we can partition our space of configurations by two concepts, whether V is within U or not, and moreover that . not simply connected – so there is no rigid motion that can allow V to escape from being within U if V starts off trapped inside



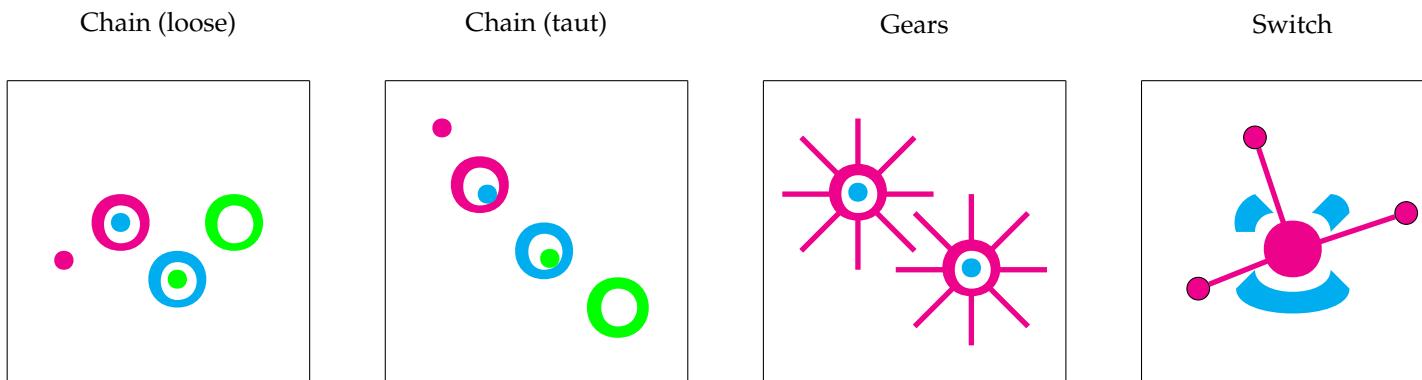
locked without being inside one another. Some potentially familiar examples are plastic models of molecular structure that we

encounter in school, metal lids in cold weather that are too tightly hugging the glass jar, or stubborn Lego pieces that refuse to come apart. The commonality of all these cases is that the two shapes must move together as one, unless deformed or broken. In other words, when two shapes are interlocked, knowing the position in space of one shape determines the position of the other, and this determination is a fixed isometry of space. So we only need to specify a range of positions S for the entire subconfiguration of interlocked shapes U and V , and we may obtain their respective positions by a fixed rigid motion ρ . Since objects may interlock in multiple ways, we may have a sum of these expressions. We additionally observe that interlocking shapes should also be touching, which translates to containment inside the touching concept. Finally, we observe that as in the case of entrapment and enclosure, rigid motions are interlocking-invariant, which translates diagrammatically to the constraint that each S, ρ expression is an entire connected component in configuration space.



CONSTRAINED MOTION

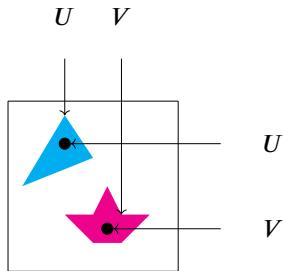
A weaker notion of interlocking is when shapes only imperfectly determine each other's potential displacements, by specifying an allowed range. Here is an understatement: there is some interest in studying how shapes mutually constrain each other's movements in this way.



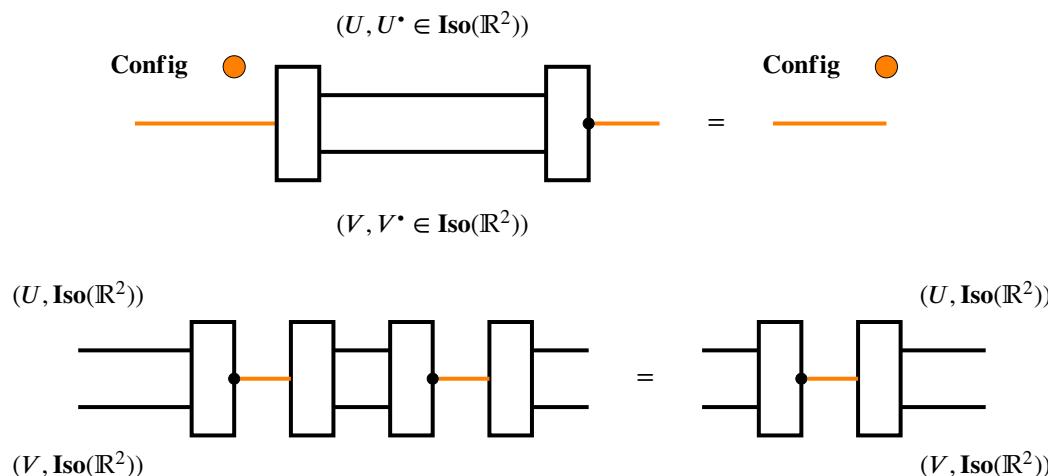
There are as many definitions to go through here as there are potential mechanical models, and among other things, there are mechanically realised clocks [], computers [], and analogues of electric circuits []. So instead, we will allow ourselves to additionally specify open sets as concepts in configuration space that correspond

please, and we assure the reader seeking rigour that blueprints exist for all the mechanisms humans have built. Of course in reality among rigid objects, and directional behaviour is provided by a source of energy, such as gravitational potential, or wound race these sources of energy by a belt that we choose to spin in one direction – our own arrow of time. We postpone discussion of d analogy for a later section.

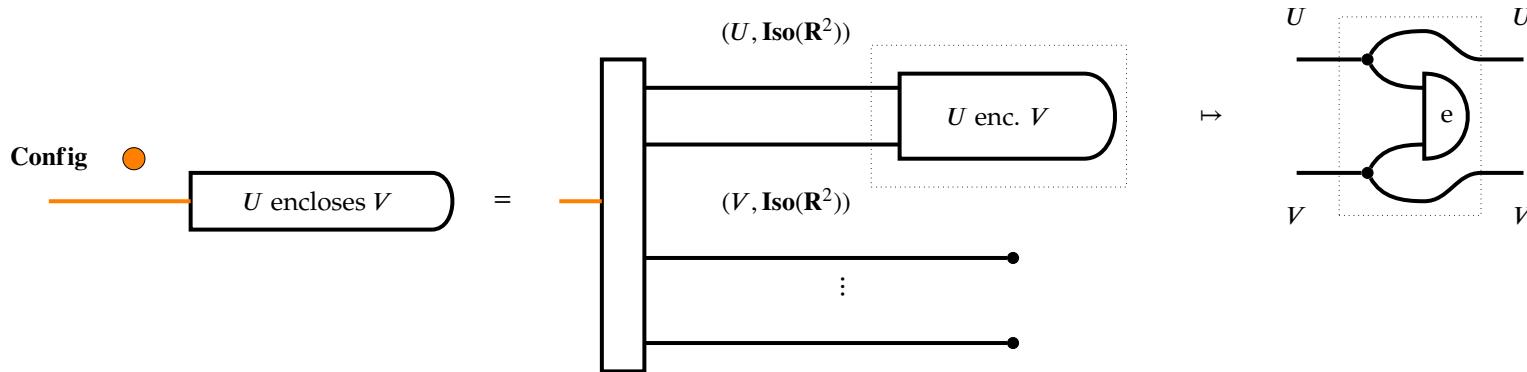
label noun wires: each wire in expanded configuration space must be labelled with the shape within the sticky spider it correctly know how to reconstruct the shapes, since each shape may have a different spatial extent.



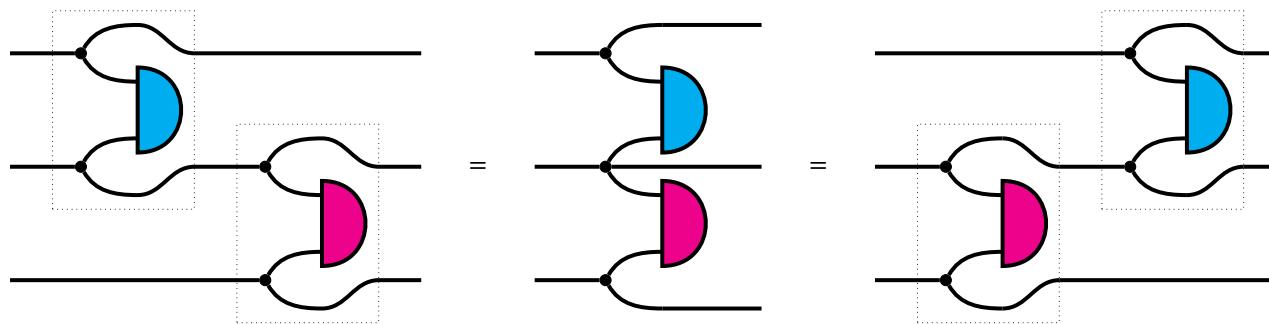
Concretely, for each shape in a sticky spider,
in order to reconstruct the shape,
the data of configuration space
only has to remember a basepoint
(which the isometries act upon)
paired with a label naming the shape
(so that the extent of the shape is reconstructible)



All of the concepts we have defined so far are open sets in configuration space – and for any concept that isn't, we are always free to take the interior of the set; the largest open set contained within the concept. Passing through the split idempotent, we can recast each as a circuit gate using copy maps.



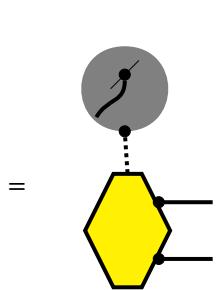
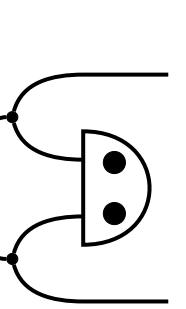
Going forward, we will just label the wires with the names of each shape when necessary. We notice that one feature of this procedure to get gates from open sets is that all gates commute, due to the commutativity of copy.



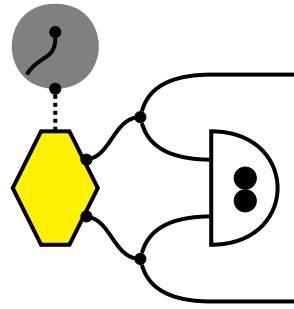
Moreover, since each gate of this form is a restriction to an open set, the gates are idempotent. So the concepts we have defined so far behave as if describing *states* of affairs in space, as if we adding commuting adjectives to space to elaborate detail. For example, *fast red car*, *fast car that is red*, *car is (red and fast)* all mean the same thing. As we add on progressively more concepts, we get diminishing subspaces of configurations in the intersection of all the concepts. So the natural extension is to ask how states of affairs can change with motion. A simple example is the case of *collision*, where two shapes start off not touching, and then they move

to touching.

A particular collision trajectory

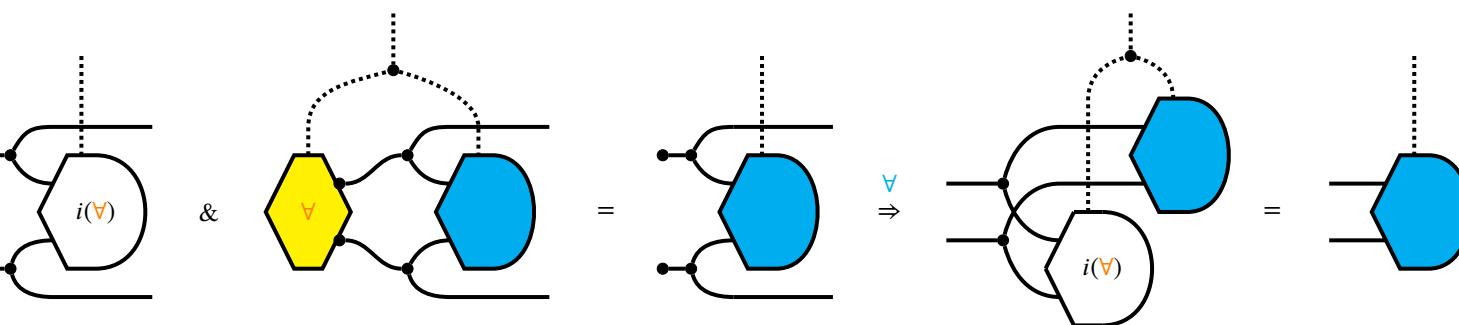


At $t = 0$, the shapes do not touch

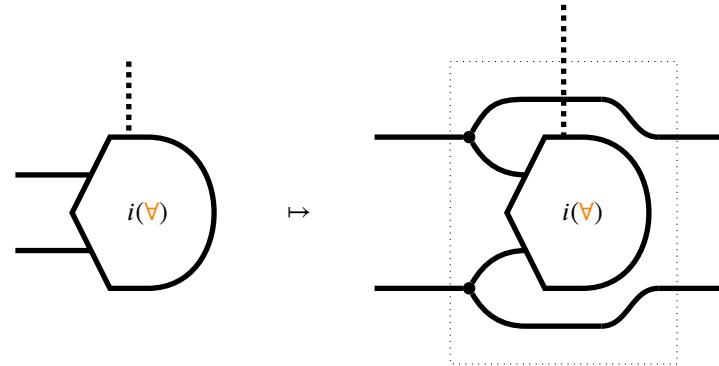


At $t = 1$, the shapes touch

Collisions are the unions of homotopies between maps, we have a homotopy that is the union of all collision trajectories, which we call interior $i(\nabla)$ as the concept of collision; the expressible collection of all particular collisions. But this is not just an open set on the space of shapes, it is a collection of open sets parameterised by homotopy.

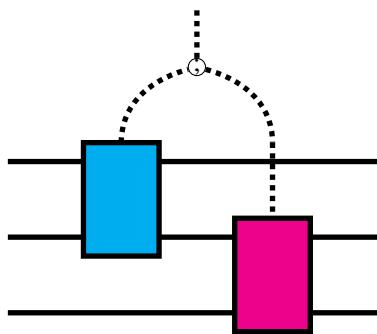


Once we have the open set $i(\nabla)$ that corresponds to all expressible collisions, we have a homotopy-parameterised gate. Following a similar procedure, we can construct gates of motion that satisfy whatever pre- and post-conditions we like.

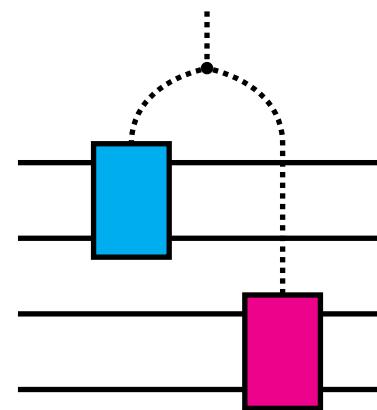


We can compose multiple rigid motions sequentially by a continuous function ; that splits a single unit interval into two: ; $\vdash := x \mapsto \begin{cases} (2x, 0) & \text{if } x \in [0, \frac{1}{2}) \\ (1, 2x - 1) & \text{if } x \in [\frac{1}{2}, 1] \end{cases}$. The effect of the map is to splice two vignettes of the same length together by doubling their speed, then placing them one after the other. We can achieve the same thing without resorting to units of measurement, because recall by Theorem 5.1.7 and by construction that we have access to a map that selects midpoints for us; we will revisit a string-diagrammatic treatment of homotopy and tenses in a later section. We can also compose multiple motions in parallel by copying the unit interval, allowing it to parameterise multiple gates simultaneously.

Sequential composition of motions

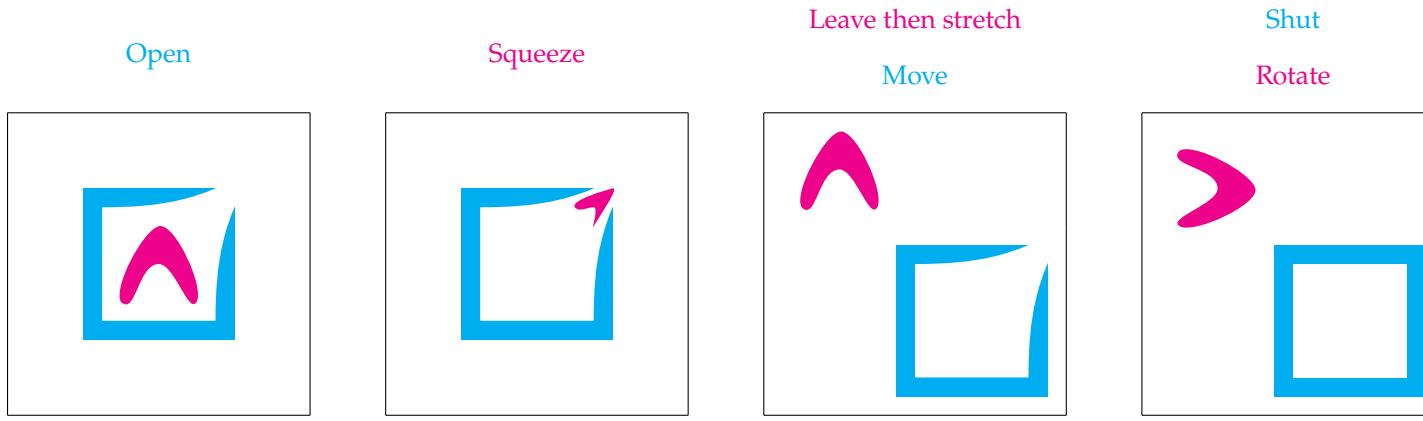


Parallel composition of motions

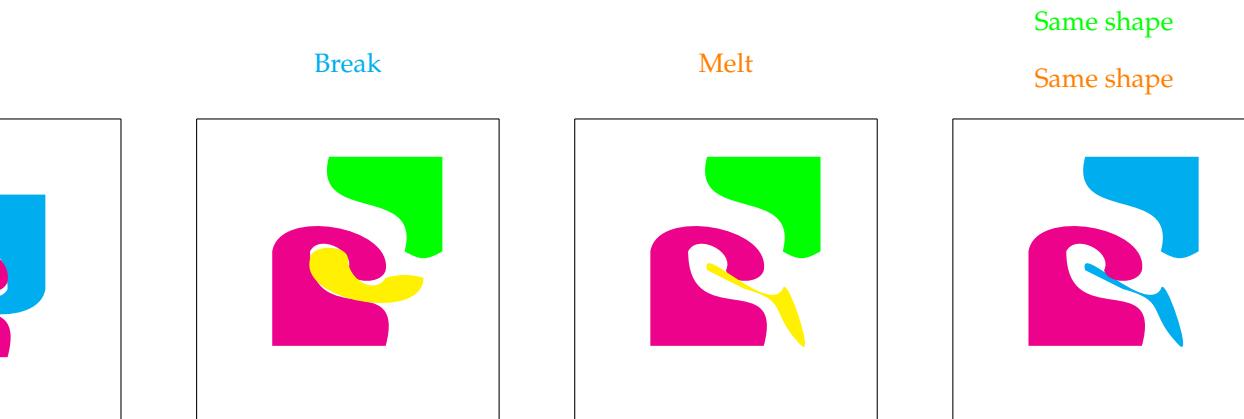


It is easy to see that the gates can always be rewritten to respect the composition order given by ; and copy, since for any input point at the unit interval the gates behave as restrictions to open sets. These new gates do not generally commute; consider comparing the situation where a tenant moves into one apartment and then

tenant reverses the order of the apartments. These are different paths, as the postconditions must be different. So now we have *ns*, or verbs. What kinds of actions are there? In our toy setting, in general we can define actions that arbitrarily change states of rigid motions. The trick to doing this is the observation that arbitrary homotopies allow deformations, so our verb gates allow in the process of a homotopy, as long as at the end they arrive at a rigid displacement of their original form.

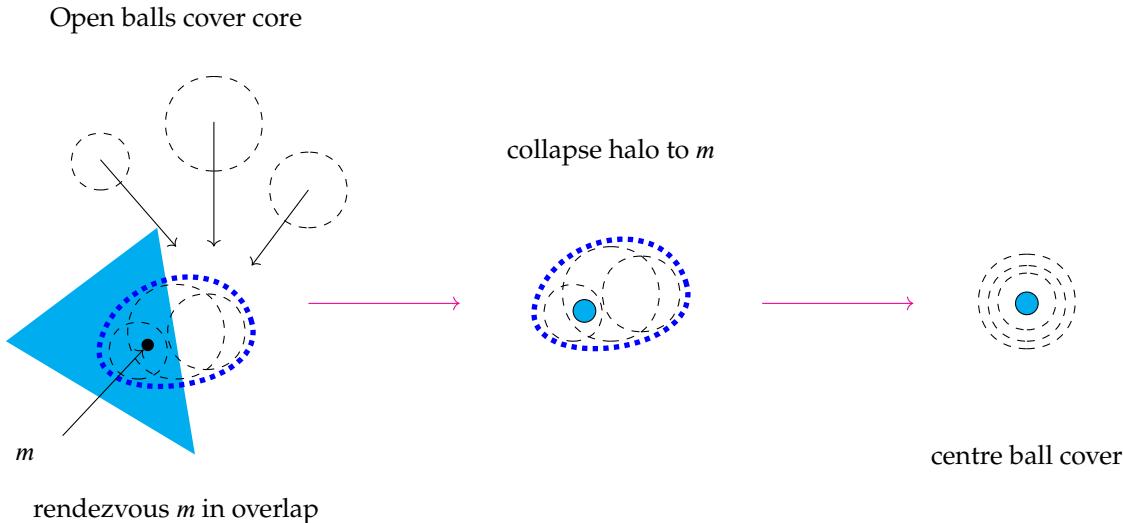


that completely different spiders can be related by homotopy, so we can model a situation where there is a permanent bend, or



as a general recipe to construct homotopies between spiders.

Construction 5.1.11 (Morphing sticky spiders with homotopies). We aim to construct homotopies relating (almost) arbitrary sticky spiders. For now we focus on just changing one shape into another arbitrary one. The idea is as follows. First, we need a cover of open balls $\cup \mathcal{J} = T^0$ and $\cup \mathcal{K} = T^1$ of the start and end cores T^0 and T^1 such that each $k \in T^1$ is expressible as a rigid isometry of some core $j \in \mathcal{J}$; this is so we can slide and rearrange open balls comprising T^0 and reconstruct them as T^1 . As an intermediate step to eliminate holes and unify connected components, we gather all of the balls at a meeting point m (to be determined shortly.) Intuitively we can illustrate this process as follows:



Second, in order to perform the sliding of open balls, we observe that, given a basepoint to act as origin (which we assume is provided by the data of the split idempotent of configuration space) we can express the group action of rigid isometries $\text{Iso}(\mathbb{R}^2)$ on \mathbb{R}^2 as a continuous function:

$$\begin{array}{ccc} \text{Iso}(\mathbb{R}^2) & & \\ \xrightarrow{\quad} & \boxed{\quad} & \xleftarrow{\quad} \\ \mathbb{R}^2 & & \mathbb{R}^2 \end{array} \quad ((\mathbf{a}, \theta), \mathbf{b}) \mapsto \begin{bmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{bmatrix} \mathbf{b} + \mathbf{a}$$

Third, before we begin sliding the open balls, we must ensure that the halo of the shape cooperates. We observe that a given shape i in a sticky spider may be expressed as the union of a family of constant continuous partial functions in the following way. Given an open cover \mathcal{J} such that $\cup \mathcal{J} = T_i$, where T_i is the core of the shape i , each function is a constant map from some $T_j \in \mathcal{J}$ to some point $x \in S_i$, where S_i is the halo of the shape i . For each $T_j \in \mathcal{J}$ and every point $x \in S_i$, the

, to x is in the family.

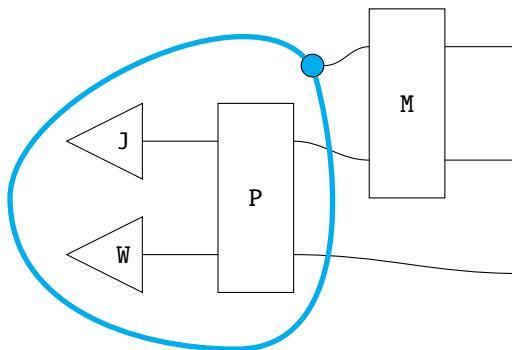
$$\begin{array}{c} \text{Diagram showing two shapes } T_i \text{ and } S_i \text{ connected by a blue dot labeled } x. \\ = \\ \text{Diagram showing a large shape } T_j \text{ containing a smaller shape } S_i, with a blue dot labeled } x \text{ at a point on the boundary of } T_j. \end{array}$$

$$\begin{gathered}
 \text{Diagram showing two shapes } T_i \text{ and } S_i \text{ connected by a blue dot labeled } x. \\
 = \\
 \text{Diagram showing a large shape } T_j \text{ containing a smaller shape } S_i, with a blue dot labeled } x \text{ at a point on the boundary of } T_j. \\
 \text{The top diagram has labels } T_i \text{ and } S_i \text{ near their respective shapes. The bottom diagram has labels } T_j \text{ and } S_i \text{ near their respective shapes. The blue dot } x \text{ is located at the junction between the two shapes in both diagrams.} \\
 \text{The top diagram is followed by an equals sign, then the bottom diagram. Above the top diagram is a union symbol } \bigcup_{T_j \in J_i \subseteq \tau : \cup J_i = T_i} \text{ and below the bottom diagram is a union symbol } \bigcup_{x \in S_i}.
 \end{gathered}$$

must exist some point m that is in both the core and the halo: we pick such a point as the rendezvous for the open balls. For each a homotopy that varies only the image point x continuously in the space to finish at m . Now we can slide the open balls to the reversible by the continuous map $t \mapsto (1 - t)$ on the interval, we can perform the above steps for shapes T^0 and T^1 to finish at ccess for T^1 and composing sequentially to obtain a finished transformation. The final wrinkle to address is when dealing with ion conditions ?? for shapes, it may be that parts of one shape are enclosed in another, so the processes must be coordinated ample, the enclosing shape must be first opened, so that the enclosed shape may leave. I will keep it an article of faith that such e up with a proof that all spiders \mathbf{R}^2 are mutually transformable by homotopy in this (or any other) way, so that will remain a deal of spiders are mutually transformable; almost certainly any we would care to draw. So this will just be a construction for

5.2 On entification, general anaphora, computers, and lassos.

ENTIFICATION IS THE PROCESS OF TURNING WORDS AND PHRASES THAT AREN'T NOUNS INTO NOUNS. We are familiar with morphological operations in English, such as *inflections* that turn the singular *cat* into the plural *cats*, by adding a suffix *-s*. Another morphological operation, generally classed *derivation*, turns words from one category into another, for example the adjective *happy* into the noun *happiness*. With suffixes such as *-ness* and *-ing*, just about any lexical word in English can be turned into a noun, as if lexical words have some semantic content that is independent of the grammatical categories they might wear. I'll call this process *entification*, which extends beyond morphology towards more complicated constructions such as a prefix the fact that that converts sentences into noun-like entities, insofar as these entities can be referred to by anaphora: for example, in the sentence *Jono was paid minimum wage but he didn't mind it*, it may be argued that *it* refers to the fact that *Jono was paid minimum wage*. Graphically, we might want to depict the gloss as a circuit with a lasso that gives another noun-wire:



A MATHEMATICAL MODELLING PROBLEM FOR SEMANTICISTS ARISES WHEN ANYTHING CAN BE A NOUN WIRE. The problem at hand is finding the right mathematical setting to interpret and calculate with such lassos. In principle, any meaningful (possibly composite) part of text can be referred to as if it were a noun. For syntax, this is a boon; having entification around means that there is no need to extend the system to accommodate wires for anything apart from nouns, so long as there is a gadget that can turn things into nouns and back. For semantics this is a challenge, since this requires noun-wires to "have enough space in them" to accommodate full circuits operating on other noun-wires, which suggests a very structured sort of infinity.

COMPUTER SCIENCE HAS HAD A PERFECTLY SERVICEABLE MODEL OF THIS KIND OF NOUN-WIRE FOR A LONG TIME. What separates a computer from other kinds of machine is that a computer can do whatever any other kind of machine could do – modulo church-turing on computability and the domain of data manipulation – so long as the computer is running the right *program*. Programs are (for our purposes) processes that manipulate variously formatted – or typed – data, such as integers, sounds, and images. They can operate in sequence and in parallel, and wires can be swapped over each other, so programs form a process theory, where we can reason about the extensional equivalence of different programs – whether two programs behave the same with respect to mapping inputs to outputs. This aim of reasoning about programs in an implementation-independent fashion contributed to the birth of computer science from programming, which was attended by the independent discovery of proto-string-diagrams in the form of flowcharts.

SPECIAL IS THAT ON REAL COMPUTERS, THEY ARE SPECIFIED BY CODE. Code is just another format of data. Programs that are equivalent have many different implementations in code: for example, there are many sorting algorithms, though all of them map the inversely, every possible program in a process theory of programs must have some implementation as code. Diagrammatically, we this: for every pair of input formats and output formats (A, B) , there is a computer for that format $ev\{A_B : A \otimes \Xi \rightarrow B\}$, which denote Ξ going forward) as an additional input, and for every possible program $f : A \rightarrow B$, there exists a state $\lceil f \rceil : I \rightarrow \Xi$

$$\forall A, B \in Ob(C) \exists ev\{A_B : A \otimes \Xi \rightarrow B\} \forall f : A \rightarrow B \exists \lceil f \rceil : I \rightarrow \Xi \quad (5.1)$$

$$(5.2)$$

erises computers as code-evaluators, provides a plan of attack for the semantic modelling problem of entification: if we take monoidal computer, we have restricted the candidate symmetric monoidal categories to model text-circuits in a way that allows

we could have made is that since computers really just manipulate code, every data format is a kind of restricted form of the to be a mathematical consequence of the above equation (and the presence of a few other operations such as copy and compare (ra), demonstrated in Pavlovic's forthcoming monoidal computer book [], itself a crystallisation of three monoidal computer pa Cockett's work on Turing categories []. Both approaches to a categorical formulation of computability theory share the common closure (monoidal closure in the case of monoidal computer and exponentiation in Turing categories) where rather than having or $B\{A$, there is a single "code-object" Ξ . They differ in the ambient setting; Pavlovic works in the generic symmetric monoidal restriction categories, which generalise partial functions. I work in Pavlovic's formalism because I prefer string diagrams to com

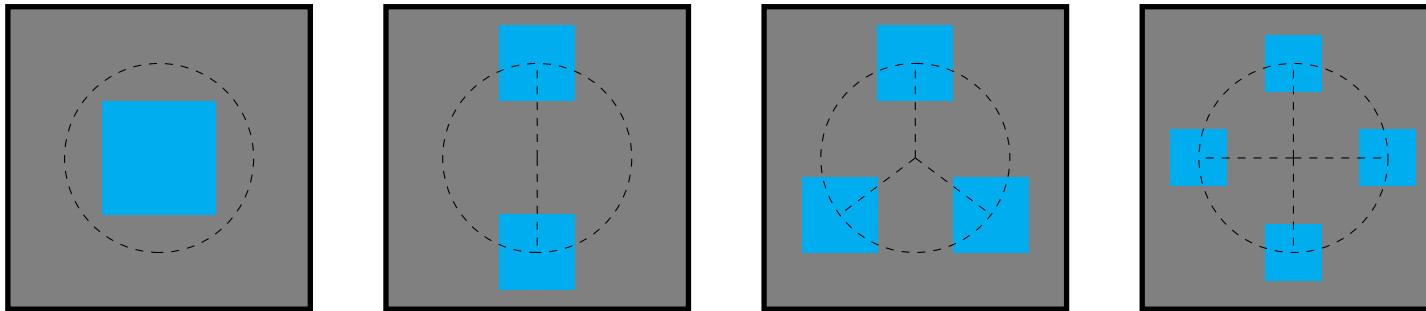
conclude that any programming language is a model for text circuits, using the code data format as the noun wire. Since semanti the following construction.

the open unit square model the category relations between countable sets equipped with a code object). Using the open unit code object, there is a subcategory of **ContRel** which behaves as the category of countable sets and relations equipped with

through any countable set X). For any countable set X , the open unit square \blacksquare has a sticky spider that splits through X^* .

We'll assume the sticky-spiders to be mereologies, so that cores and halos agree. Then we only have to highlight the copiable axis-aligned open squares evenly along them, one for each element of X . The centres of the open squares lie on the circumfer-

ence of the circle, and we may shrink each square as needed to fit all of them.



□

Definition 5.2.4 (Morphism of sticky spiders). A morphism between sticky spiders is any morphism that satisfies the following equation.

$$\text{Diagram: } \begin{array}{c} \text{---} \bullet \text{---} \square \text{---} \bullet \text{---} \text{---} \\ = \\ \text{---} \square \text{---} \text{---} \end{array}$$

Proposition 5.2.5 (Morphisms of sticky spiders encode relations). For arbitrary split idempotents through A^* and B^* , the morphisms between the two resulting sticky spiders are in bijection with relations $R : A \rightarrow B$.

$$\begin{array}{ccc}
 A^* & & \\
 \text{Diagram: } \begin{array}{c} \text{---} \bullet \text{---} \square \text{---} \bullet \text{---} \text{---} \\ = \\ \text{---} \square \text{---} \text{---} \end{array} & & \\
 \forall & : \mathbf{Rel}(A, B) \ni R \leftrightarrow & \text{Diagram: } \begin{array}{c} \text{---} \bullet \text{---} \square \text{---} \square \text{---} \bullet \text{---} \text{---} \\ \simeq \\ \text{---} \bullet \text{---} \square \text{---} \bullet \text{---} \text{---} \\ = \\ \text{---} \square \text{---} \square \text{---} \end{array} \\
 B^* & &
 \end{array}$$

(\Leftarrow) : Every morphism of sticky spiders corresponds to a relation between sets.

$$\begin{array}{c}
 \text{Diagram 1: } \text{A horizontal line with a blue dot at the left end and a pink dot at the right end, passing through a black square labeled } R' \text{.} \\
 = \quad \bigcup_{\text{blue diamond} \cup \text{orange diamond}} \text{Diagram 2: } \text{A horizontal line with a blue dot at the left end and a pink dot at the right end, passing through two blue triangles (one pointing up, one pointing down) and a black square, followed by two orange triangles (one pointing up, one pointing down).} \\
 \\
 = \quad \bigcup_{\text{blue diamond} \cup \text{orange diamond}} \text{Diagram 3: } \text{A horizontal line with a blue dot at the left end and a pink dot at the right end, passing through a blue triangle (pointing up) and an orange triangle (pointing down).} \\
 = \quad \boxed{\quad}
 \end{array}$$

Since (co)copiables are distinct, we may uniquely reindex as:

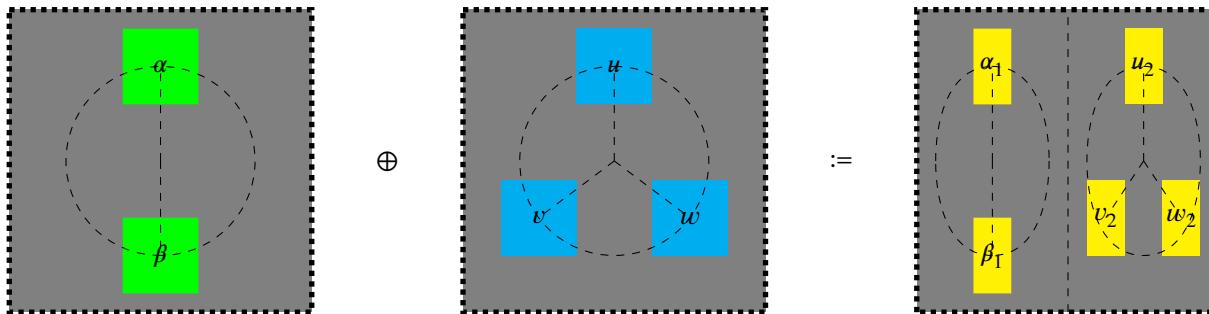
$$= \bigcup_{(a, b) \in R \subseteq A \times B} \text{Diagram 4: } \text{A horizontal line with a blue dot at the left end and a pink dot at the right end, passing through a blue triangle labeled } a \text{ and an orange triangle labeled } b.$$

By idempotence of (co)copiables, every relation $R \subseteq A \times B$ corresponds to a morphism of sticky spiders.

$$\boxed{\quad} = \bigcup_{(a, b) \in R} \text{Diagram 5: } \text{A horizontal line with a blue dot at the left end and a pink dot at the right end, passing through a blue triangle labeled } a \text{ and an orange triangle labeled } b.$$

□

Construction 5.2.6 (Representing sets in their various guises within $\boxed{\cdot}$). We can represent the direct sum of two $\boxed{\cdot}$ -representations of sets as follows.



The important bit of technology is the homeomorphism that losslessly squishes the whole unit square into one half of the unit square. The decompressions are partial continuous functions, with domain restricted to the appropriate half of the unit square.

$$(x, y) \mapsto (\frac{x}{2}, y)$$

$$(x, y) \mapsto (\frac{x+1}{2}, y)$$

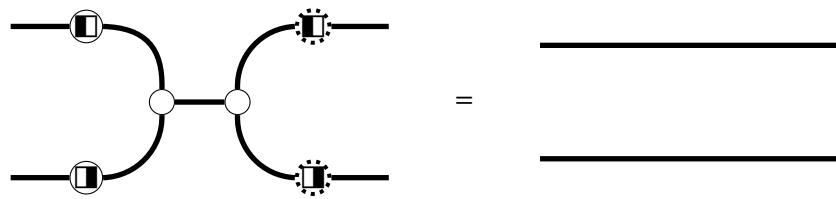
$$(x, y)|_{x < 1} \mapsto (2x, y)$$

$$(x, y)|_{x > 1} \mapsto (2x - 1, y)$$

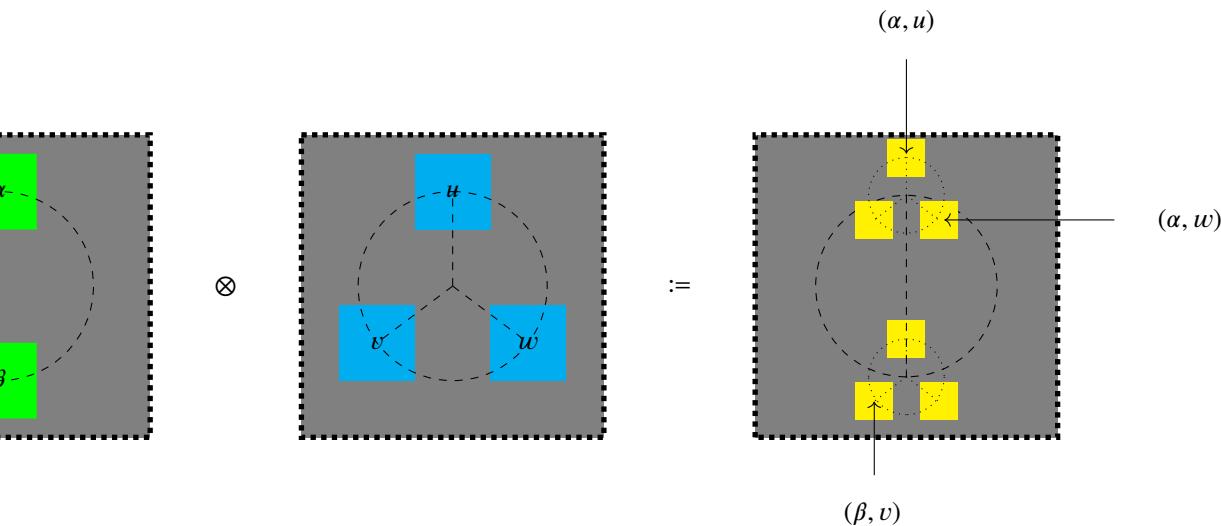
We express the ability of these relations to encode and decode the unit square in just either half by the following graphical equations.

Now, to put the two halves together and to take them apart, we introduce the following two relations. In tandem with the squishing and stretching we have defined, these will behave just as the projections and injections for the direct-sum biproduct in **Rel**.

The following equation tells us that we can take any two representations in $\boxed{\cdot}$, put them into a single copy of $\boxed{\cdot}$, and take them out again. Banach and Tarski would

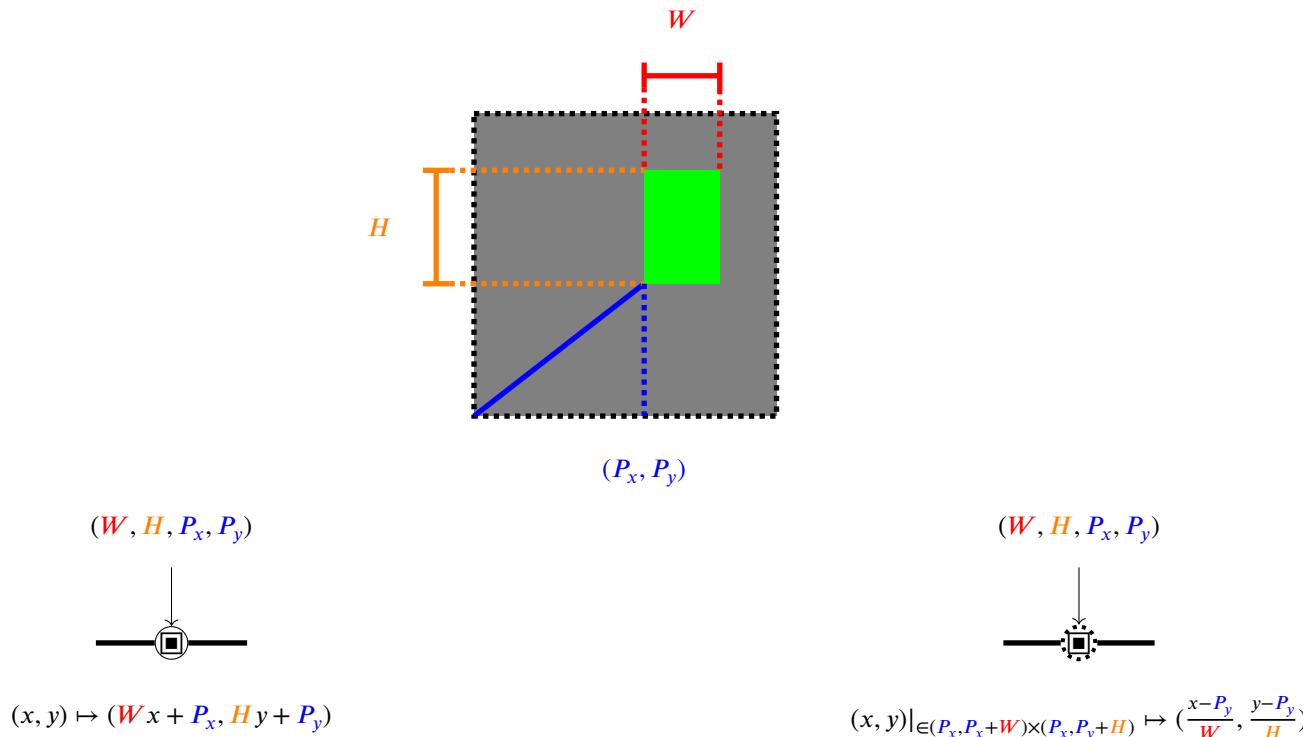


of representations by placing copies of B in each of the open boxes of A .



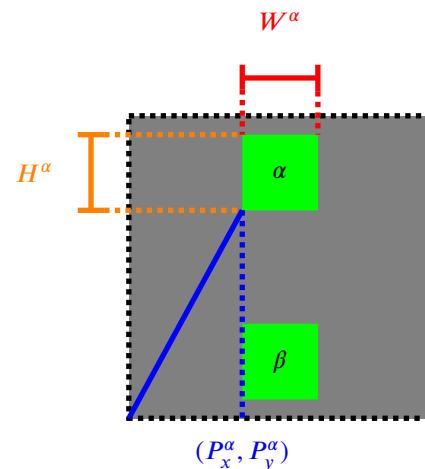
is a family of homeomorphisms of \blacksquare parameterised by axis-aligned open boxes. We depict the parameters outside the body of

the homeomorphism for clarity. The squish is on the left, the stretch on the right.



Now, for every representation of a set in \blacksquare by a sticky-spider, where each element corresponds to an axis-aligned open box, we can associate each element with a

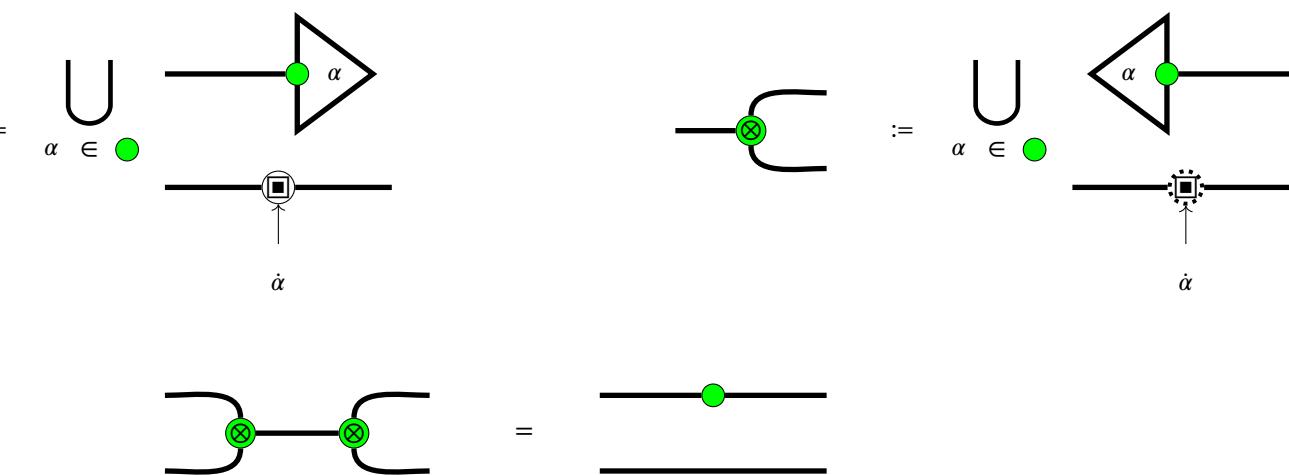
the parameters of the open box, which we notate with a dot above the name of the element



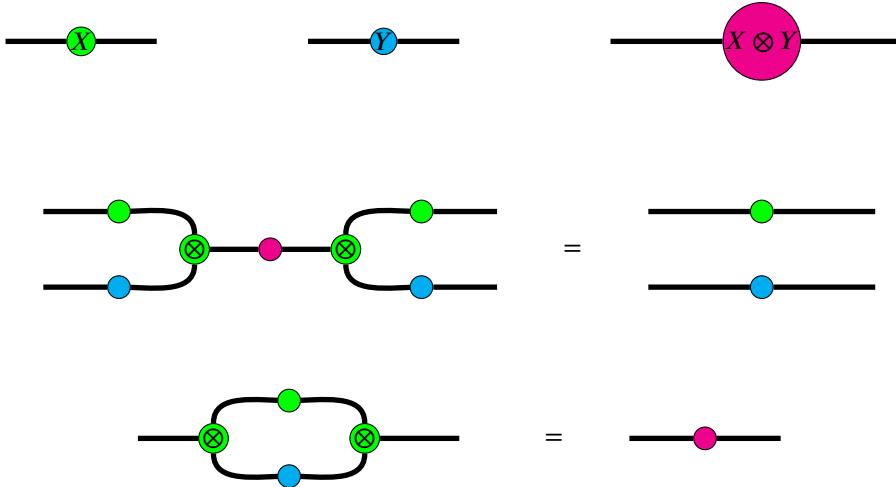
$$\dot{\alpha} = (W^\alpha, H^\alpha, P_x^\alpha, P_y^\alpha)$$

$$\dot{\beta} = (\textcolor{red}{W^\beta}, \textcolor{orange}{H^\beta}, P_x^\beta, P_y^\beta)$$

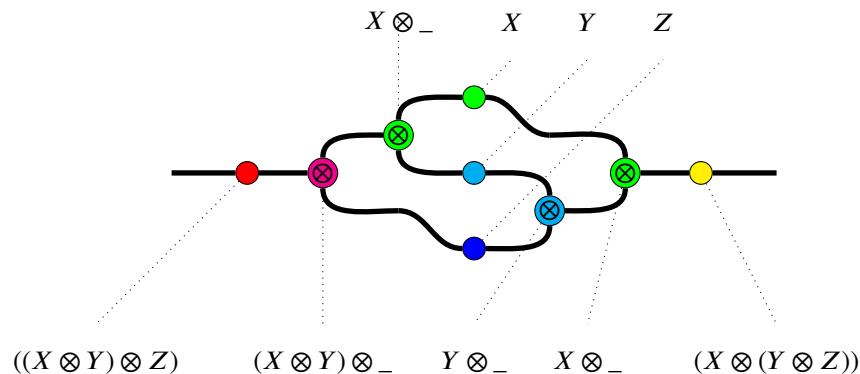
"the left" relation $_ \rightarrow X \otimes _$ and its corresponding cotensor



The tensor and cotensor behave as we expect from proof nets for monoidal categories.

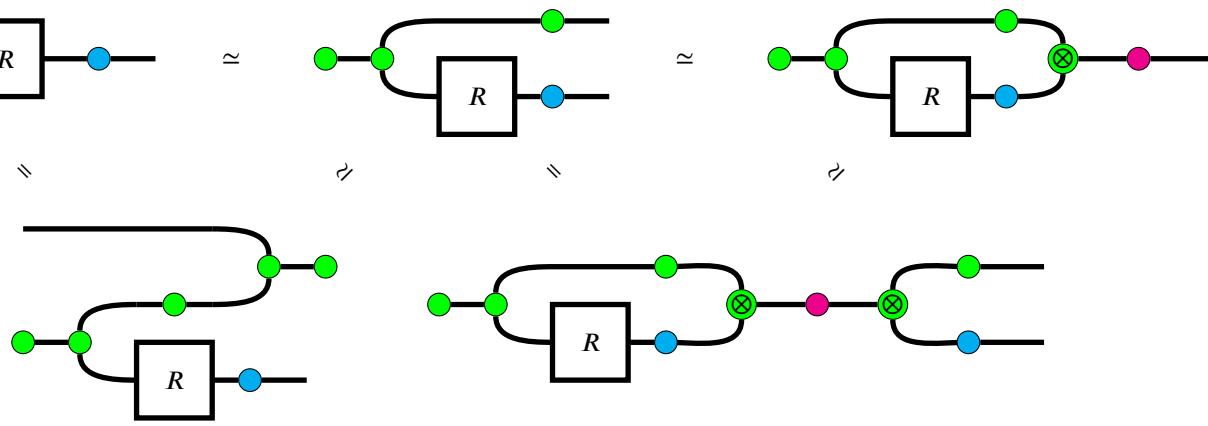


And by construction, the (co)tensors and (co)pluses interact as we expect, and they come with all the natural isomorphisms between representations we expect. For example, below we exhibit an explicit associator natural isomorphism between representations.

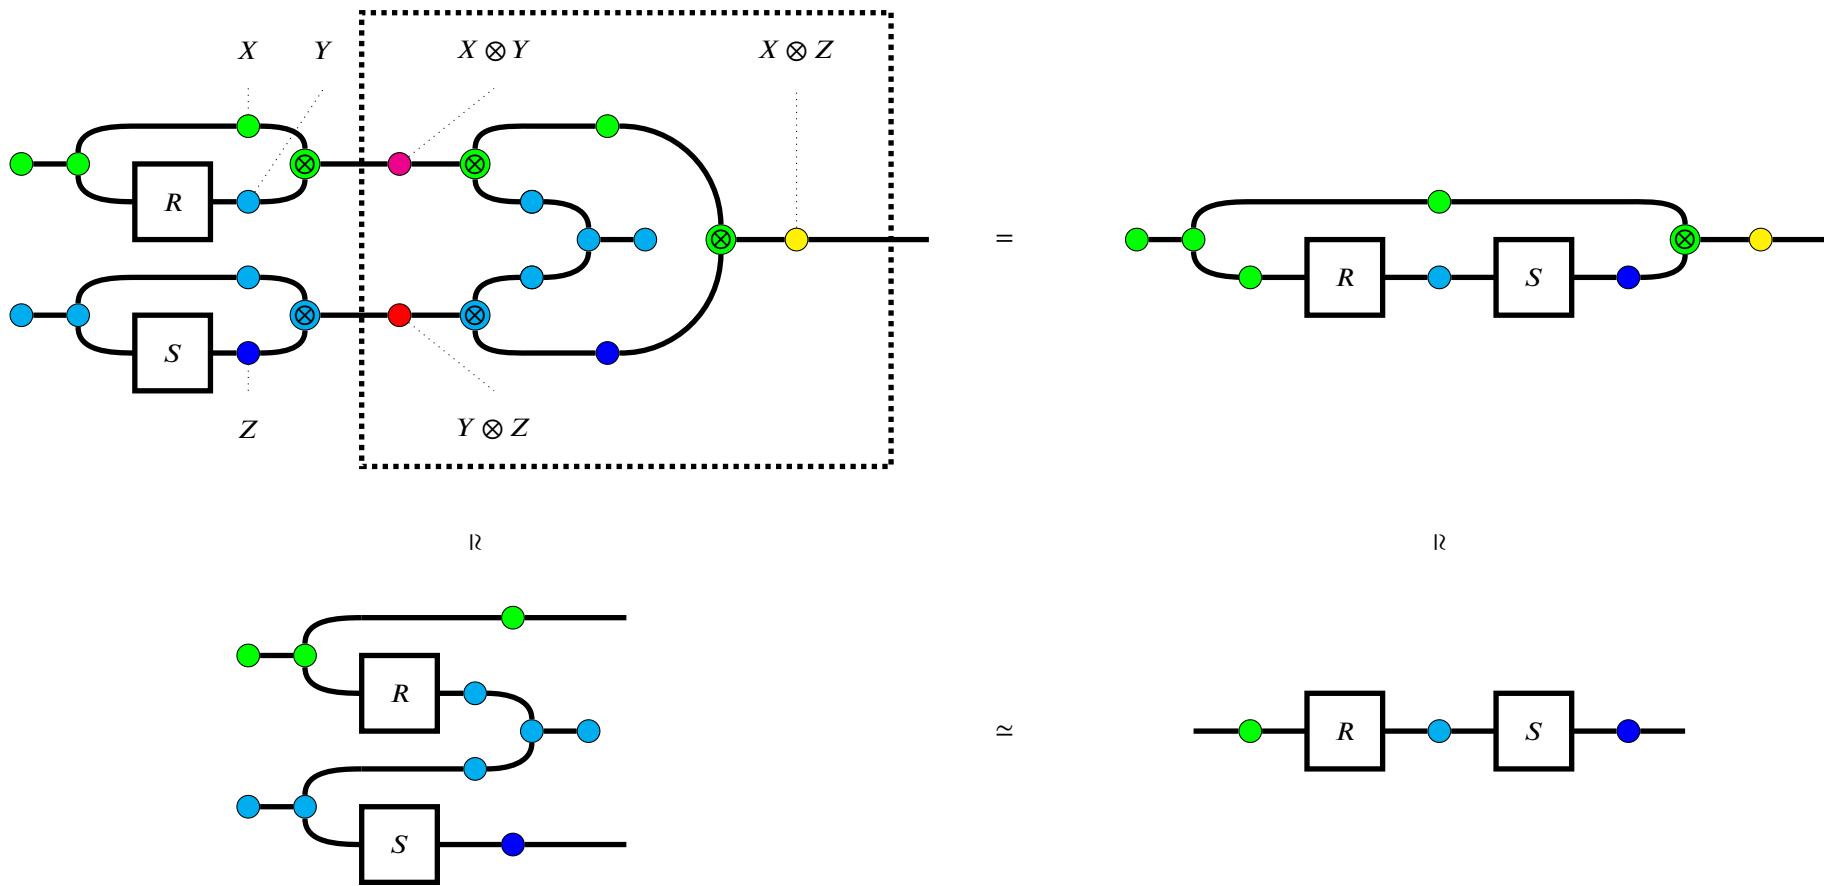


Construction 5.2.7 (Representing relations between sets and their composition within $\boxed{\quad}$). With all the above, we can establish a special kind of process-state duality;

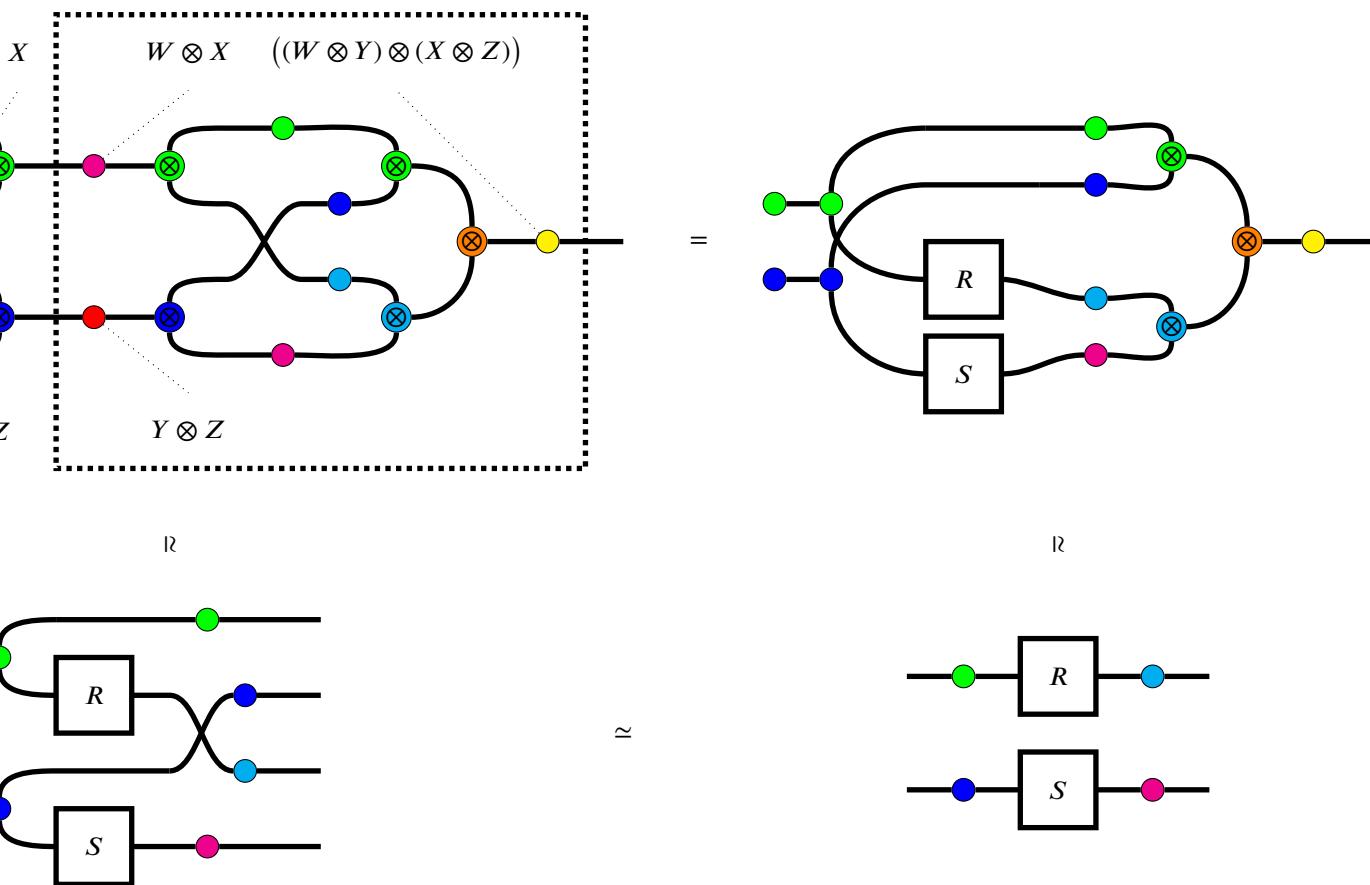
to states of $\blacksquare\blacksquare$, up to the representation scheme we have chosen.



Moreover, we have continuous relations that perform sequential composition of relations.



parallel composition of relations by tensors.



5.3 Neural nets, and (im)possibility results for learning text circuits from data

So far, we have been working process-theoretically, using equations between processes to specify their behaviour. It is natural to ask whether it is possible to realise each process in a process theory as a neural net, using the equations as training criteria so that the neural nets jointly model a process theory. This approach is worth pursuing to combine the benefits and ease of data-driven learning with the modularity and explainability benefits of process theories. Moreover, the onus is on us to demonstrate that text circuits can be learnt in this way, or else we would be no better off in terms of a practical theory of language for the age of big data.

In this sketch, we introduce neural nets diagrammatically, along with the *Universal Approximation Theorem*, which, along with variants for different architectures, states that for any dimension m and any $\epsilon > 0$, there exists a neural net that approximates any continuous function $\mathbb{R}^m \rightarrow \mathbb{R}$ on a compact subset of the domain \mathbb{R}^m within a discrepancy of ϵ . Then we introduce the notion of approximability for PROPs, and we observe that not all PROPs are approximable in terms of smooth functions of the form given by the universal approximation theorem. So we restrict our attention to PROPs for basic text circuits, which we demonstrate are suitable for certain learning tasks. We prove that basic text circuit PROPs of bounded depth and width – a notion we will define – are approximable; in other words, that text circuits work in principle alongside data-driven techniques. We close with a discussion of limitations and extensions. We give a corollary that finitely generated subcategories of **FinSet** are realisable as ensembles of deterministic neural nets, and we show how introducing probabilistic states extends the situation to **FinRel**. We posit a no-go conjecture on the basis of an observed tension between the space-resource demands of deterministic representations and unbounded compositionality.

5.3.1 Deterministic Neural Nets and Closed Monoidal Categories

For those unfamiliar, it is worth understanding the fundamentals of how neural nets work, which illuminates how they are able to transmute large amounts of input data into strong and "theory-free" [] performance at tasks. The ability to convert data (in which we are drowning), along with the universal approximation theorem, explain why data-driven learning methods have recently outperformed handcrafted rules-based approaches to artificial intelligence. I will focus only on the basic deep neural net – so the reader is warned that the state of the art in terms of learning architectures is far more sophisticated, such as transformers [] used in large language models. There are only two important takeaways from this section. First is that the reader must understand why it is that data-driven learning methods outperform human axiomatising when dealing with complex problem domains, and are therefore (with the usual caveats of explainability) preferable. Second is an understanding of the statement of the universal approximation theorem, which will reappear in Section 5.3.5, where it will be of service in an attempt to introduce compositionality via interacting ensembles of neural nets.

NEURAL NETS ARISE FROM A TOY MODEL OF BIOLOGICAL NEURONS. At a glance, biological neurons have many receptors and one output, and the neuron fires a signal out if its combined inputs exceed an activation threshold. As a simplification, McCulloch-Pitts neurons are a sum of n inputs passed through an activation function $\sigma : \mathbf{R}^n \rightarrow \mathbf{R}$ that is permitted to be nonlinear, but traditionally monotone increasing and sigmoidal, which bounds the range of the function $\exists a_{\mathbb{R}} b_{\mathbb{R}} \forall x_{\mathbb{R}} : a \leq \sigma(x) \leq b$, and asks that σ approaches the lower and upper bounds in the limit as x goes to $-\infty$ and ∞ respectively. Using the diagrammatic calculus for linear algebra [] equipped with a nonlinear activation function (interpretable in sets and functions or euclidean spaces and smooth maps), we can immediately grasp a visual resemblance between the designs of nature and man:

placeholder

THE FIRST USE OF NEURAL NETS WAS IN APPLICATION TO THE PROBLEM OF MACHINE VISION. These first, single-layer neural nets were called *perceptrons*. Mimicking the neuronal organisation of the visual cortex, it was a sensible idea to stack these layers on top of one another [] – these layers are the original reason for the word "deep" in "deep learning", but words change in meaning over time.

placeholder

THE MODERN UBIQUITY OF NEURAL NETS IS DUE TO SEVERAL FACTORS. First is Hinton's backpropagation algorithm [] (which may be obsolete when you are reading this by Hinton's forward-forward second salvo []). Observe that even after one has decided on the shape of the neural net in terms of neuronal connectivity, there are still many degrees of freedom in the parameters of the activation functions, in particular their horizontal shift (bias) and vertical stretching (weights). Borrowing diagrammatic notation for parameters as orthogonal wires from [], we can depict the degrees of freedom for a single neuron like this:

placeholder

THERE IS A MASSIVE SPACE OF PARAMETERS TO SET FOR EVEN A MODERATELY SIZED NEURAL NET. So how do we set the parameters in such a way that the neural net computes something useful? Backpropagation solves this problem by leveraging the shape of a neural net. There are many easily searchable resources that cover backpropagation for the interested reader, including category-theoretic ones []. The simple explanation goes like this. Let's just focus on the weight parameter of each neuron. By analogy each neuron is a shitty person, and their weight is how strongly they hold a binary opinion. A neural net by analogy is a shitty rigid hierarchical society with voters in the back and decision makers in the front. As a simple example, Alice and Bob each make a recommendation to Claire based on what they receive as input.

placeholder

Suppose that Claire's decision is wrong. She revises her own opinion then meets with her confidantes. Alice's recommendation was faulty, so Claire blames her; as a narcissistic defense, the viciousness of the blame is proportional to how wrong Claire was. Alice revises her own opinion proportional how mean Claire is being, and then Alice goes to seek out her confidantes to perpetuate a vicious cycle of psychological violence. Bob on the other hand was right, Claire tells him this with sheepishness proportional to her error, and he starts gloating "I told you so!" with glee proportional to how much cleverer he feels than Claire. So Bob becomes slightly more entrenched in his opinion, and then he goes to seek out his confidantes to either congratulate or belittle them, again proportional to how right he was. When all of the blame and backpatting has backpropagated throughout society, all the shitty people have adjusted their opinions, and their shitty society will be less prone to making the same mistake again. In human terms, this process is repeated for all of recorded history or longer, and then you have a neural net that can recognise handwritten digits.

placeholder

All this process needs to get started is a lot of labelled pairs of data, input along with the desired output for that input. The formal terminology for the scenario above that converts data into performance is "training", which is a computationally intensive process when lots of data is involved for big neural nets. So the second factor of the ubiquity of neural nets is Moore's law and analogues, which have overseen exponential growth in computational power and digital data storage capacity. Neural nets convert data and compute power as fuel into practical applications, and we live in an era of increasingly plentiful data and compute. Hence, the bitter lesson []; clever theories are no match for stupid methods with lots of data and a big computer. But why the hell should any of this work in the first place? Surely there are limits to what neural nets can do. Now the third factor; Moore's law and the bitter lesson might be cheated, but the third factor is a law backed by mathematics.

Theorem 5.3.1 (Universal Approximation Theorem).

ANY PROBLEM THAT CAN BE ENCODED AS A CONTINUOUS TRANSFORMATION OF LISTS OF REAL NUMBERS INTO OTHER LISTS OF REAL NUMBERS IS POTENTIAL PREY FOR A BIG ENOUGH NEURAL NET. A little creative thought will show that many practical problems can be treated in this way. The litigious can easily spot problems in neural nets outside of this law. For example, to the best of our knowledge there is no known lower bound for how much data is required – as a function of desired accuracy within a desired confidence – for a neural net to learn its target, so for all we know, any big neural net could suddenly fail on an easy input instance for no good reason. The universal approximation theorem is a double-edged sword, and the side that cuts the holder is that for complex problems, the input data cannot span the whole problem domain, so there will be

many neural nets that agree perfectly on the training data but will perform differently out-of-distribution – this common phenomenon is called "overfitting". Moreover, a major component of explainability from the human perspective is having mental representations and rules for their manipulation, whereas it is in general very difficult to determine how or even whether the neural net has internal representations and rules for its inputs. This introduces an ethical criterion that most data-driven algorithms fail: it is not a big deal to be recommended a show we do not like or have a machine play a game more innovatively than a human could, but if the stakes are higher, such as operating a vehicle or on a patient, we would prefer safety guarantees predicated upon internal representations and rules that we can understand and negotiate in human terms. In Section 5.3.5 we will try to blunt the painful edges by using the universal approximation theorem and overfitting to our advantage in search of compositional, explainable internal representations.

WHAT IS THE DIAGRAMMATIC CONTENT OF THE UNIVERSAL APPROXIMATION THEOREM?

5.3.2 Approximating Text Circuits with deterministic neural nets

There is a lot to be gained from a process-theoretic view of interacting ensembles of neural nets. For a simple example, consider that an autoencoder is precisely a pair of neural nets trained cooperatively encode a large input space into a small latent space and decode the original input from the latent space. Diagrammatically, this amounts to asking for the equations of a split idempotent to be treated as training conditions for a pair of processes.

autoencoder

If that's what we can do with a pair of equations, what can we do with an arbitrary PROP? We first need to decide what qualifies as a valid interpretation the generators of a PROP in terms of neural nets. Not just any functor will do, because we want to rule out trivial solutions that map all processes to constant functions. We also need to put in some work to interpret what equality of processes should mean in the setting of neural nets.

Definition 5.3.2 (Approximating a (coloured) PROP). An $(\epsilon^=, \epsilon^{\neq})$ -approximation of a finitely presented coloured PROP \mathfrak{P} is a strict symmetric monoidal functor \mathcal{T} that interprets \mathfrak{P} in the (cartesian) symmetric monoidal subcategory of **Top** generated by Euclidean spaces with the usual metric as wires equipped with cartesian copy and delete, along with neural nets as processes. As a nontriviality condition, \mathcal{T} must send each wire colour in \mathfrak{P} to a Euclidean space of finite positive dimension. Equality relations presented in \mathfrak{P} are interpreted as $\epsilon^=$ -closeness by \mathcal{T} , i.e. if \mathfrak{P} stipulates that $f = g$ for $f, g : A \rightarrow B$, then we have the following inequality in the metric of $\mathcal{T}B$:

$$\forall \mathbf{x}_{\in \mathcal{T}A} : d_{\mathcal{T}B}(\mathcal{T}f(\mathbf{x}), \mathcal{T}g(\mathbf{x})) \leq \epsilon^=$$

Any PROP that equates generators directly is redundant, and we can without loss of generality restrict consideration to PROPs where each generator is implicitly assumed to be distinct. We interpret inequality as ϵ^\neq -farness, i.e., for all pairs of generators f, g of the same type $A \rightarrow B$, we ask that \mathcal{T} satisfies:

$$\exists \mathbf{y} \in \mathcal{T}_A : d_{\mathcal{T}_B}(\mathcal{T}f(\mathbf{y}), \mathcal{T}g(\mathbf{y})) \geq \epsilon^\neq$$

Since the category is cartesian monoidal, states are points in euclidean space, and the above definition specialise to treating points as "equal" if they are $\epsilon^=$ -close and "inequal" if they are ϵ^\neq -far. We choose to treat the determination of equality and inequality as separate semidecidable procedures because "equality" as we have defined it is not necessarily transitive, but we can recover a form of bounded transitivity by making $\epsilon^=$ very small compared to ϵ^\neq , so that equality is testable within a tolerance of ϵ^\neq , granting $\frac{\epsilon^\neq}{\epsilon^=}$ -fold transitivity. We can always recover decidable "equality" at the expense of transitivity by setting $\epsilon^= = \epsilon^\neq$. With that out of the way, we observe that since the target category of deterministic neural nets is cartesian monoidal, not all PROPs are approximable.

Example 5.3.3 (Not all PROPs are approximable). We take the snake equations as an example. The PROP generating the snake equation is as follows:

snakeprop

Since we are dealing with a cartesian monoidal category, the cup can only be interpreted as a pair of points, and the cap can only be a pair of deletes $[]$. The only Euclidean space in which the identity is equal to a constant map is the singleton zero-dimensional space.

trivialityproof

So nondeterminism is a necessary but possibly insufficient condition for the realisation of general PROPs. Not all is lost; if we restrict our consideration to well-behaved PROPs, such as those of simple text circuits, then we can get somewhere eventually.

Definition 5.3.4 (Basic Text Circuit PROP). A *basic text circuit PROP* has two colours of wires, N for "noun" and A for "answer". The generators fall under four main families. *Nouns* have type $1 \rightarrow N$. *Gates* have type $\bigotimes^k N \rightarrow \bigotimes^k N$ for some positive k . *Queries* have type $\bigotimes^k N \rightarrow A$ for some positive k . *Answers* have type $1 \rightarrow A$.

states, gates, queries, answers

The relations of a text circuit fall under three families. *Axioms* are equations between pairs of nonempty composites of gates; the only kind we disallow is an equality between two generators.

axiom

Instances are equations between a composite of nouns and gates and a single query on the left – a *datum* – and an answer on the right – a *label*.

instance

In addition, we ask for a special generator with a non-finite family of rules to enforce a coherence condition; we decide that distinct nouns should maintain their identity no matter what relations they participate in. The generator is *Name*, of type $N \rightarrow N$, and its relations are such that applying *Name* to any noun-wire that traces back to a noun will return that noun.

name

Example 5.3.5 (How basic text circuits PROPs may be used in practice).

A limitation that arises from the interaction of the definition of approximability and the universal approximation theorem is that any compact subset of Euclidean space can be covered by finitely many ϵ -balls. This means that the universal approximation theorem is unable to provide us guarantees if we want potentially unboundedly large text circuits, but we might reasonably expect compositional behaviour up to some notion of text circuits with bounded width and depth, which we state as follows.

Definition 5.3.6 (Bounded width and depth). We say that a basic text circuit PROP \mathfrak{T} is *terminating* if all of its axiom relations can be equipped with directions so that applying directed equality rewrites to any diagram (necessarily finite) yields a finite set of equal diagrams. As an easy example, a basic text circuit PROP with a single idempotent gate is terminating when we equip the idempotence relation with a direction that reduces the number of gates; we don't want to deal with cases where equalities explode to infinity.

idempotentreduce

Observe that if \mathfrak{T} is terminating, then applying axioms to the datum of any instance relation will also yield a finite set of equal diagrams, and each instance diagram may be rewritten by isotopies so that parallel gates are displaced so that each gate occupies a distinct level, sandwiched by a layer of nouns and query, which we also count as layers. We say that \mathfrak{T} has *bounded depth* $d \in \mathbb{N}$ if the maximum depth in layers obtained in this way from any datum in \mathfrak{T} is bounded above by d . The case of width is simpler, because axioms cannot change the number of wires in a datum, which is the same as the number of nouns; we say that \mathfrak{T} has *bounded width* $w \in \mathbb{N}$ if the maximum number of nouns that occur in any datum is bounded above by w .

Even this is not enough. Another issue arises from the interaction of approximability with the nature of cartesian monoidal categories.

Theorem 5.3.7 (Fox's Theorem).

A consequence of Fox's theorem is that in any cartesian monoidal category, we can equip every wire with canonical cocommutative comonoids (copy and delete), and every morphism in a cartesian monoidal cat-

egory is a cohomomorphism with respect to copy and delete; recall from Section 2.1.5 that this is the diagrammatic definition of deterministic maps. This consequence means that for some text circuit PROPs, approximable-equalities may hold in *any* of their interpretations as deterministic neural nets that are not equalities licensed by the original PROP.

Example 5.3.8.

The deeper essence of this issue is that in a cartesian monoidal category, every wire can at most carry data about its diagrammatic causal past, since equalities of the following sort always hold.

placeholder

This conflicts with the nature of updating representations with text, where later information may force revisions of earlier representations. For a crude example, consider this transcription of a meme about a morning routine:

Wake up. Take a shit. Eat. Get out of bed. Have breakfast.

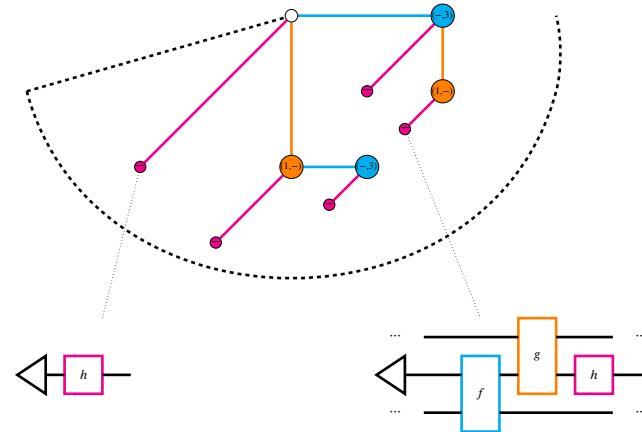
Now I will kill the joke by overthinking it. What happens in our mind's eye if we are constructing a little vignette of events? In this example, the first three clauses take a pedestrian interpretation – Taking a shit normally happens in toilets, and the inferred object of Eat is breakfast. Clauses four and five require belief revisions. The internal representation of the sequence of events up to those points must be retroactively changed in a manner that is not representable as gates updating entities locally:

cartoon

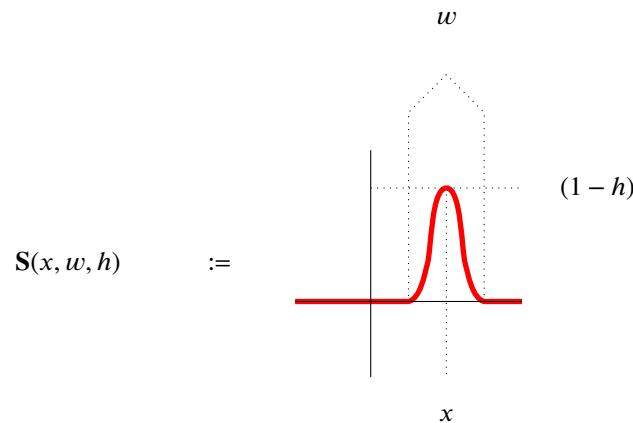
So in the deterministic setting, after every local update, we require a *global* coordination gate that gives all representations access to each other's data and a chance for them to revise themselves. To summarise our restrictions so far, we are only dealing with text circuit PROPs, of bounded depth to remain within the guarantees of the universal approximation theorem, and of fixed width as a diagrammatic consequence of having a single global coordinator gate.

Theorem 5.3.9 (Basic Text Circuit PROPs of bounded depth and fixed width are approximable using a global coordinator). *Proof.* We sketch a construction. The essential idea is to have each noun represent all of the gates it has passed through. The representation is (the monoid analogue of) a Cayley graph of a free group in $\mathbb{R}^{|\mathbf{G}|}$, where if $w \in \mathbb{N}$ is the fixed width, \mathbf{G} is the set of gates where every wire but one is assigned: $\{g(i, j, \dots, -, \dots, k) \mid g \in \text{Gates}, 1 \leq i, j, \dots, k \leq w\}$. In the illustration below, we can for instance record that a noun wire has passed through a unary gate h as the only argument, or that it has passed through three gates, first a binary f gate as the first argument, sharing the gate with the third wire, then a binary g gate as the second argument, shar-

ing the gate with the first wire, and then a unary h gate. Each distinct generator has its own dimension in Euclidean space.

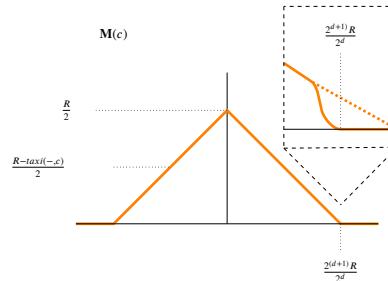


By asking for each node to be an ϵ -ball, we can test for equality by a spike function.



We don't want equal steps in each direction in the graph, or else all pairs of steps would commute. To reflect the order in which steps occur, we want subsequent steps to be smaller than earlier ones. So for each colour of node in the graph, we add the value of a move function in the desired dimension, parameterised by the initial position of the noun, so that each subsequent step is half the distance of the previous one. To make life easier, we have each move function use the taxicab metric. By bounded depth, there is some $d \in \mathbb{N}$ after which we don't need to carry representations, so we modify the ends of the move function to hit zero early. Setting $R >> 2^d \epsilon$, each move function acts to translate all graph nodes within an R -ball to their destination graph nodes after uniformly applying some generator. We ask that each wire is initialised at least $2R$ -far

apart, so each wire has their own R -ball neighbourhood in which to encode data as a graph.



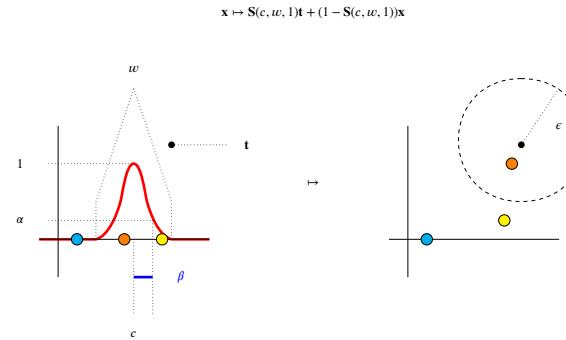
Now we can jointly specify representations of nouns and gates. For example, suppose f is a binary gate, and that $w = 3$. Then there are six shift functions

$$\mathbf{M}_f(-, 1), \mathbf{M}_f(-, 2), \mathbf{M}_f(-, 3), \mathbf{M}_f(1, -), \mathbf{M}_f(2, -), \mathbf{M}_f(3, -)$$

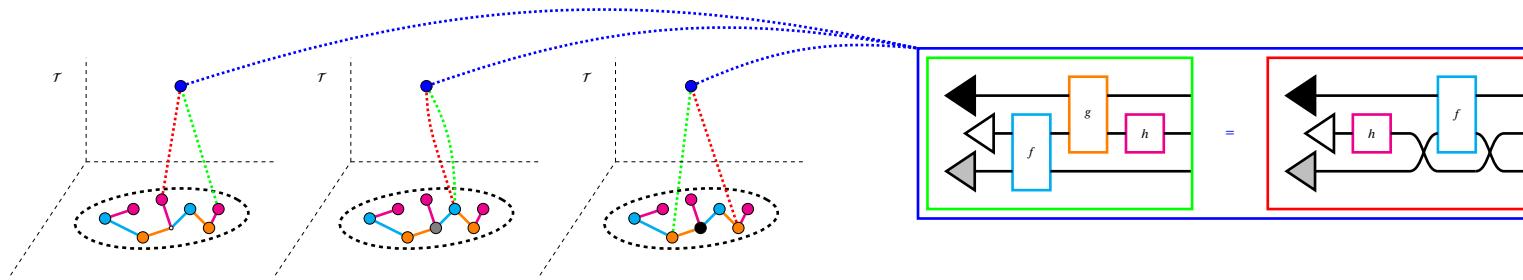
Set an ansatz dimension \mathcal{A} aside, which distinguishes noun wires by indicator-spike functions $\mathbf{1}_1, \mathbf{1}_2, \mathbf{1}_3$ – a continuous analog of one-hot encoding – then let f be:

$$f(x, y) := \begin{cases} x \mapsto x + \sum_{i^x, i^y=1}^w (\mathbf{1}_{i^x}(x) \mathbf{1}_{i^y}(y) \mathbf{M}_f(-, i^y)) \\ y \mapsto y + \sum_{i^x, i^y=1}^w (\mathbf{1}_{i^x}(x) \mathbf{1}_{i^y}(y) \mathbf{M}_f(i^x, -)) \end{cases}$$

In prose, the coefficients obtained by the indicator functions of each summand behave as a case guard, identifying the wire-numbers of x, y by sending all other cases to zero. This generalises to representations of gates with higher arity. At this point, we have obtained representations for a free theory of just the gates of bounded-depth and fixed-width text circuits – we have not yet incorporated the potential equalities of the theory. First we deal with the axiom-type equalities. Here is where we make use of a global coordinator gate; where width is w and each noun is represented in some space \mathbb{R}^J , the global coordinator is typed $\prod_w^w \mathbb{R}^J \rightarrow \prod_w^w \mathbb{R}^J$. The global coordinator looks at the local data recorded by each wire, and when it spots an equality in the theory, it pinches free graph nodes together. The pinch function is a modification of the spike and shift functions, moving only points within a tolerance β into a target ϵ -ball.



The global coordinator waits to see if the product state of all the wires satisfies an equality, and when it does, it pinches the agreeing free representations on each wire to a new node along a new dimension – we ask for one dimension for each equivalence class of expressions quotiented by equality in the theory, and here the bounded-depth and fixed-width conditions mean that we only have to add finitely more dimensions in this way.



The target location of each pinch can be chosen to be compatible with the shift functions that generate the graphs – i.e. each equality-class is treated as if it were its own generator with its own shift function. The global coordinator is the sum of all such pinches, and the way to use it is to sandwich it between each gate.

placeholder

Finally, the query morphisms in the instance-type rules of the text circuit PROP are continuous functions that assign each ϵ -ball node a value in a Euclidean answer-space, which can mimick a finite discrete number of cases by spike-indicator functions. All of this is, by construction, an (ϵ, R) -approximation of a text circuit PROP. \square

5.3.3 Text circuits with unbounded depth and width

For a better understanding of what is possible, we might approach "from above", seeking first a mathematical setting in which we have unrestricted compositionality, and then gradually applying sensible constraints.

One such mathematical setting is **Rel**, which we show can faithfully model any finitely presented coloured PROP.

Definition 5.3.10 (Faithful models and expressive completeness). Given a coloured prop \mathfrak{P} , a symmetric monoidal category \mathcal{M} , and a symmetric monoidal functor $T : \mathfrak{P} \rightarrow \mathcal{M}$, we say that T is a *faithful model of \mathfrak{P} in \mathcal{M}* if, for all diagrams $f, g \in \mathfrak{P}$, $f =_{\mathfrak{P}} g \iff T(f) =_{\mathcal{M}} T(g)$. Given a collection of props $\{\mathfrak{P}_i\}$, we say that \mathcal{M} is *expressively complete for $\{\mathfrak{P}_i\}$* when each \mathfrak{P}_i has a faithful model in \mathcal{M} .

Theorem 5.3.11 (\mathbf{Rel}^\times is expressively complete for all finitely presented coloured PROPs). *Proof.* Our strategy has two main ideas. First is to take the Lindenbaum-Tarski algebra of diagrams for an arbitrary \mathfrak{P} , quotiented by the equivalence relation $=_{\mathfrak{P}}$; this will give faithfulness. Second is to treat the sought functor T for \mathfrak{P} as a category of elements construction, adapted to the symmetric monoidal setting. Let \mathfrak{P}^* denote the finitely presented coloured PROP \mathfrak{P} augmented with new generators that do not obey any relations: a state o and effect o for each colour C in \mathfrak{P} . Let \mathfrak{P}_{LT}^* denote the set of diagrams of \mathfrak{P}^* quotiented by the equivalence relation obtained by equality $=_{\mathfrak{P}}$ in \mathfrak{P} . Let $T(0)$ be the singleton. For each colour $C \in \mathfrak{P}$, let $T(C)$ be the subset of \mathfrak{P}_{LT}^* that selects all states on C . For each nonempty list of colours $[C_i]$, let $T([C_i])$ be the subset of \mathfrak{P}_{LT}^* that selects all states on $C_1 \otimes_{\mathfrak{P}} C_2 \otimes_{\mathfrak{P}} \dots \otimes_{\mathfrak{P}} C_i$. For $f : C \rightarrow D$ in \mathfrak{P} , let $T(f) := \{(\Gamma(c), \Gamma(c); f^\top) \mid c \in T(C)\}$, where corner notation denotes an equivalence class of states under $=_{\mathfrak{P}}$; as a particular consequence, for a state $\langle C |$ on C , $T(\langle C |) = \langle C | \subseteq T(C)$. Note that applying f after any state $\langle C | \in T(C)$ yields a state $\langle C | ; f^\top \in T(D)$, and that these states include actual states native to \mathfrak{P} and formal states from \mathfrak{P}^* where diagrams from \mathfrak{P} with an output wire typed C have their other wires "capped off" by o and o .

T is a functor; $T(1_C) = 1_{T(C)}$ since $\{(\Gamma(c), \Gamma(c); 1_C^\top) \mid \Gamma(c) \in T(C)\}$ is the identity relation on $T(C)$; $T(f) ;_{\mathbf{Rel}} T(g) = T(f ;_{\mathfrak{P}} g)$ since the relational composite is

$$T(f) ;_{\mathbf{Rel}} T(g) := \{(\Gamma(c), \Gamma(e)^\top) \mid c \in T(C), \exists \Gamma(d) \in T(D) : \Gamma(c) ; f^\top = \Gamma(d) \& \Gamma(d) ; g^\top = \Gamma(e)^\top\}$$

We observe that $\Gamma(d) \in T(D) : (\Gamma(c) ; f^\top = \Gamma(d) \& \Gamma(d) ; g^\top = \Gamma(e)^\top)$ implies that $\Gamma(c) ; f ; g^\top = \Gamma(e)^\top$, so we have $T(f ;_{\mathfrak{P}} g) \subseteq T(f) ;_{\mathbf{Rel}} T(g)$. For the other inclusion, we observe that $\Gamma(c) ; f^\top$ yields the state $\Gamma(c) ; f^\top \in T(D)$ in the bracketed expression, thus satisfying the existential quantifier.

The unitors, associators, braidings, and coherences are tedious to write but conceptually trivial, so I will skip them. The tricky part of showing that T is monoidal is providing isomorphisms $T(C) \times T(D) \simeq T(C \otimes_{\mathfrak{P}} D)$. We present them first and comment later.

$$\begin{aligned}
T(C) \times T(D) \rightarrow T(C \otimes_{\mathfrak{P}} D) &:= \{ \left(\begin{pmatrix} \Gamma \langle c | \sqcap \\ \Gamma \langle d | \sqcap \end{pmatrix}, \begin{cases} \Gamma \langle x | \sqcap & \text{if } \exists \langle x | \in T(C \otimes_{\mathfrak{P}} D) : \Gamma \langle c | ; \neg_C \sqcap = \Gamma \langle x | ; (\neg_C \otimes_{\mathfrak{P}} \neg_D) \sqcap = \Gamma \langle d | ; \neg_D \sqcap \\ \Gamma \langle c | \otimes_{\mathfrak{P}} \langle d | \sqcap & \text{otherwise} \end{cases} \right) \mid \Gamma \langle c | \sqcap \in T(C), \Gamma \langle d | \sqcap \in T(D) \} \\
T(C \otimes_{\mathfrak{P}} D) \rightarrow T(C) \times T(D) &:= \{ \Gamma \langle x | \sqcap, \begin{cases} \begin{pmatrix} \Gamma \langle c | \sqcap \\ \Gamma \langle d | \sqcap \end{pmatrix} & \text{if } \exists \Gamma \langle c | \sqcap \in T(C) \Gamma \langle d | \sqcap \in T(D) : \Gamma \langle x | \sqcap = \Gamma \langle c | \otimes_{\mathfrak{P}} \langle d | \sqcap \\ \begin{pmatrix} \Gamma \langle x | ; (1_C \otimes_{\mathfrak{P}} \neg_D) \sqcap \\ \Gamma \langle x | ; (\neg_C \otimes_{\mathfrak{P}} 1_D) \sqcap \end{pmatrix} & \text{otherwise} \end{cases} \mid \Gamma \langle x | \sqcap \in T(C \otimes_{\mathfrak{P}} D) \}
\end{aligned}$$

We walk through the construction case-by-case. For the first case top-to-bottom, if two states $\langle c |, \langle d |$ are obtainable by formally deleting the C and D outputs (using \neg) of some state $\langle x |$ on $C \otimes_{\mathfrak{P}} D$, then we send the pair $(\Gamma \langle c | \sqcap, \Gamma \langle d | \sqcap)$ to $\Gamma \langle x | \sqcap$. Note that in the first case, even if $\langle x |$ is tensor separable as $\langle c | \otimes_{\mathfrak{P}} \langle d |$, formal deletion creates new formal scalars which must also agree by the case guard. The total effect of the first case is to identify when $(\langle c |, \langle d |)$ arise as partial diagrams (where ends are truncated with \neg) of some state $\langle x |$. The second case is where $(\langle c |, \langle d |)$ are not partial diagrams of some $\langle x |$ in this way, in which case we send the pair to their tensor product. The third case is the inverse of the second, sending all $\langle x |$ that are equivalent to a tensor-separable composite state $\langle c |$ and $\langle d |$ to that pair of states. The fourth case is the inverse of the first; if $\langle x |$ has no tensor-separable equivalent, then we project $\langle x |$ to states on C and D by formally deleting the appropriate outputs. The effect of taking equivalence classes makes the first relation as a whole an injective function, and likewise for the second relation.

All that remains is to show that T is a faithful model, i.e., for all lists of colours $[C], [D]$ and for all $f, g :$
 $[C] \rightarrow [D]$, $f =_{\mathfrak{P}} g \iff T(f) =_{\text{Rel}} T(g)$. For the forward direction \Rightarrow , if $f =_{\mathfrak{P}} g$, then for any state $\langle x |$:
 $0 \rightarrow [C]$, we have that $\langle x | ; f = \langle x | ; g$. Because $T(\langle x |) = \Gamma \langle x | \sqcap \in T([C])$, we have $\Gamma \langle x | ; f = \langle x | ; g \sqcap \in T([D])$,
thus $T(f) =_{\text{Rel}} T(g)$. For the converse direction \Leftarrow , if $f \neq_{\mathfrak{P}} g$, then by the Lindenbaum-Tarski construction and the relational-freeness of the generator \neg , $\Gamma \neg ; f \sqcap \neq \Gamma \neg ; g \sqcap$, so $T(f)$ as a relation contains the pair
 $(\Gamma \neg_{[C]} \sqcap, \Gamma \neg_{[C]} ; f \sqcap)$ that is not present in $T(g)$, and symmetrically for $(\Gamma \neg_{[C]} \sqcap, \Gamma \neg_{[C]} ; g \sqcap)$, thus $T(f) \neq T(g)$. \square

Corollary 5.3.12 (FinRel will work for bounded-size compositionality). Just limit consideration to diagrams with upper bounds on the maximal number of generators among diagrams in their equivalence classes, or whatever reasonable finiteness constraint you want, I'm not a cop.

5.3.4 Text circuits of unbounded width in noncartesian settings

WHAT'S SPECIAL ABOUT **Rel** THAT DISTINGUISHES IT FROM A CATEGORY OF DETERMINISTIC PROCESSES? For one, **Rel** is not cartesian monoidal; while the monoidal product in **Rel** is the categorical product of sets, and every object in **Rel** has a copy-delete comonoid, not every relation is a cohomomorphism with respect to this comonoid (in fact, only the functions are.) It turns out that this property is what allows the existence of tests that differ from deleting. In a cartesian monoidal category, every test is equal to deleting.

placeholder

In a noncartesian monoidal category such as **Rel** or **Stoch** – the category of stochastic maps between measurable sets, we may have tests corresponding to postselection, which greatly increases our expressive capacity for composing constraints. Let's consider an example in **Stoch**, where states are probability distributions, for example a coinflip as a state on the space $\{H, T\}$. Observe that copying the outcome of a coinflip gives a probability distribution on $\{H, T\} \times \{H, T\}$ that is different from what we obtain by flipping two independent coins, as the first is perfectly correlated and the second is not. So **Stoch** is not cartesian monoidal, while it does have a copy-delete comonoid for every object.

placeholder

The scalars in **Stoch** correspond to probabilities in the range $[0, 1]$. Another way we could generate perfectly correlated coinflips is to flip two coins and throw away any outcomes that are not either $H \times H$ or $T \times T$; diagrammatically, this corresponds to copying the outcomes of the two coins and applying a test to one set of copies. Postselecting will yield outcomes will equal in distribution to copying a single coinflip, but only half of the time.

placeholder

Postselection gates are commutative, so in a circuit made up entirely of postselection gates, the only relevant size measure is the width of the circuit.

placeholder

Although it is a very stupid method, one way to achieve unbounded-width compositionality using postselection gates is to just continue sampling randomly until you find a sample that passes all the postselection tests. For example, suppose we have a group of people $\{\text{Alice}, \text{Bob}, \text{Claire}, \dots\}$ standing in a queue modelled as a unit interval, and we have one binary linguistic relation X is-in-front-of Y , which postselects with con-

dition $X > Y$ for $X, Y \in [0, 1]$. Assume every person's initial state is the uniform random distribution on $[0, 1]$. Then we can model the text Alice is in front of Bob. Claire is in front of Alice. as:

placeholder

Any postselected sample in this case is a satisfying instance of our linguistic positions: an assignment of possible positions in the queue to people such that Alice really is in front of Bob and Claire is really in front of Alice. In this case, we can have any finite number of people and it will remain true that if we manage to obtain a postselected sample, it will be a configuration that satisfies our constraints. Moreover, no globally coordinating gate is required in this example. So by introducing random variables as initial states and posts-election, we have shown that it is (in principle, and in some cases) possible to obtain approximations for text circuit PROPs using deterministic neural nets without global coordination gates and without the restriction of fixed width.

5.3.5 A value proposition for quantum machine learning

Since I am from the quantum group in my department, I should probably write something extolling quantum computers, so here is the token gesture. A practicality issue arises for the random sampling approach in the form of postselection, which is the procedure of throwing away samples that do not adhere to the constraints we have set and rerolling the dice. For very specific constraints that jointly independent samples are very unlikely to satisfy, we may have to reroll many times before we obtain a satisfying instance, which may take longer than we are willing to wait. While I am sure there are clever ways to mitigate the problem of postselection on a classical computer beyond my current knowledge, let me suggest a *possibility* to avoid postselection altogether; that is, every measurement will always return a satisfying instance. The main ingredient is a parameterised quantum circuit, which consists of unitary gates controlled by classical parameters; just as a neural net is a classically-parameterised process where the parameters determine weights and biases per neuron, classical parameters can determine the preparation of e.g. phase states within a unitary quantum gate. We only want to traffic in causal quantum processes so we never have to worry about postselection, i.e. the only test we ever want to apply is discarding. So without loss of generality we express each parameterised gate as the following composite by Stinespring dilation.

placeholder

That is, we deal with gates that consist of classically-parameterised unitaries and a classically-parameterised ansatz state on ancilla wires that we discard. Since the only test we ever apply is discarding, we do not have to postselect. If we can find classical parameters for each gate such that the equations of a text circuit PROP are satisfied, then we would have a faithful model that would in principle keep the benefit of width-

unbounded compositionality as in the probabilistic case, but without the drawbacks of postselection. To be balanced, I should qualify how big of an "if" we are dealing with. There are at least three obstacles that spring to mind. First is practical: the space of parameters for parameterised quantum circuits are not easily differentiable against a loss function unlike their classical counterparts and batching is difficult when dealing with ensembles of parameterised processes in different configurations, so training takes more time. Second is theoretical: there is more work to be done to understand what kinds of text circuit PROPs may be modelled faithfully in principle by quantum and classical gates, and in particular whether bounded-depth compositionality transfers from **FinRel** to **FdHilb** with quantum processes. Third is mystical-heuristical: we can approximate that finding a satisfying instance by random sampling on a classical computer takes something like $\mathcal{O}(e^N)$ samples where N is the number of gates, so if we believe in a complexity-theory analogue of "no free lunches", finding appropriate parameters for an ensemble of quantum or classical gates may be around that hard as well.

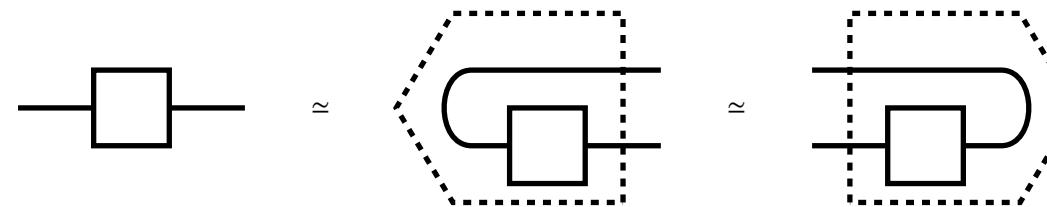
5.4 Towards learning gates that satisfy First-Order Logic specifications over boundedly finite models

Formal semanticists love logic, so I will throw them an unsatisfying bone. I will sketch a bridge between FOL over boundedly finite models and machine learning. We have seen so far how equational constraints between interacting processes can shape their behaviour in desired ways. Moreover, it is evident that when the wires are equipped with a metric, we may use equational constraints as a loss function to train gate-representations. In the worst case, we must perform raw stochastic gradient descent across a massive ensemble of gates such that they all satisfy the equations we want of them, but this is good enough for learning finite models for first-order logic using gradient descent *in principle*.

The setup is this: suppose that we are in the process of learning representations for words like Professor and Student and Likes. How do we interpret or enforce statements like Every professor likes some student? It is already known how first-order logical theories are interpreted diagrammatically [] as equations between a scalar corresponding to a closed first order logic sentence and the empty scalar, which, when applied as a constraint to a functor into **Rel**, yields a set-theoretic model. My contribution here is providing an alternative formulation that breaks up large scalars into more piecemeal equational constraints, limited to boundedly finite and typed set-theoretic models.

The first obstacle we encounter is that there are no symmetric monoidal functors from **FinRel** $\{\times\}$ to **FinVect** $\{\otimes_{\mathbb{R}}\}$, where the latter is closer to where machine learning tends to happen, once nonlinear activation functions are introduced. The reason for this is that while **FinRel** can be viewed as **FinVect** over the boolean semiring, the boolean semiring has idempotent addition and multiplication ($a \wedge a = a$), while \mathbb{R} doesn't. Moreover, **Rel** is discrete, and we would like some kind of continuous semiring so that gradient descent can do something. So we just observe that all we need in order to define matrix multiplication is a semiring, and that the boolean semiring embeds into the $(\max, \times, [0, 1])$ semiring, which is isomorphic by natural log to signed restrictions of the min-plus and max-times tropical semirings [] ($\min, +, \mathbb{R} \{ \geq 0 \cup \{\infty\} \}$) and $(\max, +, \mathbb{R} \{ \leq 0 \cup \{-\infty\} \})$). So denoting these restricted tropical semiring \mathbb{T} , we have an evidently faithful symmetric monoidal functor from **FinRel** $\{\times\}$ into **(SemiRing)FinVect** $\{\otimes_{\mathbb{T}}\}$. So what we have achieved so far is to describe a setting in which gradient descent can take place to learn equational constraints presented in **FinRel** with the \times monoidal product.

Before we get started with diagrams, just a reminder that in **Rel** or **ContRel** restricted to interpret **FinRel** with sticky-spiders, we don't have to care about the direction on the page in which wires go, because we can always use spiders to turn inputs into outputs and vice-versa.



So without loss of generality we present equational constraints taking predicates to be effects. We additionally assume that we may have different wire-types, corresponding to a type-guarded FOL. A modelling assumption we take is that the non-zero copyable states or points of a wire correspond to the individuals in the FOL-model. Now we can present some equational constraints and their corresponding interpretation in first-order logic. Let's start with some simple quantifiers.

Proposition 5.4.1.

$$\begin{array}{ccc} \text{Diagram: } & = & \exists x_A : Rx \\ \text{A blue line labeled } A \text{ enters a pink triangle labeled } R & & \end{array}$$

Proof. If there exists some point x_A such that Rx , then the equation on the left holds. If there does not exist any x_A such that Rx , the RHS must equal the zero scalar.

Proposition 5.4.2.

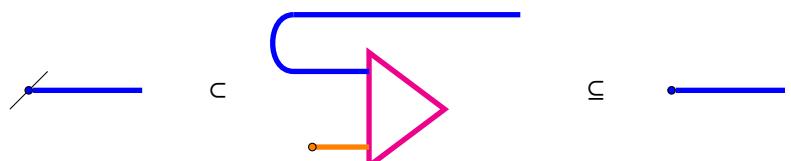


Proof. This is just how R is interpreted as a subset of A if it is the case that all x_A satisfy Rx .

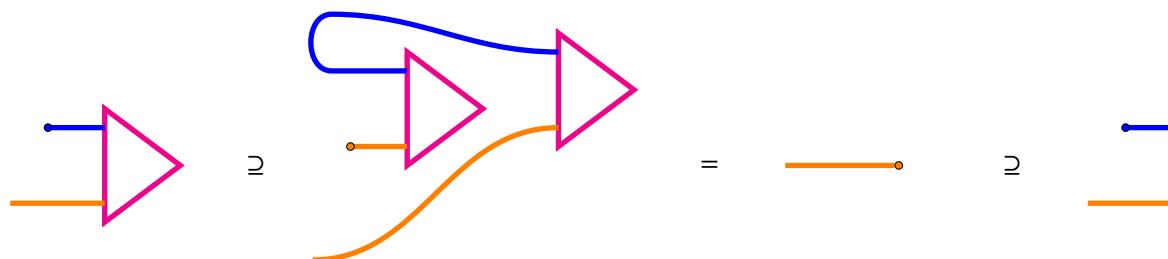
Proposition 5.4.3.



Proof. Suppose the equation holds. Then plugging in any individual state y_B yields the unit scalar, which means there is some x_A holds. Then we can diagrammatically construct the set of all x_A such that Rxy for some y . We know from the FOL this set cannot most all of A :



So starting with the LHS of the equation, we first have an inequality from above, followed by a definitional equality: since $\forall y_B \exists x_A$ of A s must be all of B . Finally, we have that "all of B " is the maximal subset of B , which altogether yields our equality.



To get at $\exists x \forall y$, we need some extra help and trickiness. Observe that since $\exists x \forall y : Rxy$ entails $\forall y \exists x : Rxy$, so we are looking for that will force R to assign at least one x for which all $y Rxy$. Recall that we're essentially dealing with one-hot encodings of subsets.

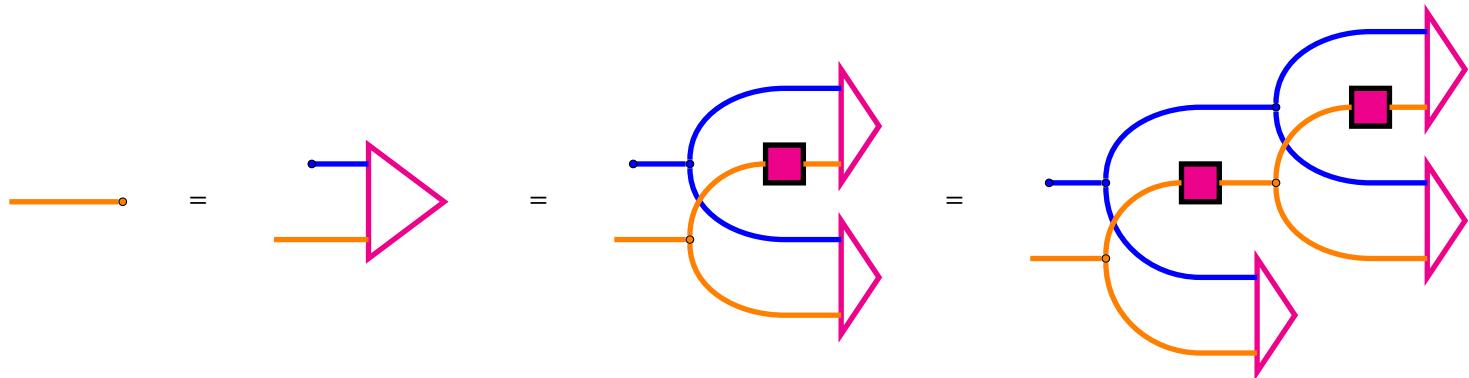
corresponds to an individual, and that bases may be permuted. Let's suppose we are granted a "cycle" function for every finite collection of wires, the job of which is to permute all of the copyable individuals or basis vectors, such that if there are N copiables, or equivalently that states on the wire represent subsets of a set of size N , then we want the following diagrammatic equations to hold:

$$\begin{array}{c} \text{Top Row: } \\ \text{Left: } \text{A wire with a square node branching into two wires} = \text{A wire with a square node followed by a cycle node} \\ \text{Right: } \text{A sequence of } N \text{ wires with square nodes} = \text{A single horizontal line} \\ \text{Bottom Row: } \\ \text{Left: } \text{A wire with a square node followed by a dot} = \text{A single horizontal line with a dot} \\ \text{Right: } \text{A sequence of } N \text{ wires with square nodes, where the } i\text{-th and } j\text{-th wires are swapped} = \text{A sequence of } N \text{ wires with square nodes where the } i\text{-th and } j\text{-th wires are swapped} \\ \qquad \qquad \qquad \forall i \neq j \in [0, N) \end{array}$$

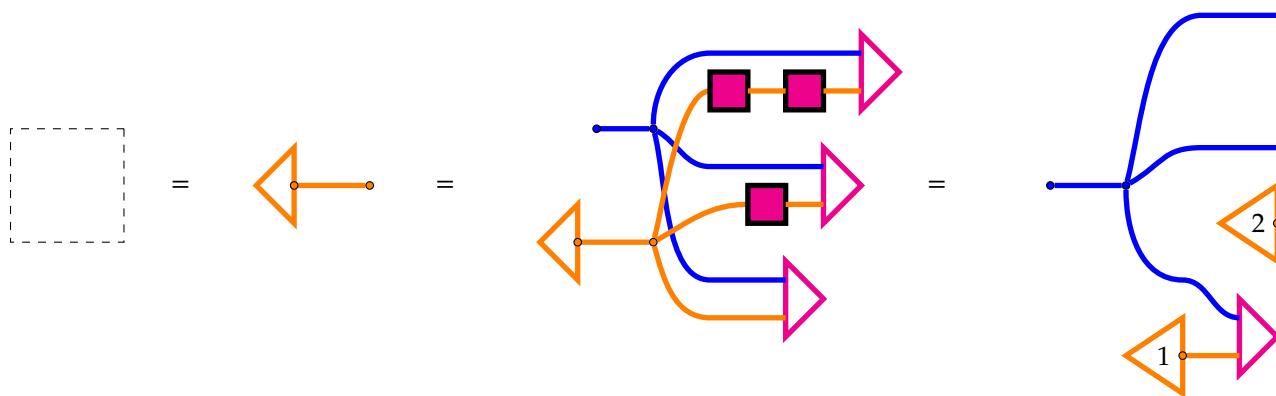
Proposition 5.4.4.

$$\begin{array}{c} \text{Left: } \text{A wire with a square node and a blue wire} = \text{A wire with a dot} \\ \text{Middle: } \text{A wire with a square node and an orange wire} \& \text{A wire with a square node and a blue wire} = \text{A wire with a dot} \\ \text{Right: } \text{A complex loop structure involving blue and orange wires} \Rightarrow \exists x_A \forall y_B : Rxy \end{array}$$

Proof. Note that we only need to prove one direction this time! Suppose the two equations hold. Then we have the following:



We can keep on going beyond 3 copies of R to eventually get N copies of R , where each copy-branch of the B wire has a distinct y_B . We'll keep it at 3 for diagrammatic simplicity. If we plug in an individual y_B into this N -fold copy expression, by design of the cycle



From left to right, the first equality holds because deleting a non-zero copiable gives the unit scalar. The second equation follows by copying and cycling. The final diagram states that for all 3 distinct elements of the set B , there is some commutativity precisely the FOL.

Proposition 5.4.5. We can obtain diagrammatic equational constraints for any quantifier ordering $\mathbf{Q}_1 x_A, \dots, \mathbf{Q}_k x_Z : \Phi(x_A, \dots, x_Z)$ $\Phi(x_A, \dots, x_Z)$ that only contains \wedge and positive atoms.

Proof. Proposition 5.4.3 applies for any quantifier prefix that has all \forall s followed by \exists s, by duplicating wires and counits. Proposition 5.4.4 applies for any quantifier prefix that has all \exists s followed by \forall s, by moving quantified variables past a universally quantified one in the quantifier order; in the case of nonparticipating wires, we cap them off. We then apply the equation of Proposition 5.4.3 then apply additional equations using Proposition 5.4.4 as appropriate to obtain the desired quantification.

All that's left to deal with is negation. Unfortunately negation – or complementation in the language of sets – is not linear by any means.

negation by learning a new relation for whatever it is we want to negate.

Proposition 5.4.6.

$$\begin{array}{c}
 \text{---} \quad \oplus \quad \text{---} \\
 \diagup \quad \diagdown \\
 P \qquad Q
 \end{array}
 = \quad \text{---} \bullet \quad \& \quad \text{---} \quad \text{---} \quad \text{---} \quad \Leftrightarrow \quad Q = \bar{P}$$

Proof. We just have to explain notation. For the first equation, set union in $\text{FinRel}\{\times\}$ maps to the direct sum in $(\text{SemiRing})\text{FinVect}\{\otimes_{\mathbb{T}}\}$, so we are just asking that the union of the original (possibly composite) predicate P along with a new predicate Q gives the whole set, which is one half of complementation. For the second equation Composing two effects by copy map gives their intersection, so we are asking that the intersection of the original relation P and the new predicate Q is empty, and together the two equations characterise what complementation is. \square

Corollary 5.4.7 (Learning FOL). Provided cycling functions and the described one-hot encoding of type-guarded FOL over boundedly finite models in $(\text{SemiRing})\text{FinVect}\{\otimes_{\mathbb{T}}\}$, equational constraints between diagrams suffice to capture FOL-conditions on learnt predicates.

5.4.1 Discussion

I named this section *towards* learning gates, and now it's time to unpack this qualifier and explain some technical choices. First, the Kronecker product of two finite dimensional vector spaces yields a vector space with dimension equal to the product of the input dimensions, rather than the sum of dimensions in the case of direct sums. Rather misleadingly for category theorists and physicists, the "tensor" in TensorFlow doesn't refer to this kind of generic tensor with possibly exponential dimension blowup, so to the best of my knowledge there are no out-of-the-box methods to put what I've written into practice.

Second, because the Kronecker product is so thicc and training is hard enough as-is, it seemed necessary to me to split up the single-scalar representations of FOL from [] into smaller pieces, where there is at most one box to have weights to be adjusted for each equation. For linguistically useful but logically-redundant things such as conditionals, two boxes have to be learnt at once, which is possibly technically hard, but graphically easy, since $x \Rightarrow y \iff x \wedge y = x$, and we know already how to express intersection by composing effects with copy. While I'm on the topic of logical presentations, I'll also remark for logicians that I'm dealing with non-empty finite models.

Third, the Kronecker product requires certain concessions and opens certain possibilities. Computer scientists may have been confused at the roundabout and costly way of defining negation, when it is surely perfectly reasonable to sandwich a function $x \mapsto 1 - x$ around the relation to be negated. The reason we couldn't just do this is because introducing such a nonlinearity breaks the categorical definition of tensor product, and the diagrammatic syntax would no longer be safe: it is not clear how to unambiguously define the Kronecker product outside of a fundamentally linear setting. So on the one hand we've reduced the problem of learning finite models of FOL to linear regression – the older cousin of ML – but on the other hand we still have to deal with potentially very high-dimensional vector spaces. Here is purportedly where quantum computers have a potential advantage over classical computers, since qubits compose in parallel by the kronecker product. However, even if there were some way to translate this story wholesale to the realm of ZX-calculus, it is still extremely difficult in complexity-theoretic terms just to test that quantum circuits are equal []; intuitively the reason is that the input space is exponentially large in the number of qubits, and in order to check whether two circuits are equal requires sampling to cover the whole input space and comparing the outputs. This seems like a roadblock, but I don't know enough about quantum

computing or machine learning to make an informed assessment of feasibility.

Fourth, it is unclear how vector representations in \mathbb{R} will behave when compressed into \mathbb{T} . Optimistically, Tropical mathematics in a machine learning context [], and rather nicely, tropical polynomial curves are precisely piecewise-linear approximations to convex regions. There is a latent connection to some ongoing work that seeks to bring together Gärdenfors's work on concepts with machine learning []. Nevertheless, this needs to be done in order to import vector-representations for words into the tropical and kronecker settings.

Fifth, the initial constraint of seeking a functor from **Rel** to some variant of **Vect** may be misguided; I obeyed this constraint for consistency, but string diagrams could remain central. It could very well be that there are far better ways to learn FOL-representations in practice than neural nets, and it could be that there are fine string diagrams somewhere that allow for nonlinear operators to behave as if they had linear ones. I suspect that there is something about FOL that necessitates this kind of space-costly representation, but it's more likely that the real problem is that I'm not really that brainy and I don't really know of (or care about) any alternative routes.

That is all the dirty laundry I could think of for this sketch, so there you have it: a highly impractical sketch bridge between FOL and string diagrams. FOL is passé for formal semantics, it's a start. If nothing else, we might conjecture that diagrammatic equations are all you need.

5.5 Modelling metaphor

I will take a *metaphor* to be text where systematic language for one kind of concept is used to structure meanings for another kind of concept where literal meanings might not apply. This may subsume some cases of what would otherwise be called *similes* or *analogies*.

THE IDEA: First, we observe that we can model certain kinds of analogies between conceptual spaces by considering structure-preserving maps between them. For example, Planck's law gives a partial continuous function from part of the positive reals measuring Kelvin to wavelengths of light, and the restriction of this mapping to the visible spectrum gives the so-called "colour temperature" framework used by colourists.

Second, we observe that we can use simple natural language to describe conceptual spaces, instead of geometric or topological models. Back to the example of colour temperature, instead of precise values in Kelvin, we may instead speak of landmark regions that represent both temperature and colour such as "candle", "incandescent", "daylight", which obey both temperature-relations (e.g. candle is a lower temperature than incandescent) and colour-relations (daylight is bluer than incandescent).

Third, we observe that we can also use simple natural language to describe more complex conceptual schemes with interacting agents, roles, objects, and abilities. Organising this linguistic data in the concrete structure of a text circuit allows us to formally specify what it means for one conceptual scheme to structure another by describing structure-preserving maps between the text circuits. This will allow us to reason about and calculate with metaphors such as "Time is Money".

THE MOTIVATION:

5.5.1 Orders, Temperature, Colour, Mood

5.5.2 Complex conceptual structure

TIME IS MONEY

"Do you have time to look at this?"

"This is definitely worth the time!"

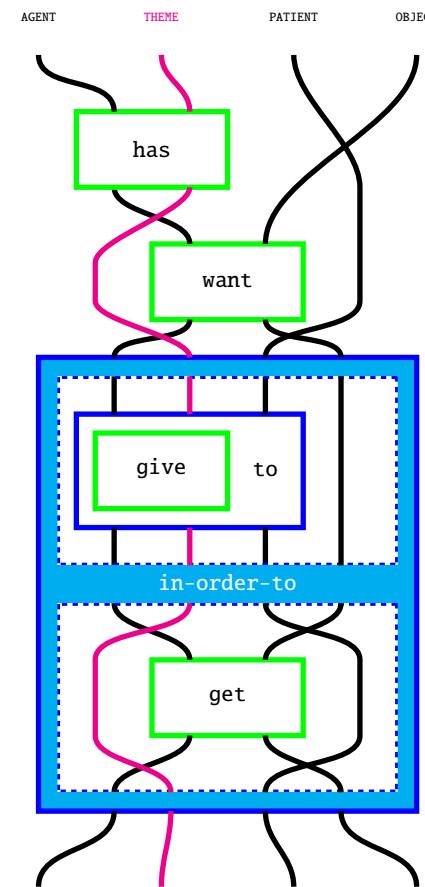
"What a waste of time."

I will work through an example of partially using the concept of Money to structure that of Time. Part of the concept of Money is that it can be *exchanged* for something. The concept of exchange can be glossed approximately as the following text, with variable noun-entries capitalised.

AGENT has THEME.

AGENT wants OBJECT.

AGENT gives THEME to PATIENT in-order-to get OBJECT.



The

5.5.3